

September 12, 2024

6.S894

Accelerated Computing

Lecture 2: Throughput Processors

Jonathan Ragan-Kelley 

What is a processor?

a programmable computer

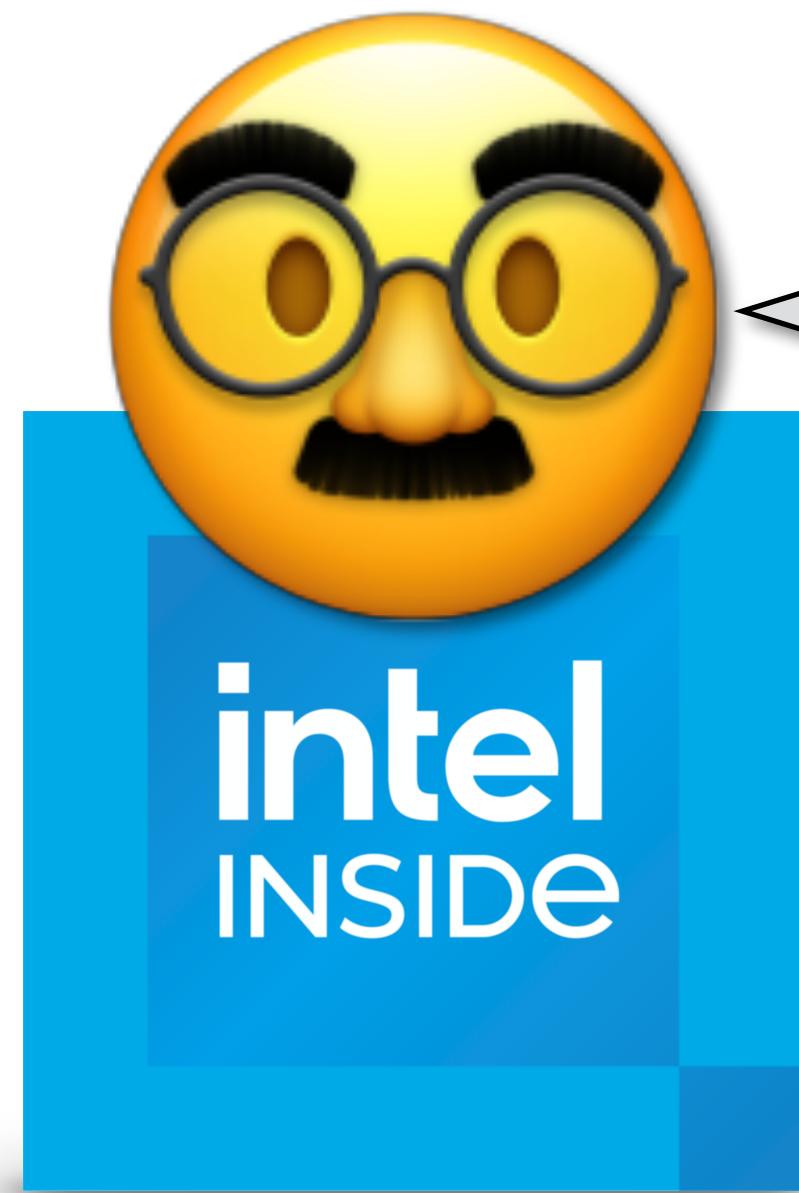
that runs a sequence of
instructions over time

including control flow,
computation & state updates

What is a processor?



1970s



2020s

Why yes,
I am still
basically a
PDP-11

A simple processor

1. Instr. fetch
2. Decode
3. Operand fetch
4. Execute
5. Write back

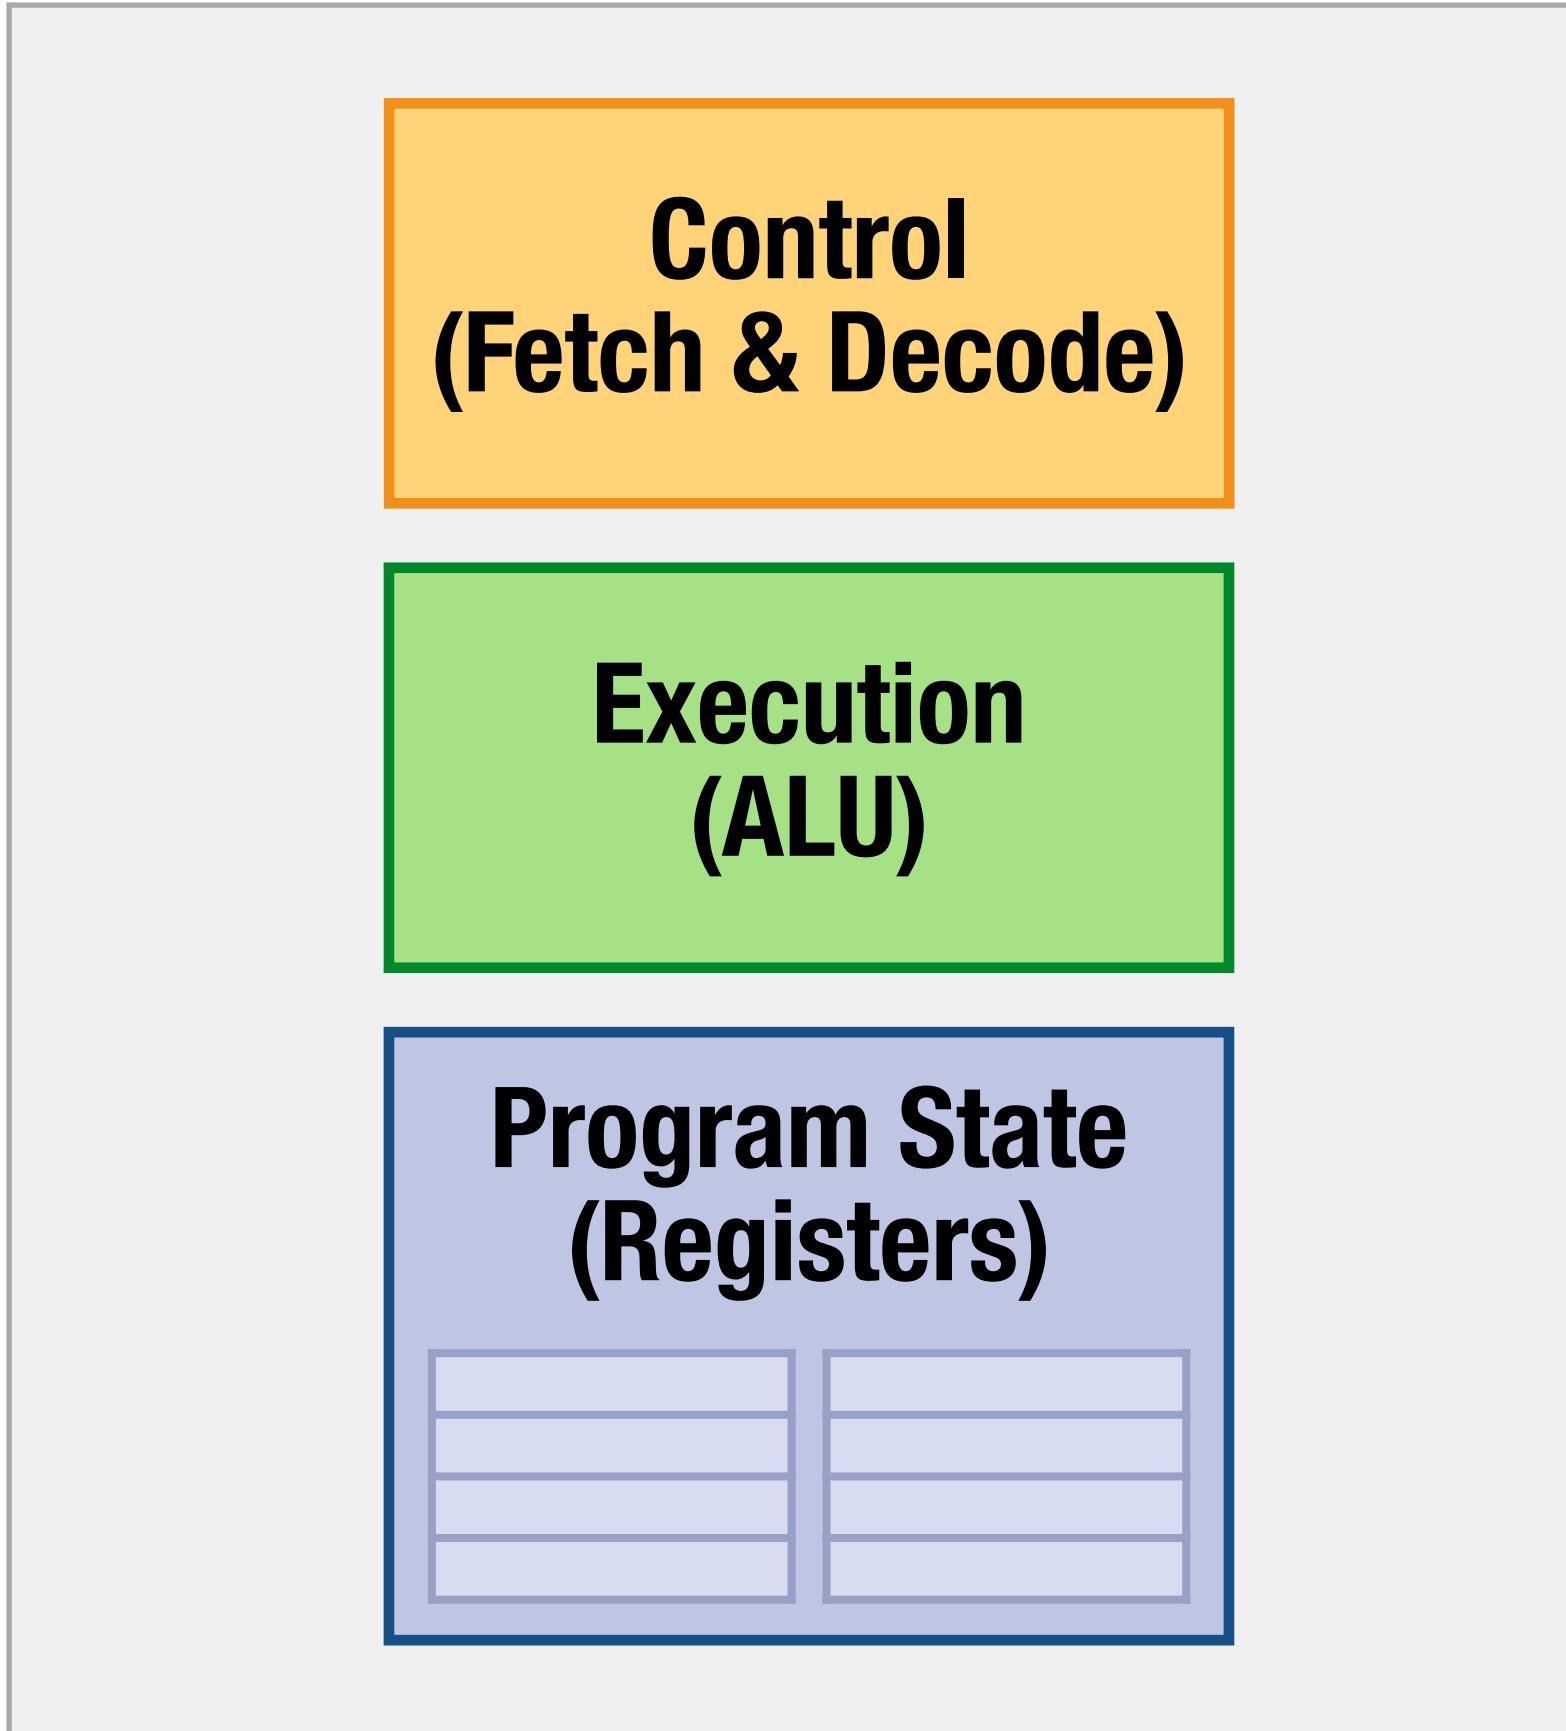
Processor

an interpreter for instructions!

ld	r3, mem[r0+r2]
ld	r4, mem[r1+r2]
mul	r3, r3, r4
add	r5, r5, r3
addi	r2, r2, 4
blt	r2, \$400, LOOP
st	addr[r6], r5
...	

Program

A simple processor



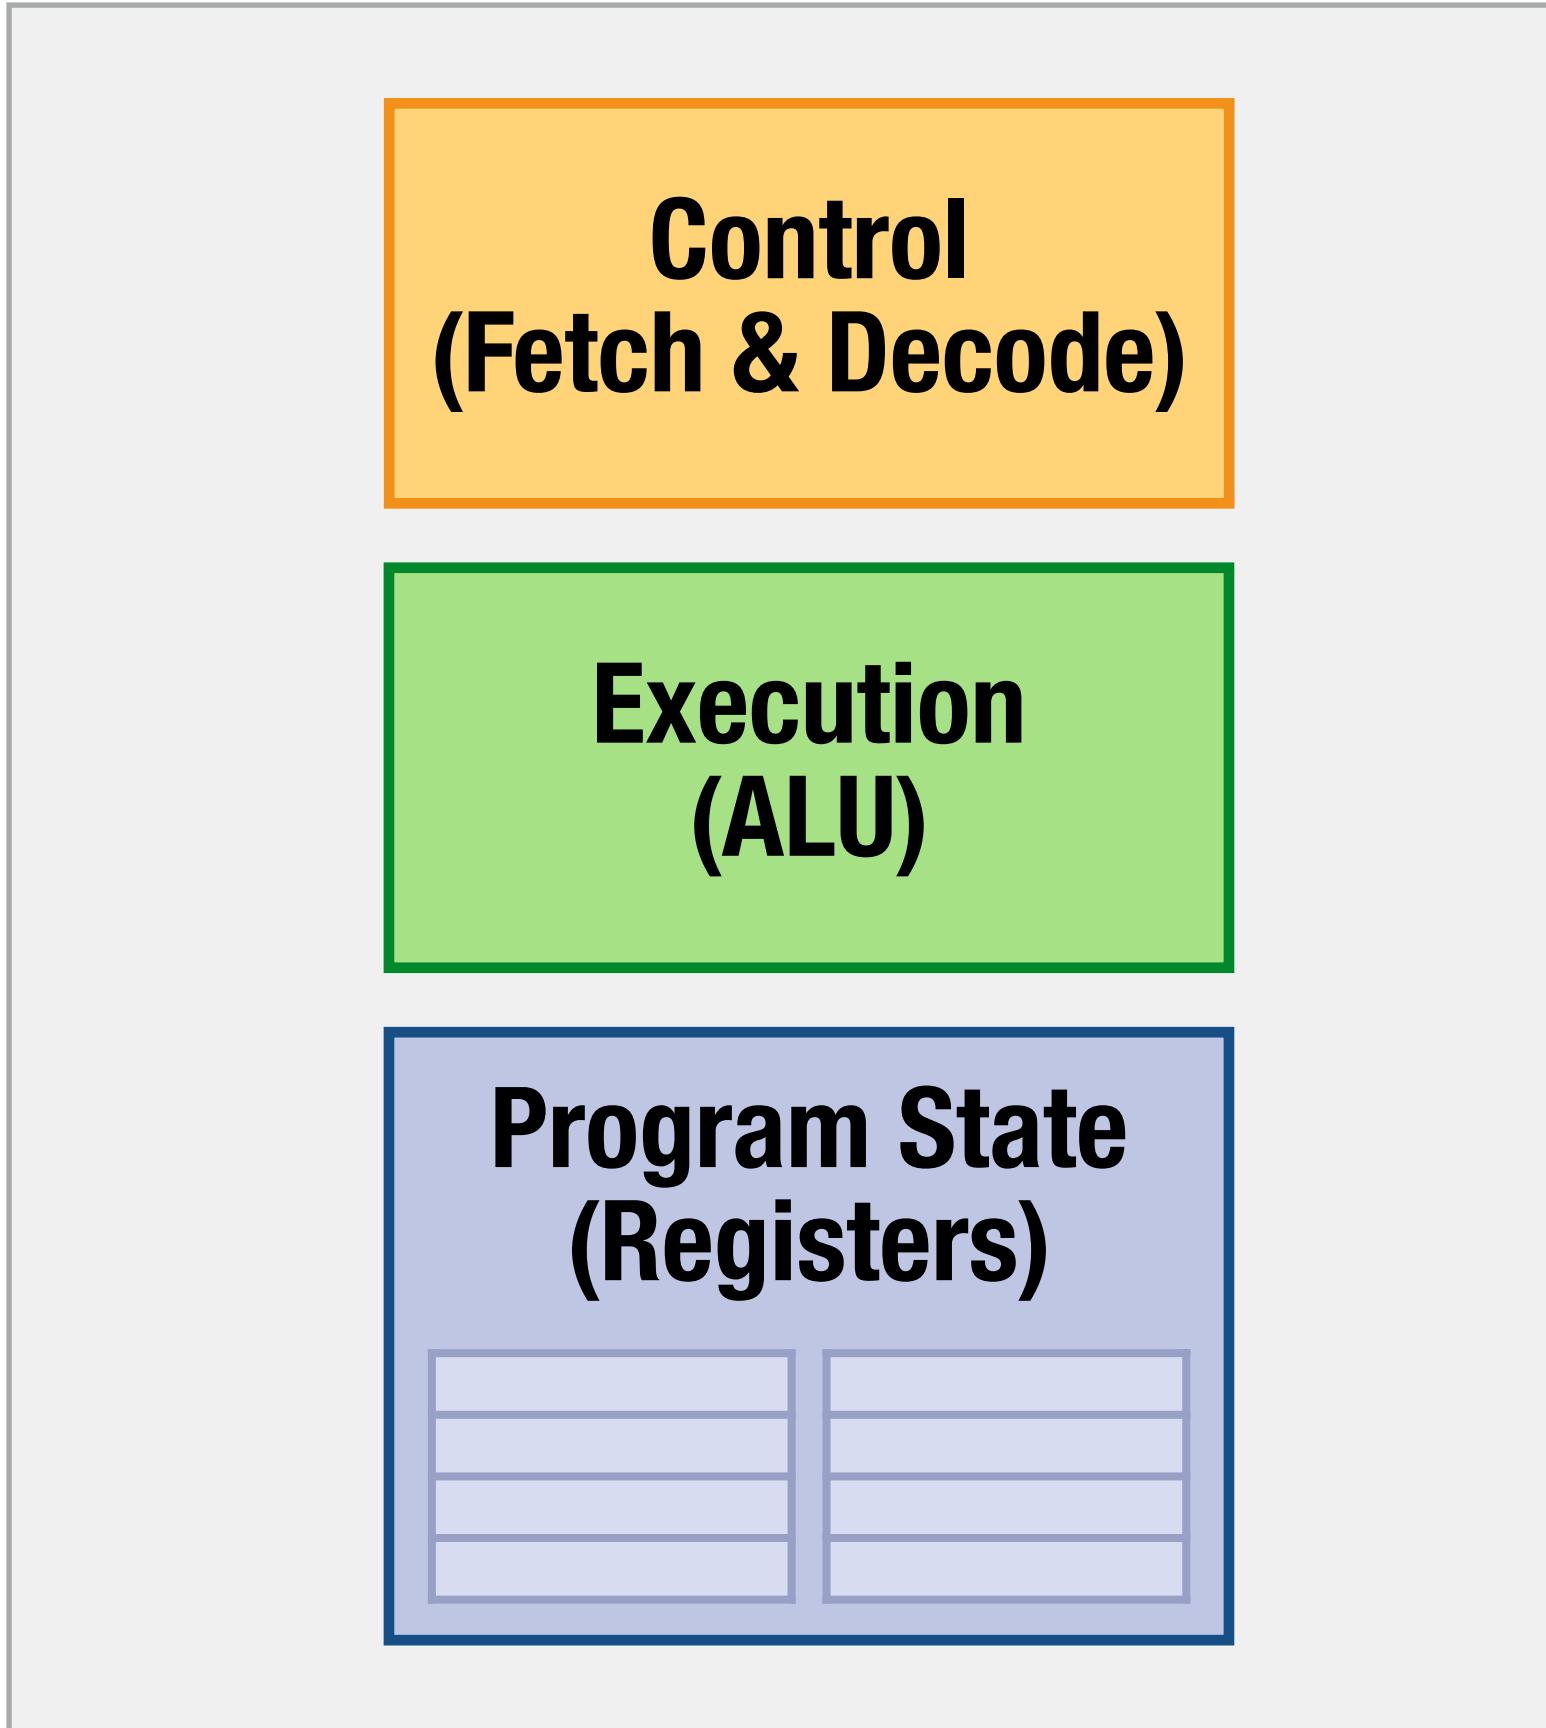
Processor

an interpreter for instructions!

ld	r3, mem[r0+r2]
ld	r4, mem[r1+r2]
mul	r3, r3, r4
add	r5, r5, r3
addi	r2, r2, 4
blt	r2, \$400, LOOP
st	addr[r6], r5
...	

Program

Optimization 1: Increase the clock speed



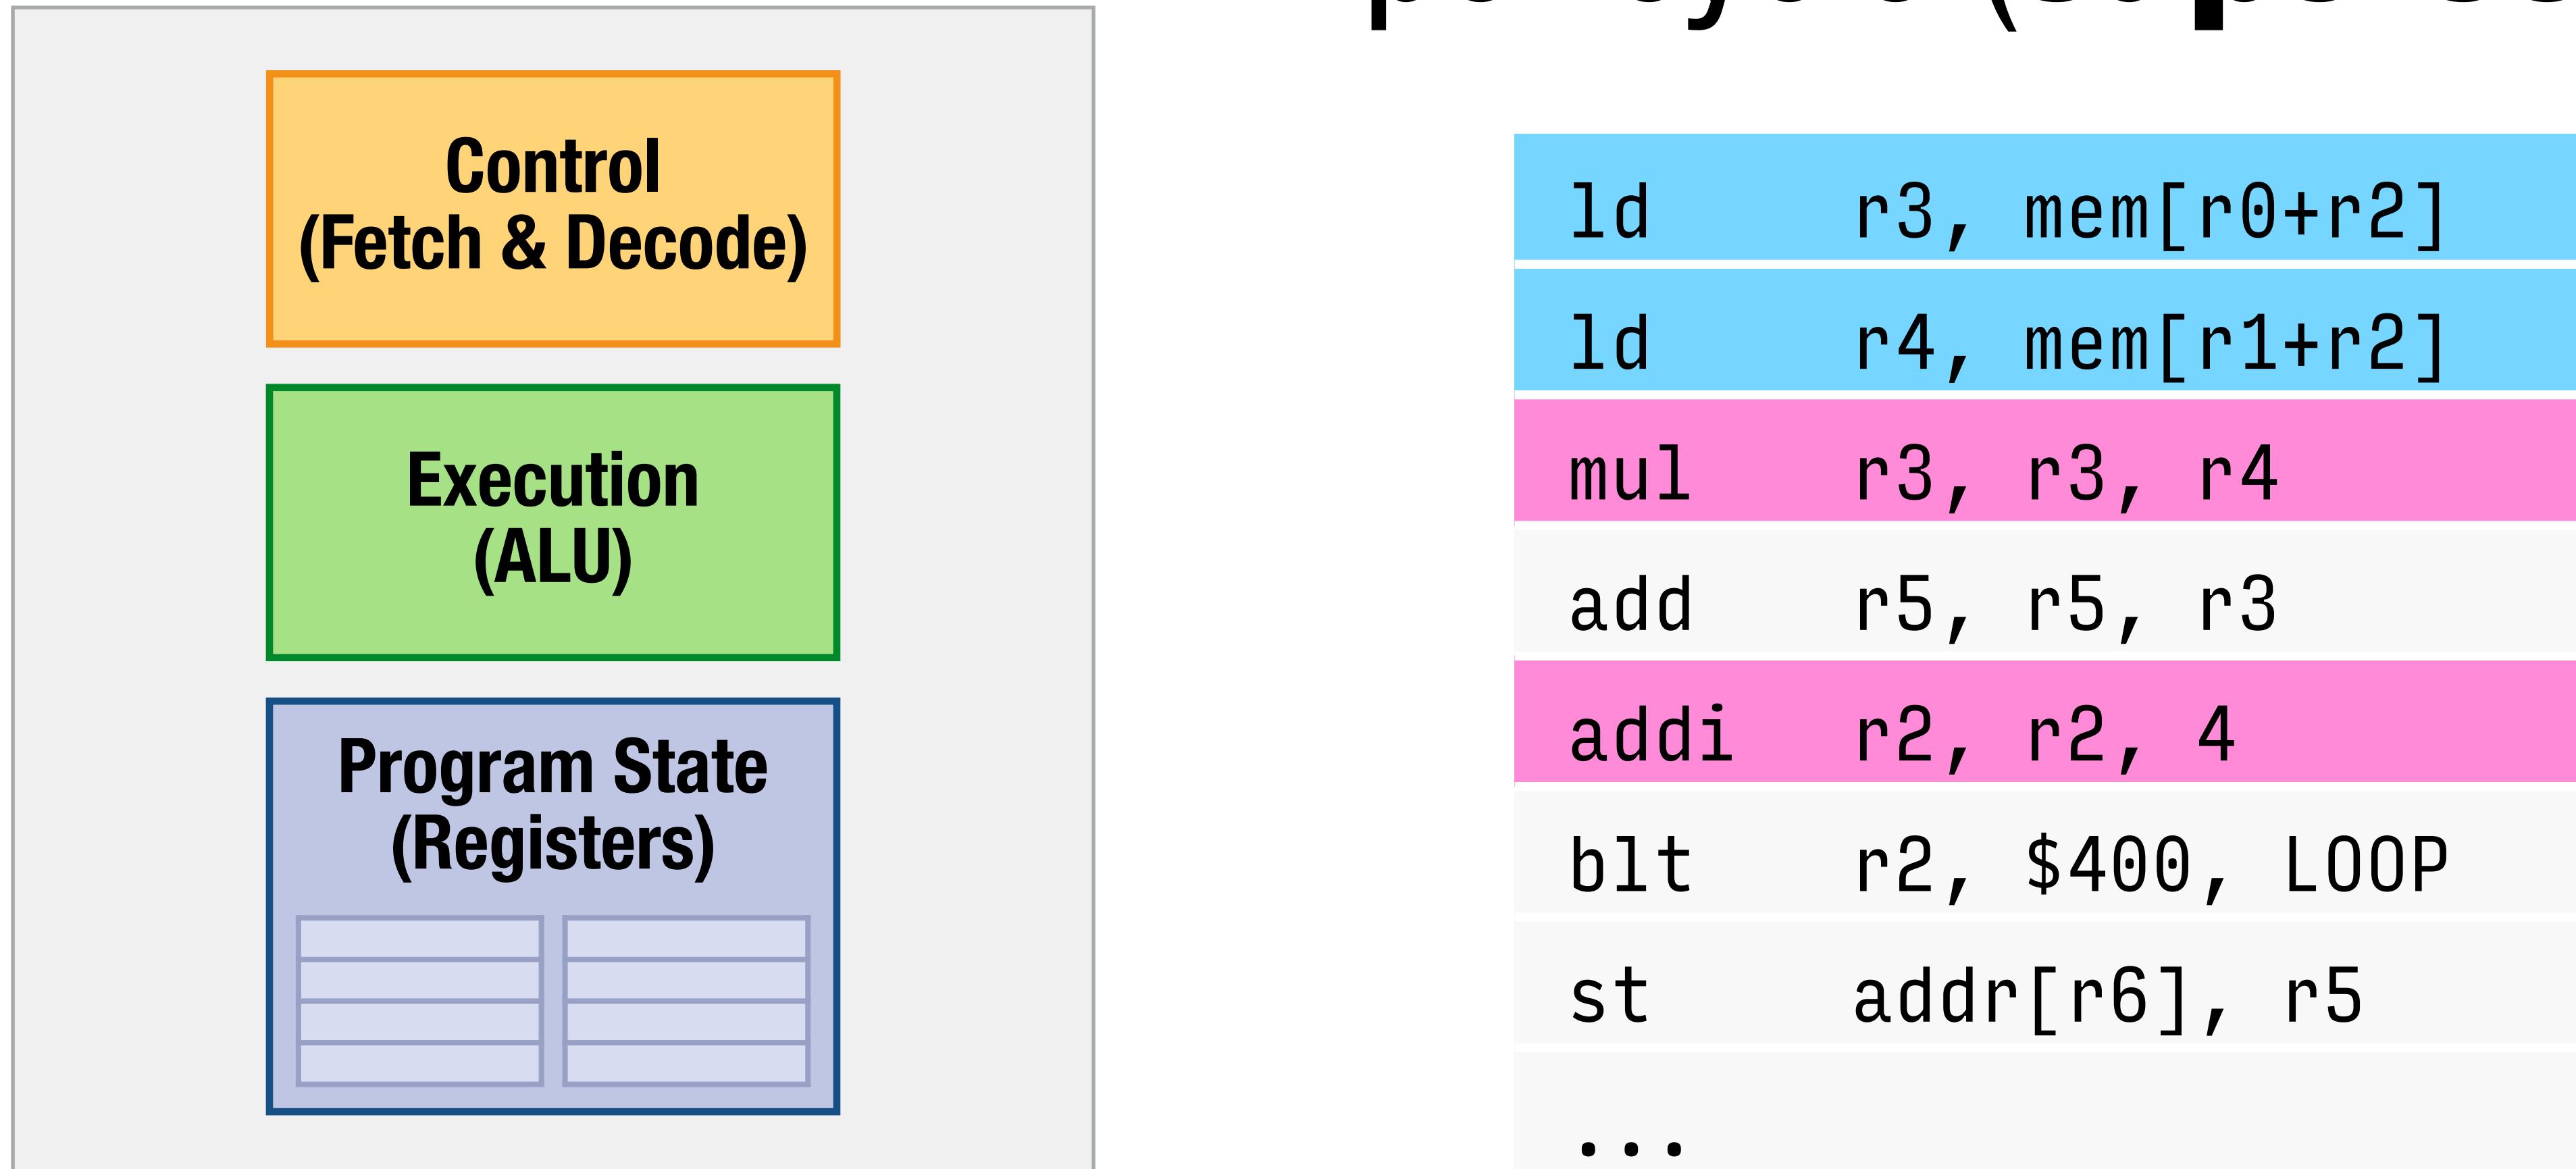
Processor

Higher voltage
↳ power grows with v^2

Faster transistors
↳ higher leakage

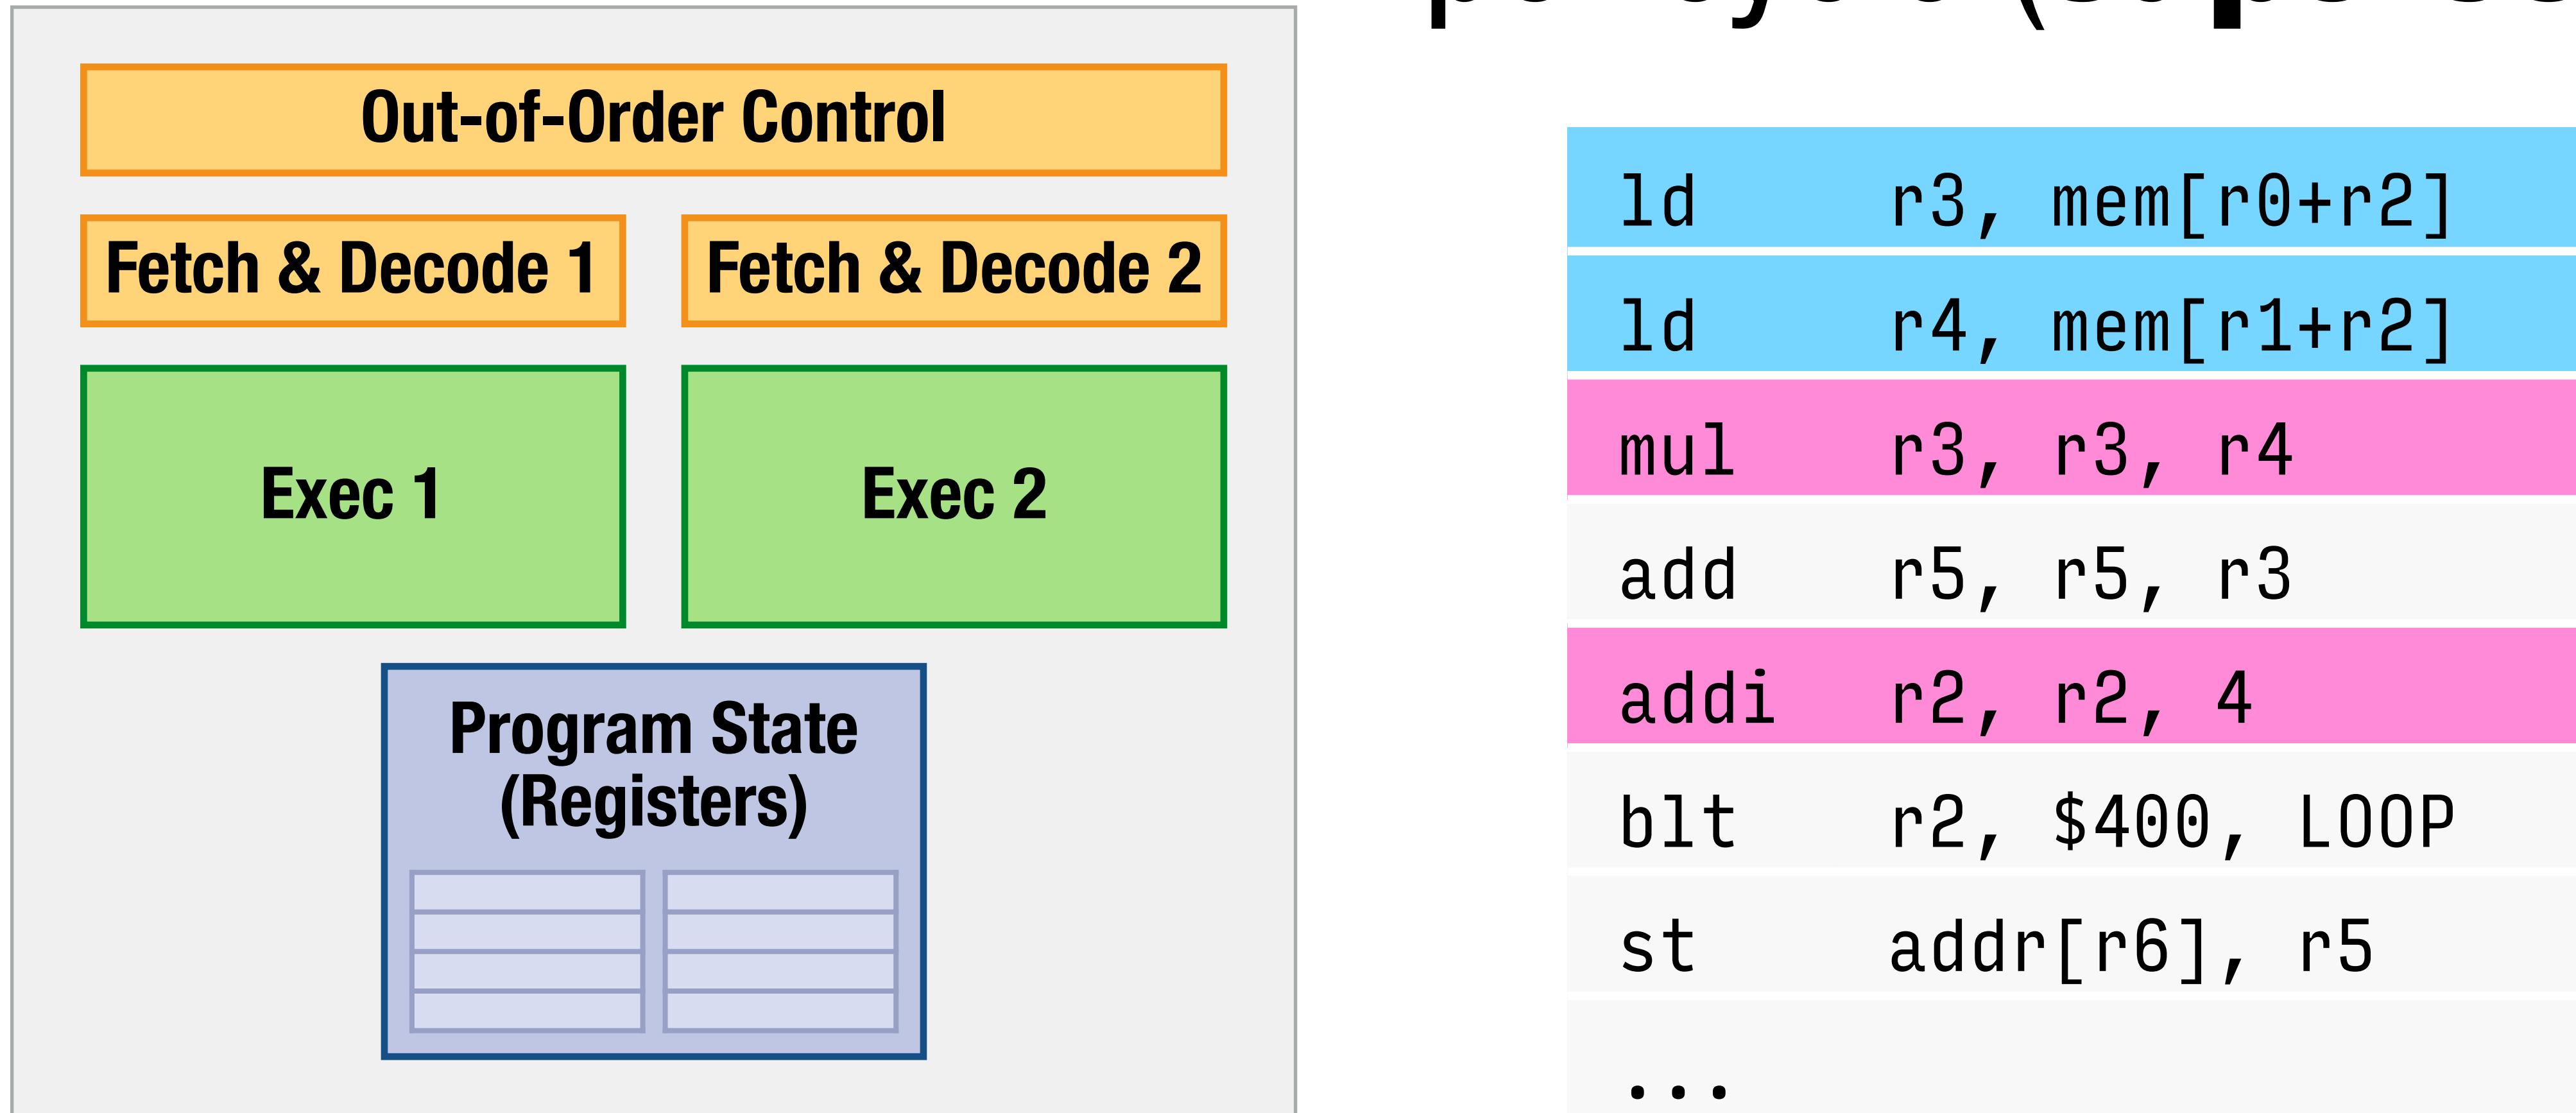
Deeper pipelining
↳ lower average IPC

Optimization 2: Execute multiple instructions per cycle (superscalar)



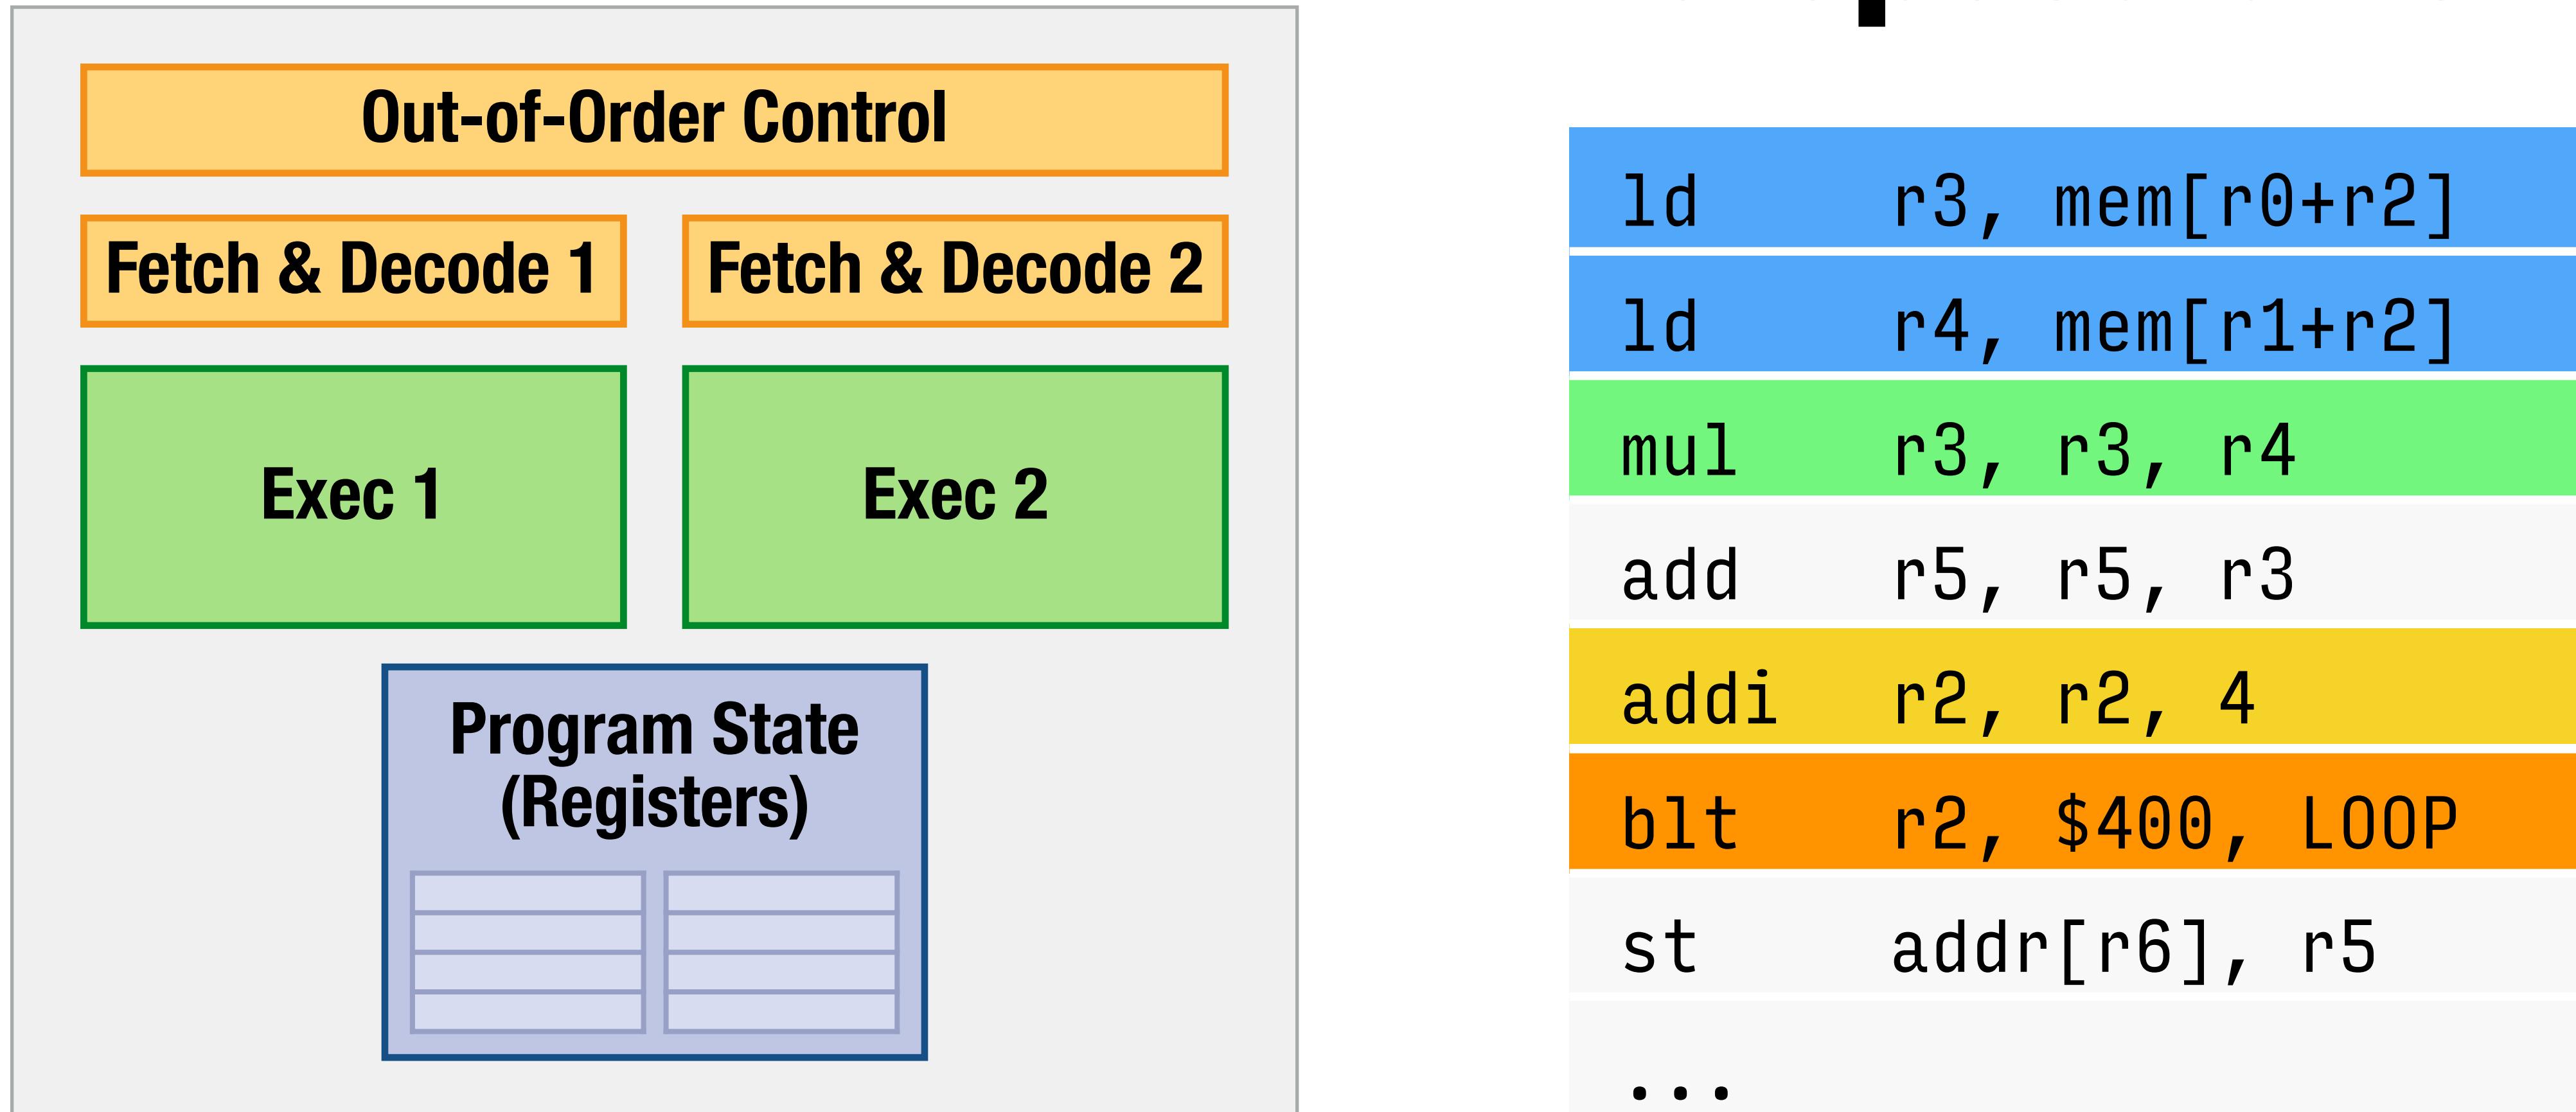
Processor

Optimization 2: Execute multiple instructions per cycle (superscalar)



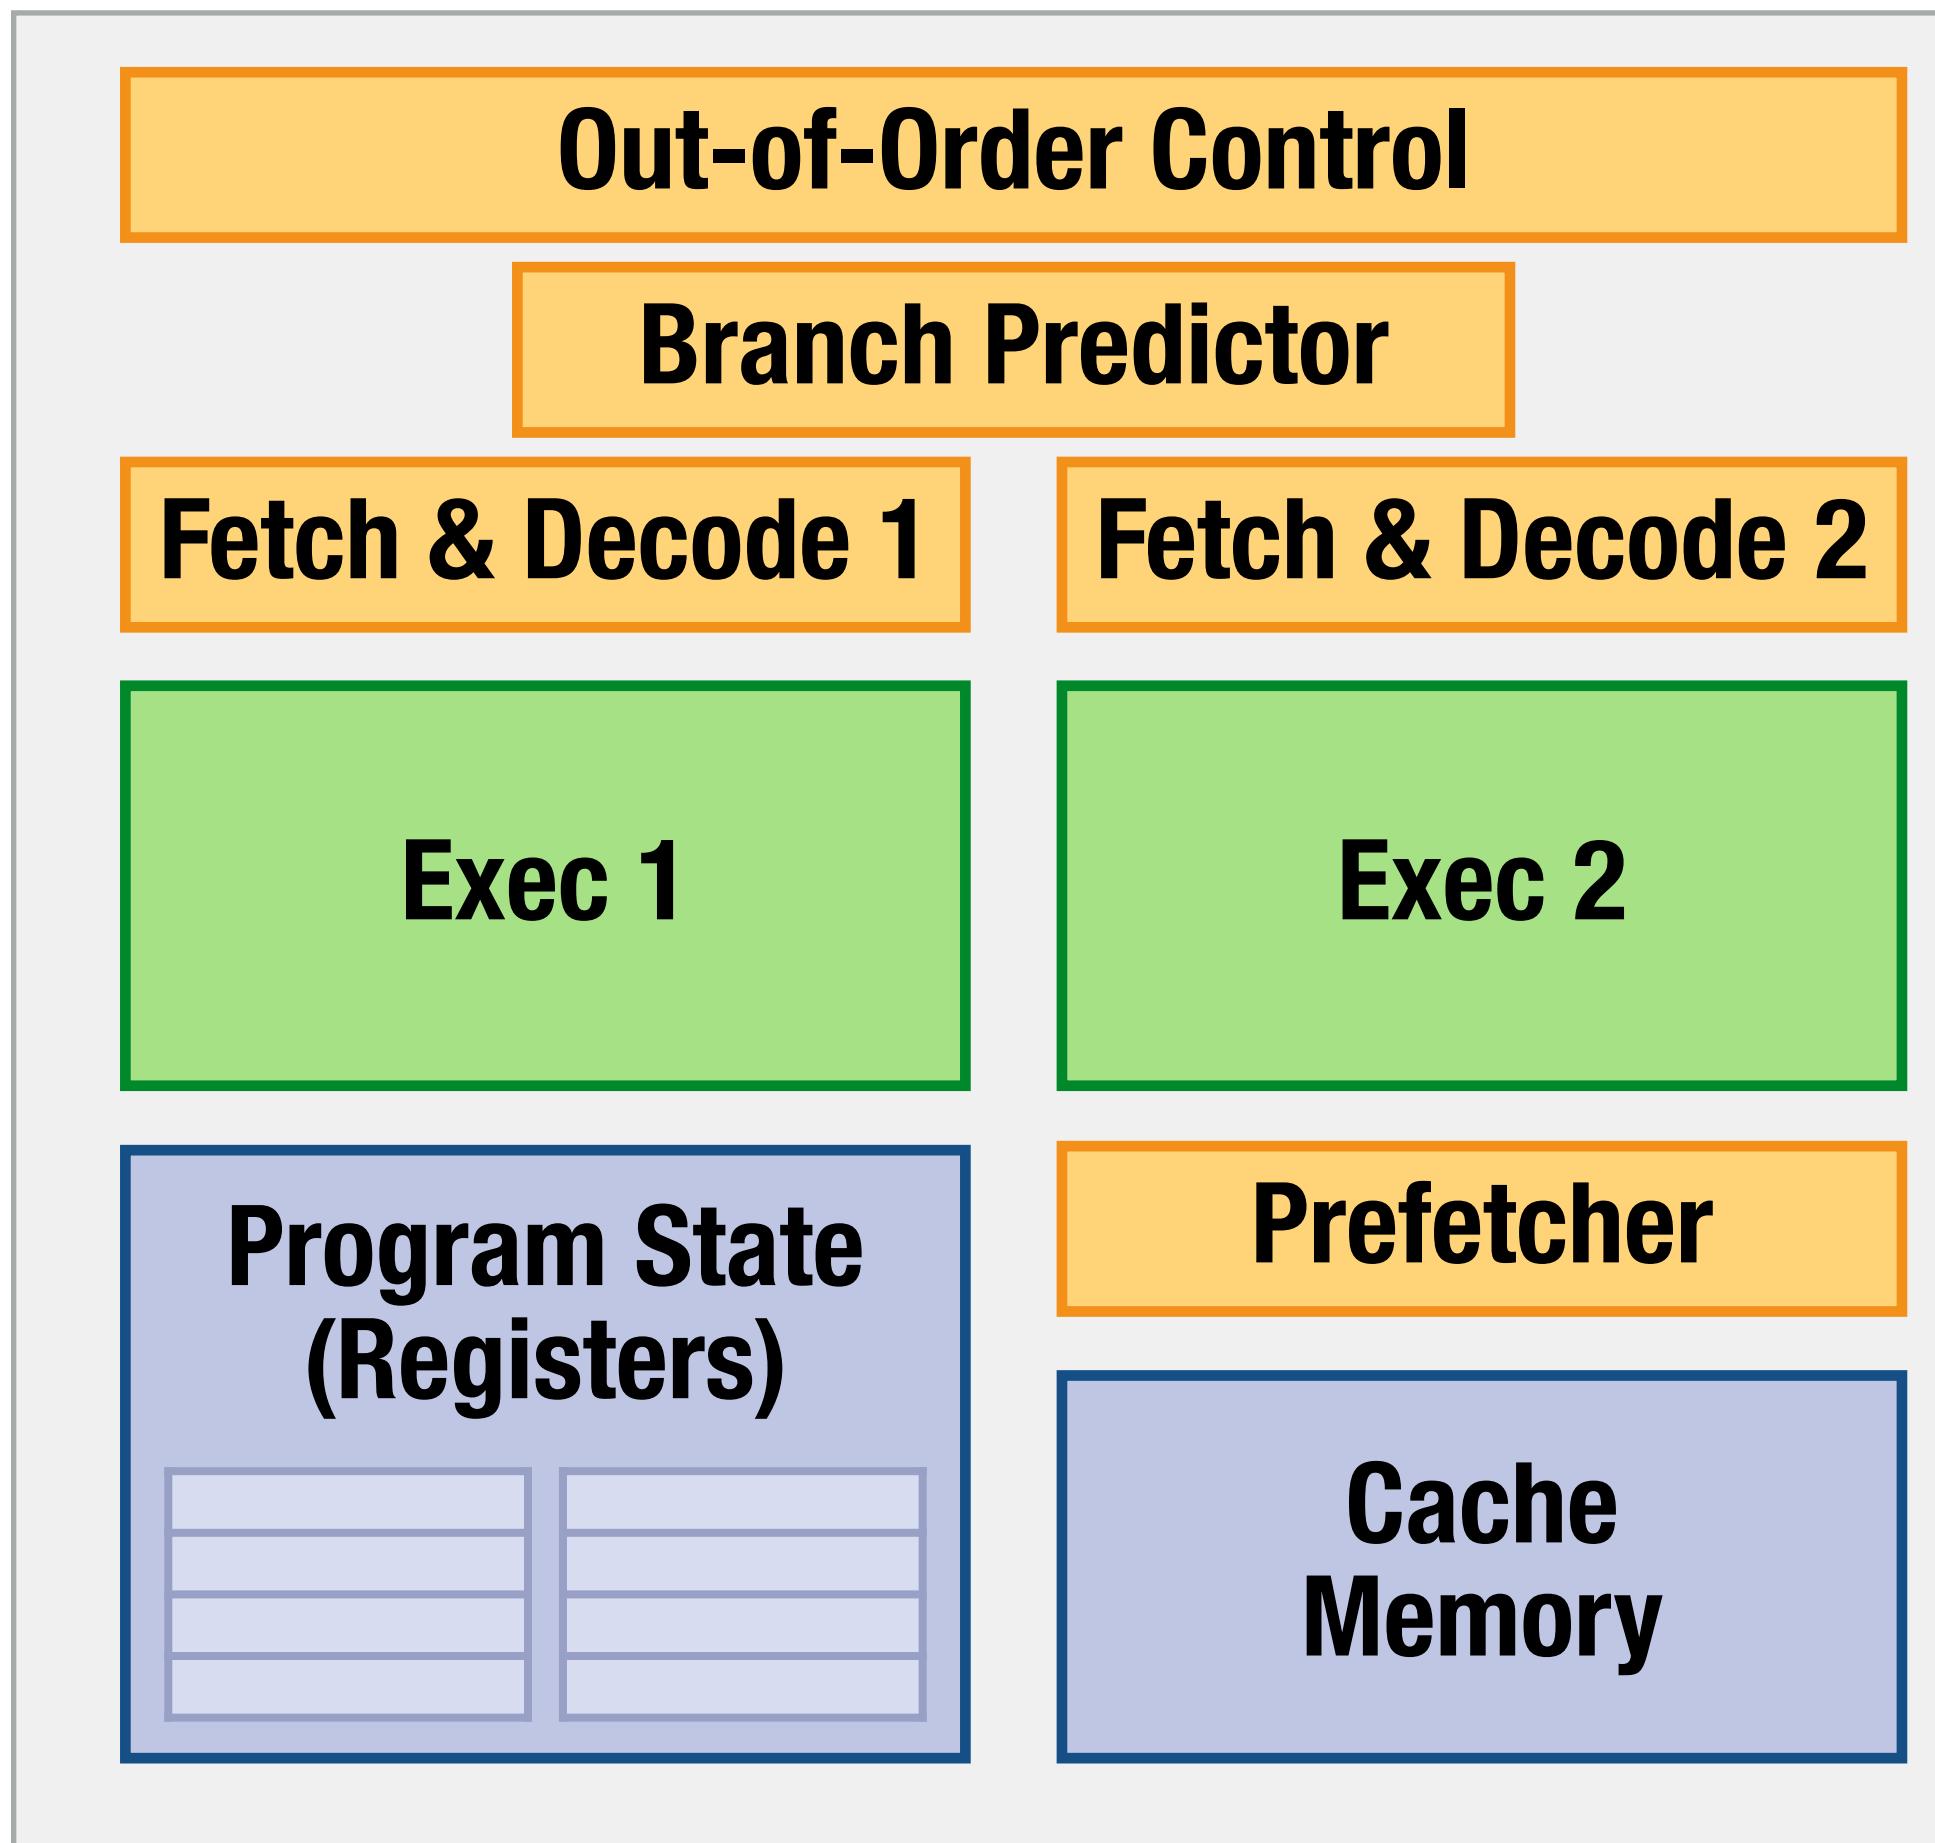
Processor

Optimization 3: Avoid stalls through ILP & speculation



Processor

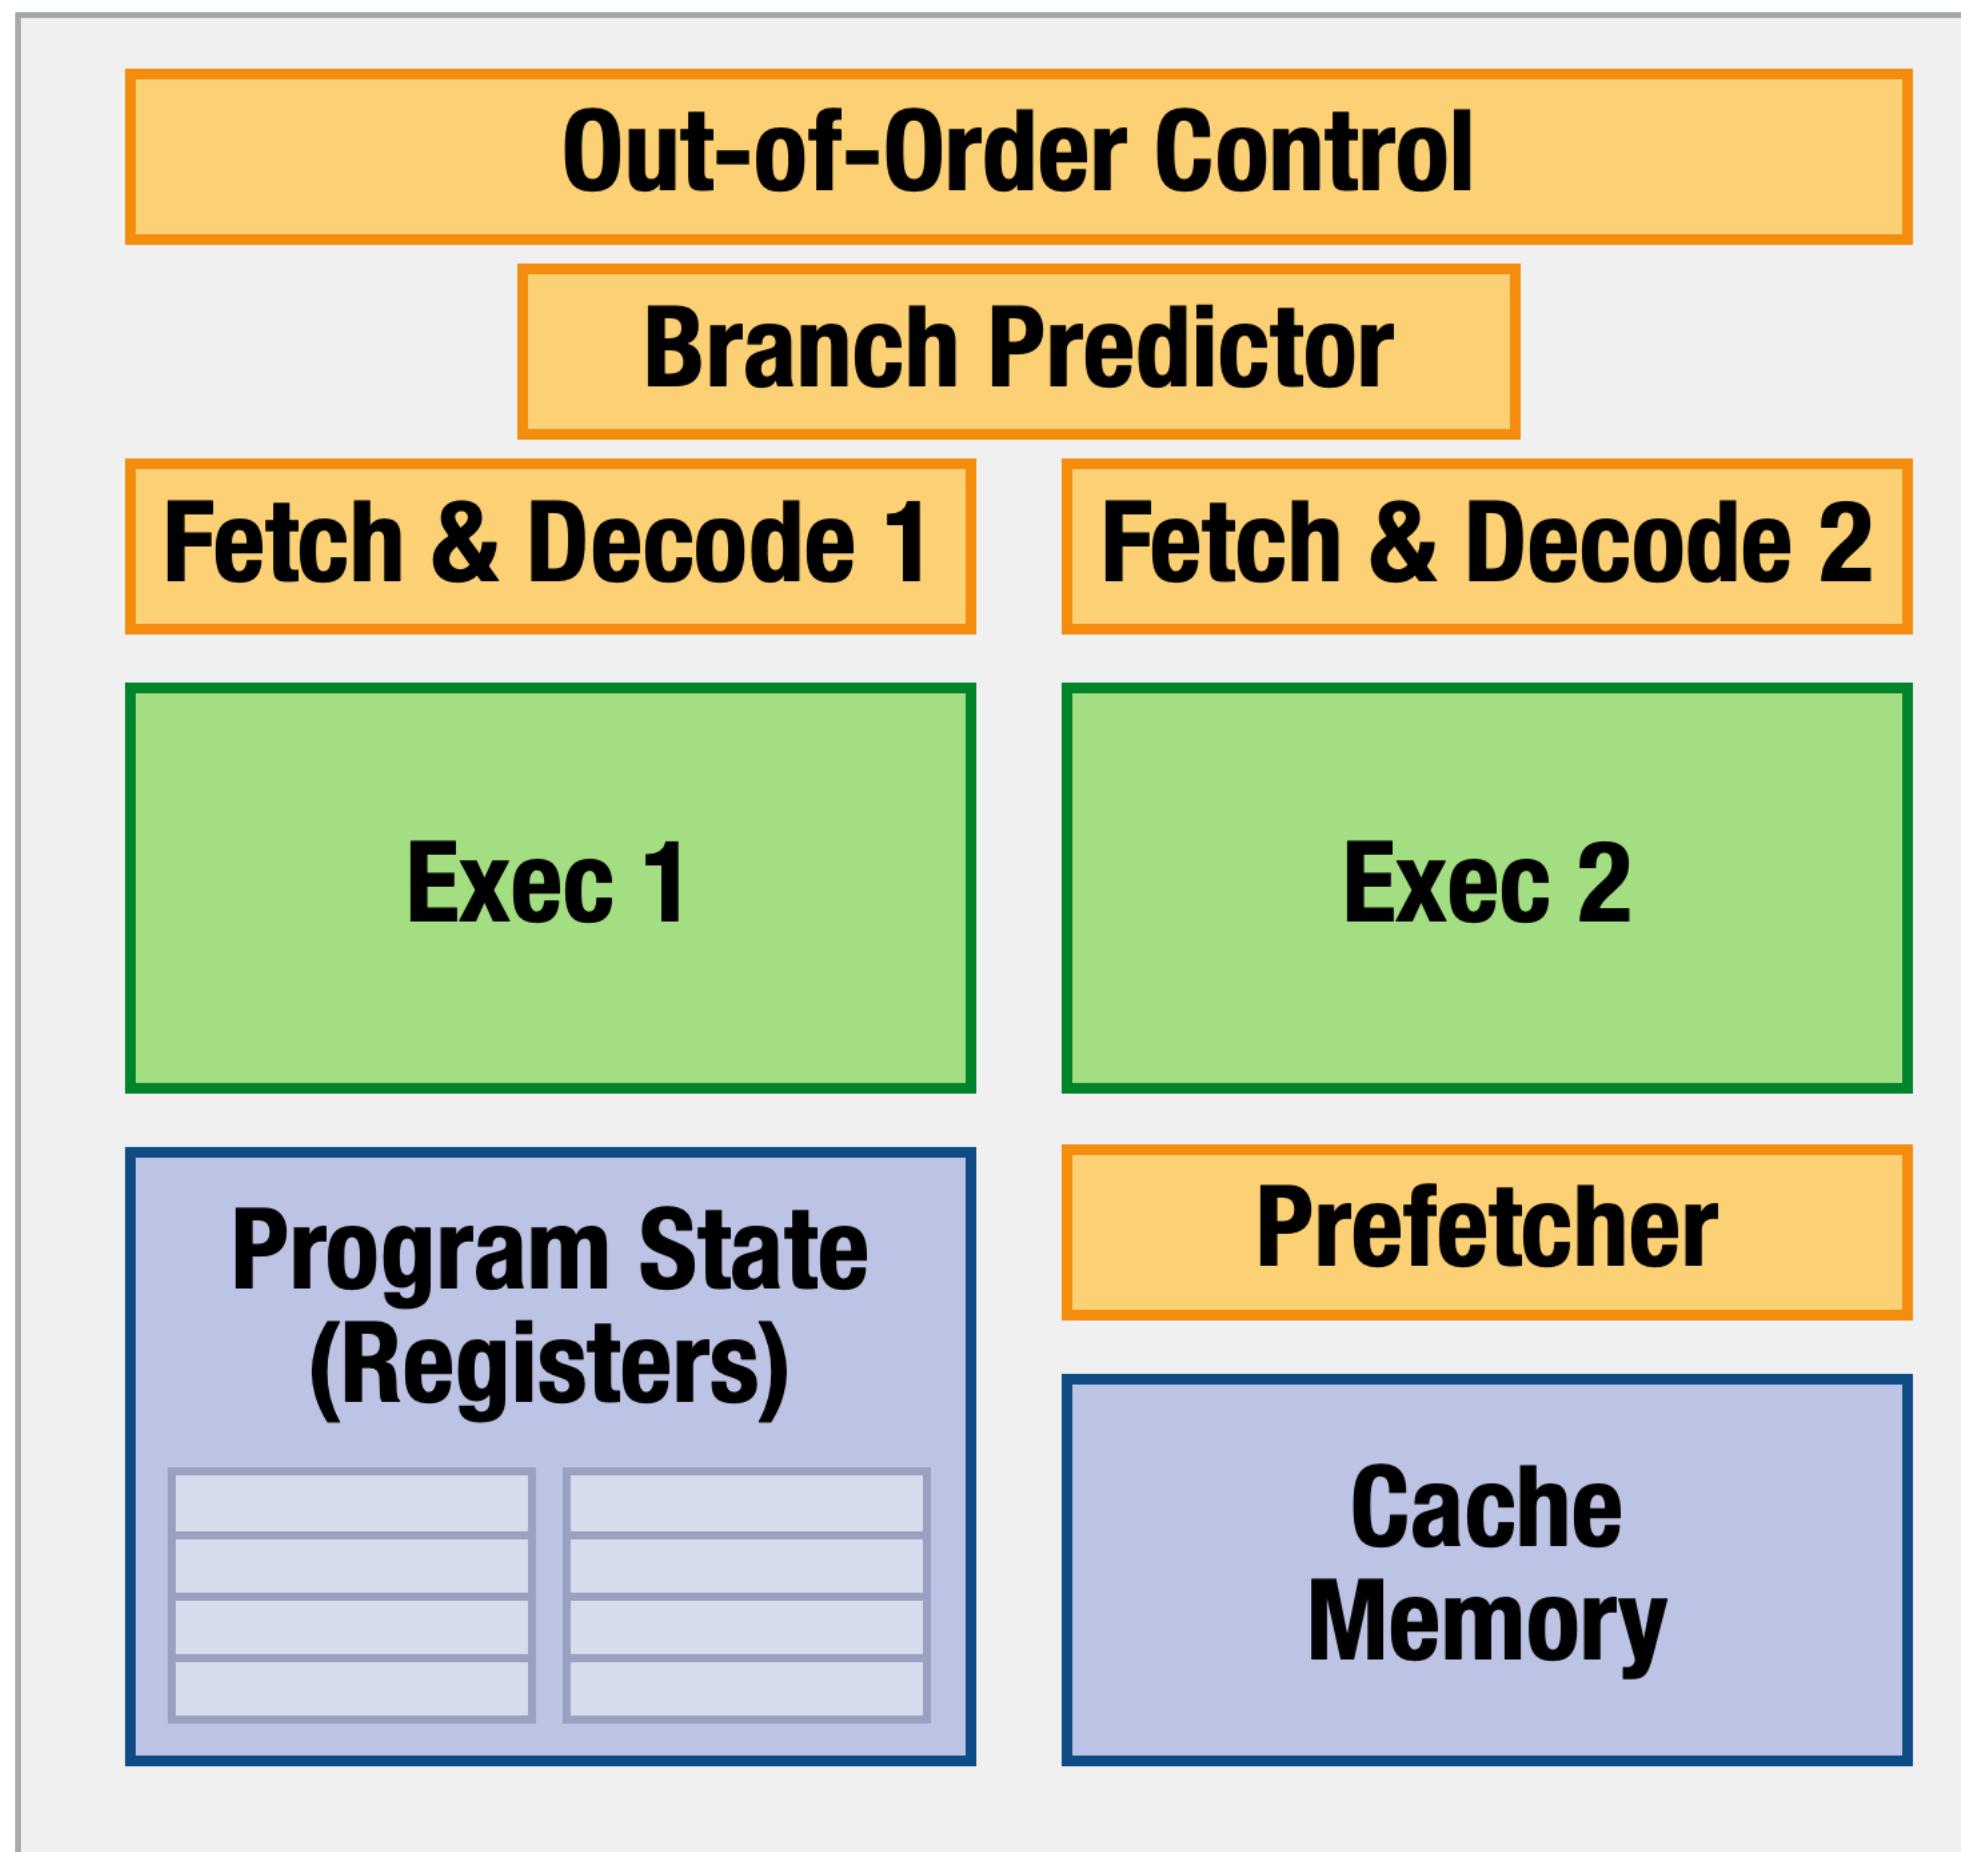
Optimization 3: Avoid stalls through ILP & speculation



Processor

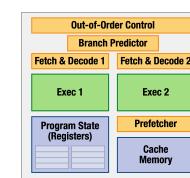
ld	r3, mem[r0+r2]
ld	r4, mem[r1+r2]
mul	r3, r3, r4
add	r5, r5, r3
addi	r2, r2, 4
blt	r2, \$400, LOOP
st	addr[r6], r5
...	

This makes sense when you have only one processor



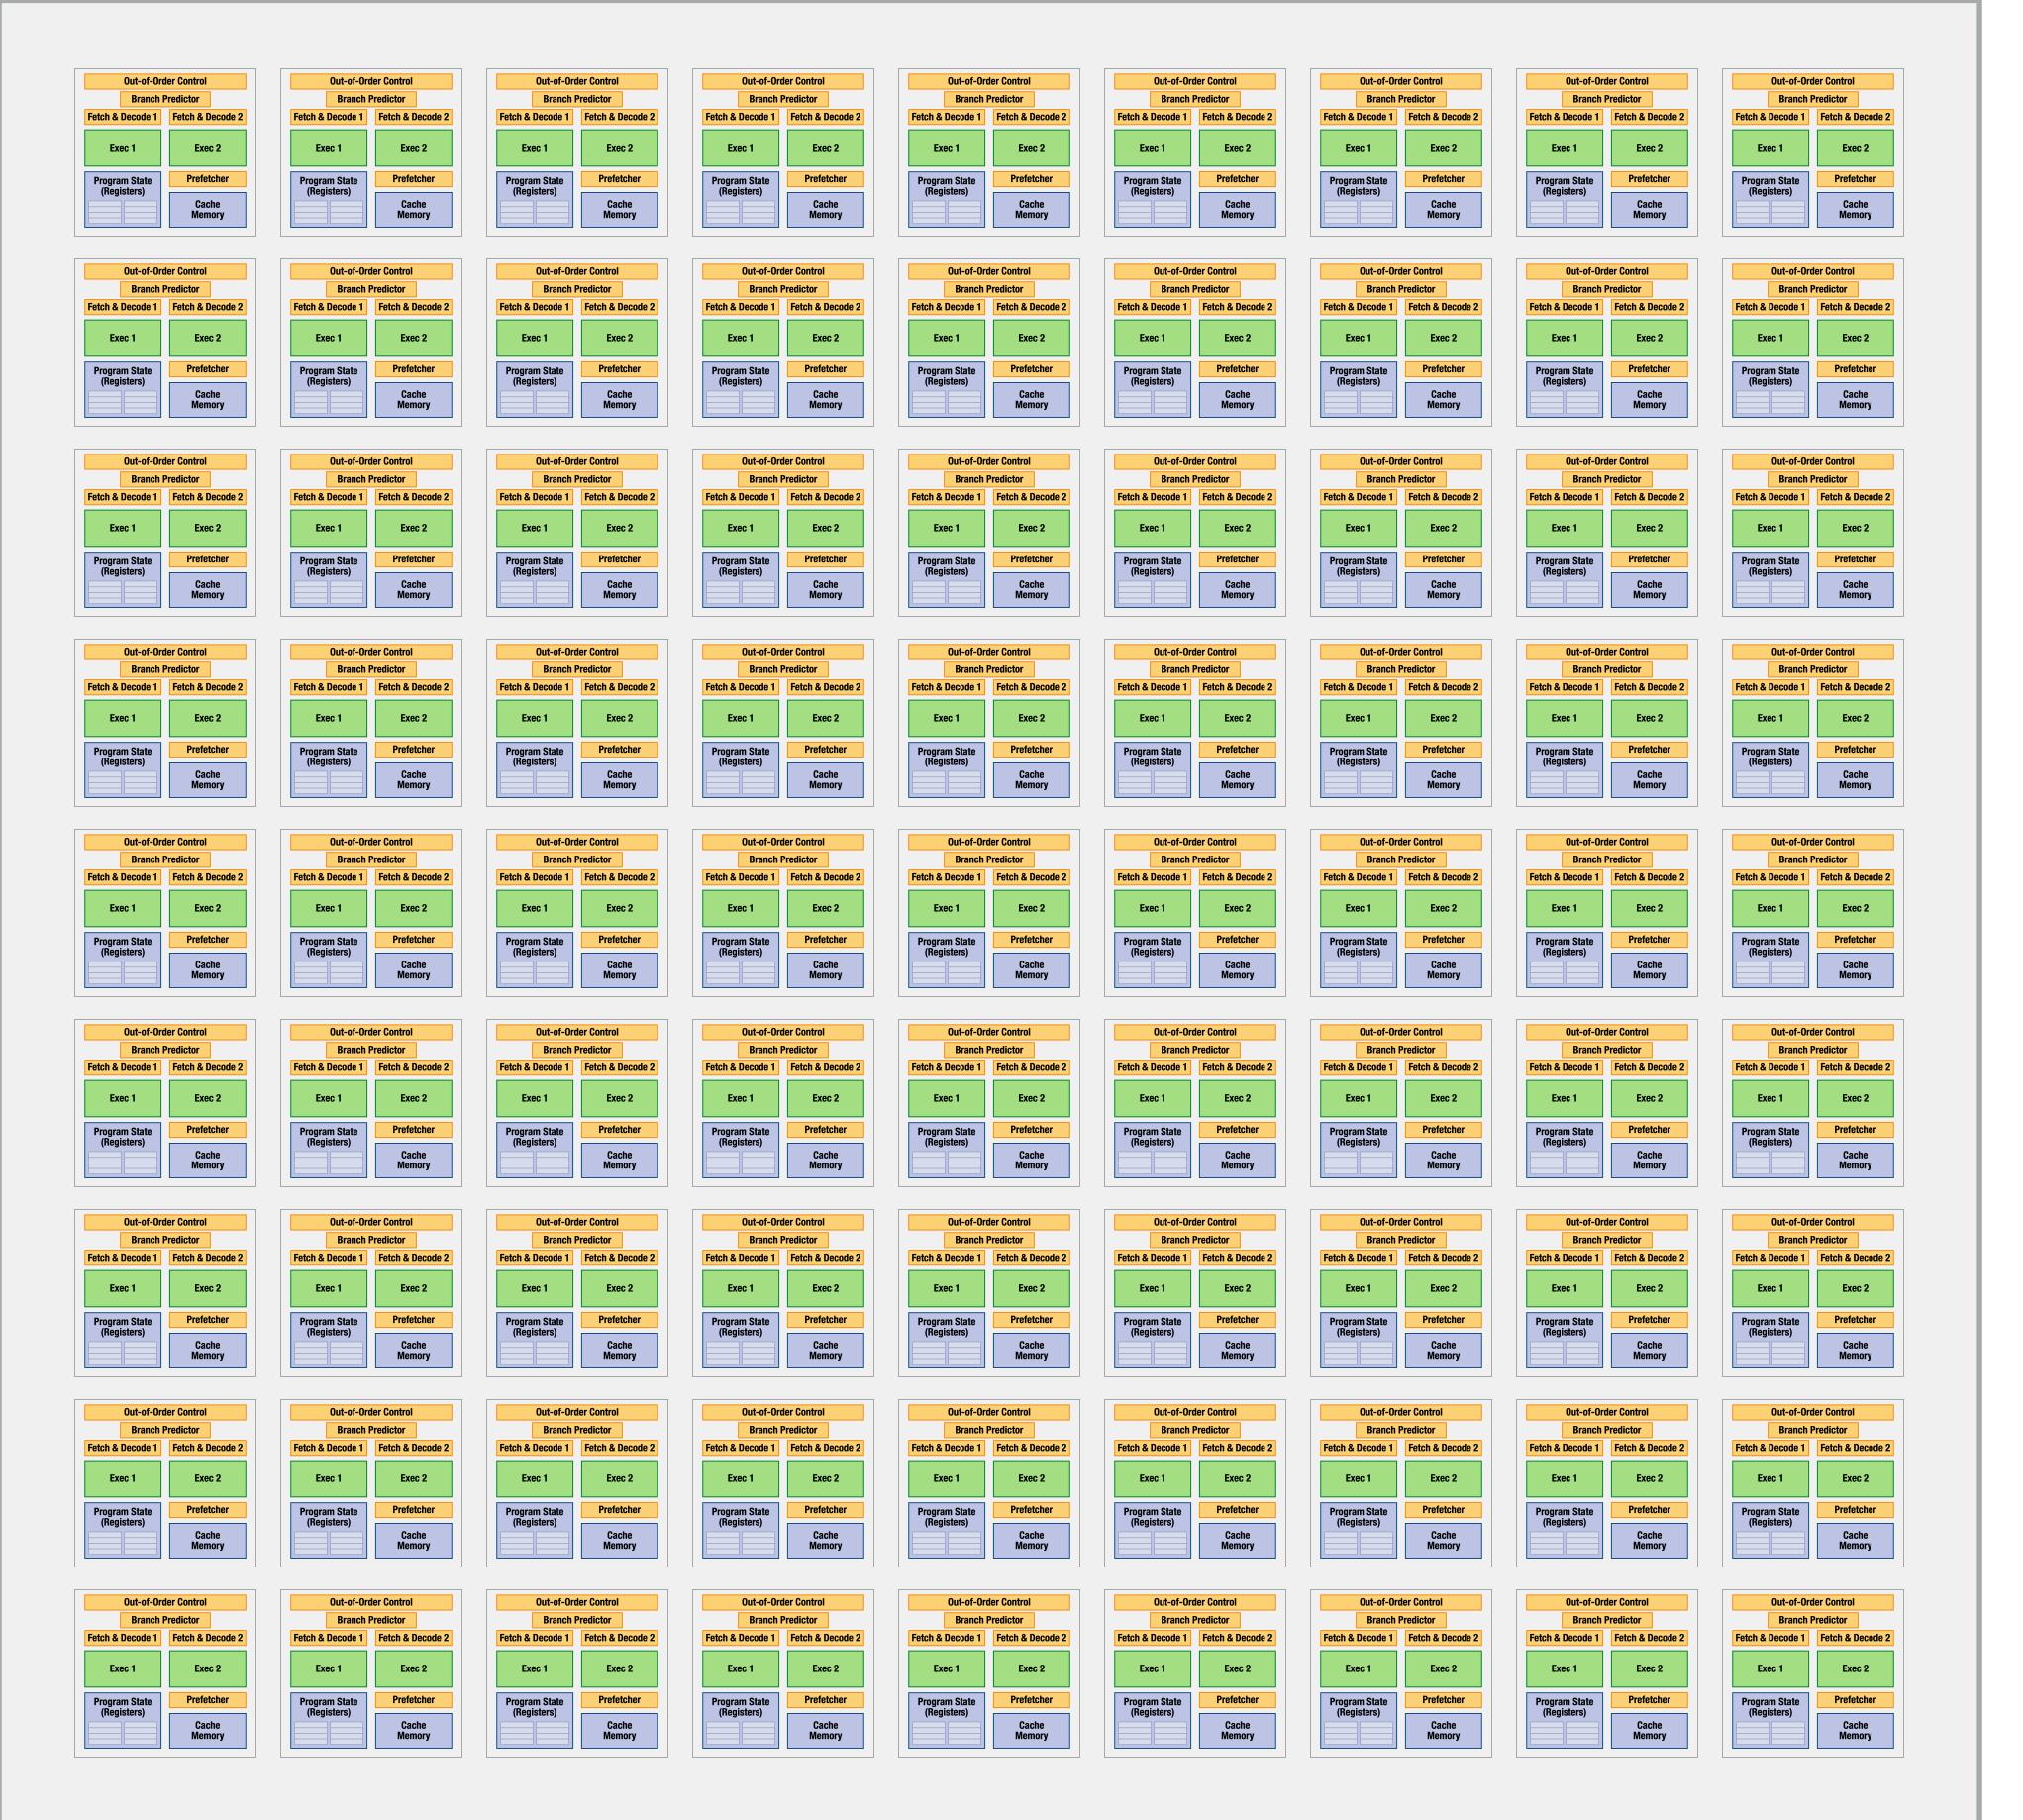
Processor

This makes sense when you have only one processor



Processor (to scale)

The reality today:



At the same time:

We most need
performance when
processing lost of stuff

↳ abundant data-
parallelism

Silicon Chip

A **throughput-oriented**
view of performance:

Optimize aggregate rate
of processing many items

Three things drive throughput:

1

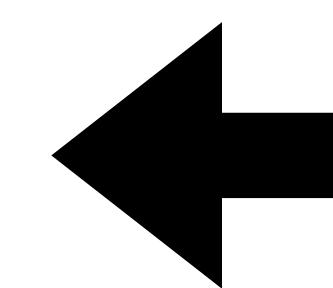
Amount of **work**
to be done

2

Amount of **resources**
to be applied (silicon, energy)

3

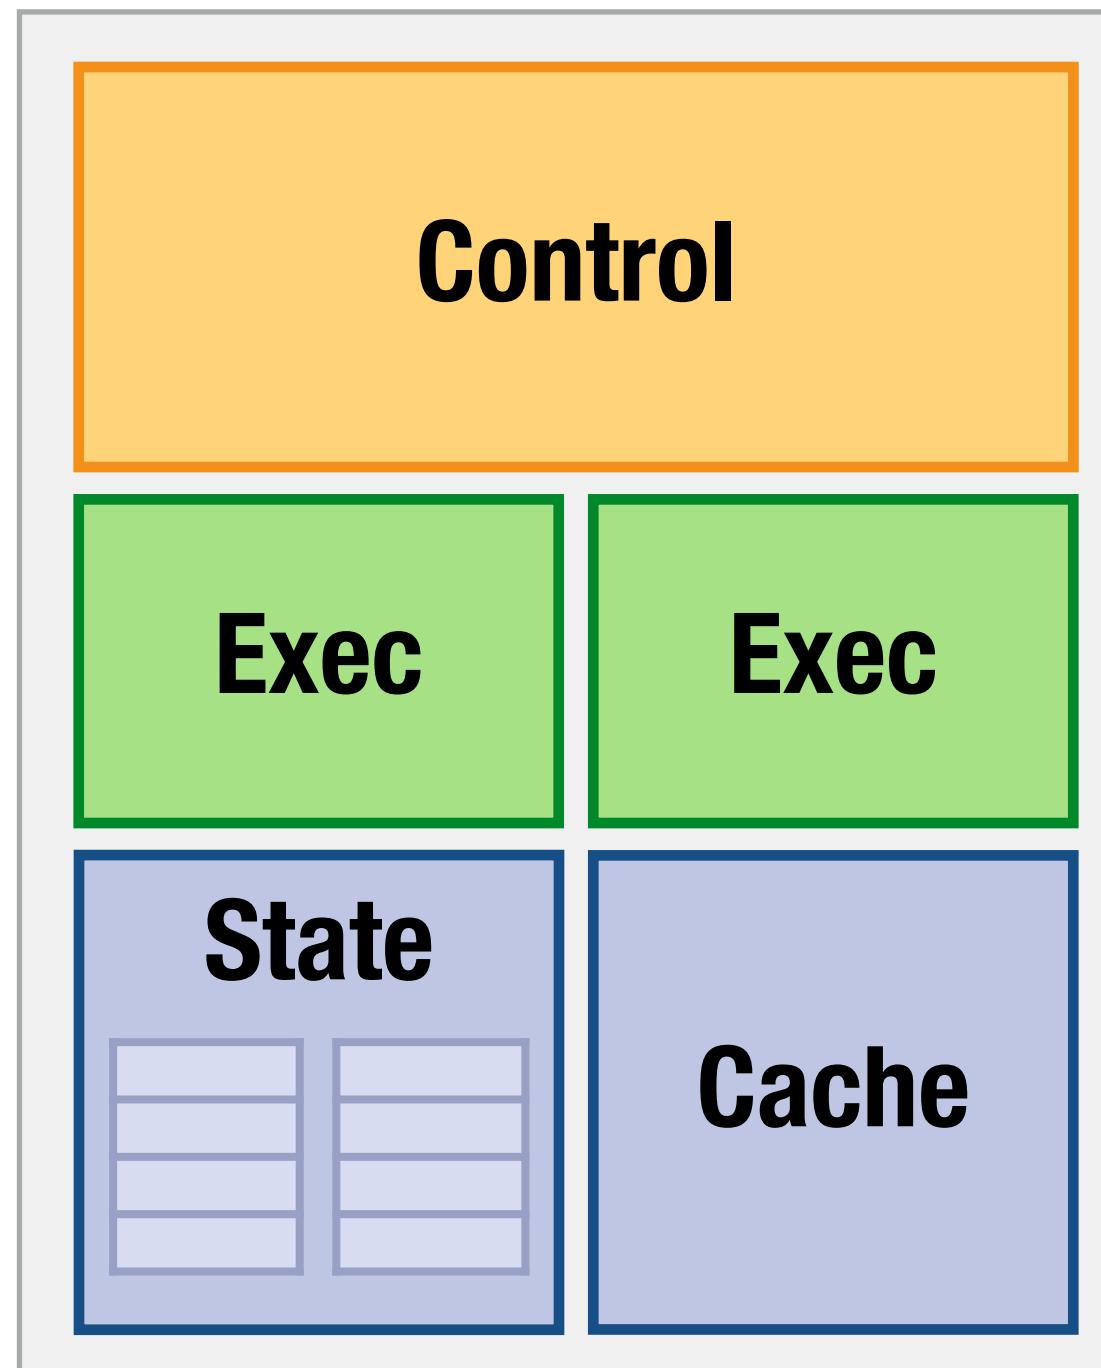
Efficiency of applying
them to useful work



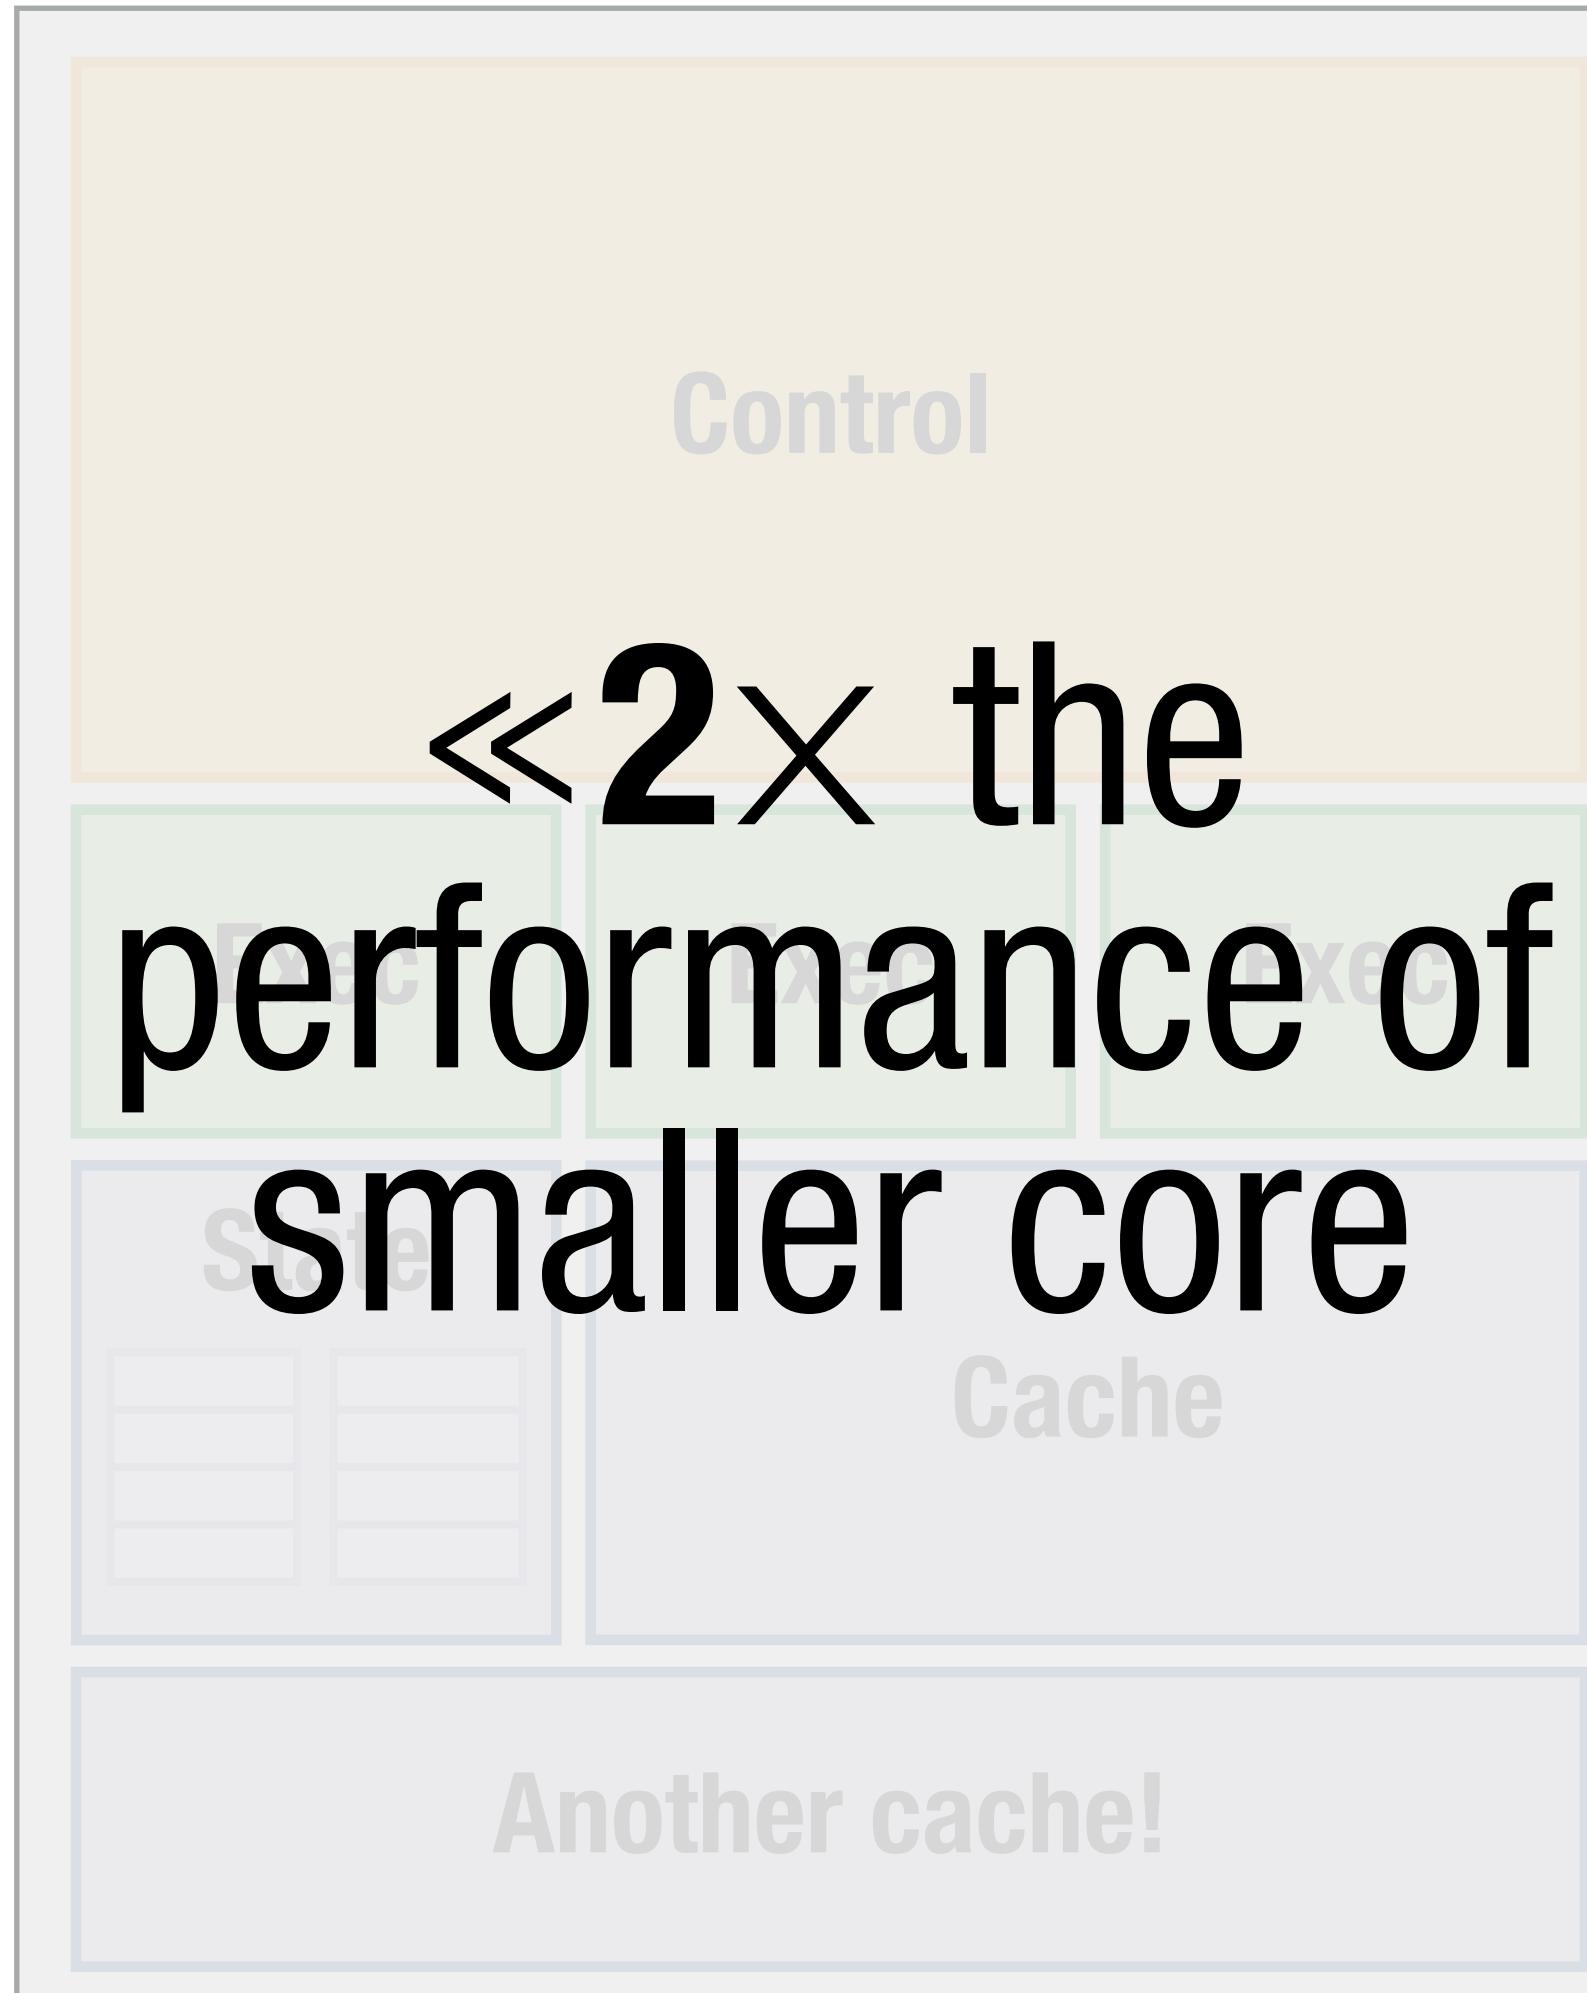
Constraints set
by application,
Si process node

Goal:
optimize this!

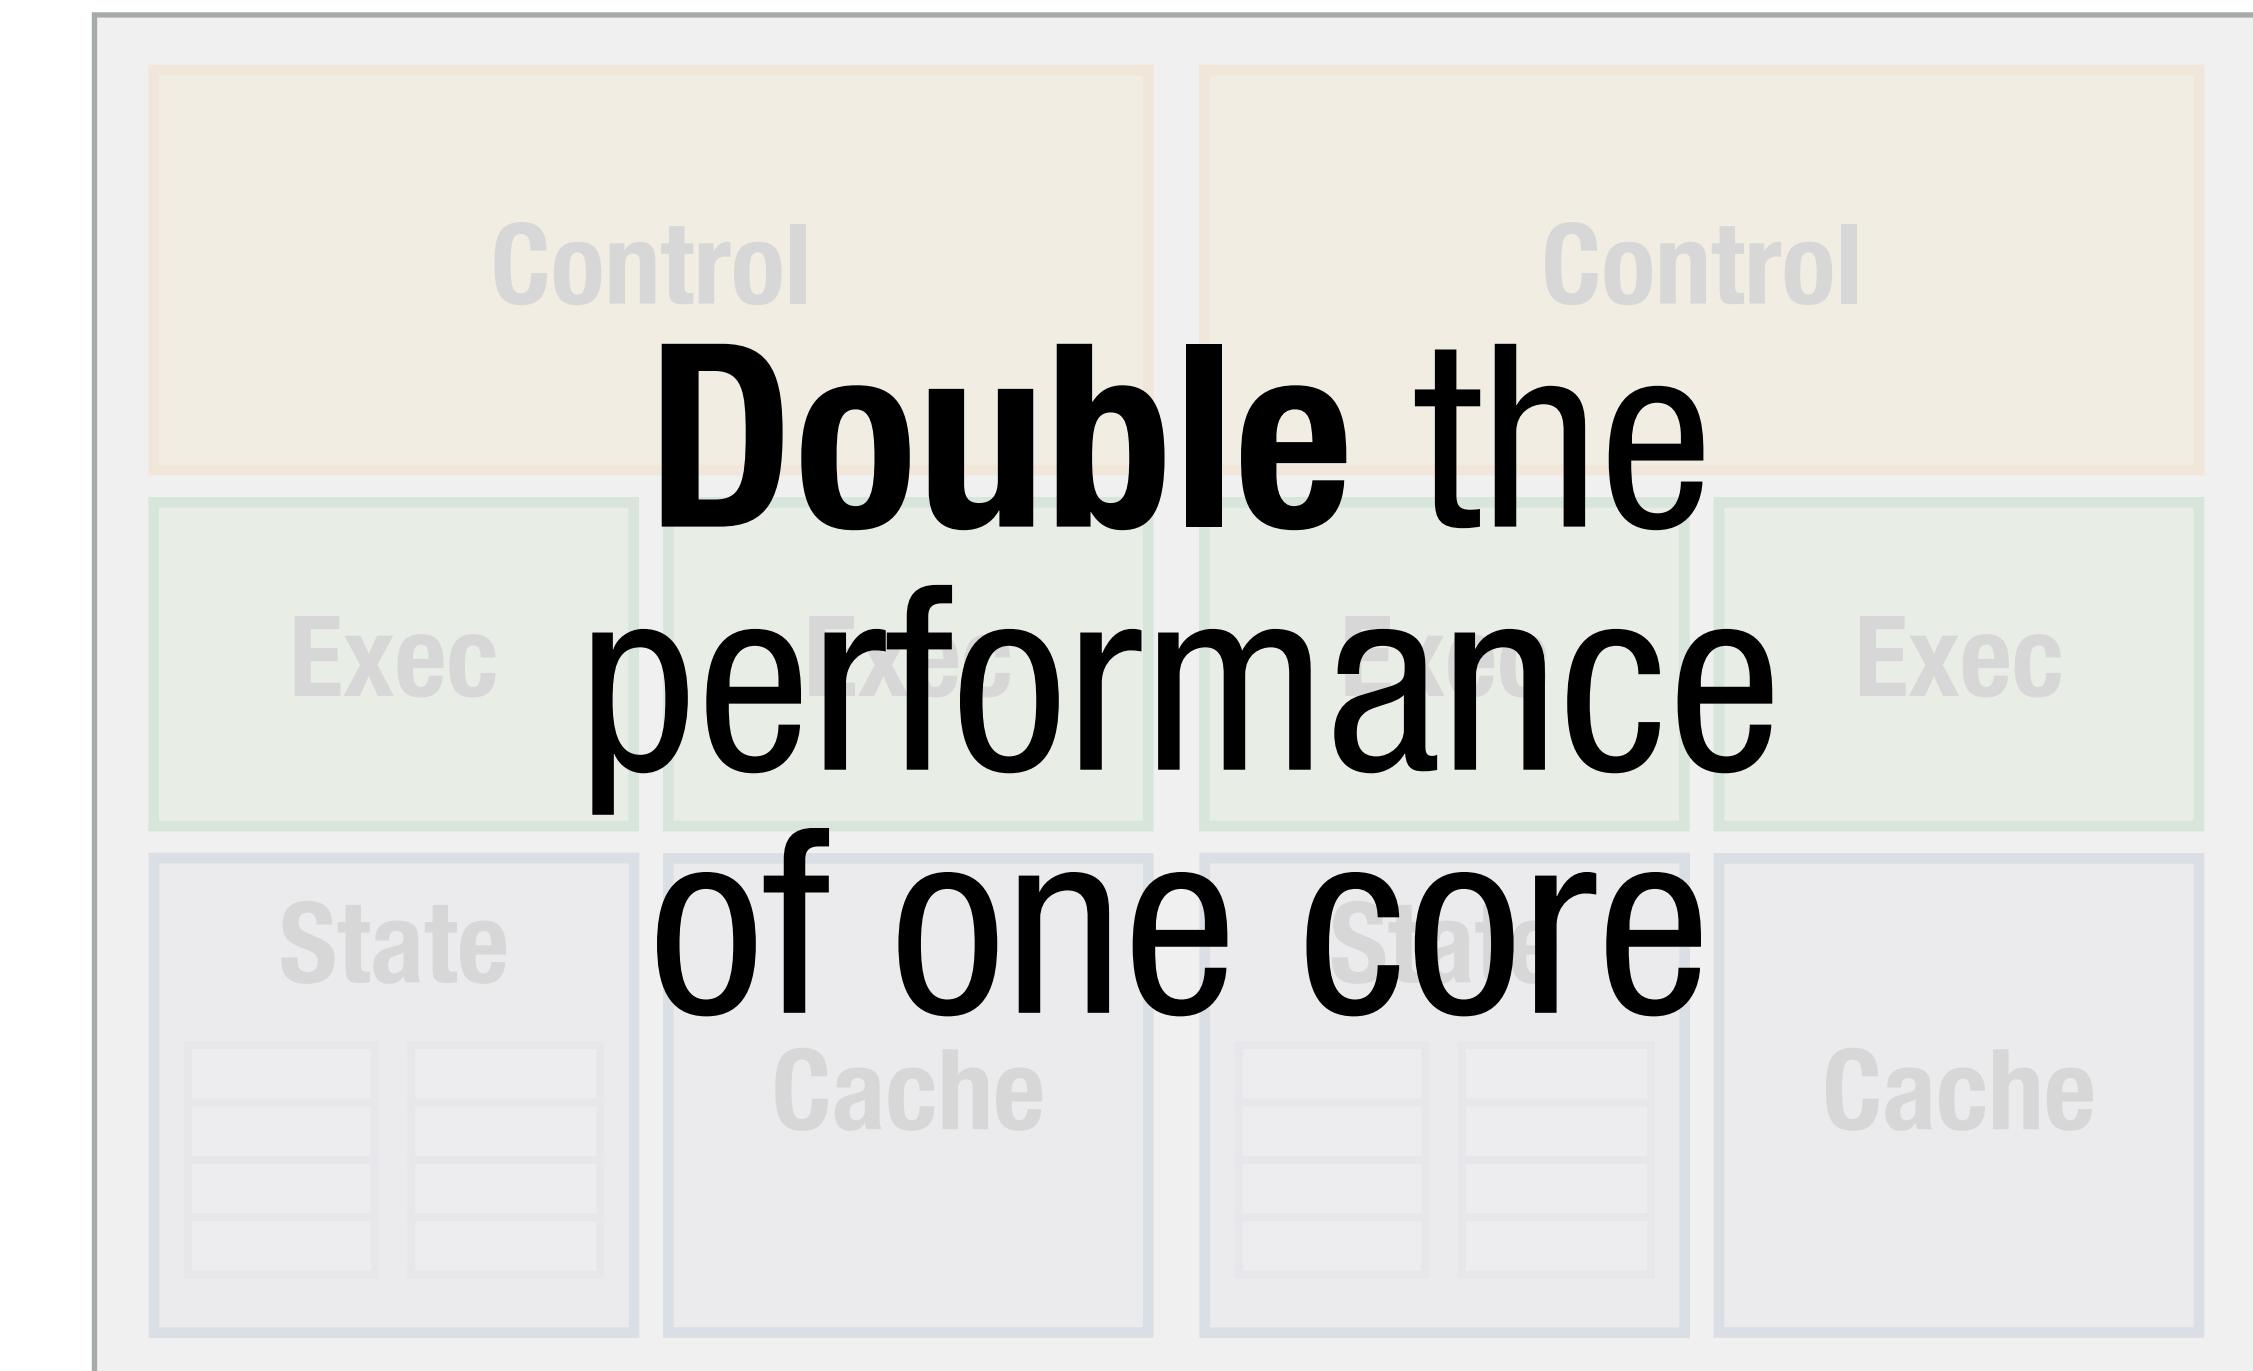
What to do with twice the silicon?



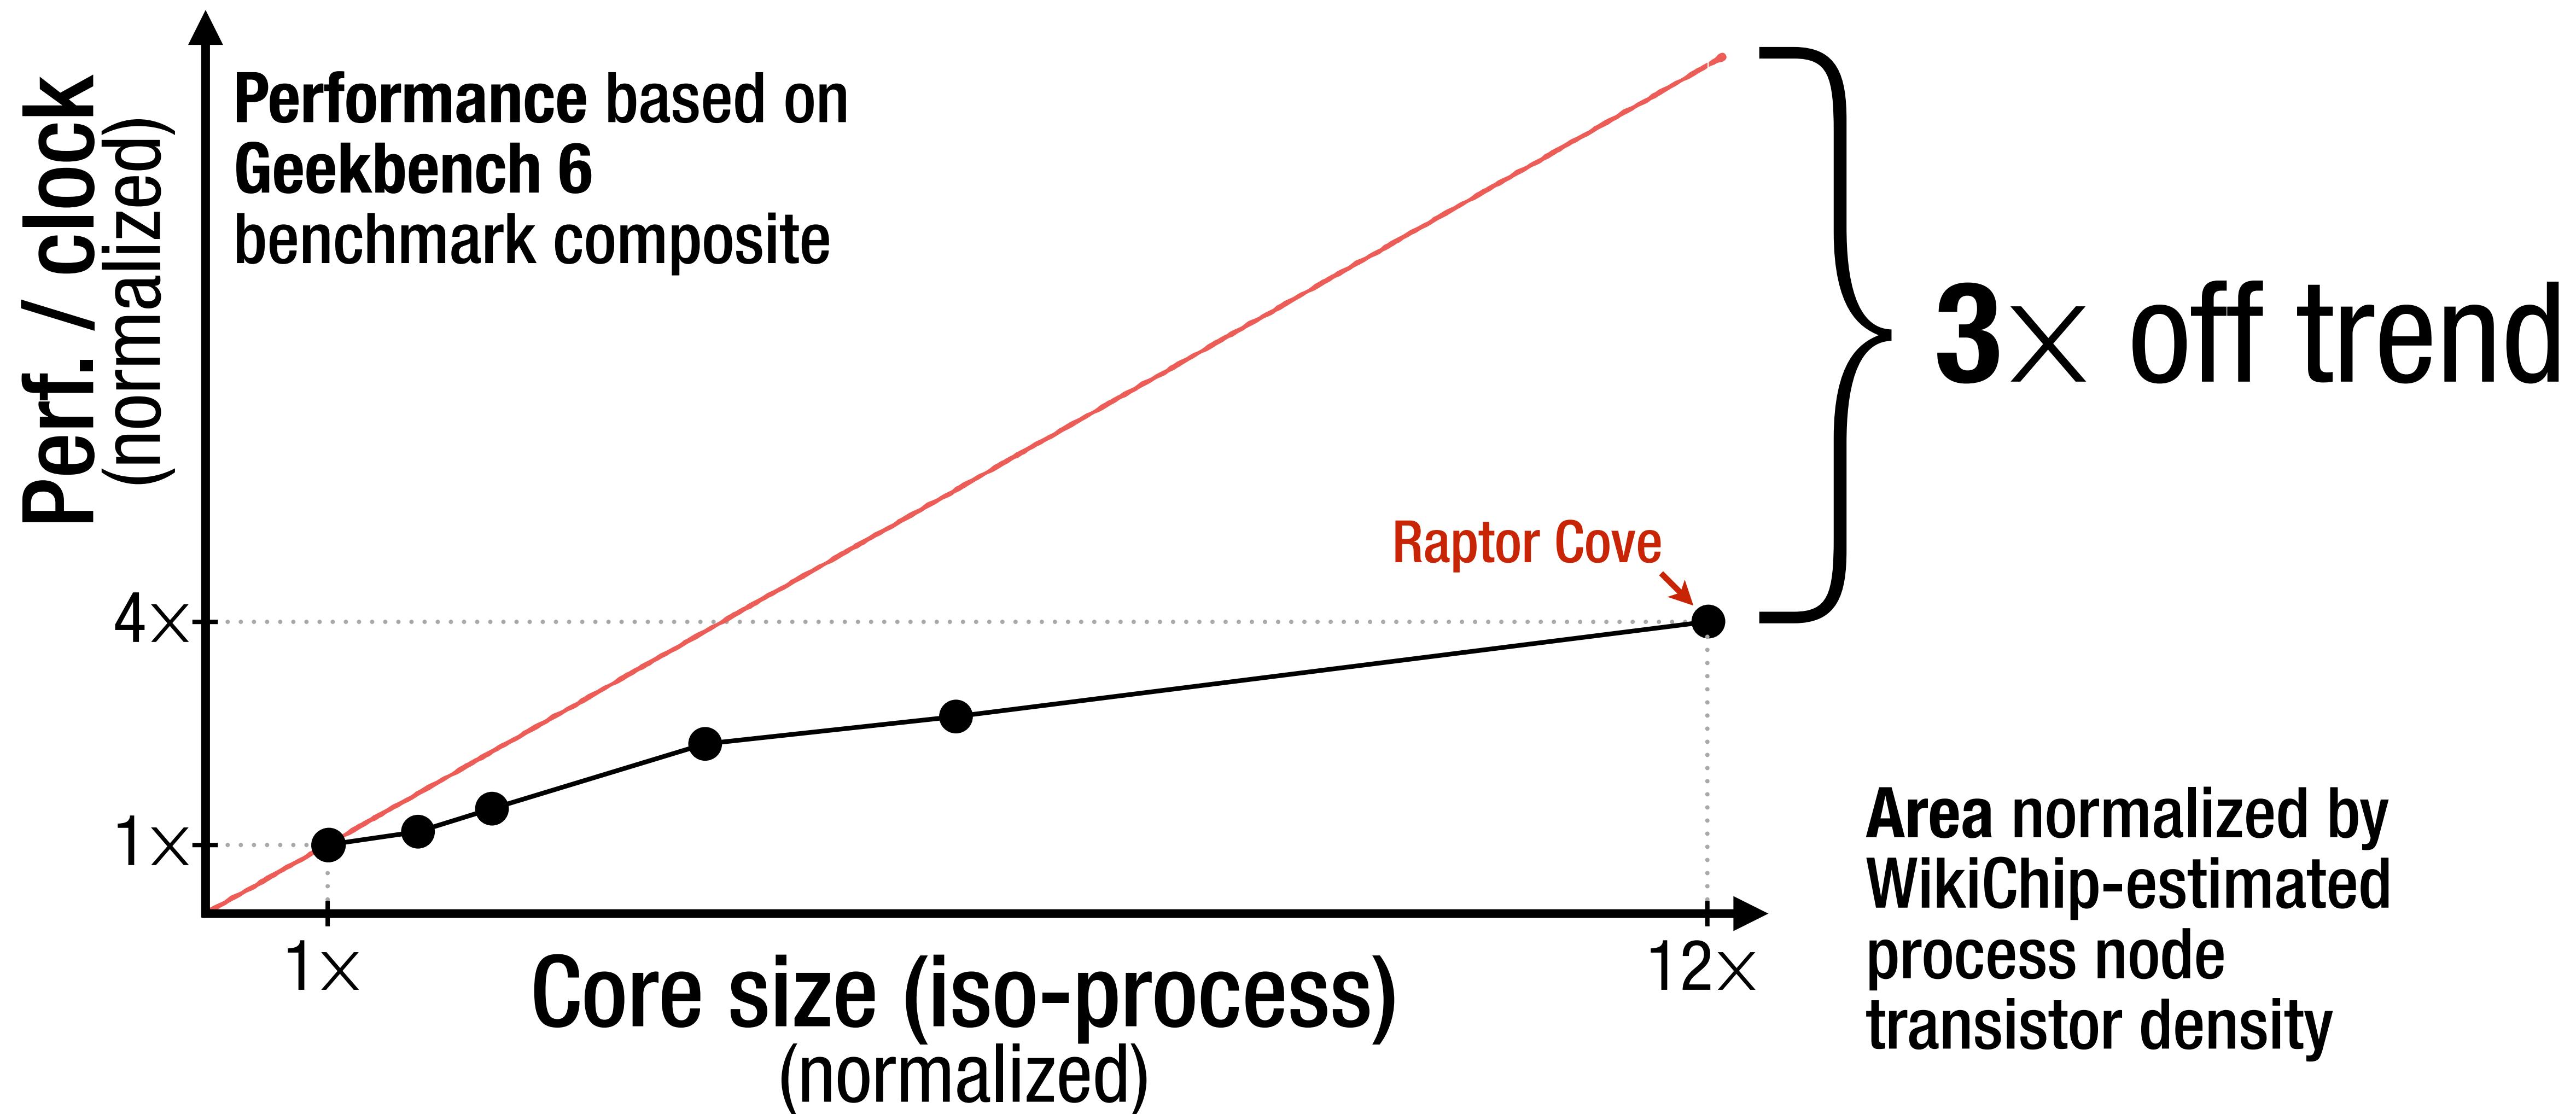
What to do with twice the silicon?



=
equal
area

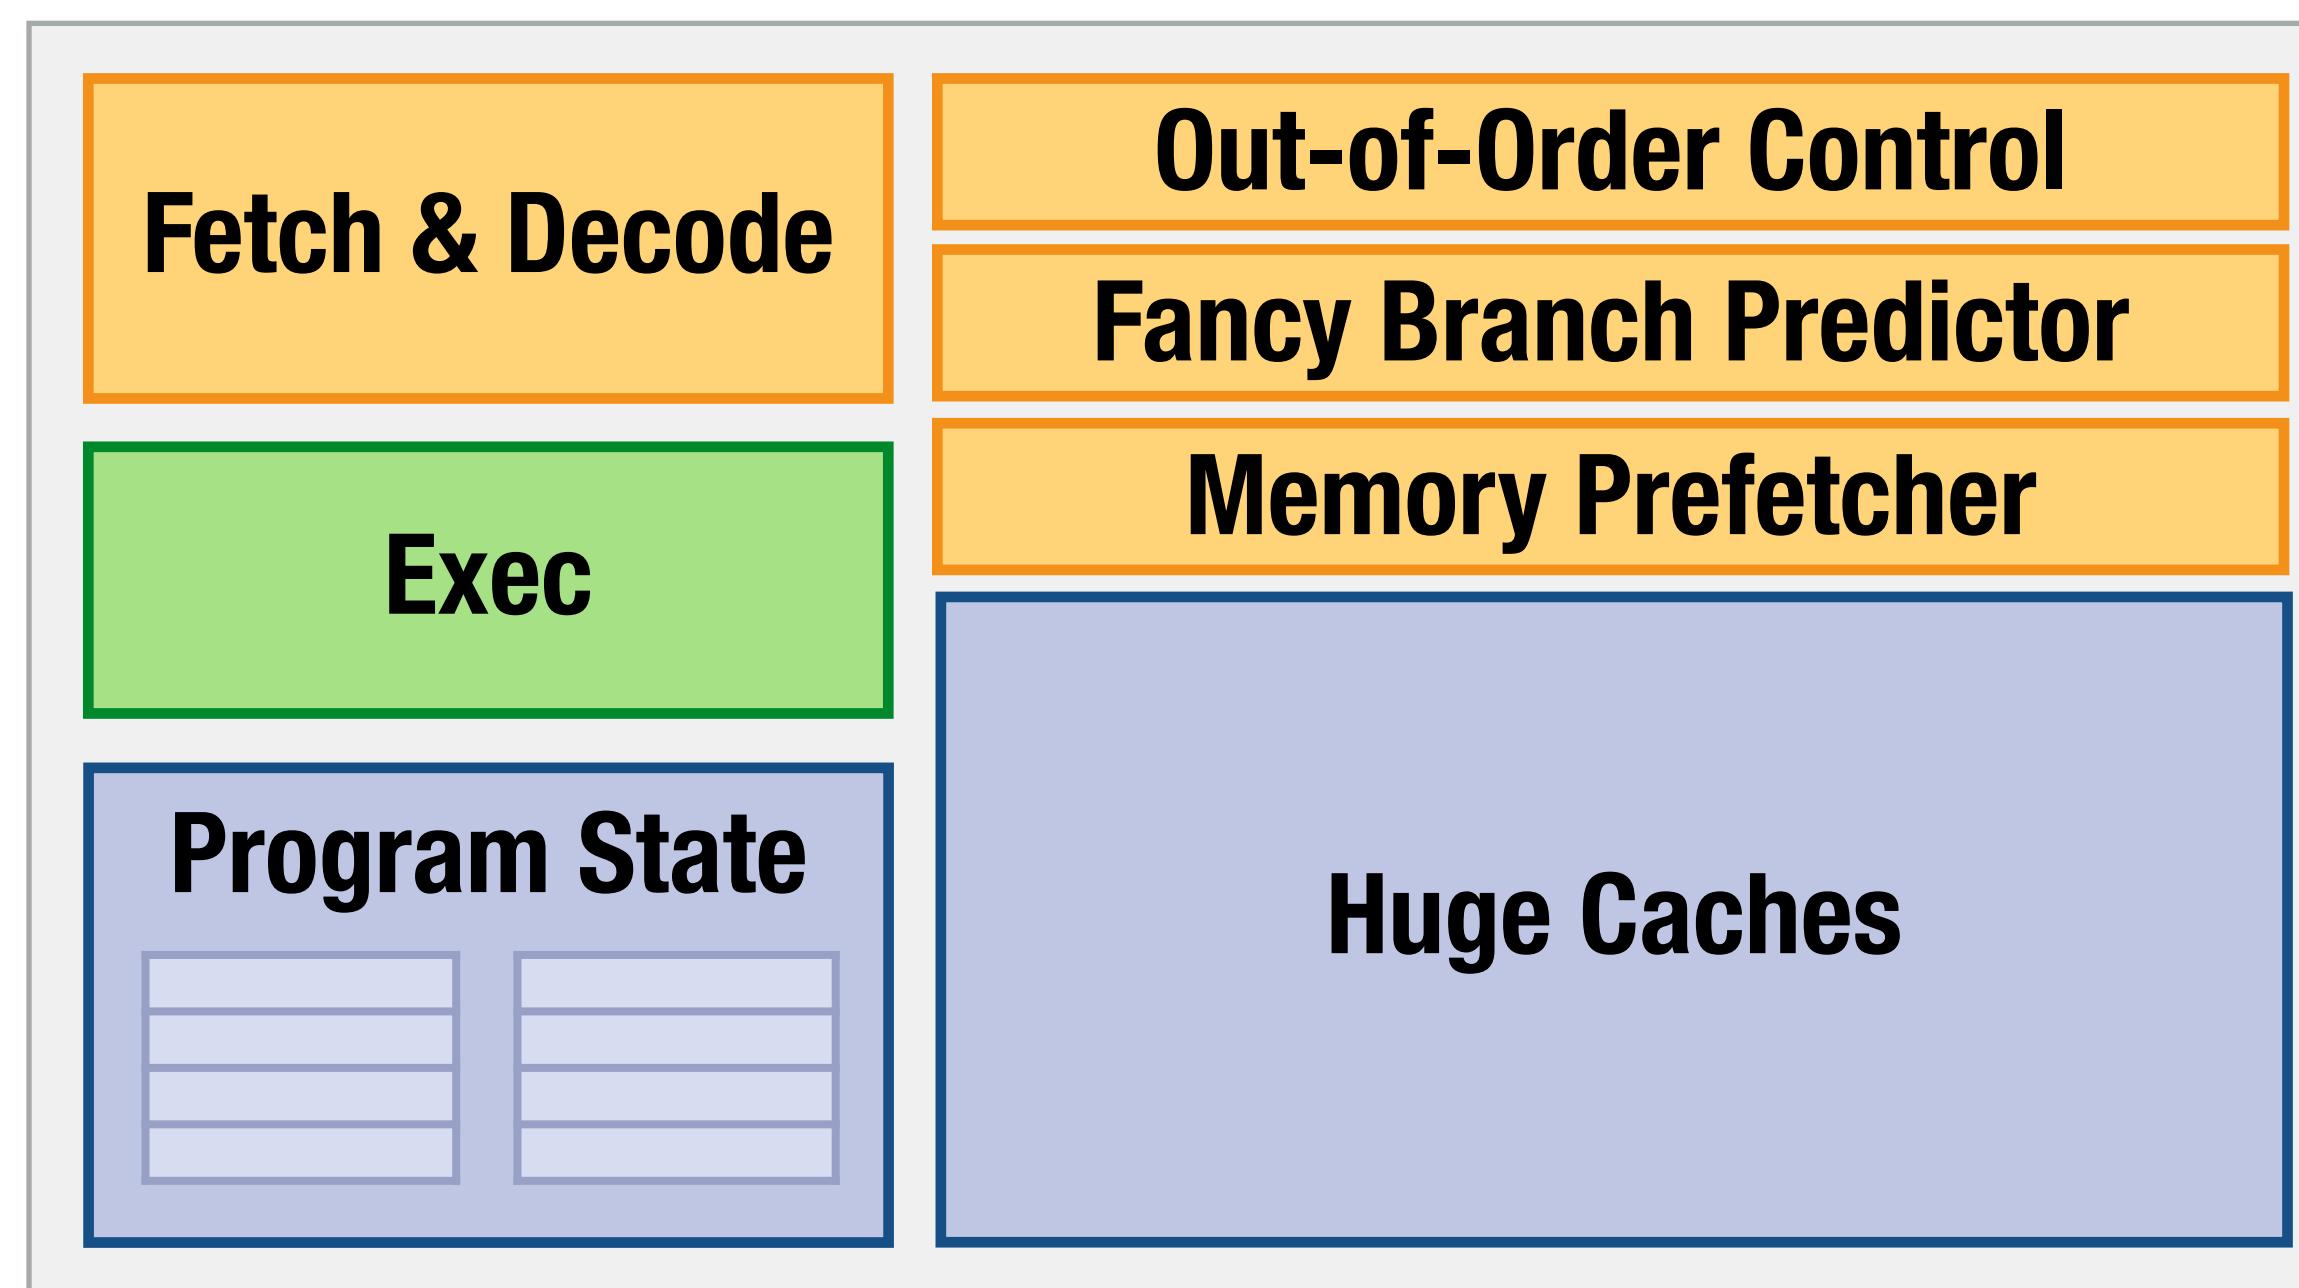


Diminishing returns to scaling single-core performance

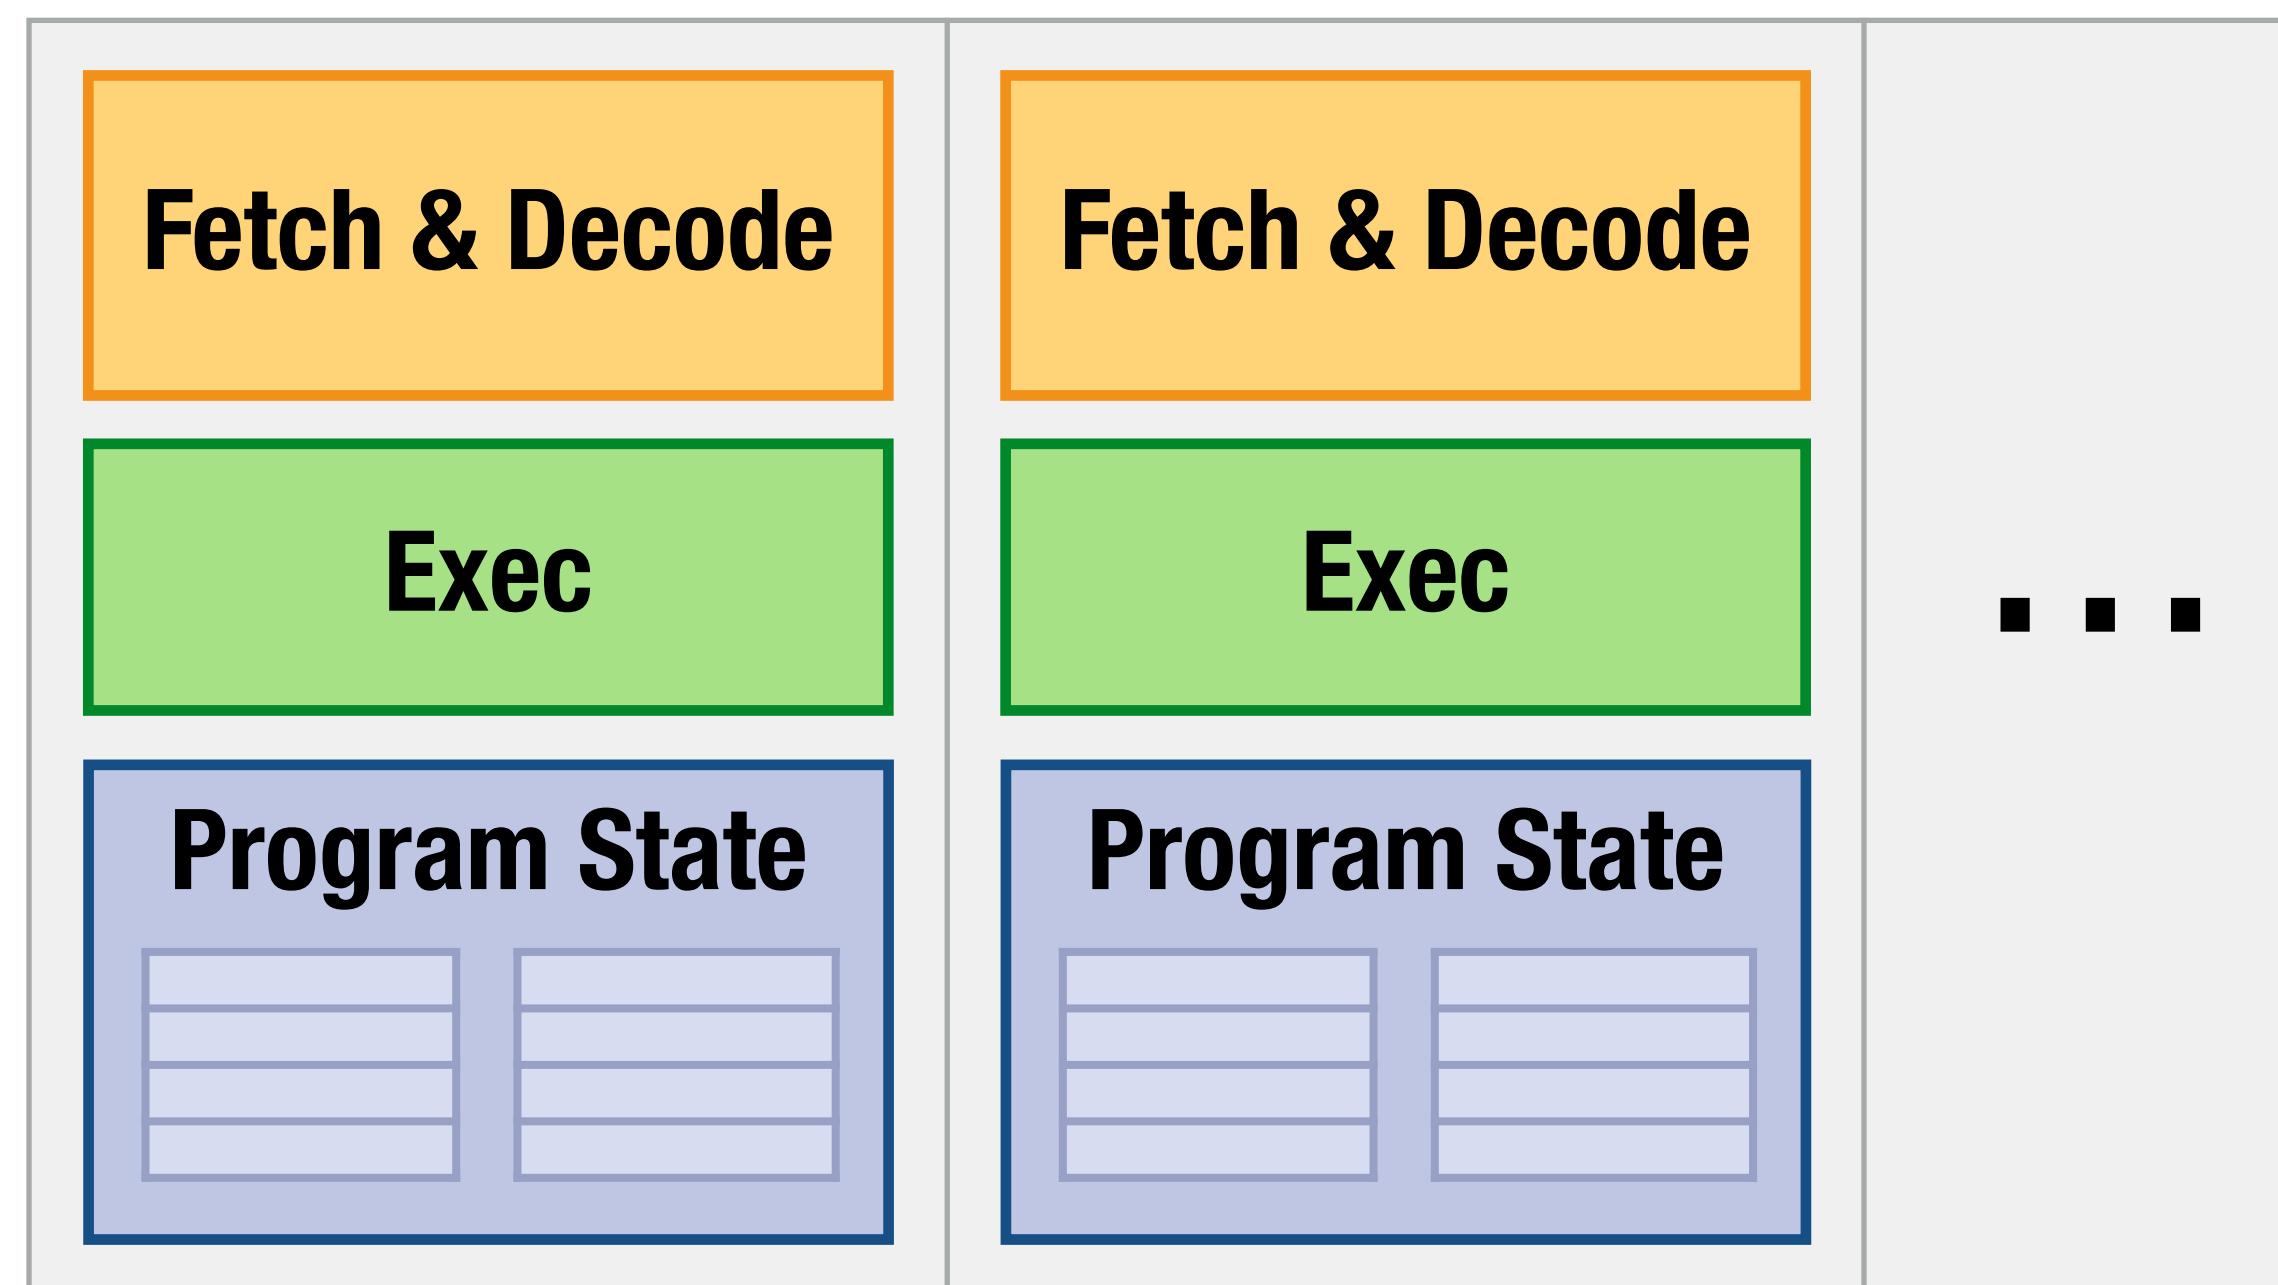


If we want to optimize **throughput**,
is there a **better way** to scale
performance?

Idea 1: remove hardware to optimize single-thread performance



Idea 1: remove hardware to optimize single-thread performance



Invest savings in parallelism

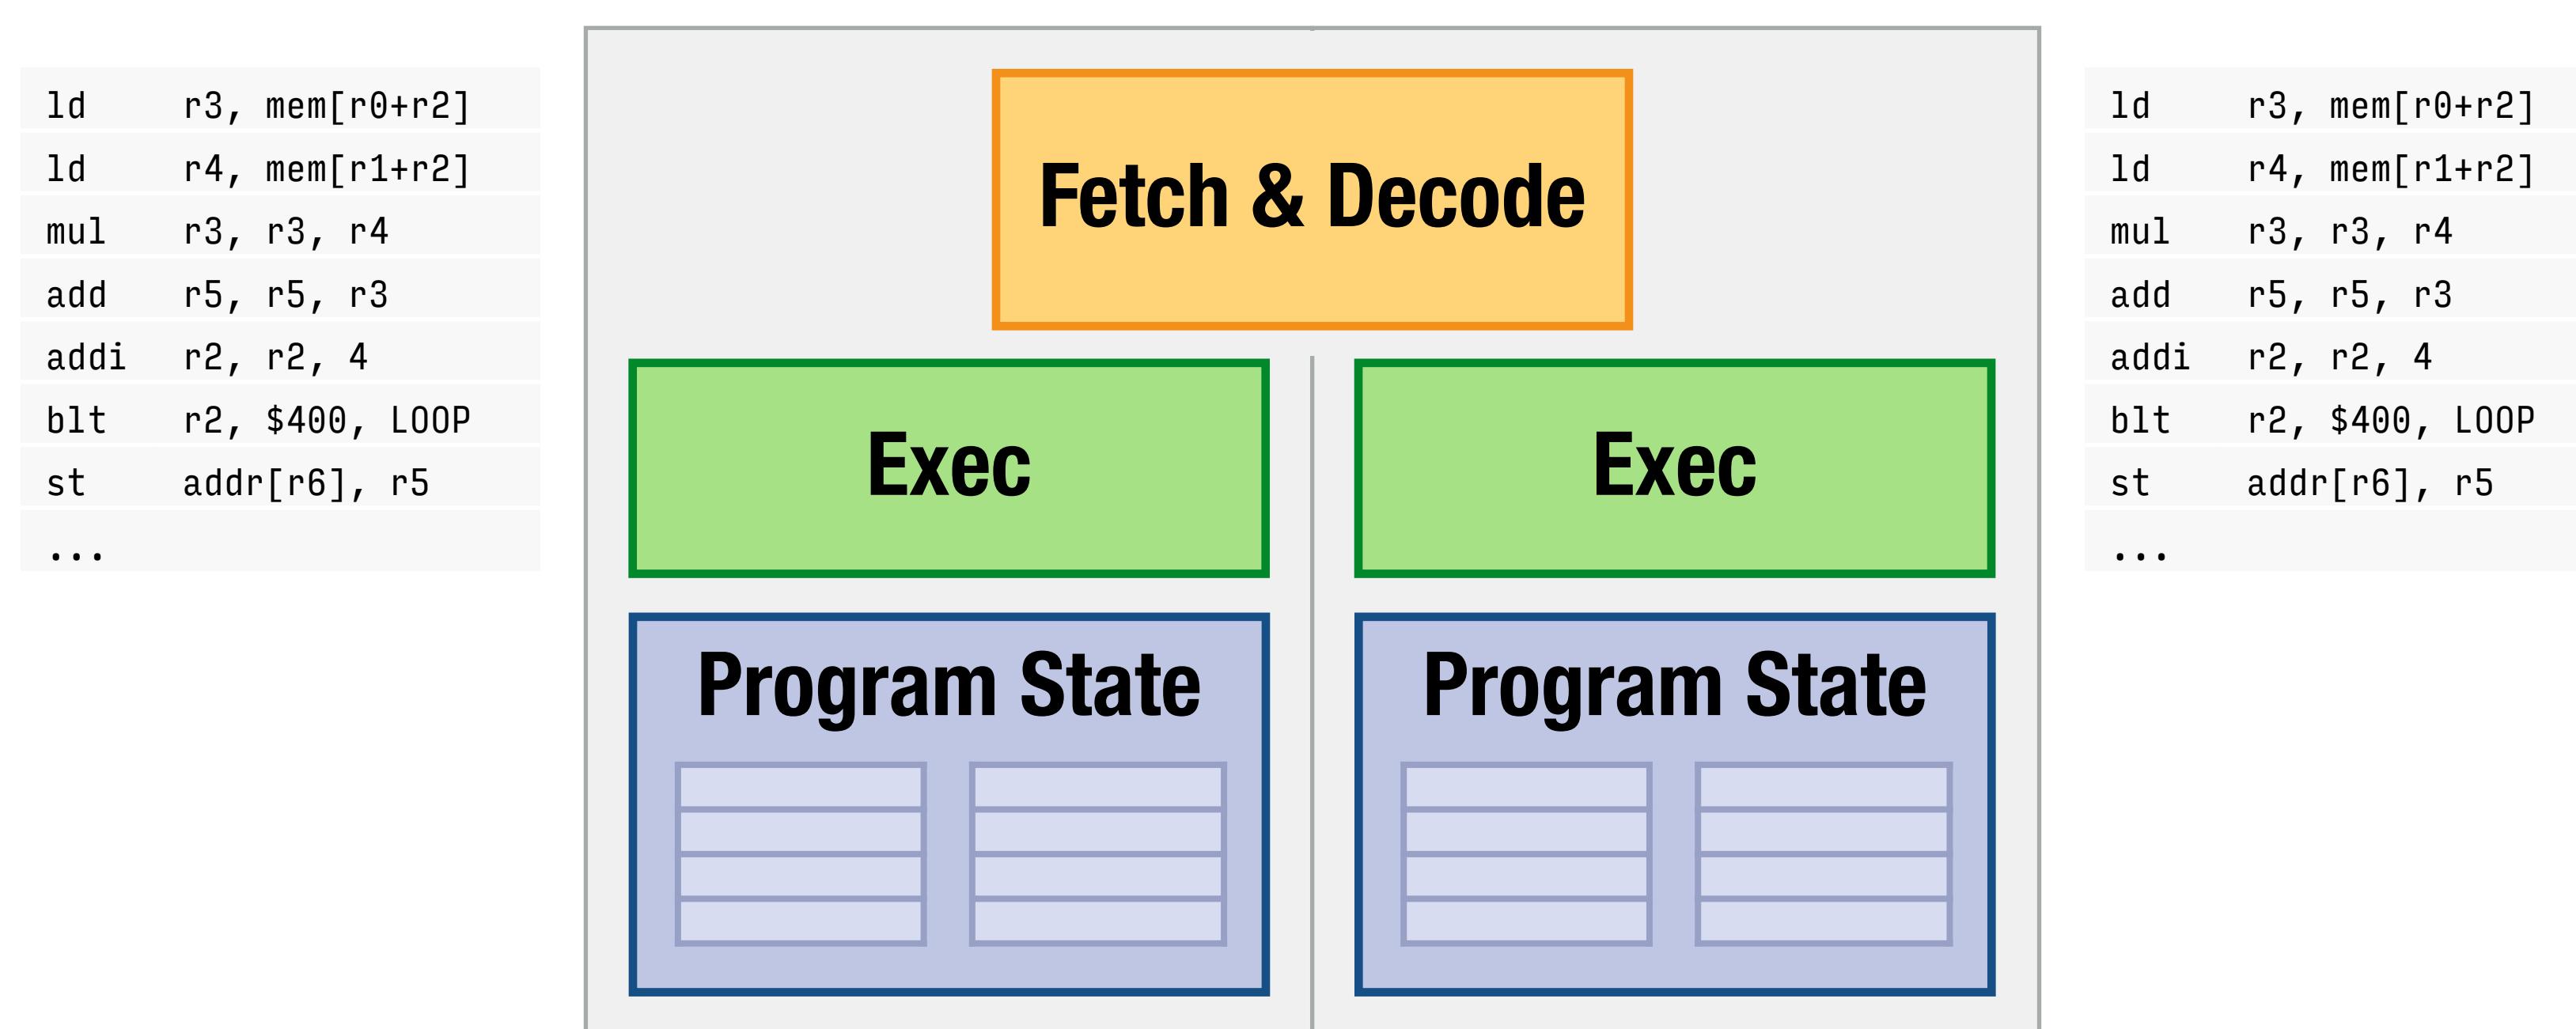
```
for (int i = 0; i < N; ++i)
    for (int k = 0; k < N; ++k)
        for (int j = 0; j < N; ++j)
            C[i*N + j] += A[i*N + k] * B[k*N + j];
```

```
for (int j = 0; j < N; ++j)  
    C[i*N + j] += A[i*N + k] * B[k*N + j];
```

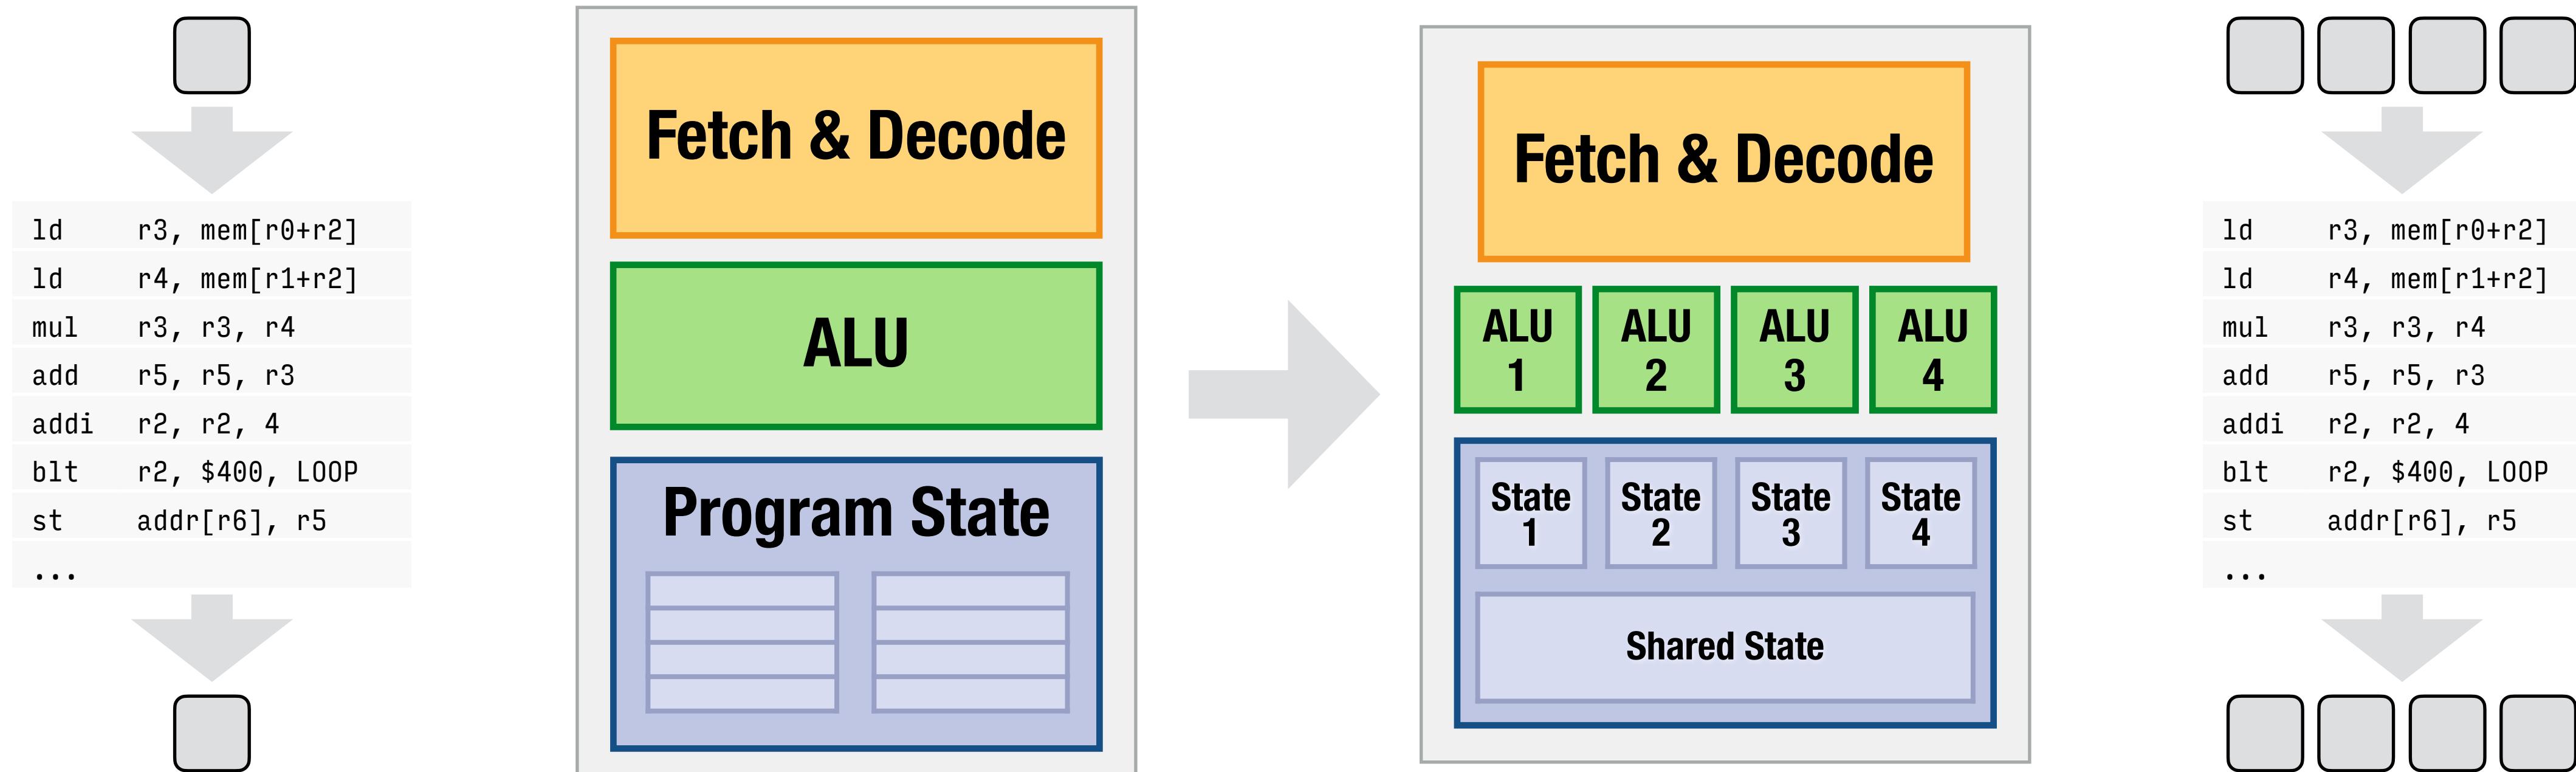
```

for all (int j = 0; j < N; ++j)
    C[i*N + j] += A[i*N + k] * B[k*N + j];

```



Idea 2: amortize control overhead with SIMD execution

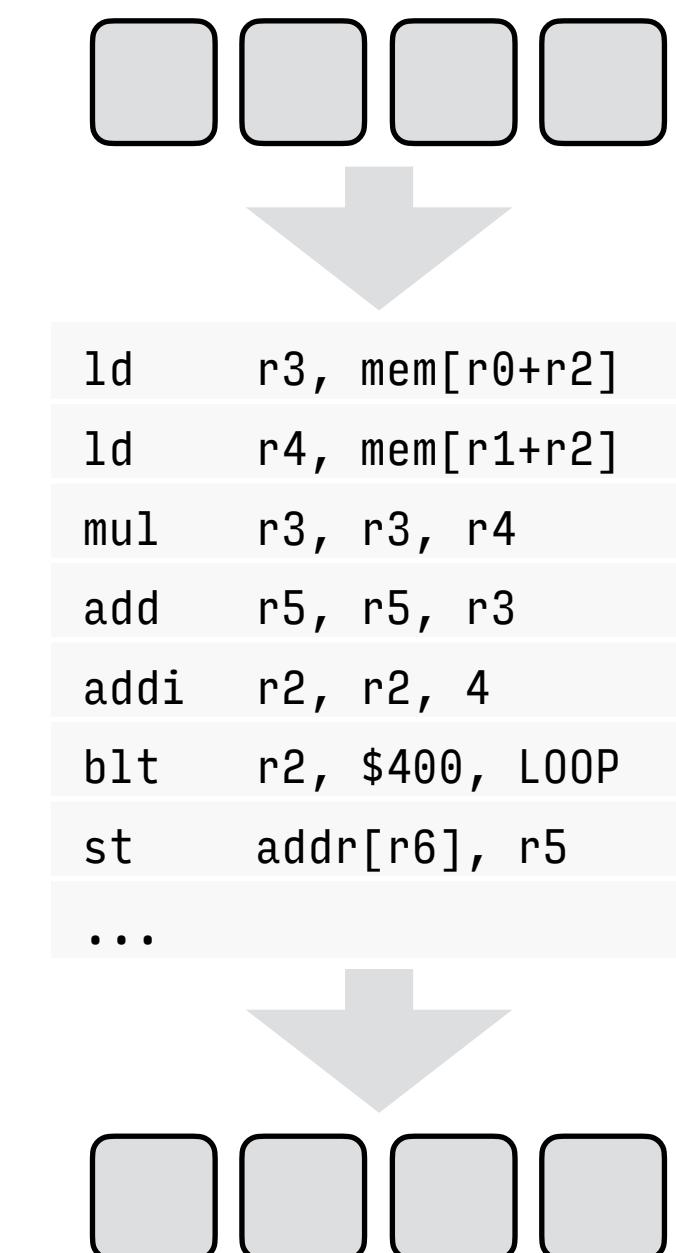
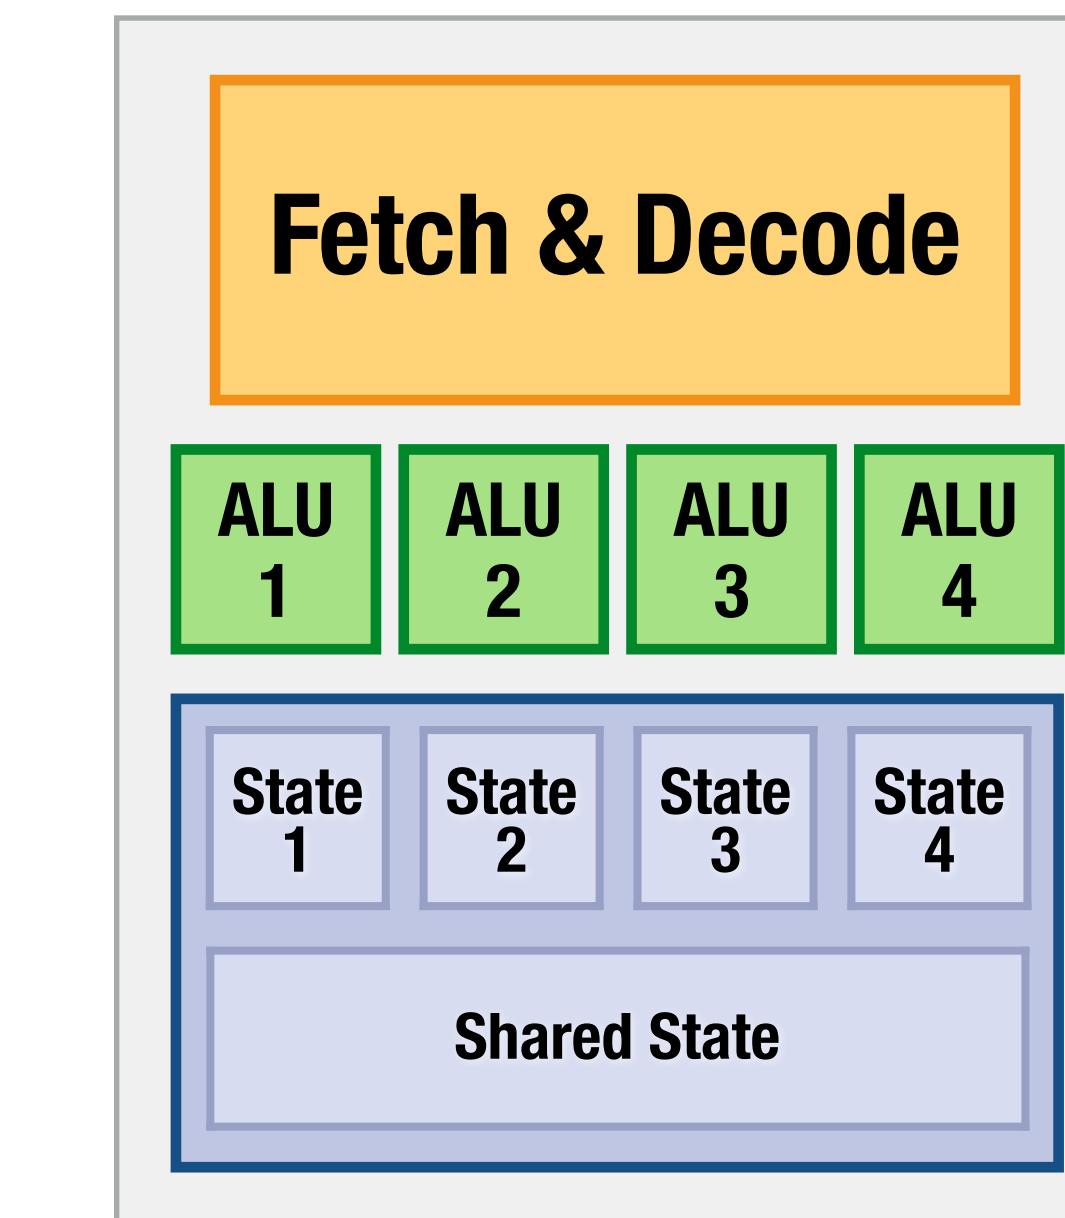


Idea 2: amortize control overhead with SIMD execution

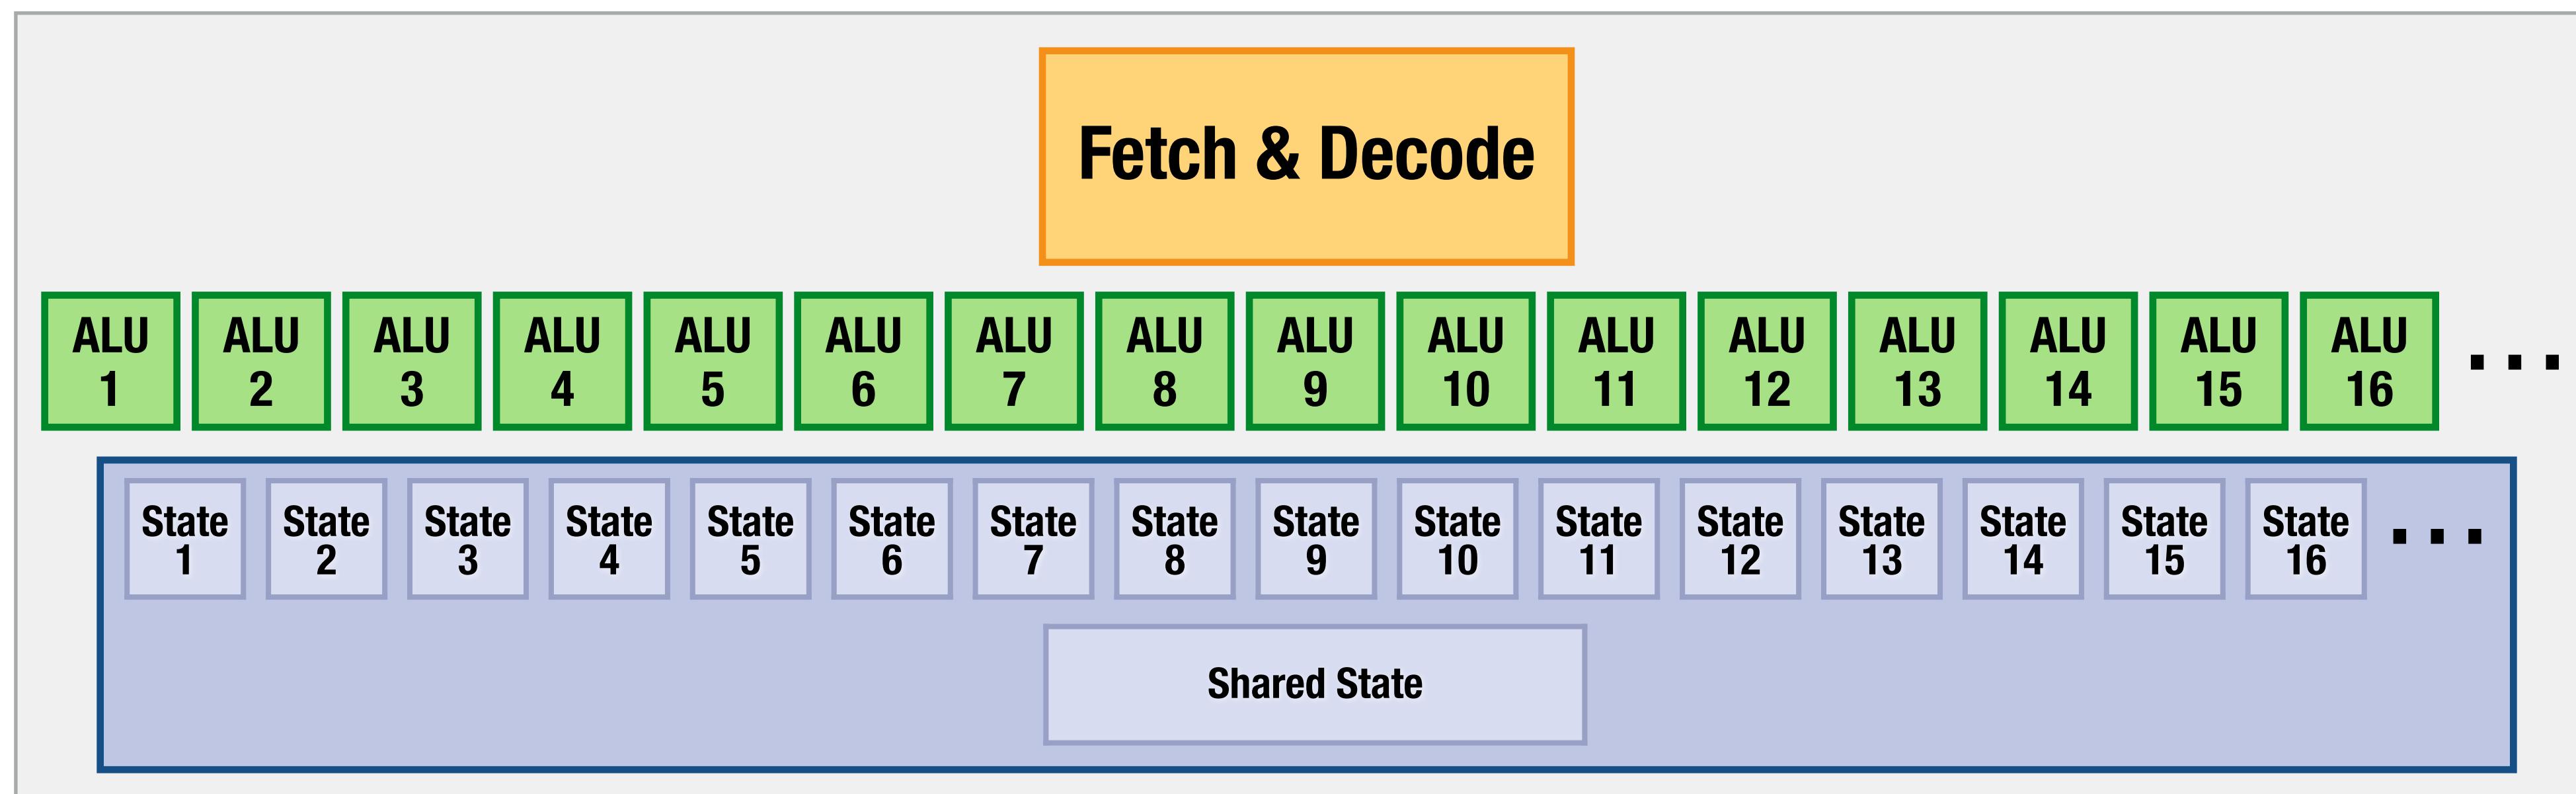
SIMD can be
explicit in ISA
or **implicit in hardware**

Intel AVX,
ARM NEON,
etc.

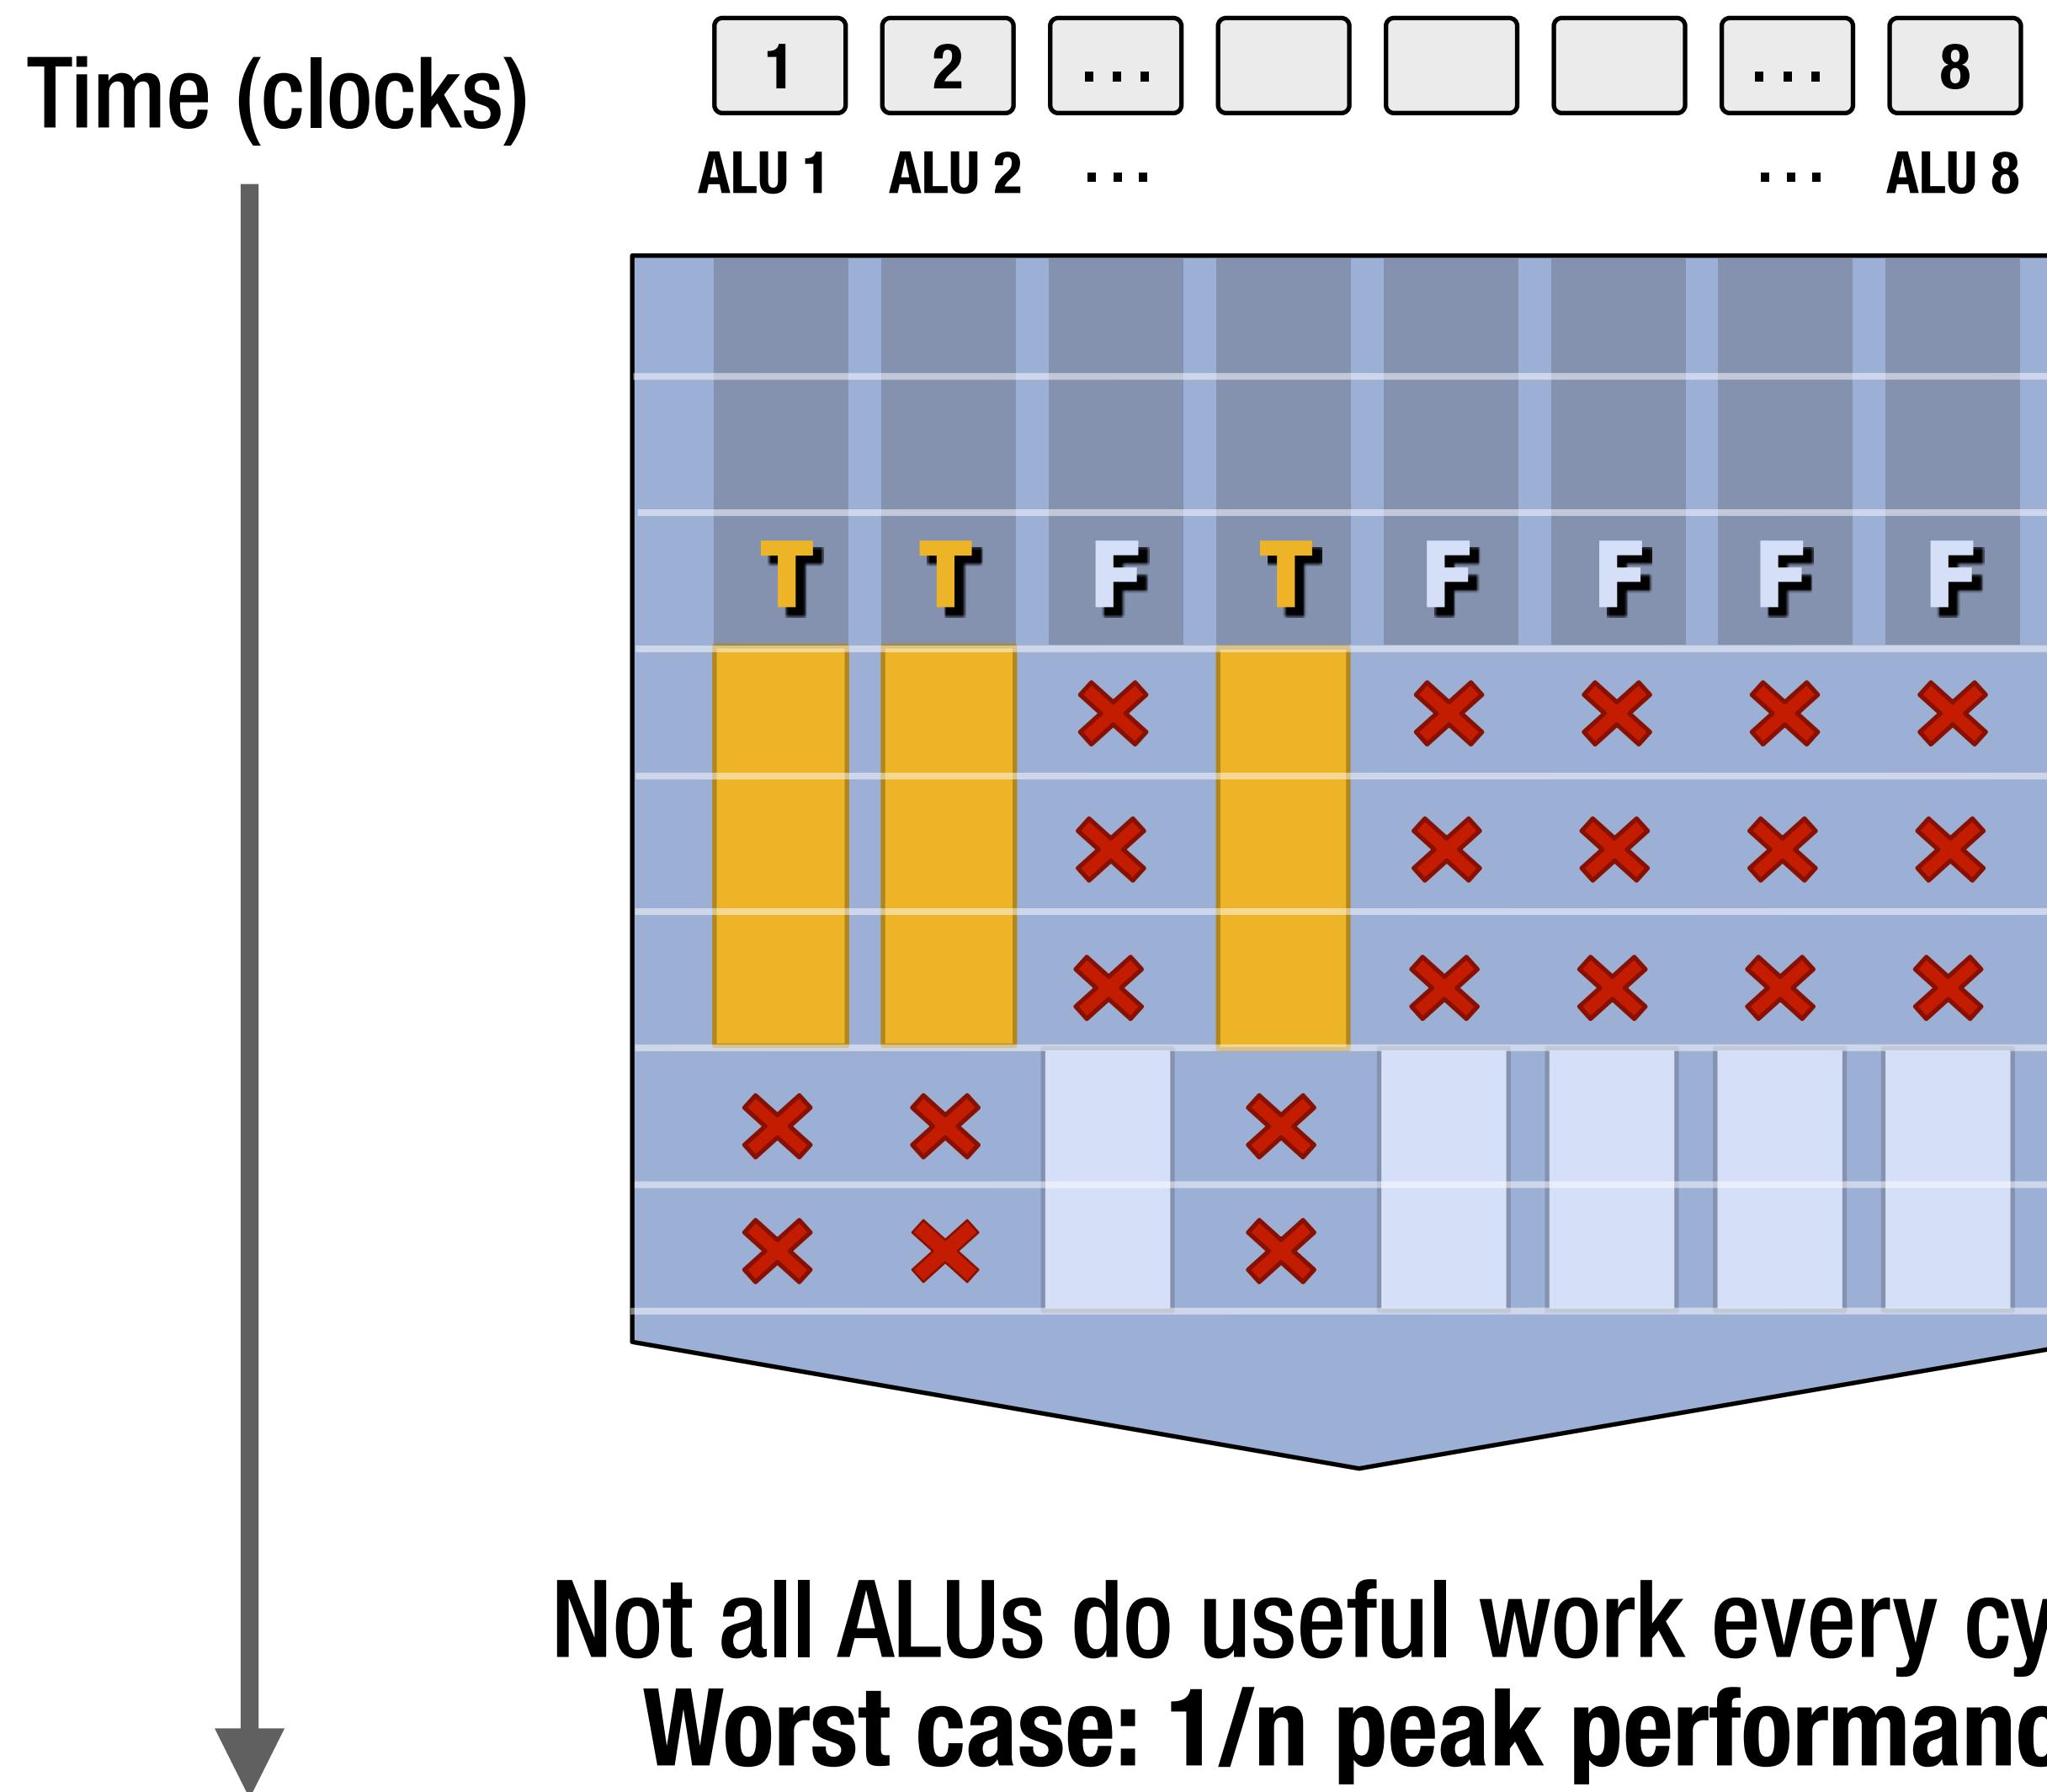
many GPUs



SIMD execution requires coherent control



SIMD execution requires coherent control

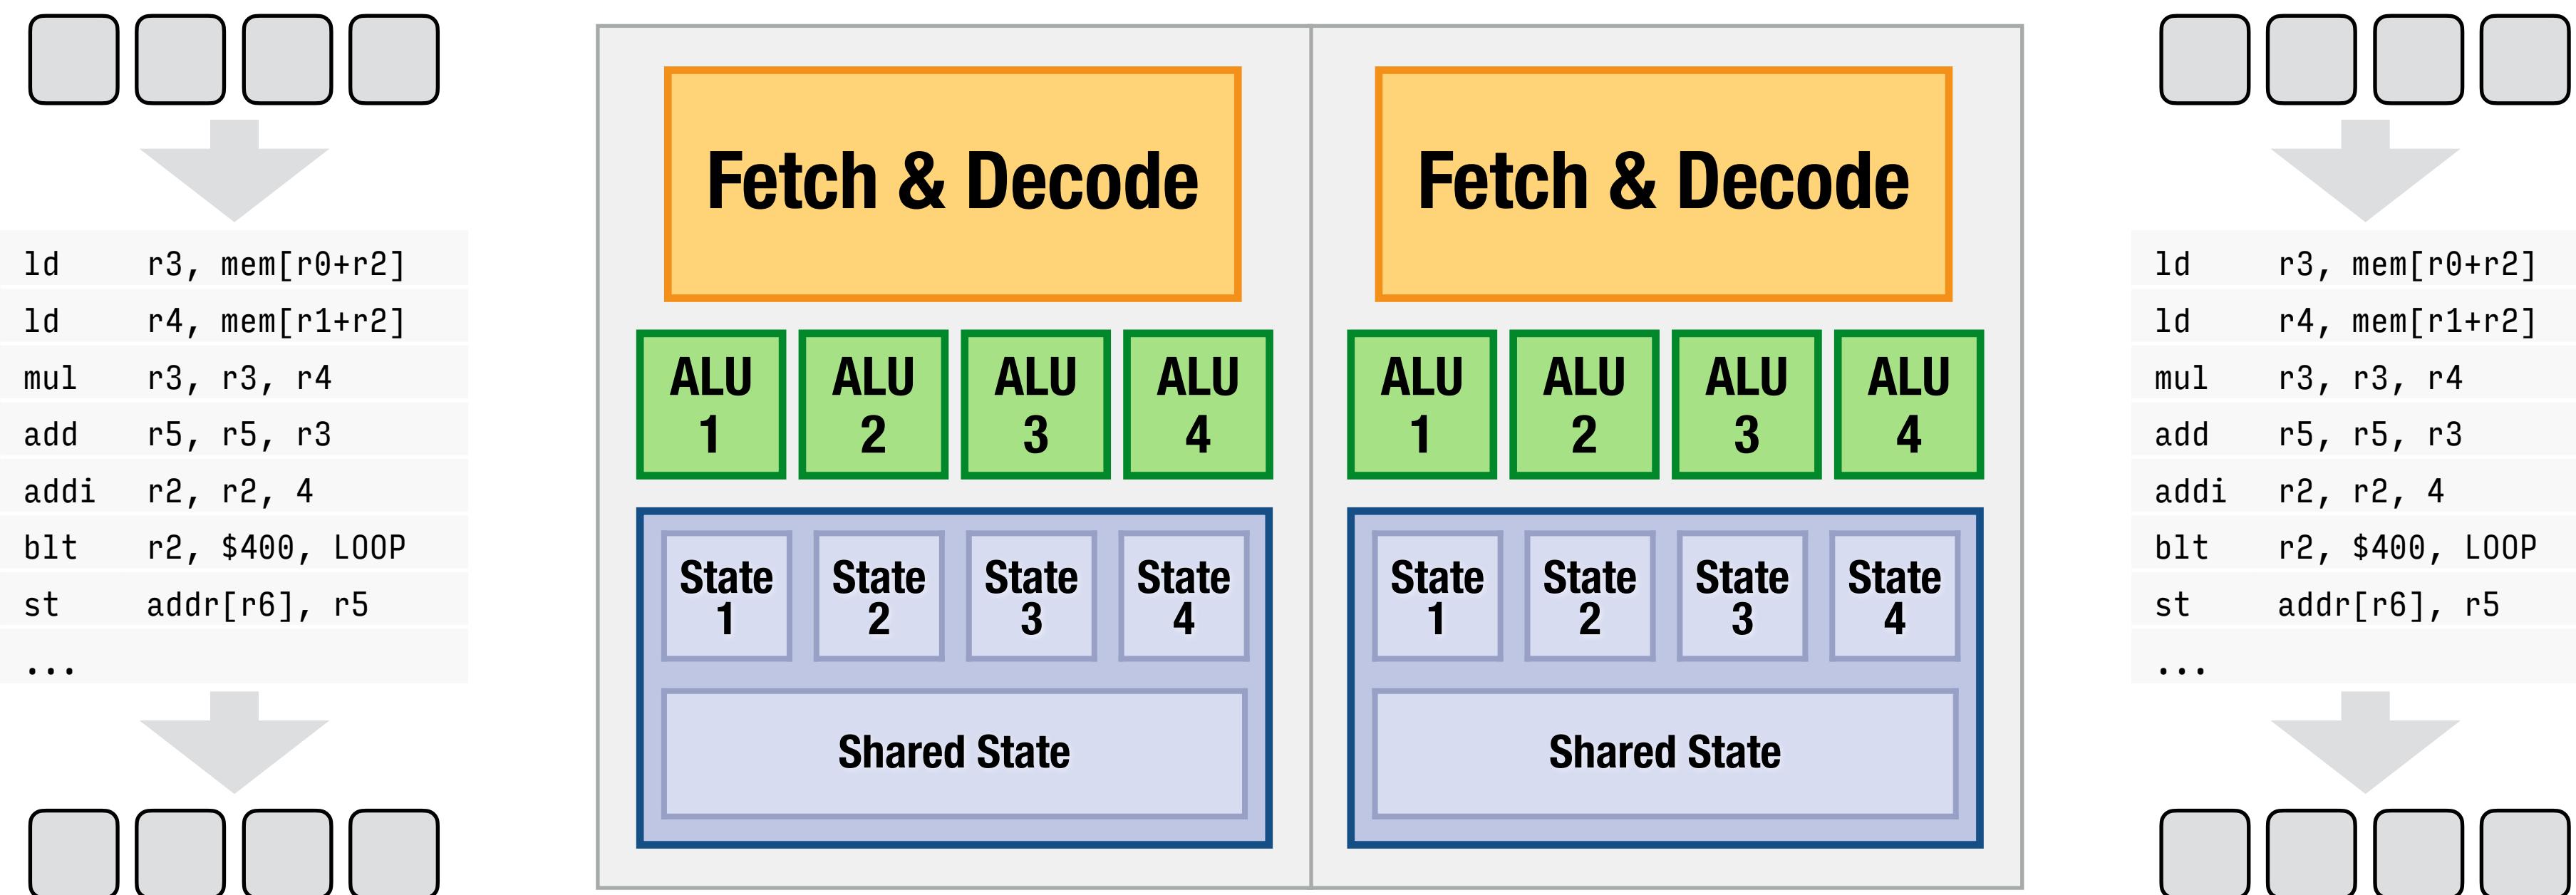


```
<unconditional code>

if (x > 0) {
    y = pow(x, exp);
    y *= Ks;
    refl = y + Ka;
} else {
    x = 0;
    refl = Ka;
}

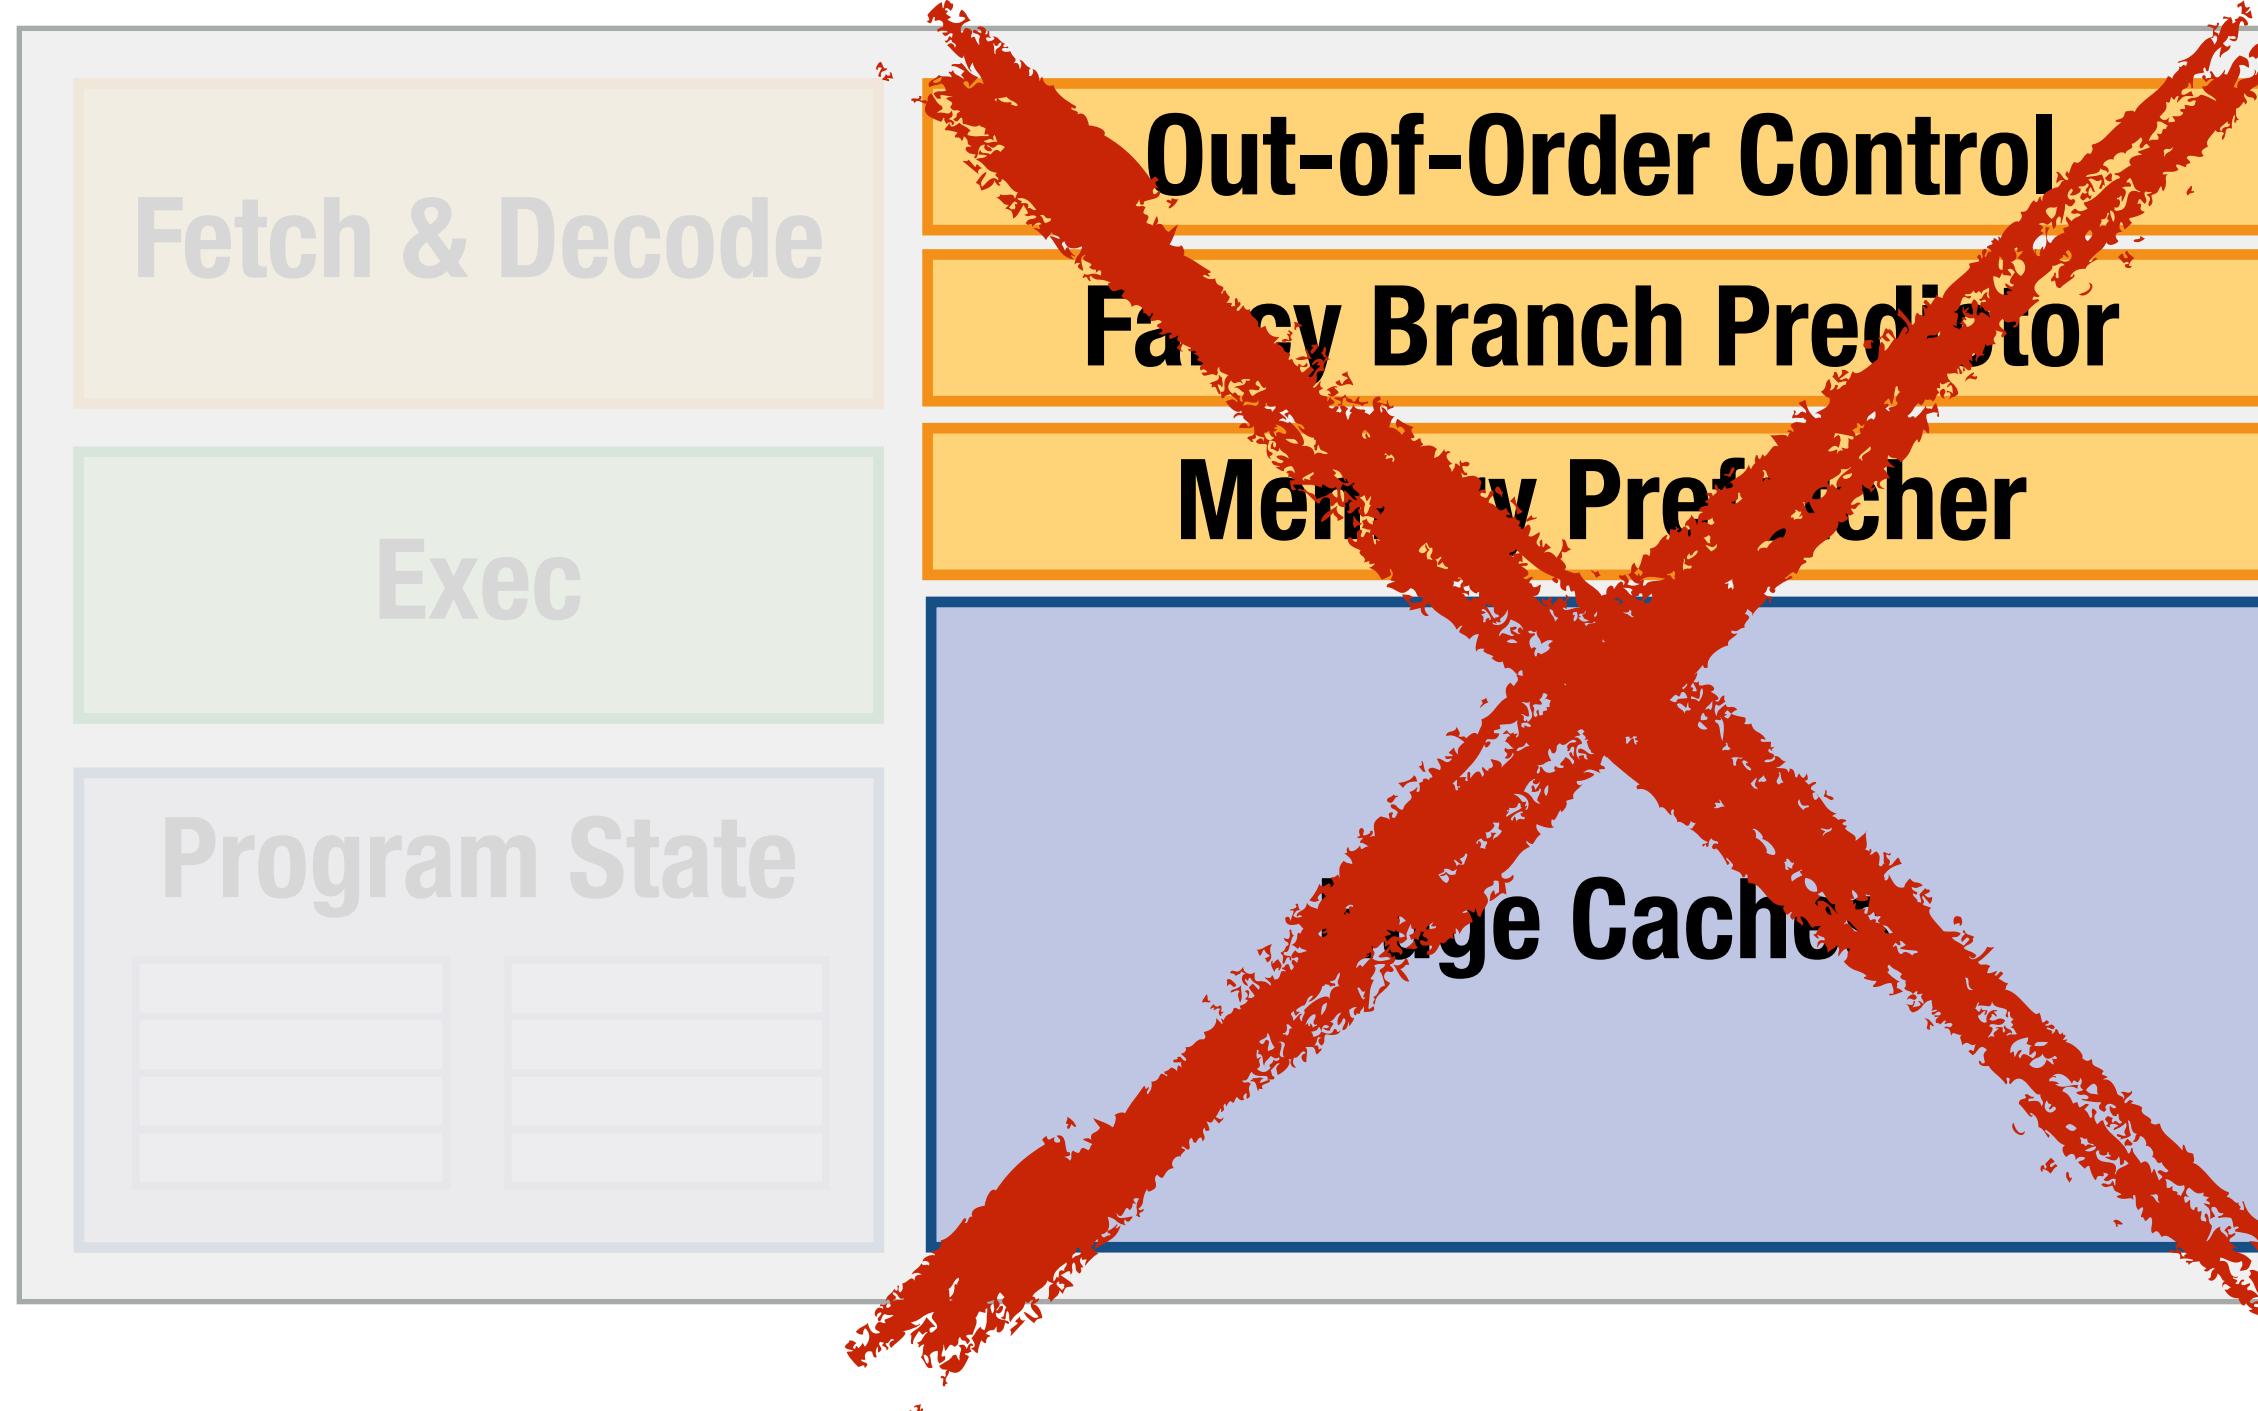
<resume unconditional code>
```

Diminishing returns, scale further with multicore



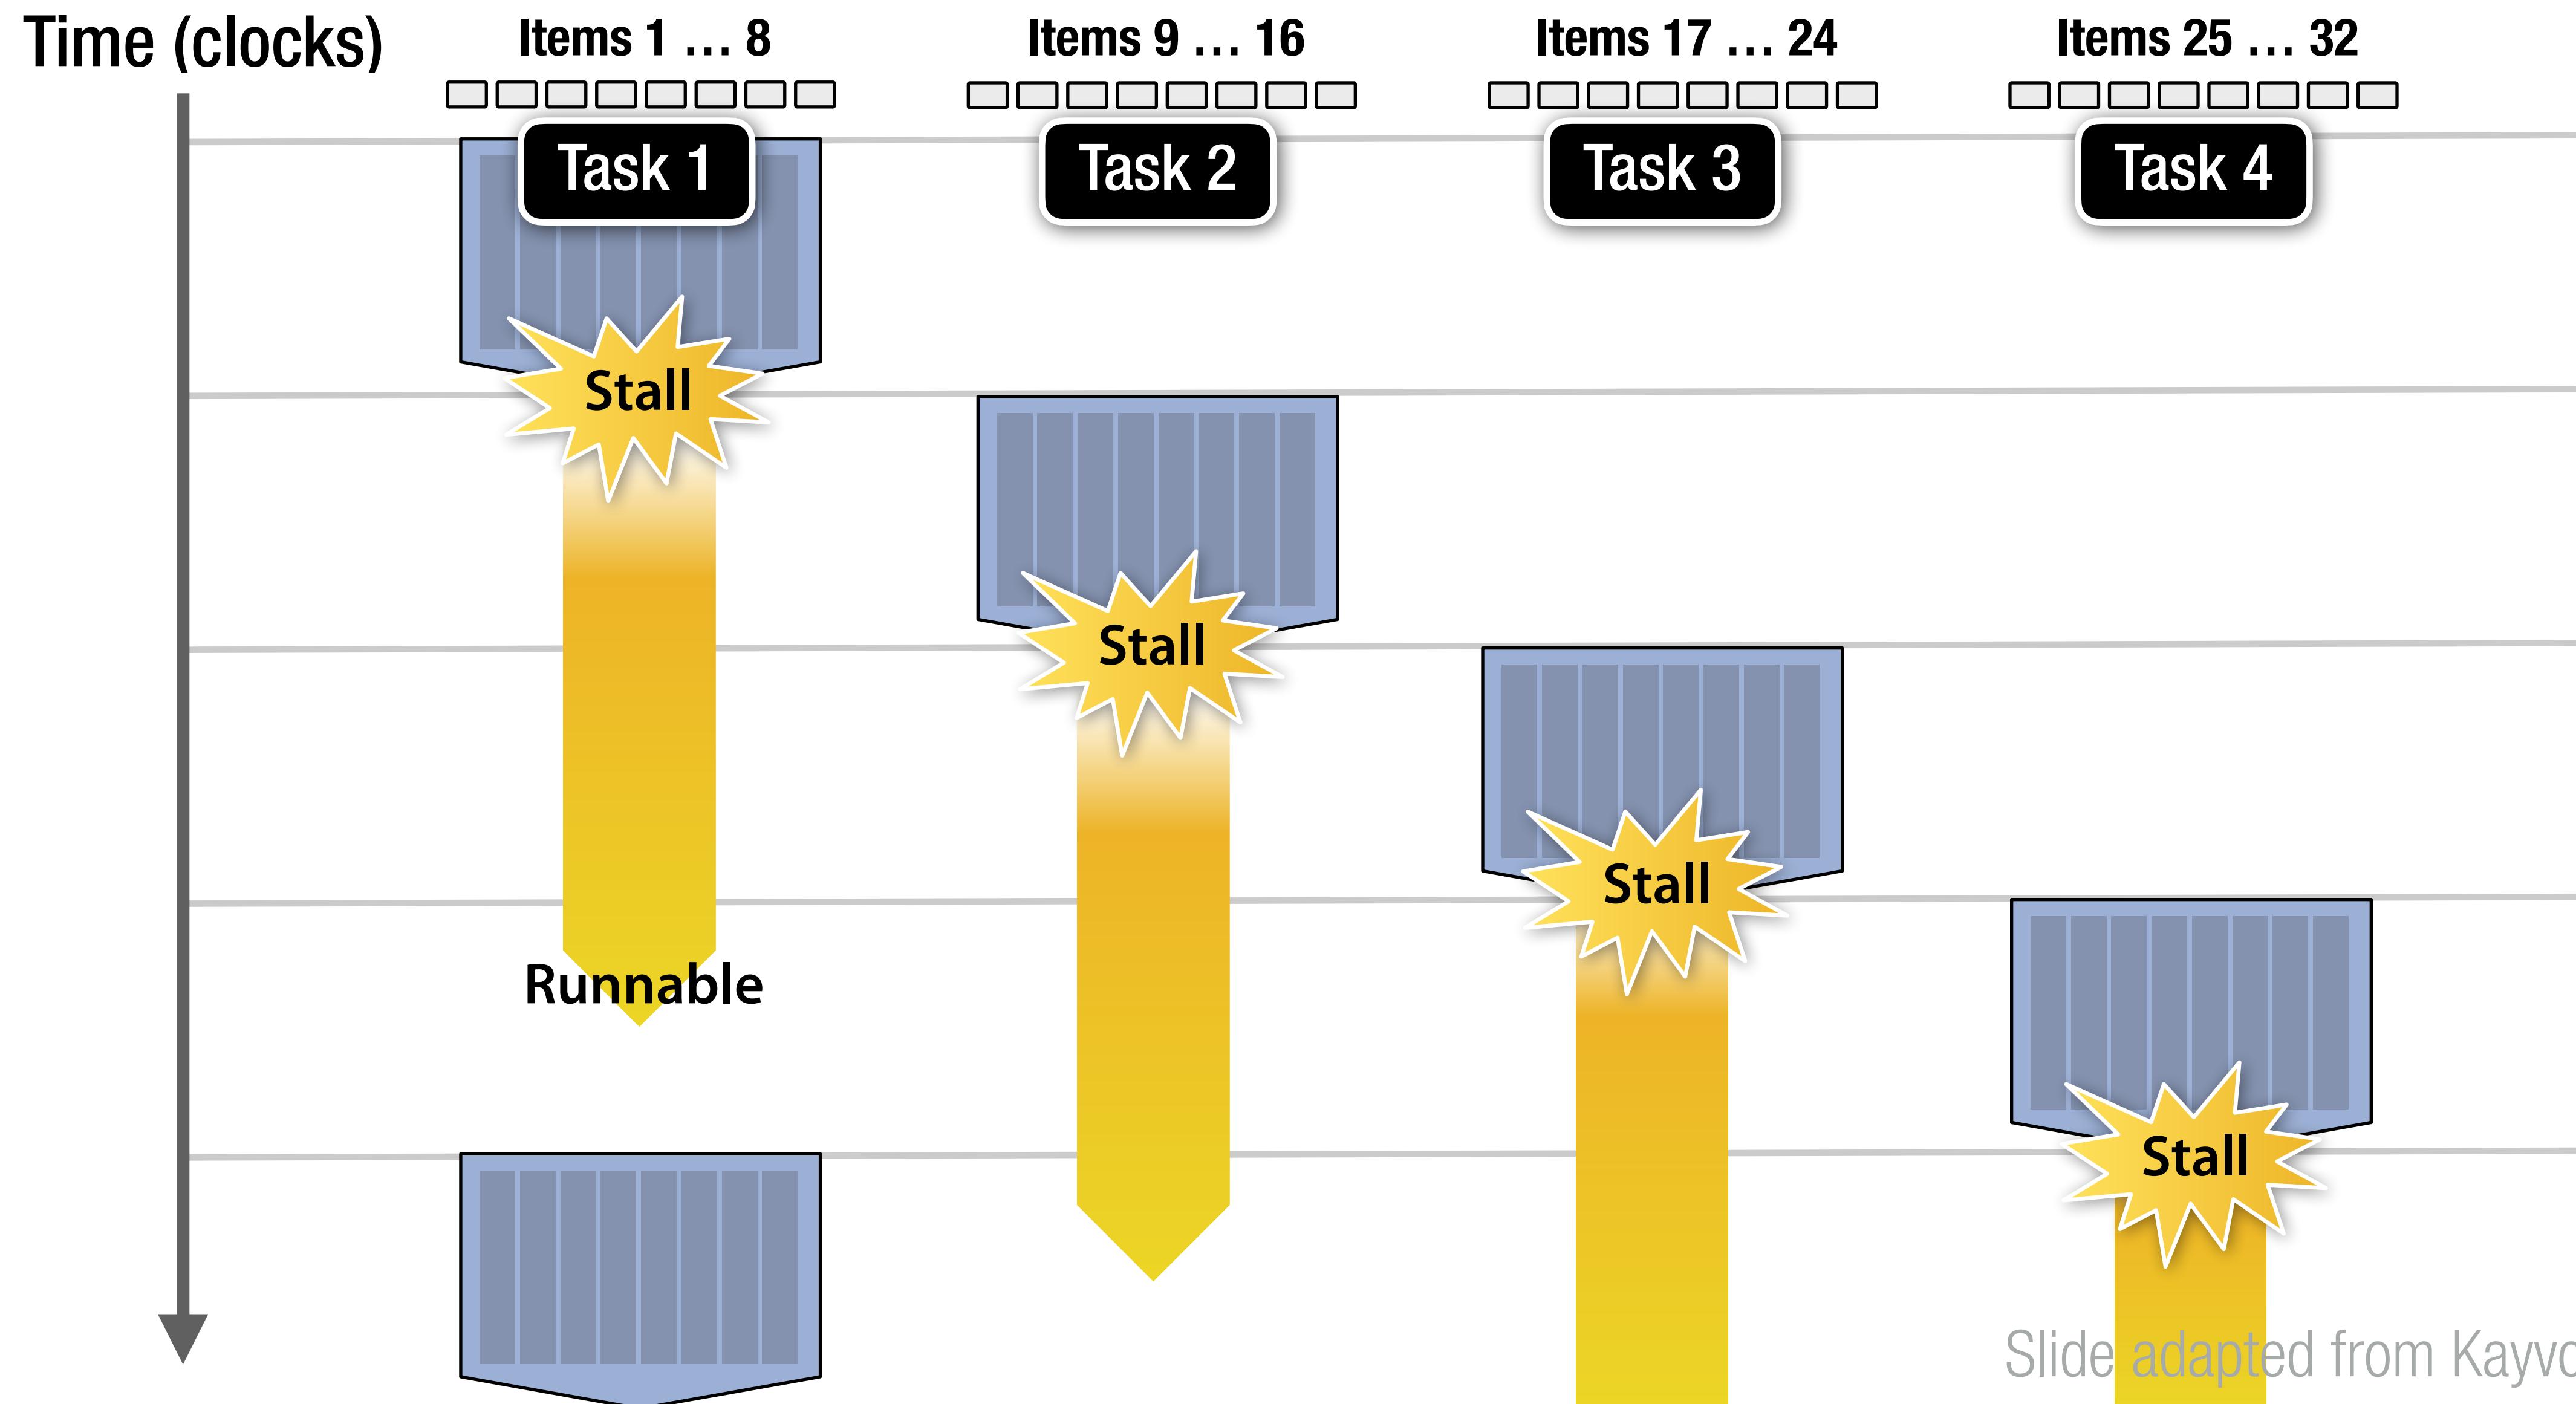
Typical SIMD width: 8 ~ 64

How to deal with latency without stalling



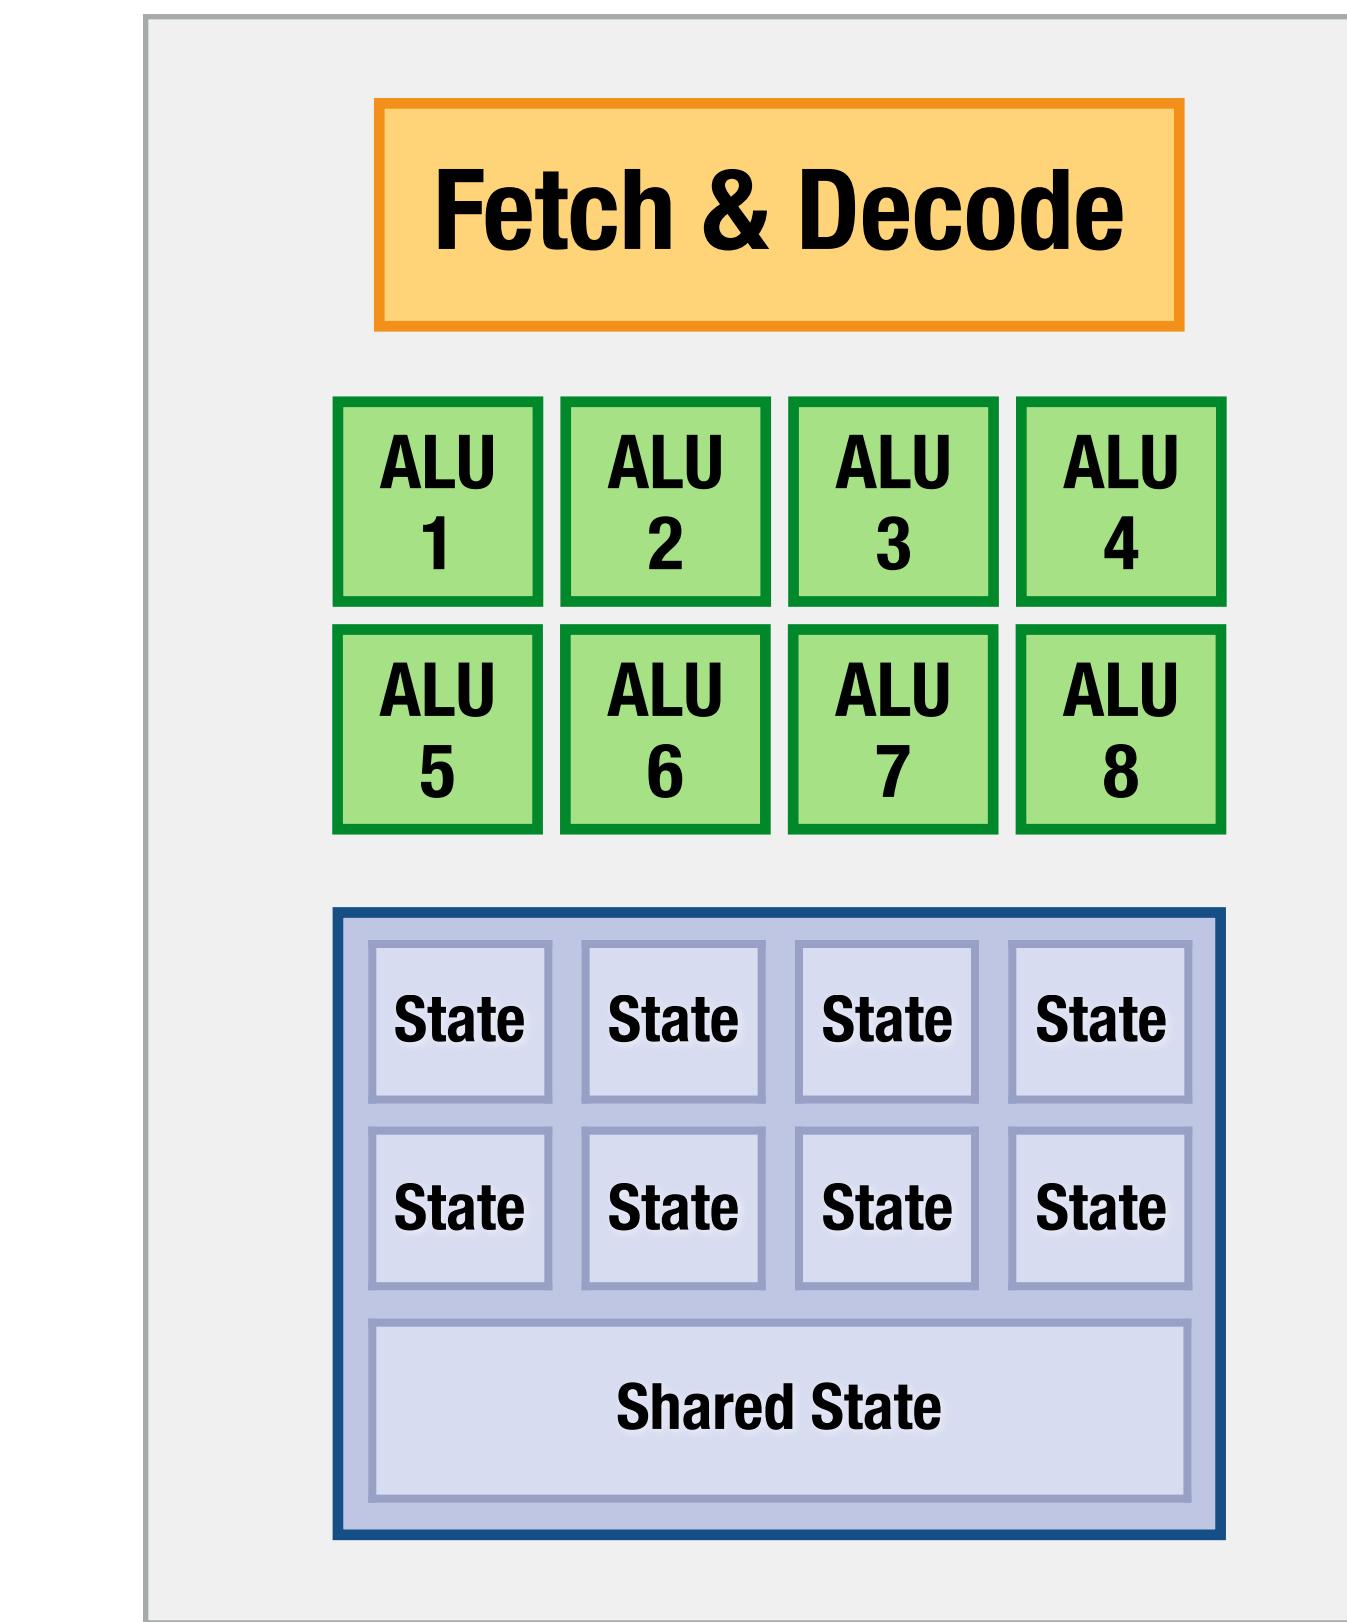
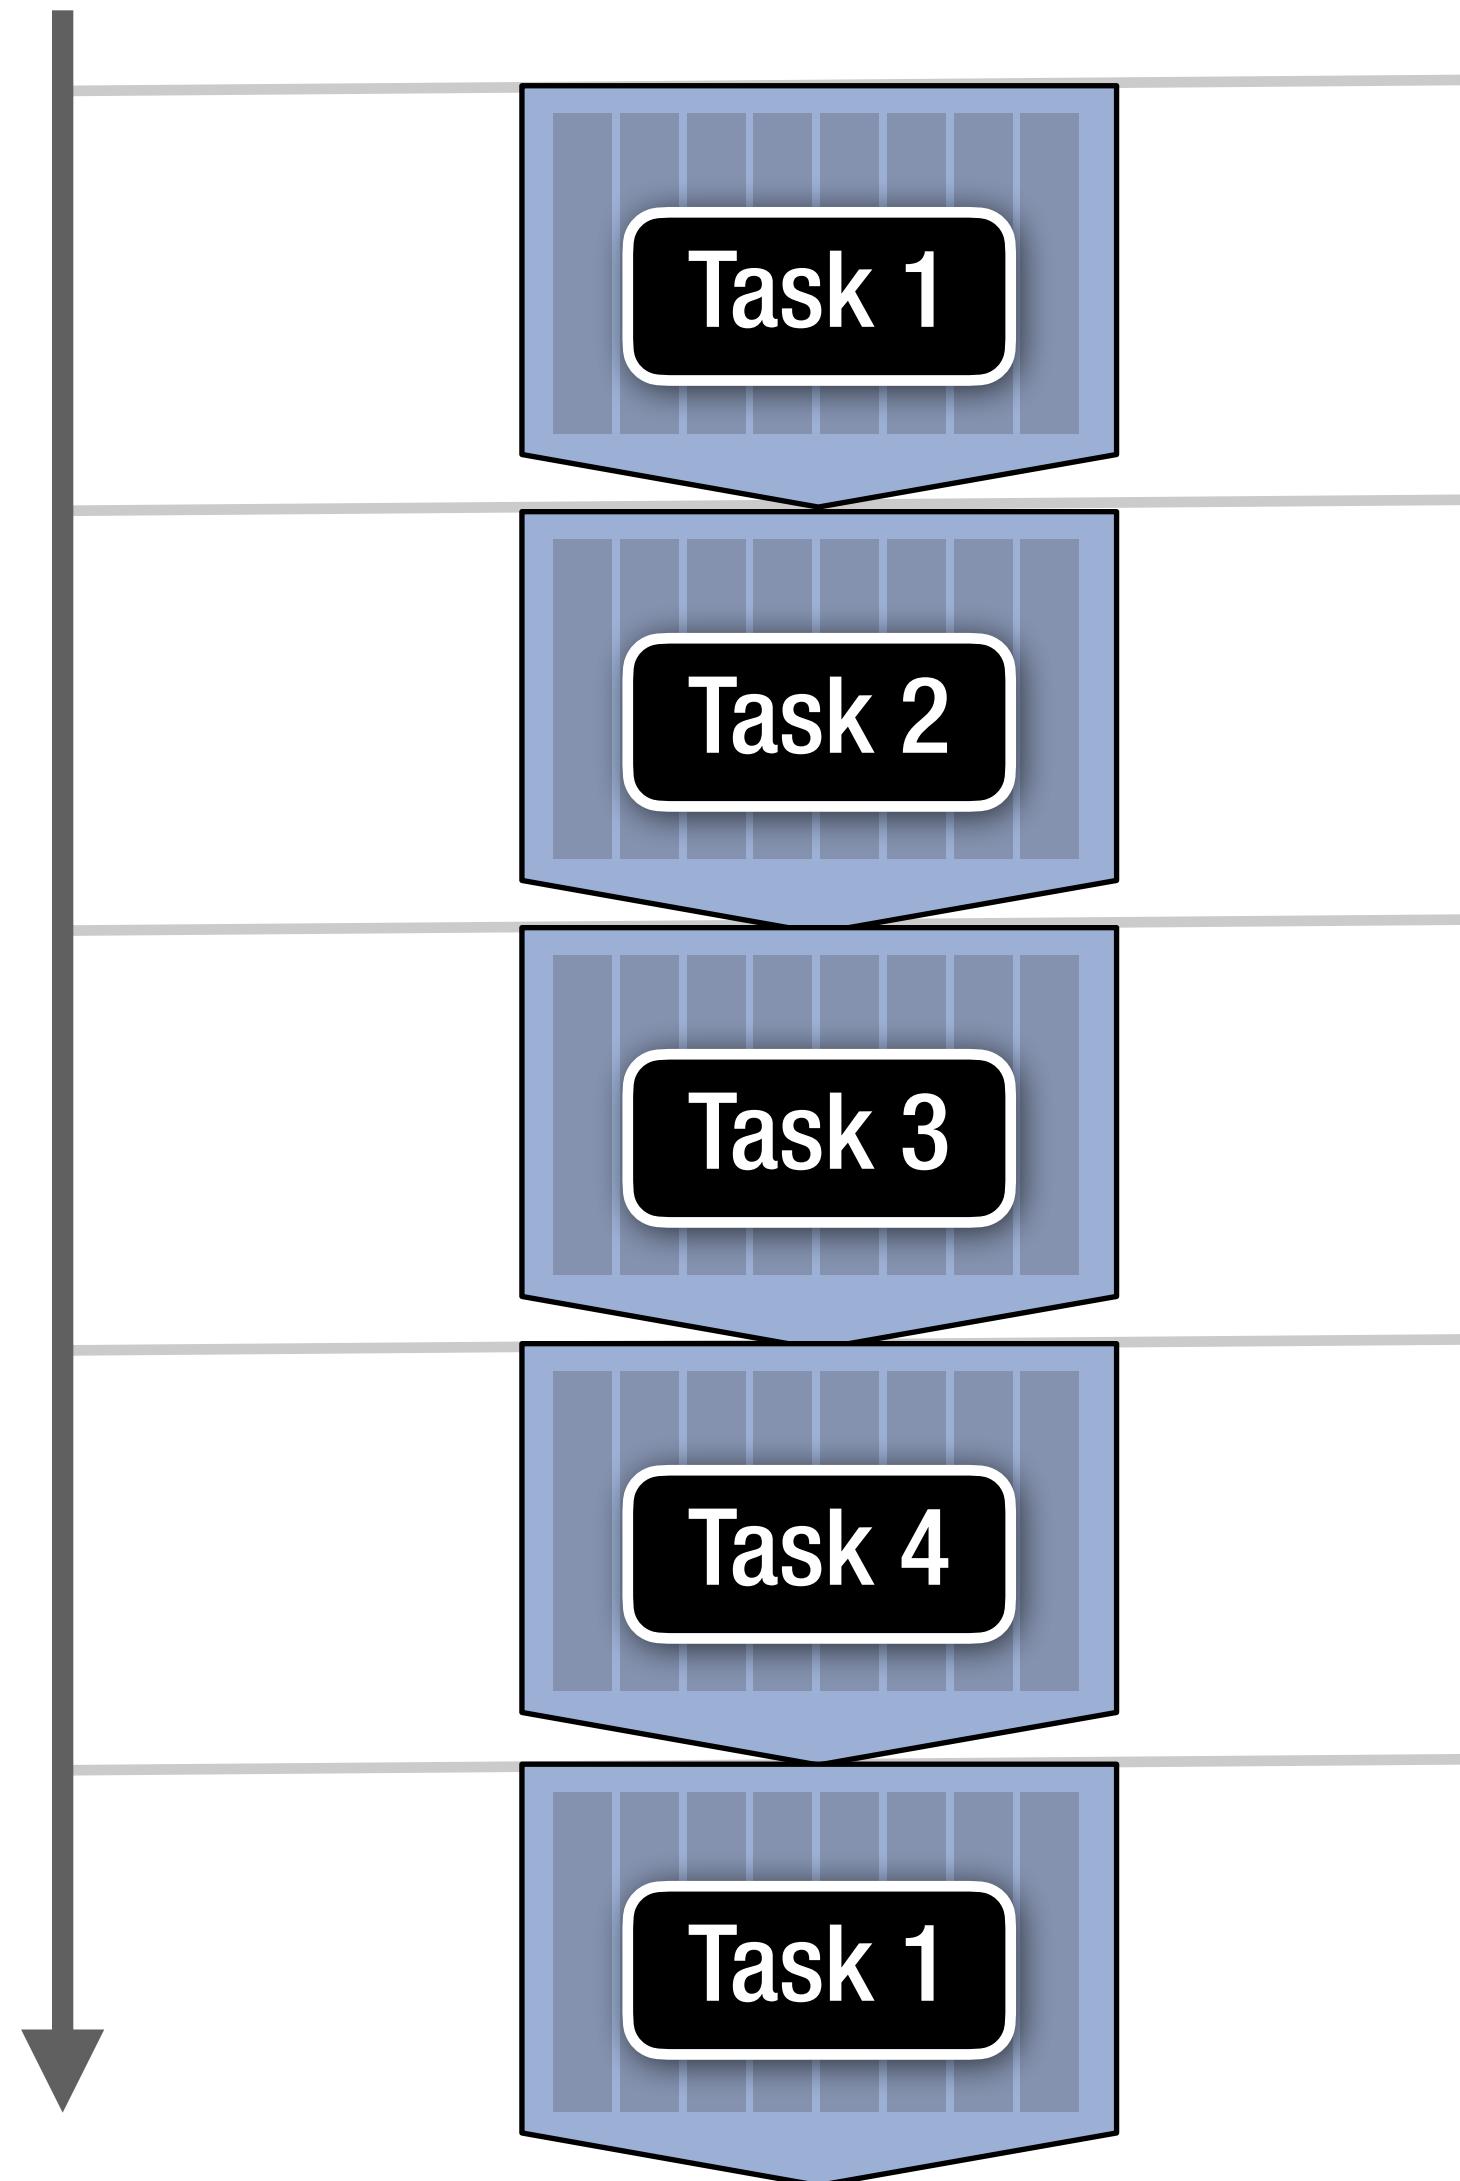
Opportunity:
exploit parallelism
to hide latency

Idea 3: Interleave parallel tasks to hide latency



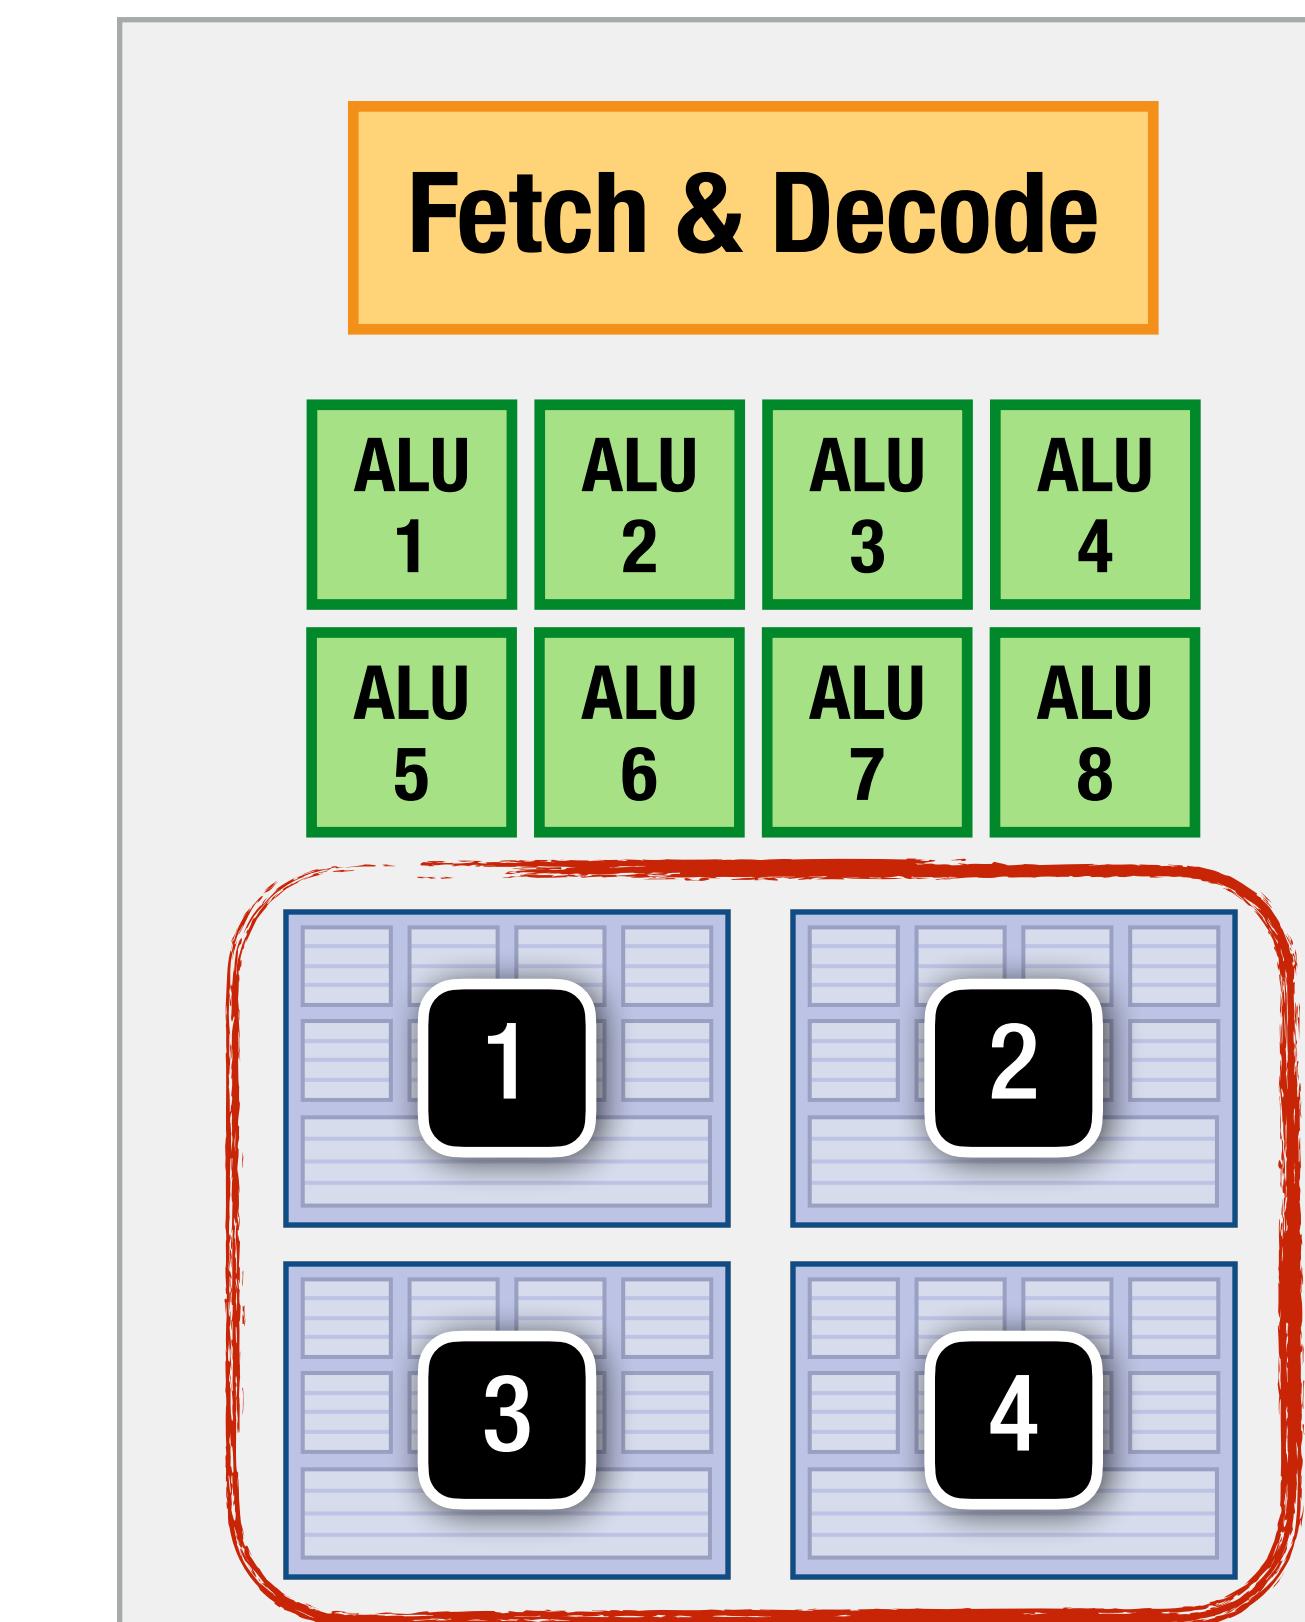
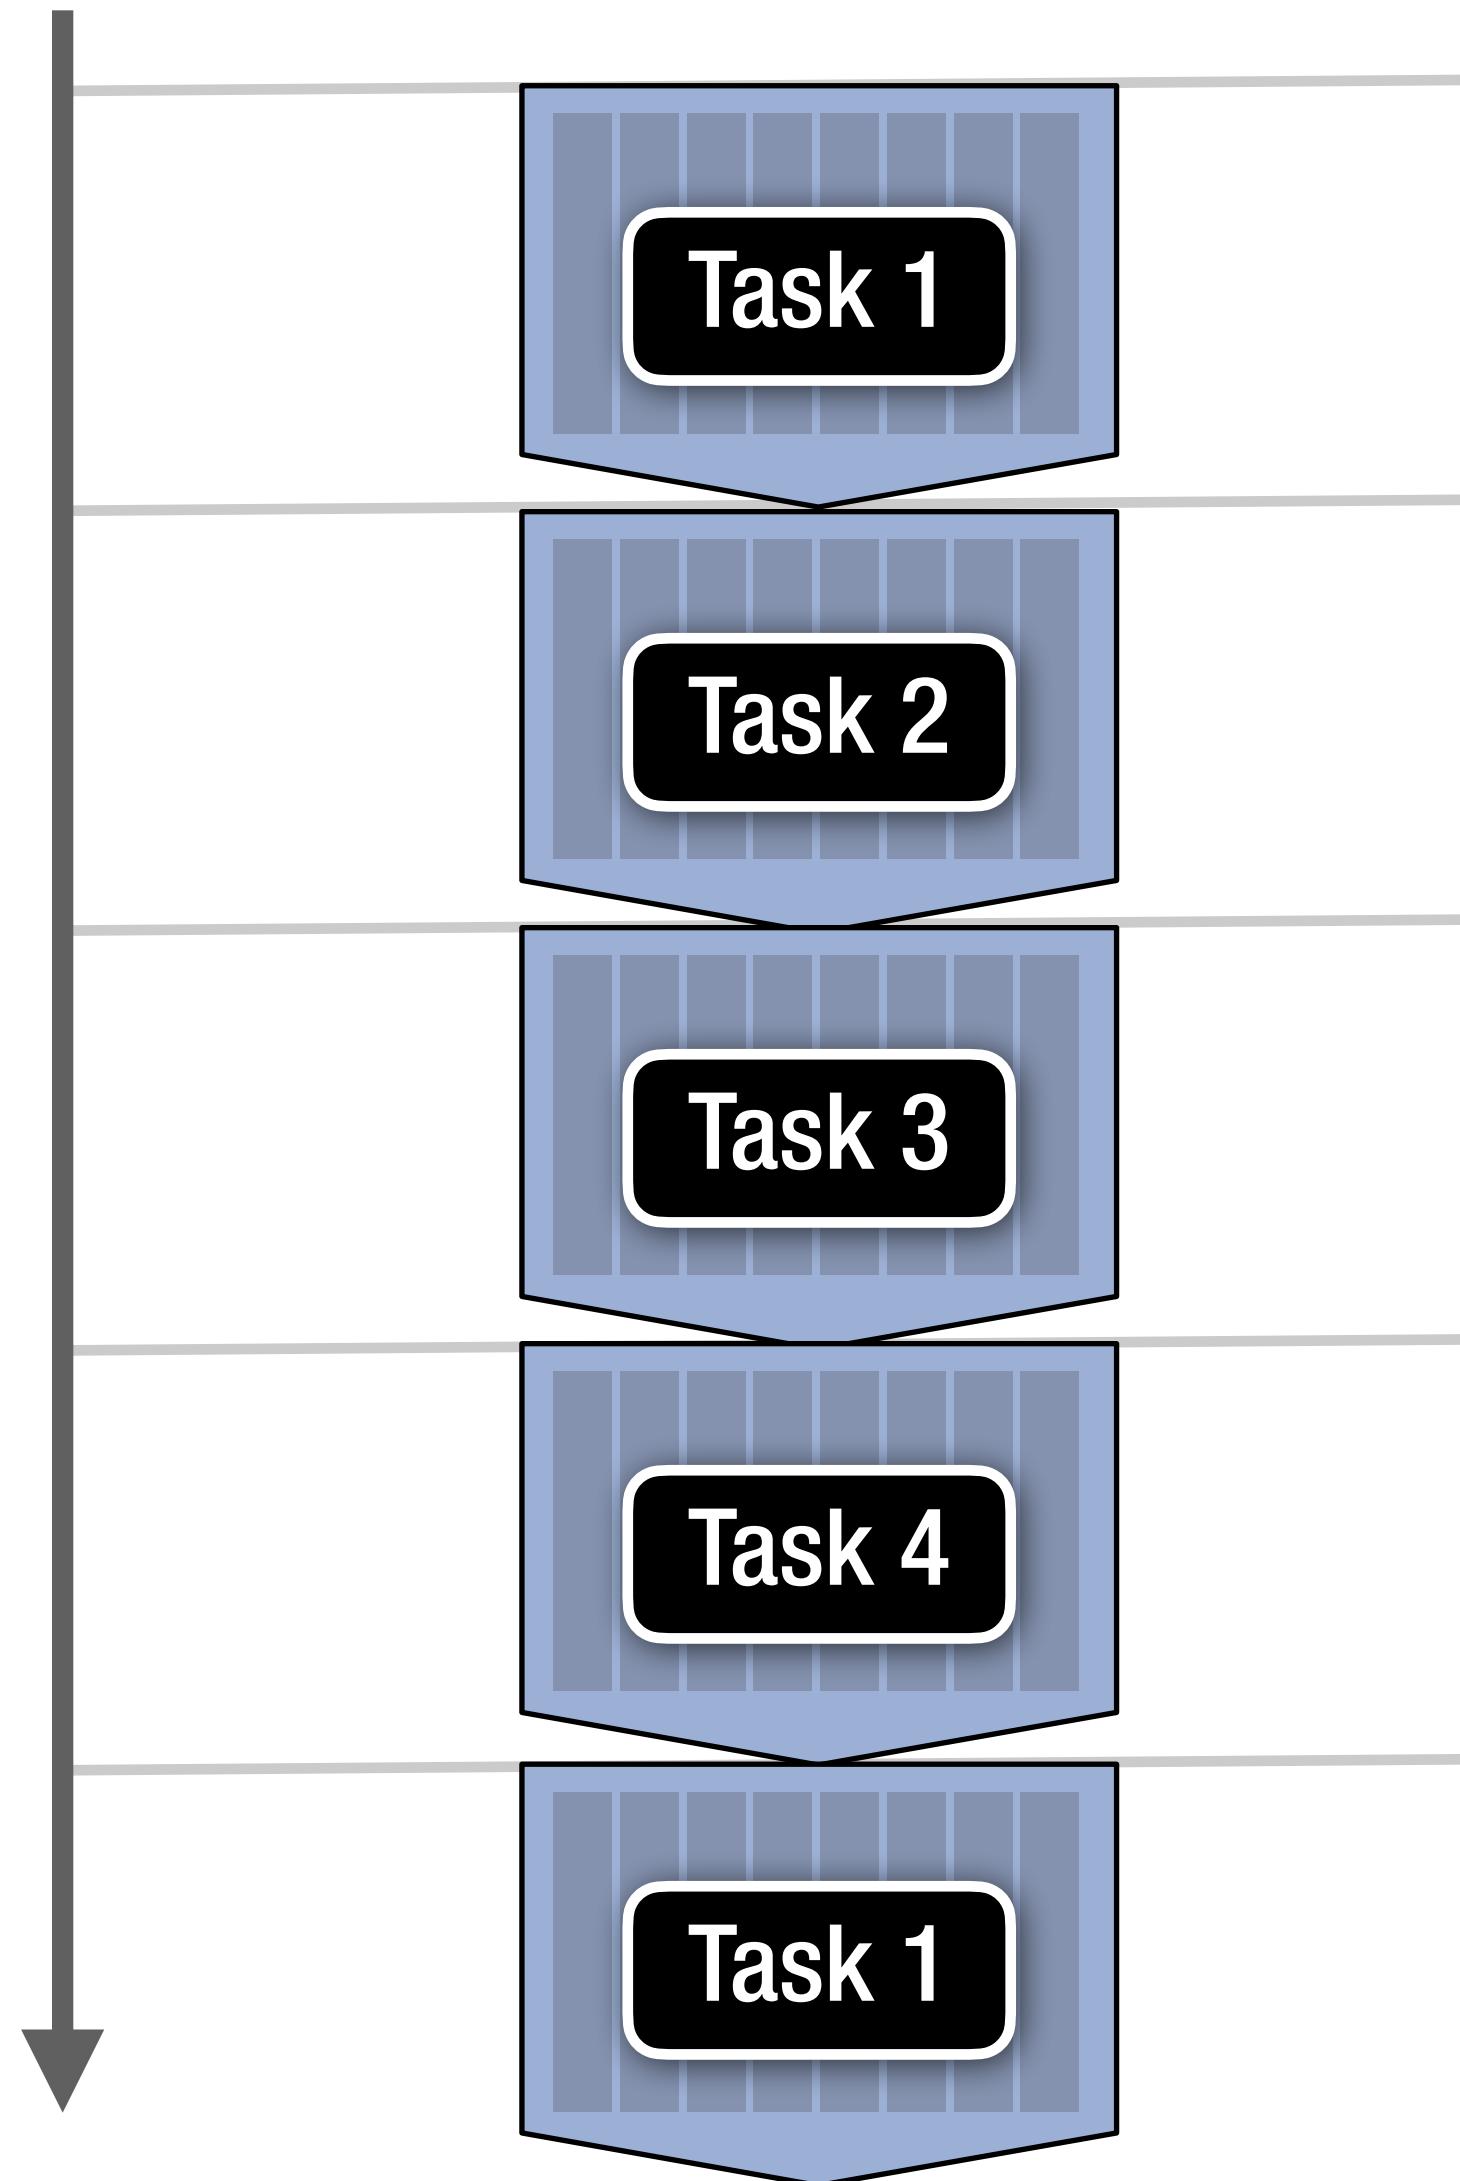
Idea 3: Interleave parallel tasks to hide latency

Time (clocks)



Interleaving requires more state storage

Time (clocks)



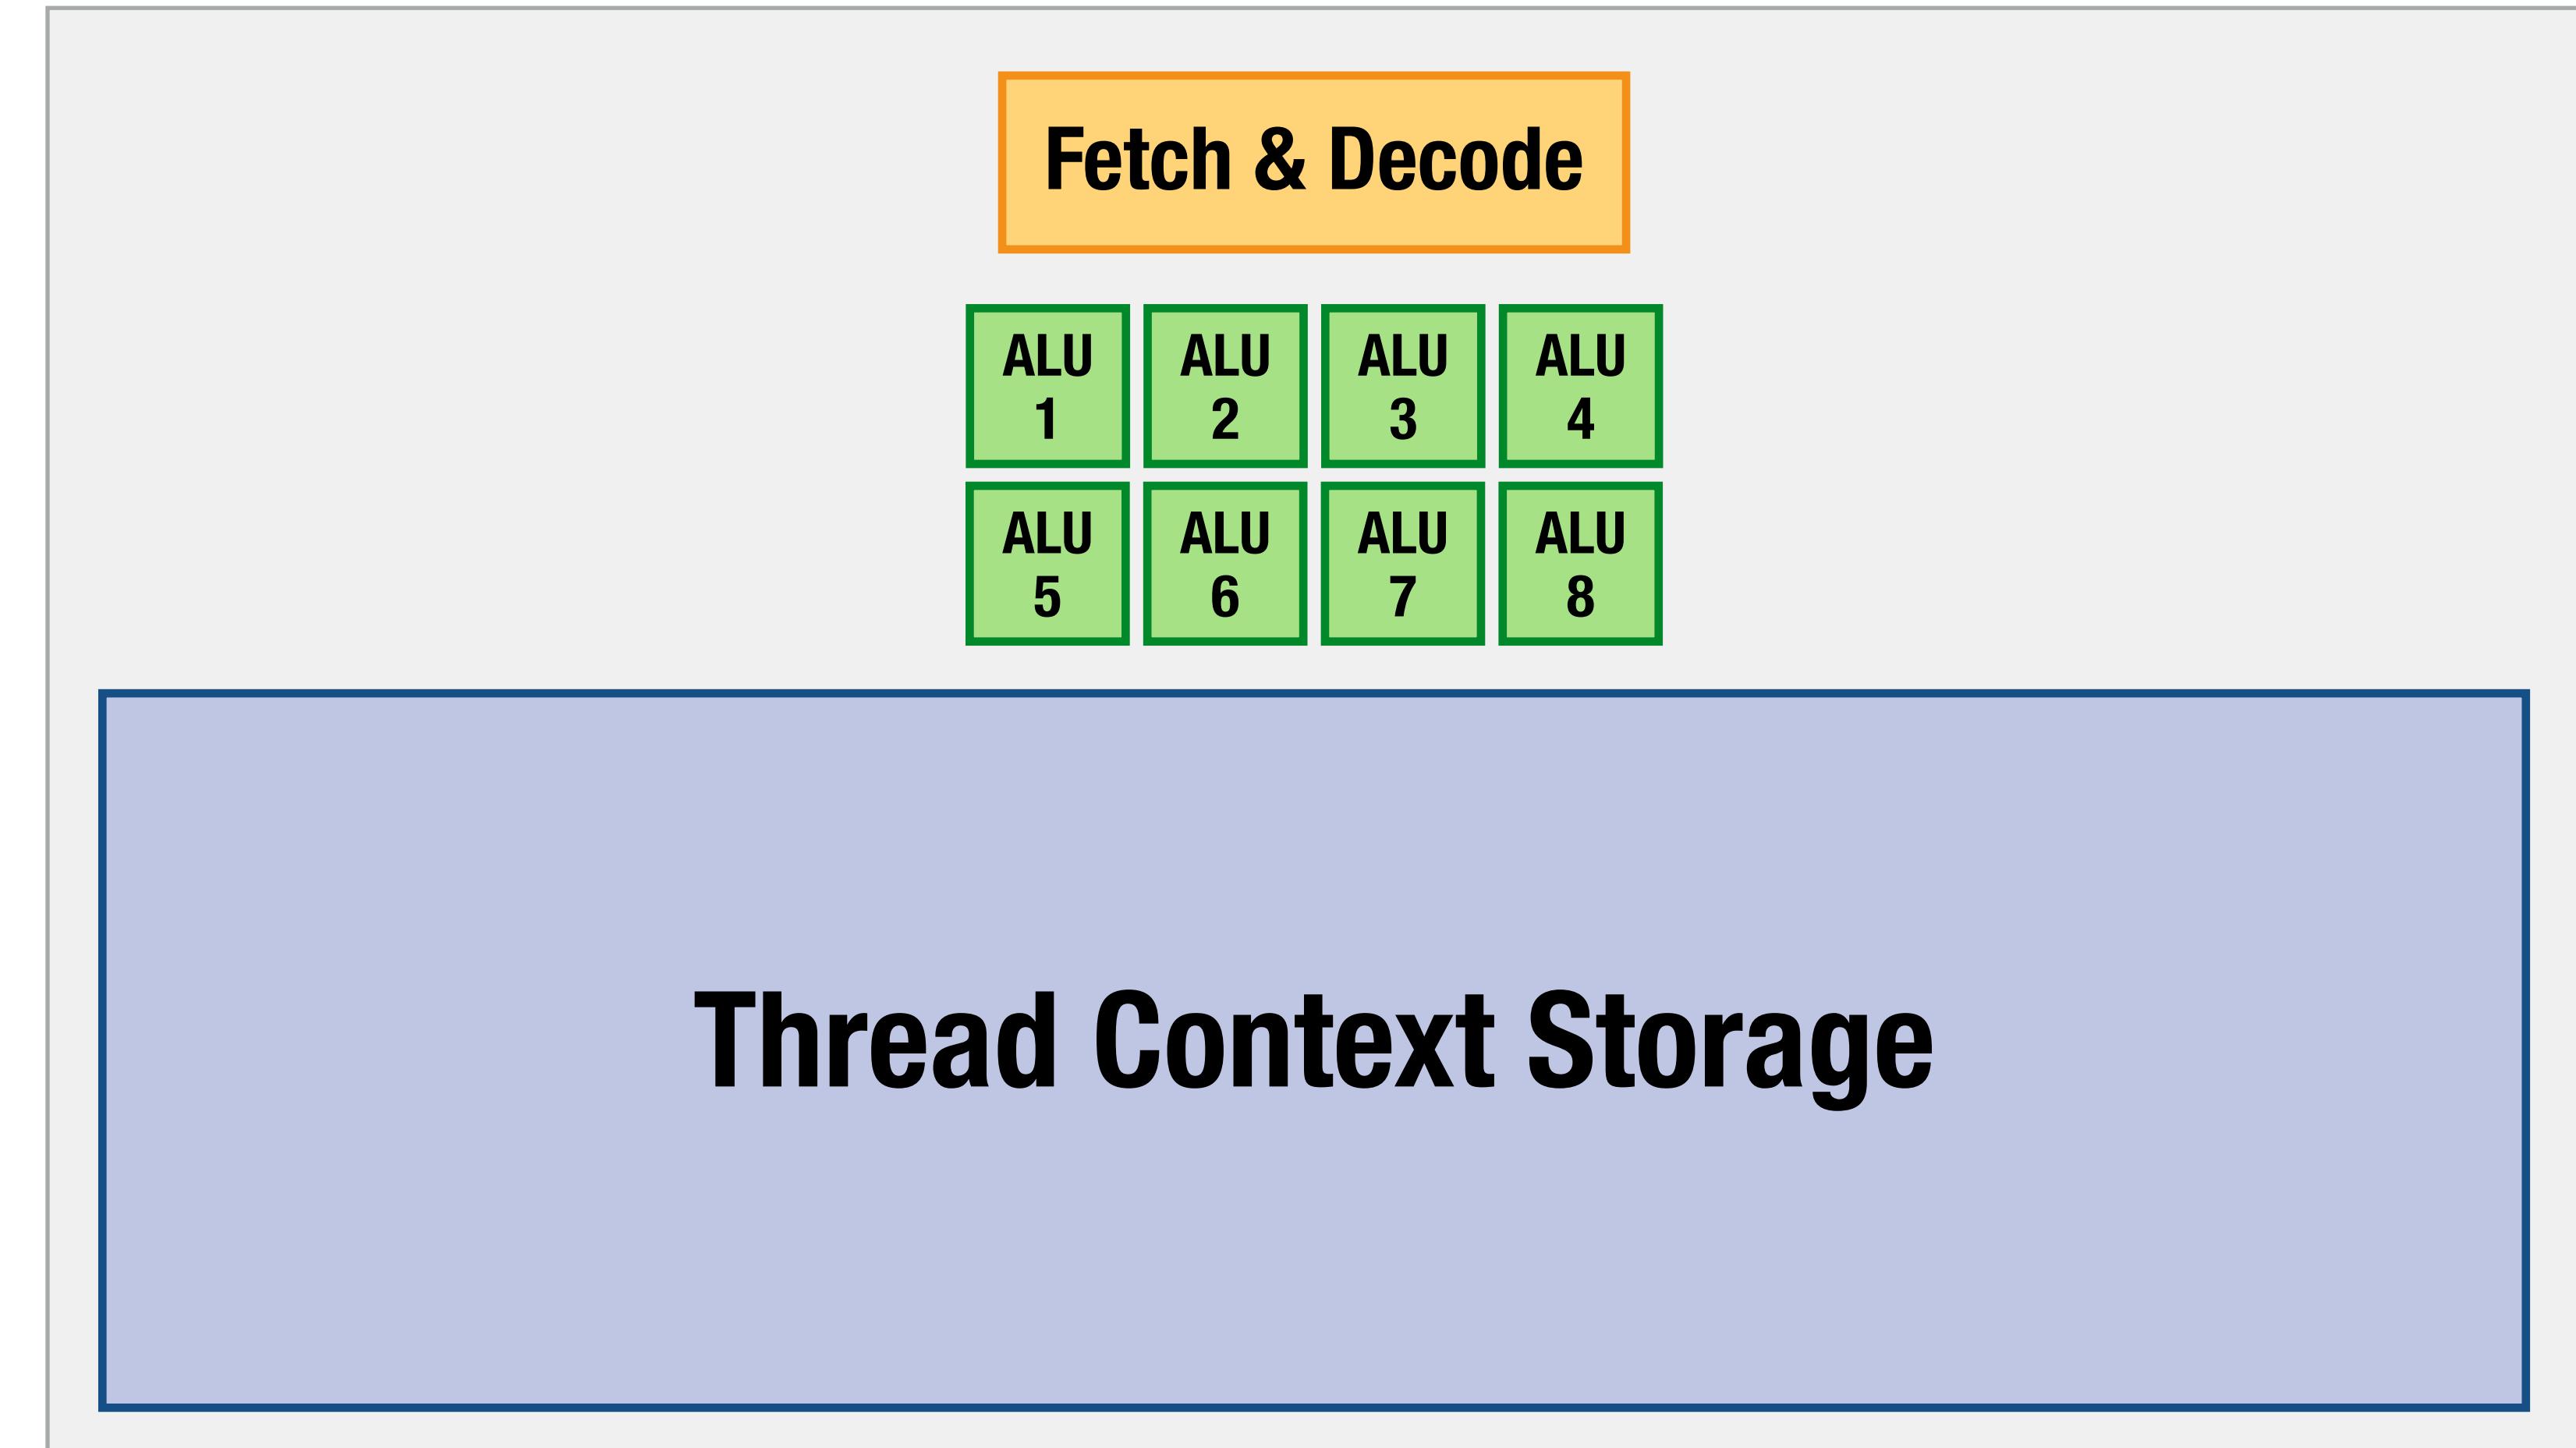
Multithreading
requires **extra**
state storage
for idle threads

Interleaving requires more state storage

DRAM latency:
100s of cycles

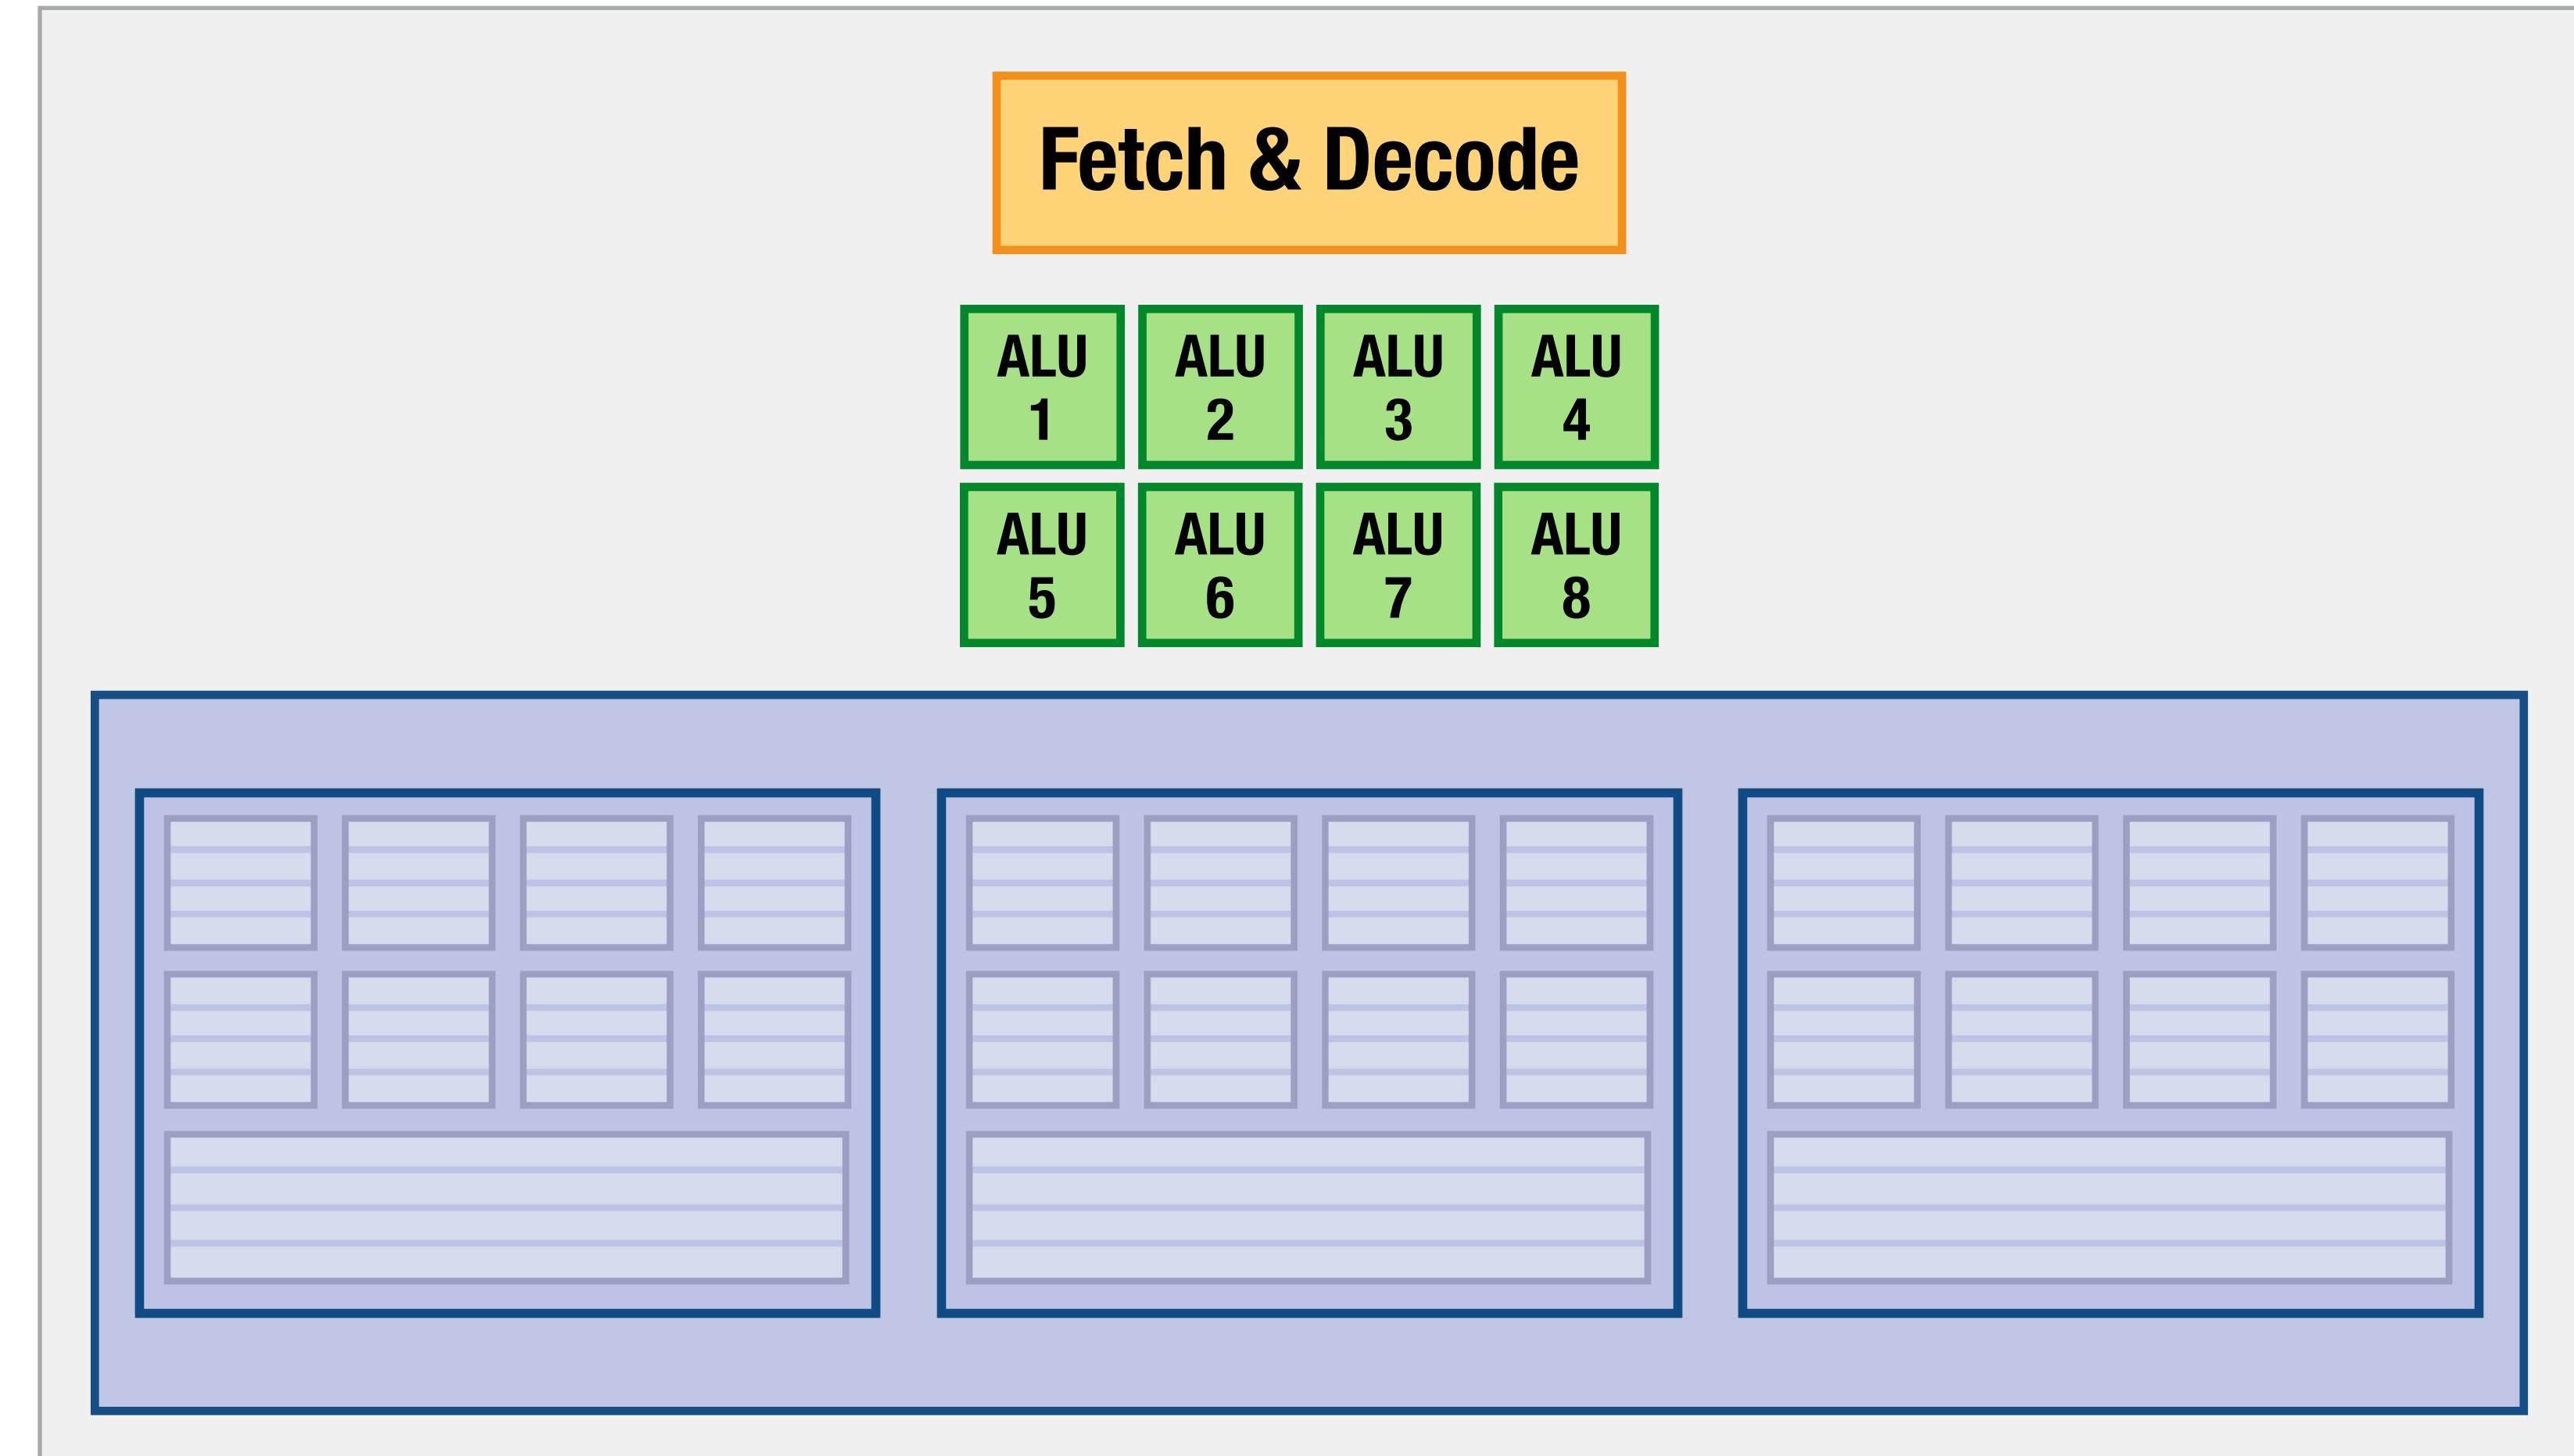
AMD/Intel Hyperthreading:
2 threads

NVIDIA H100:
16k 32-bit words
per warp scheduler
↳ **32MB** (core)
per GPU



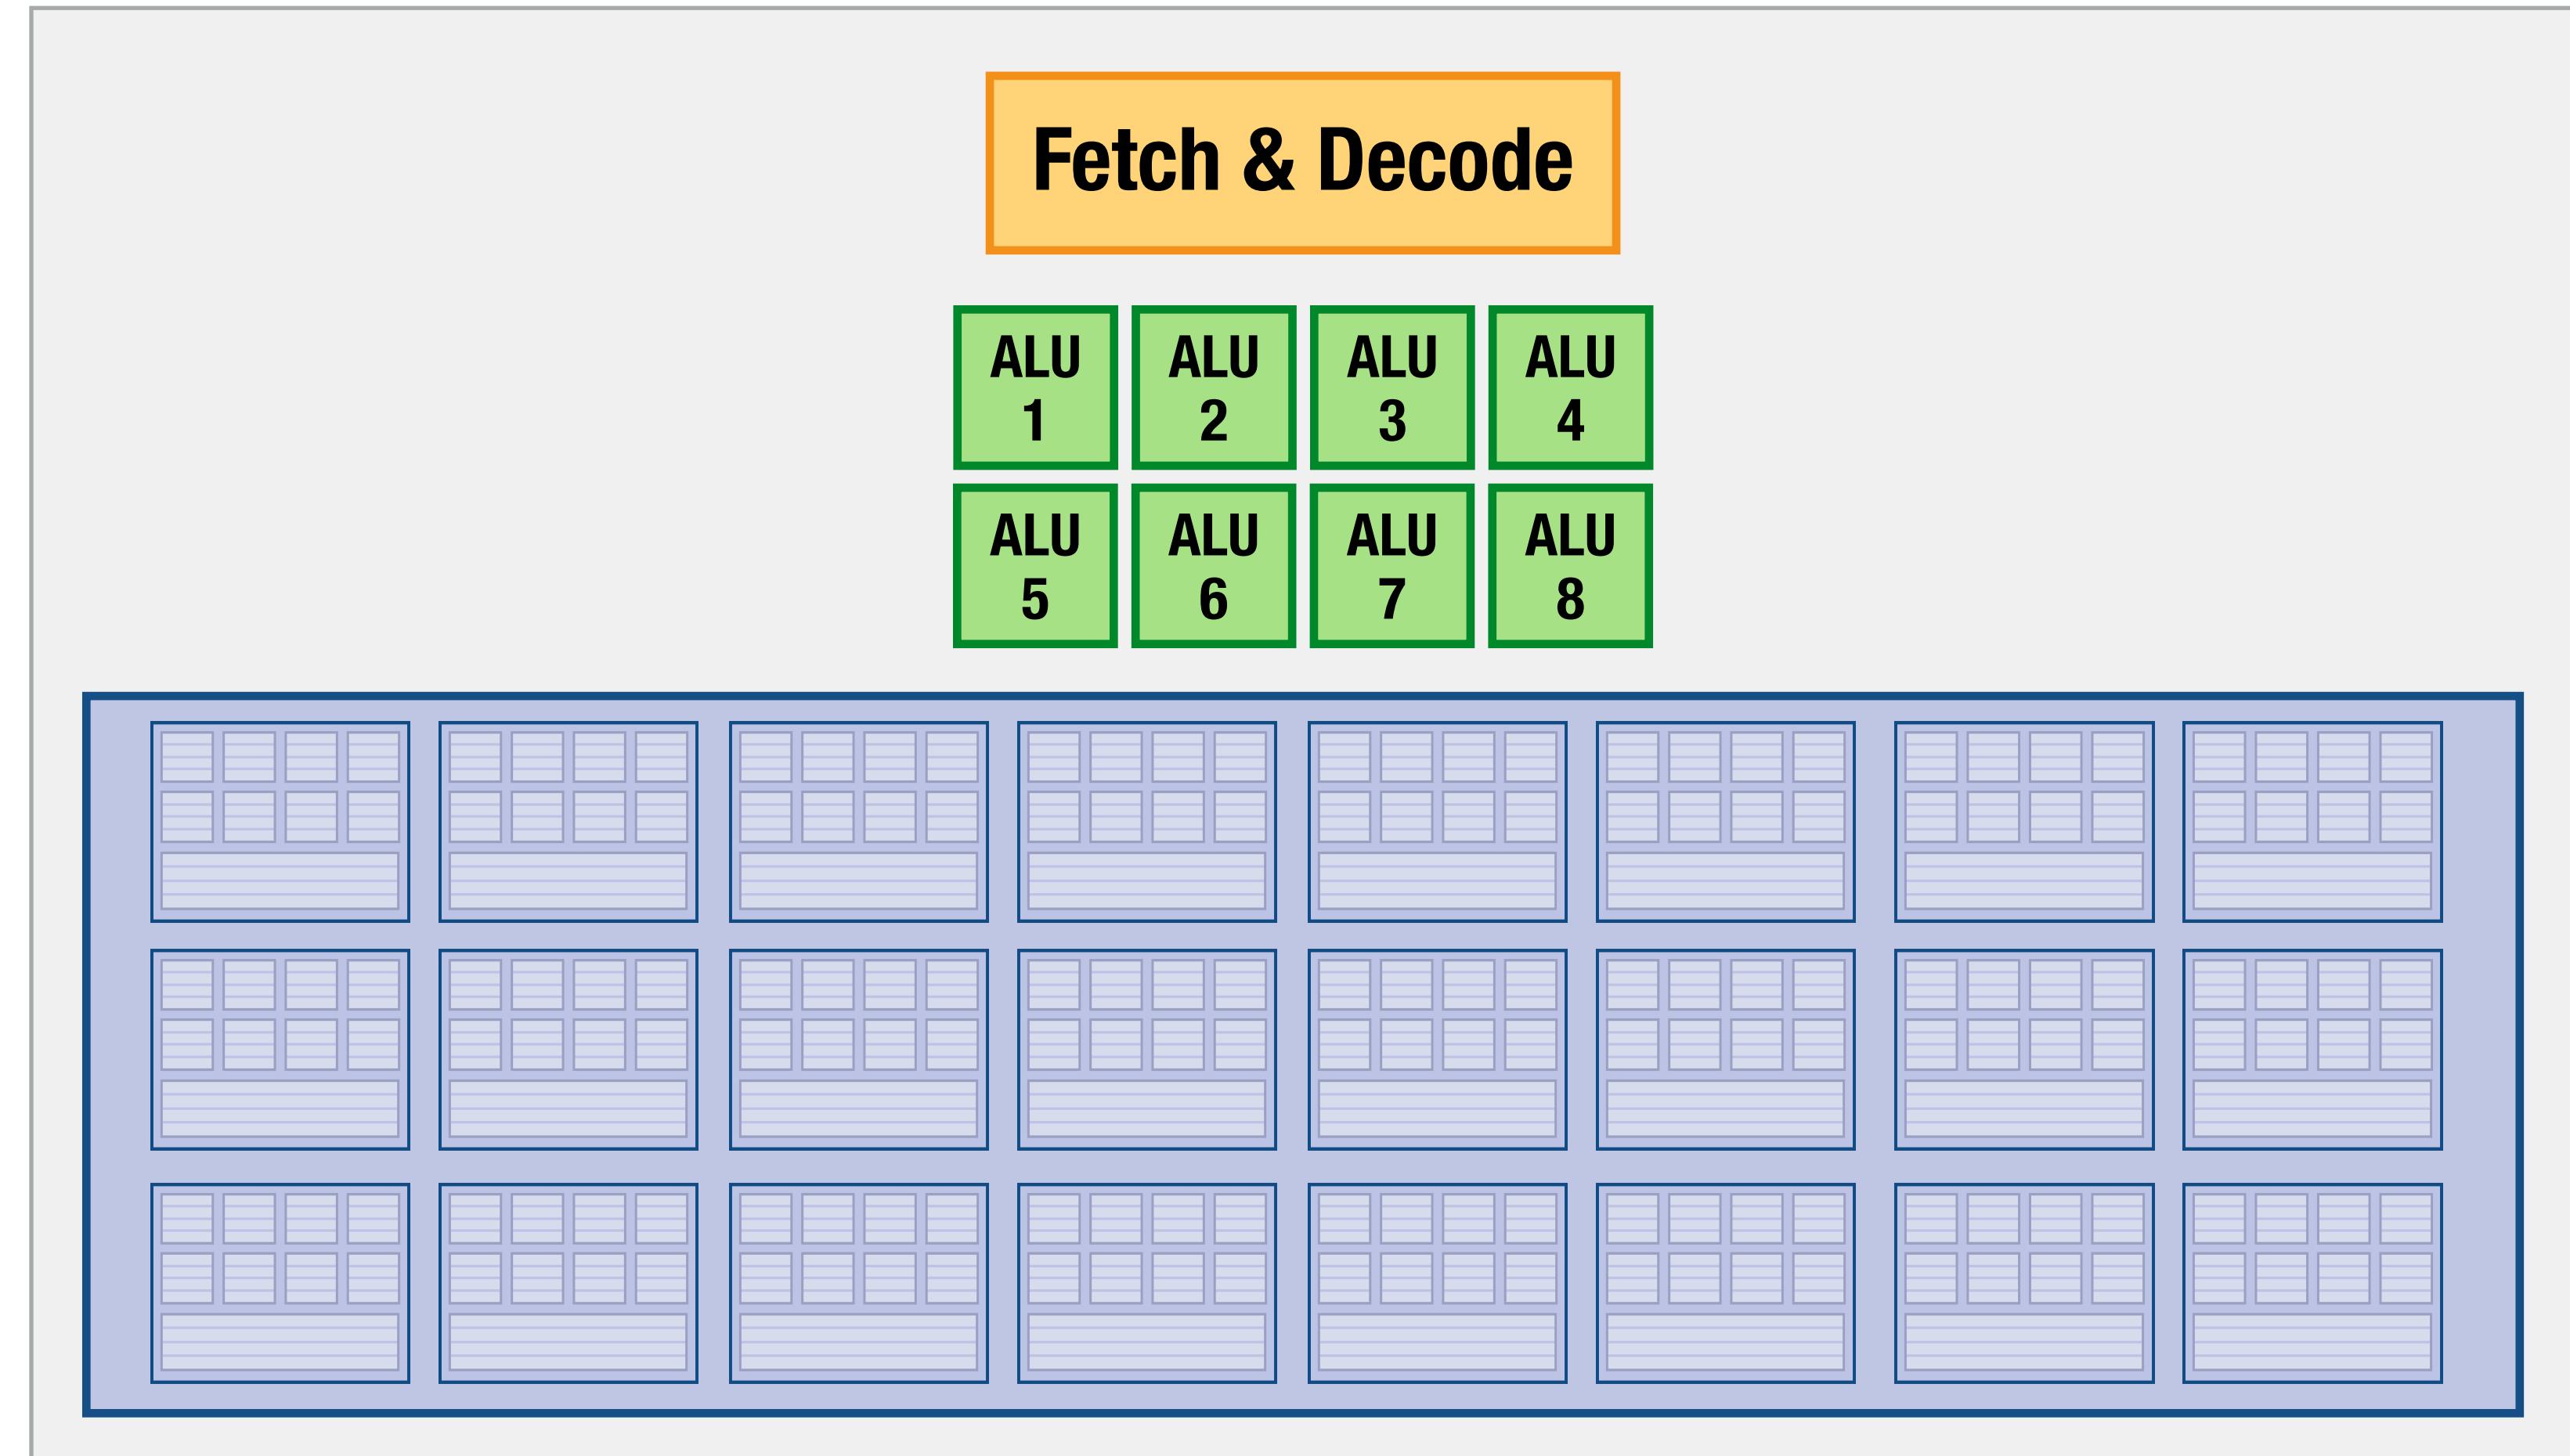
Tradeoff: per-thread state vs. latency hiding

Few, large contexts:
limited latency hiding



Tradeoff: per-thread state vs. latency hiding

Many, small contexts:
maximal latency hiding



A throughput-oriented processor exploits abundant parallelism for efficiency

- 1 Scale performance with **multicore**,
not instruction-level parallelism
- 2 Amortize **control overhead**
with **SIMD execution**
- 3 Hide **latency** with **concurrent
threads**, not speculation

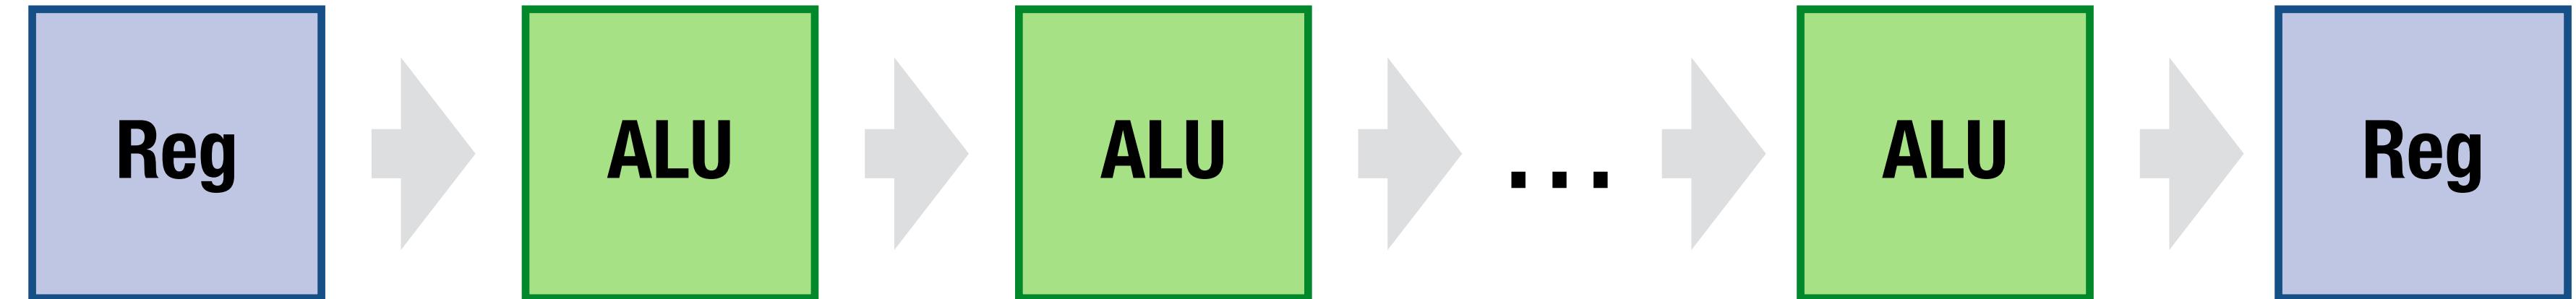
Idea 4: amortize instruction overheads with more complex instructions

Primitive op:
("RISC")



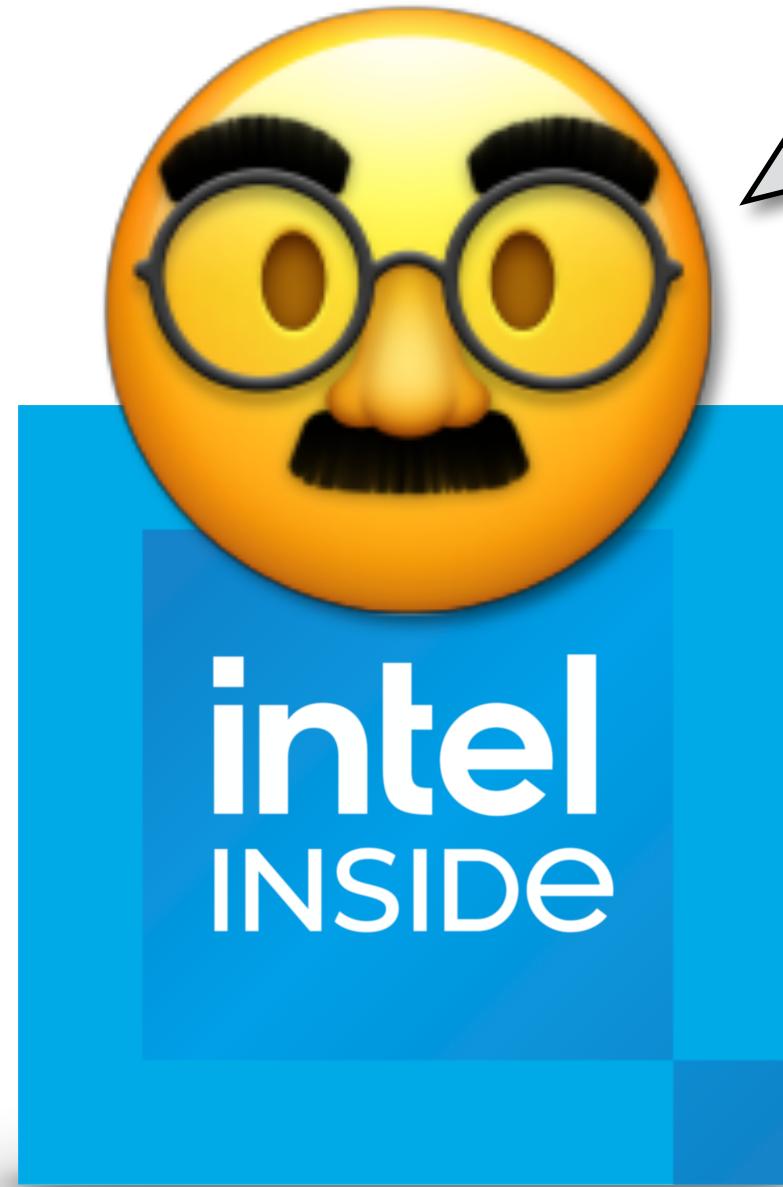
"ASIC-in-an-Instruction"

Complex op:
("CISC++")



e.g., AES, video encode/decode,
DSP, texture filtering, ...

and especially
matrix multiply!



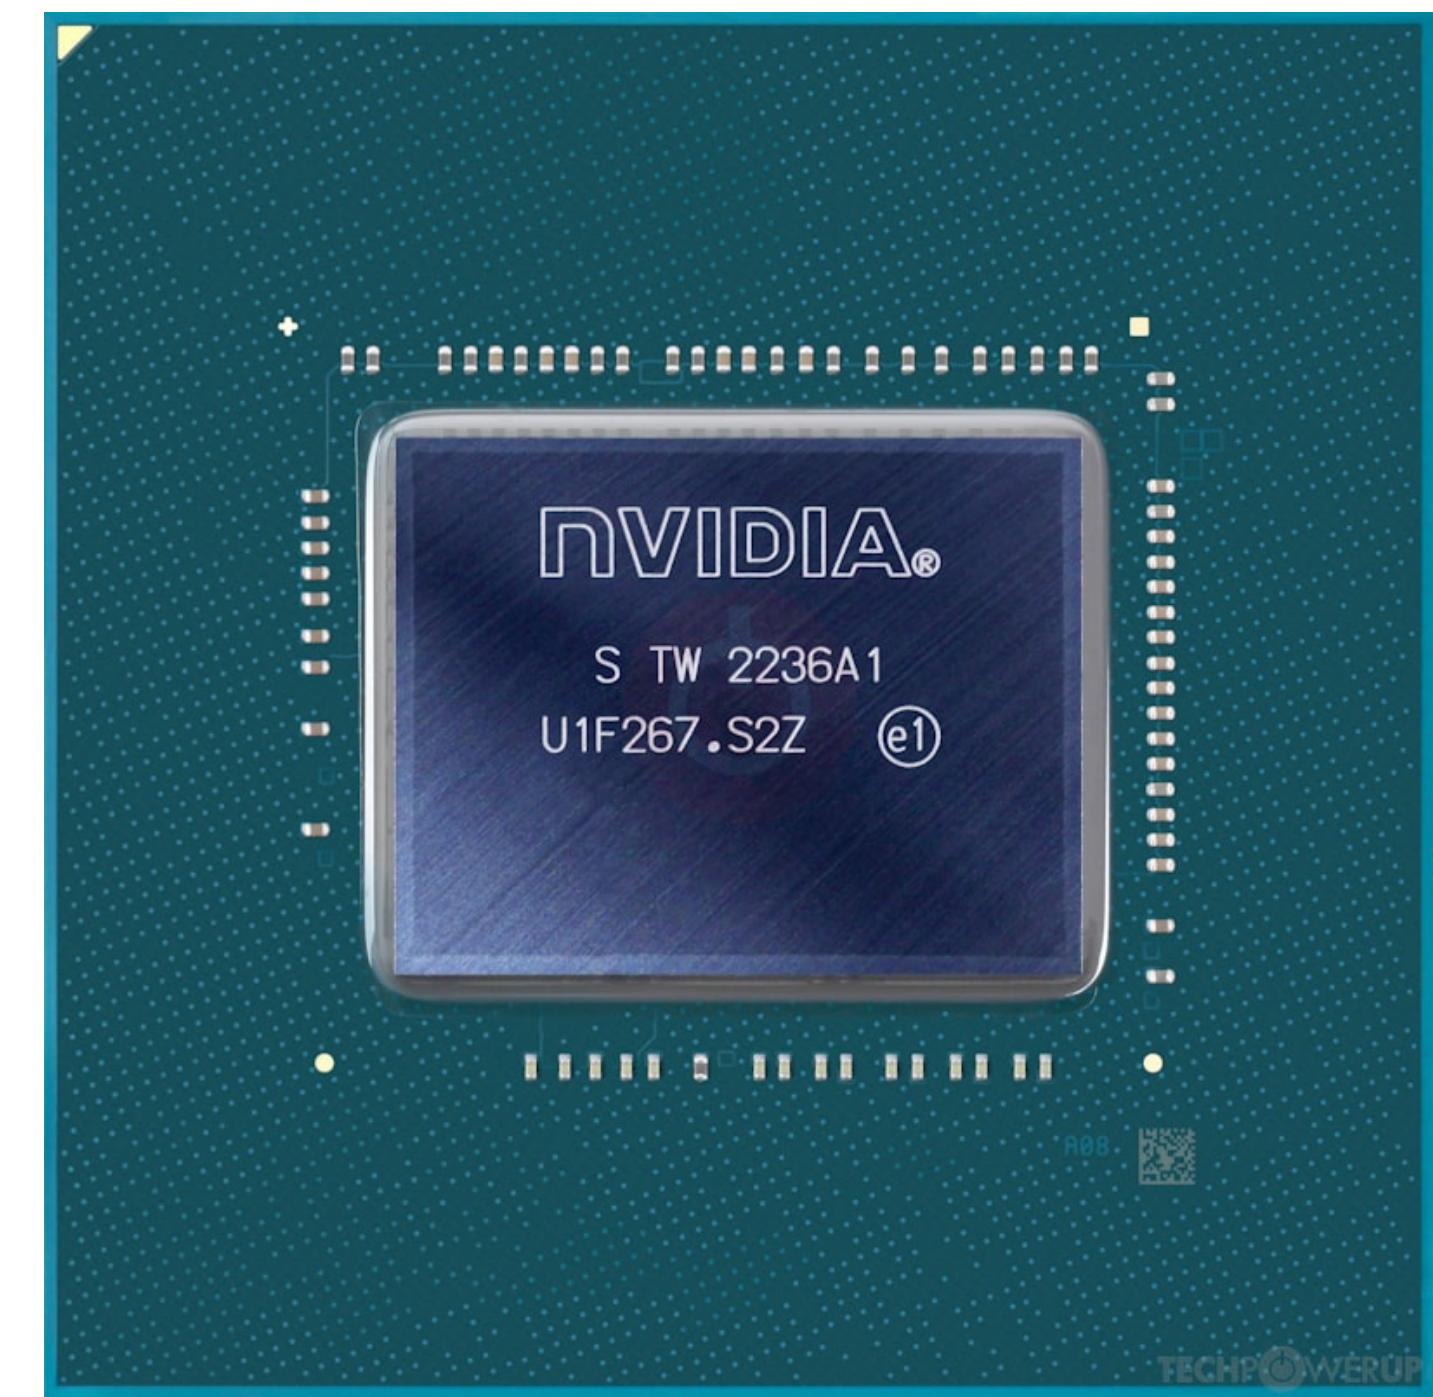
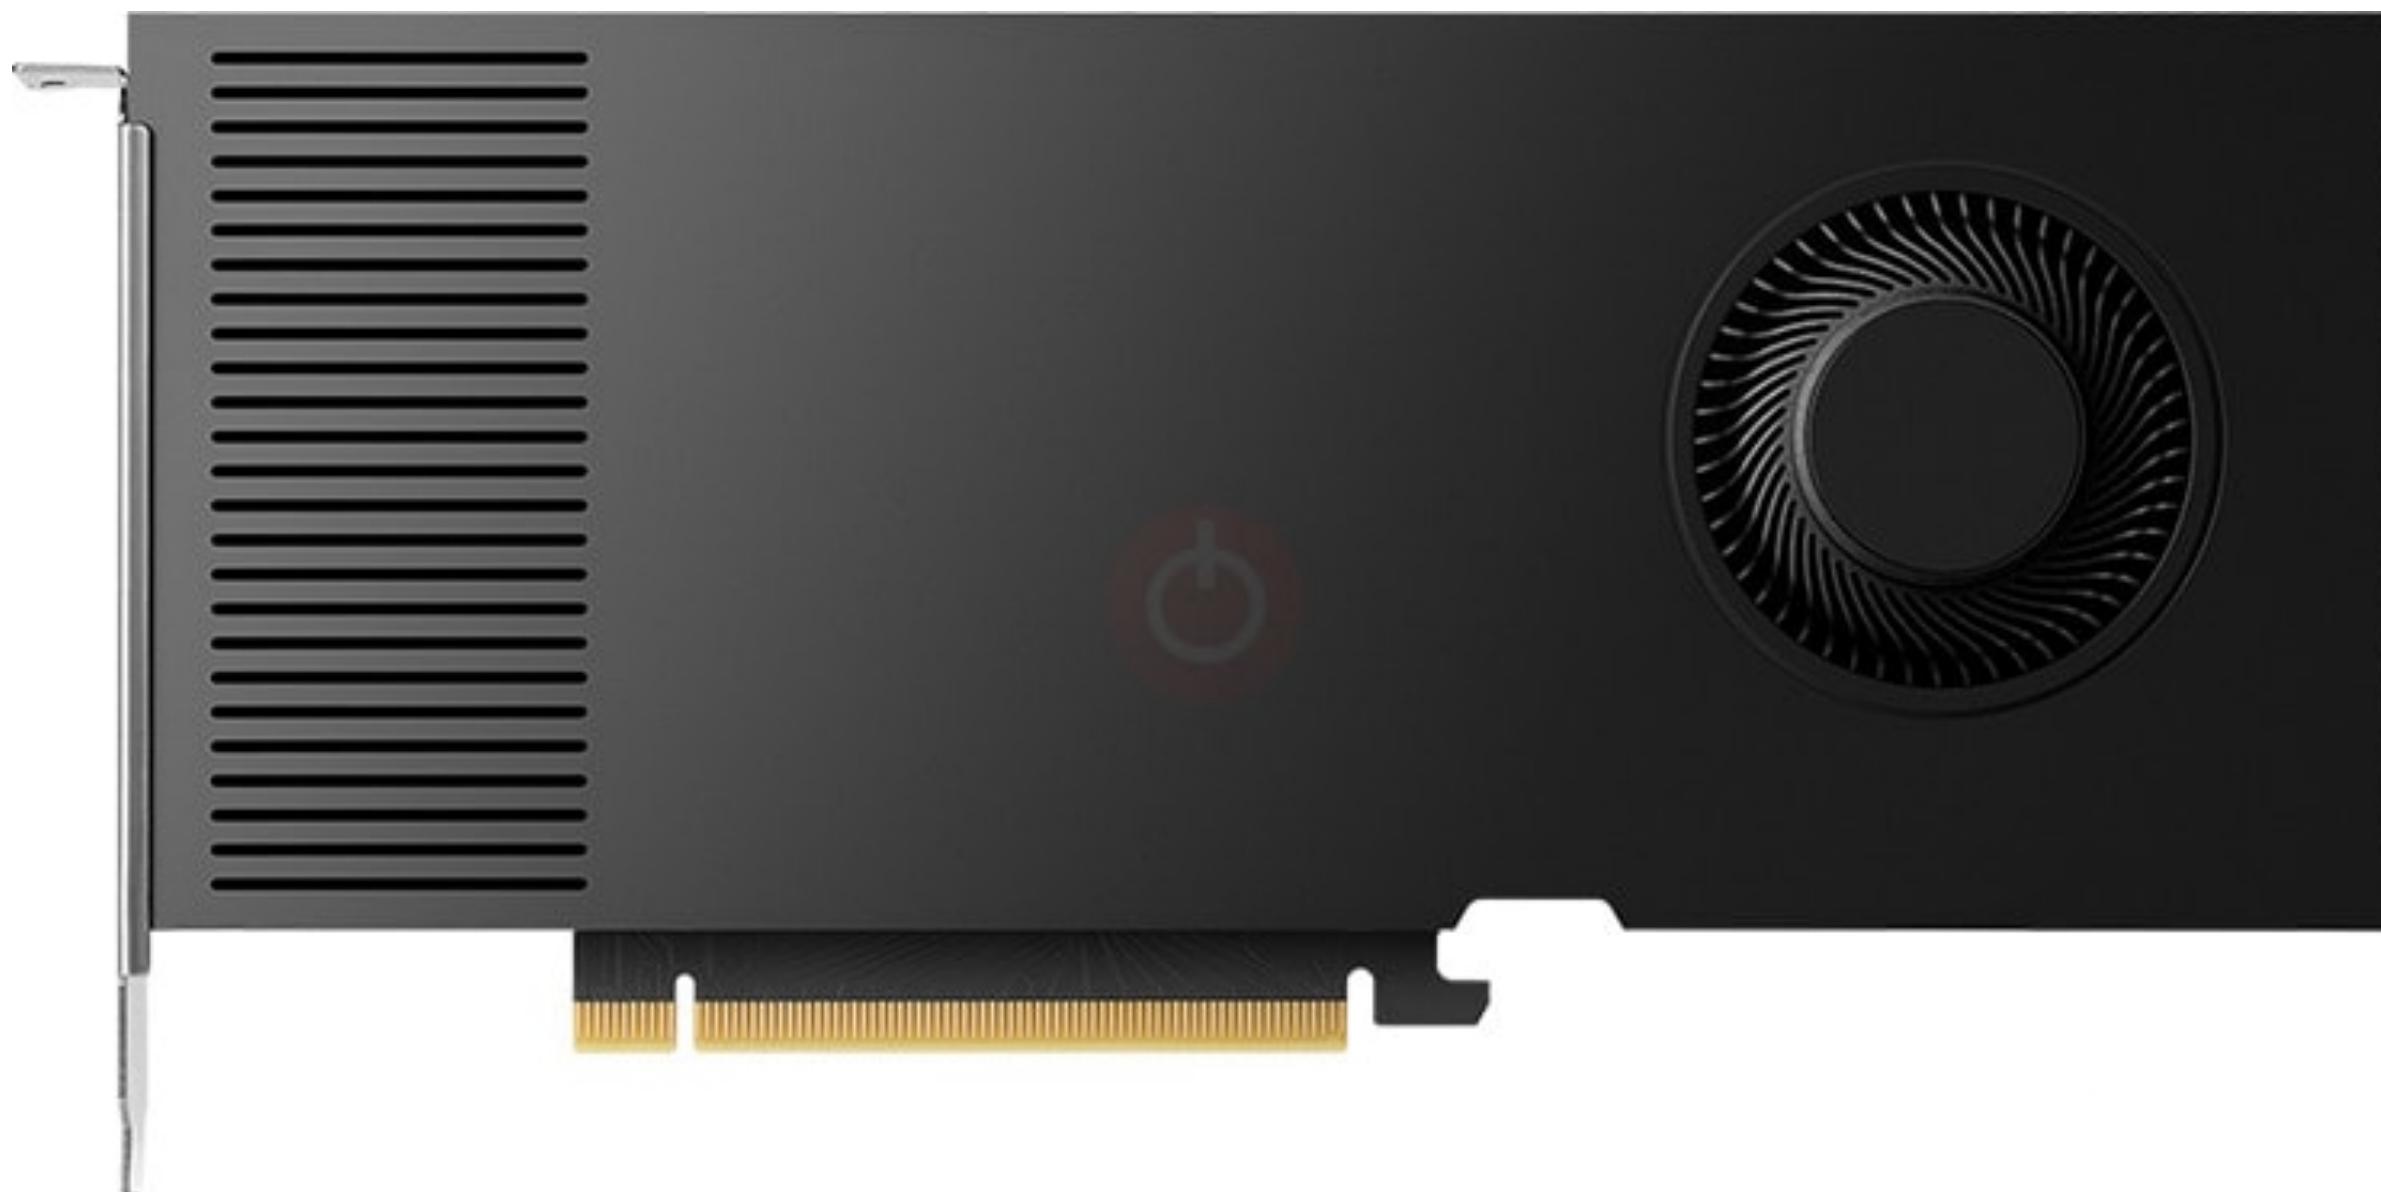
Why yes,
I am still
basically a
PDP-11

You will **rewrite**
all your code in
CUDA . . .
and pray I don't
alter the deal
any further!



**How do these ideas appear
in real hardware?**

Our GPU: NVIDIA RTX 4000 Ada (AD104 chip)



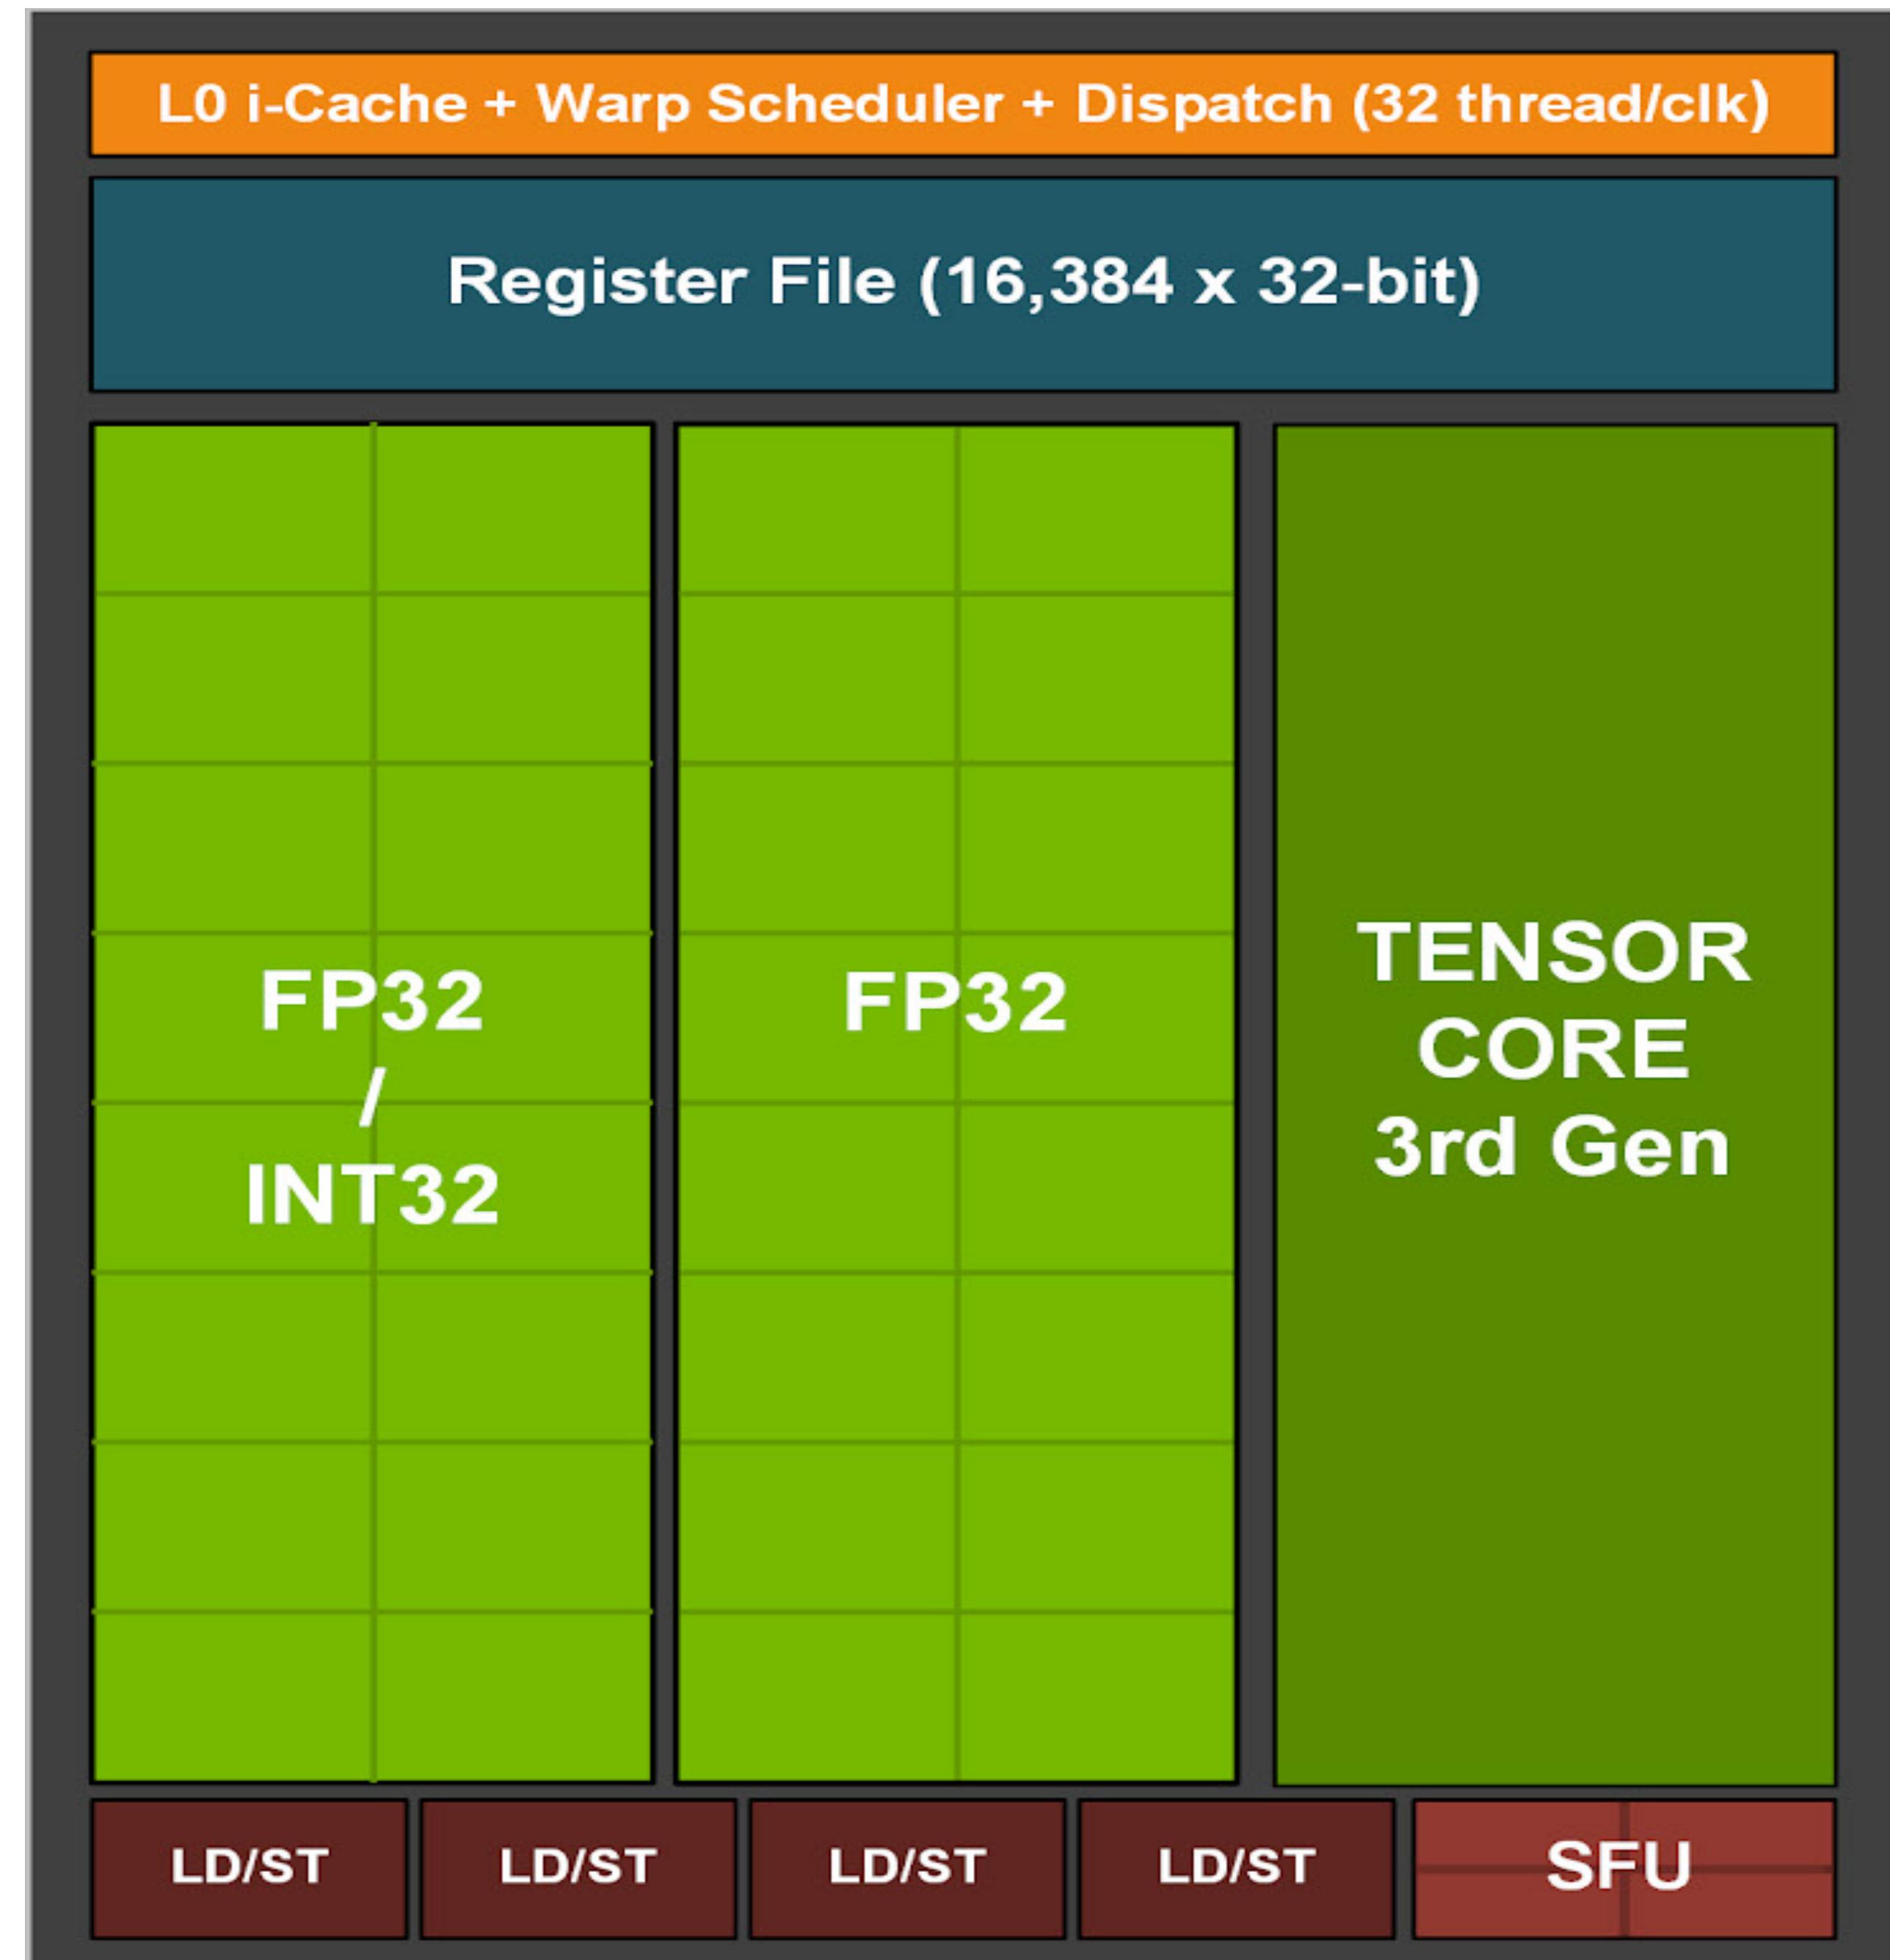
One “core” of our GPU (“warp scheduler”)

32 x 32-bit Exec. Units
(ALUs / vector lanes)

512 x 32 x 32-bit regs

1 warp instruction / clock
(32 lanes)

Up to 12 live threads
(independent warps)



One “SM” is a cluster of 4 warp schedulers



Our whole GPU

48 SMs

$x 4 = 192$ cores

$x 32 = 6144$ exec. units

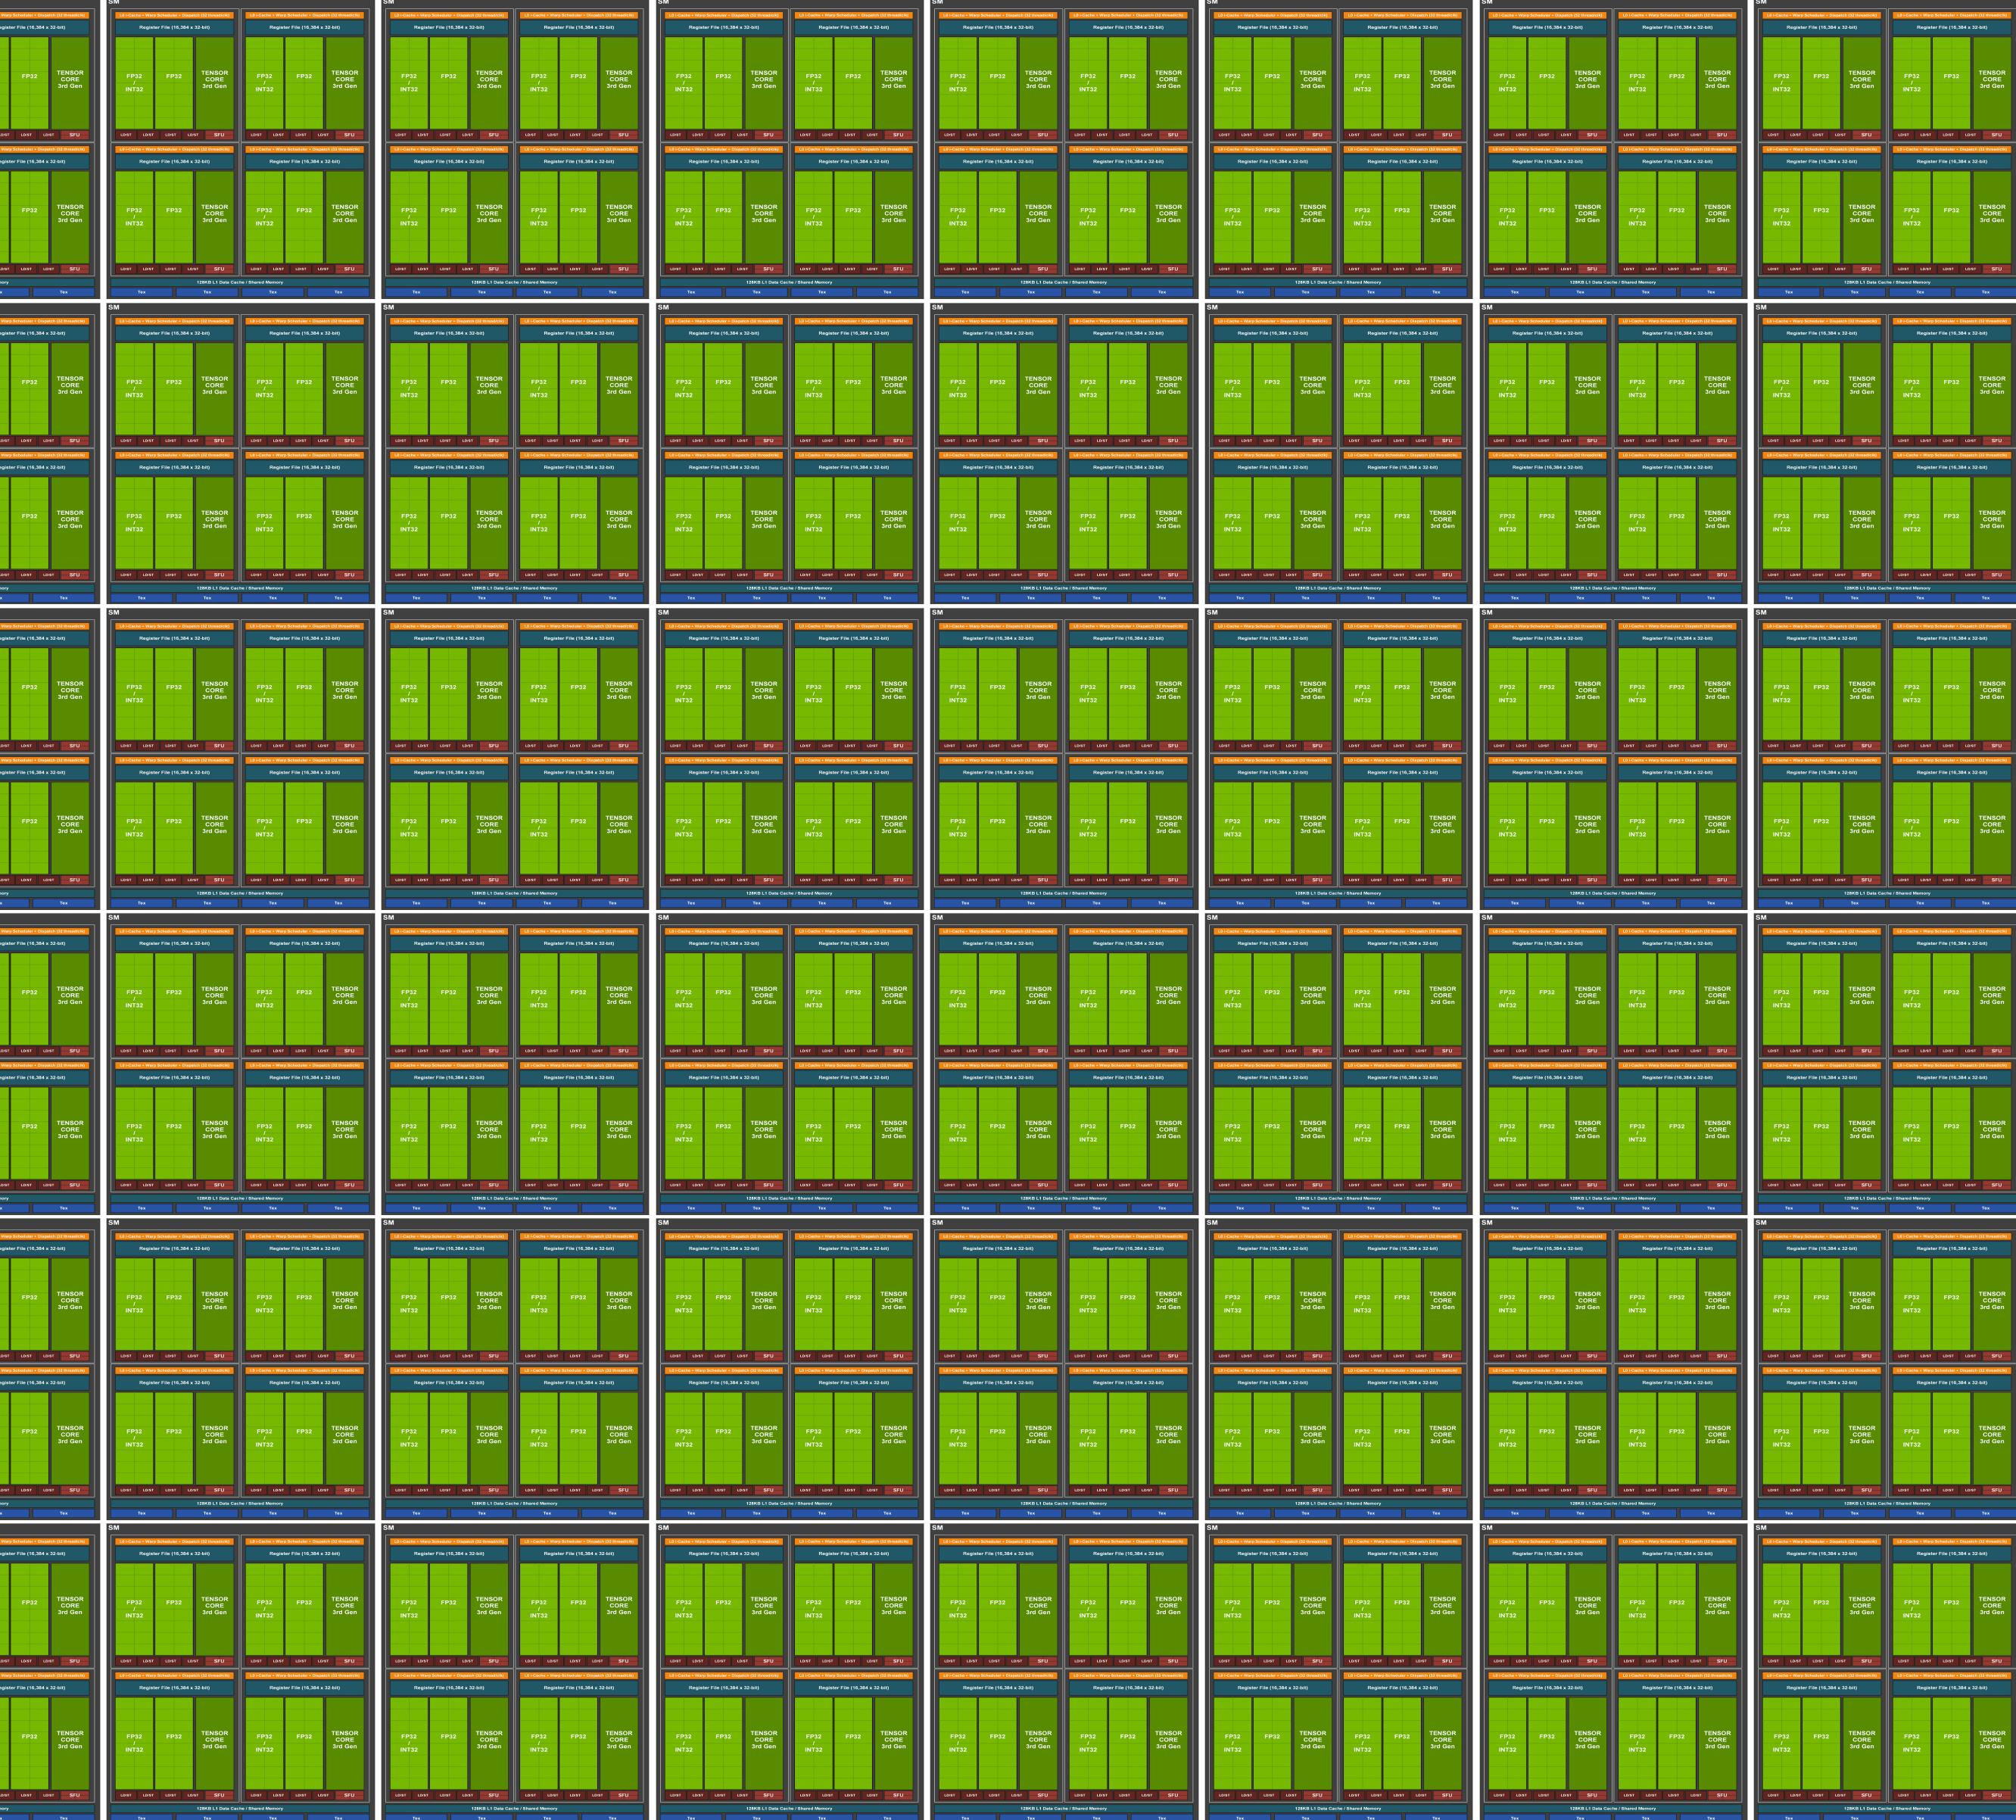
@ 2.175 GHz

= 26.7 TFLOPS

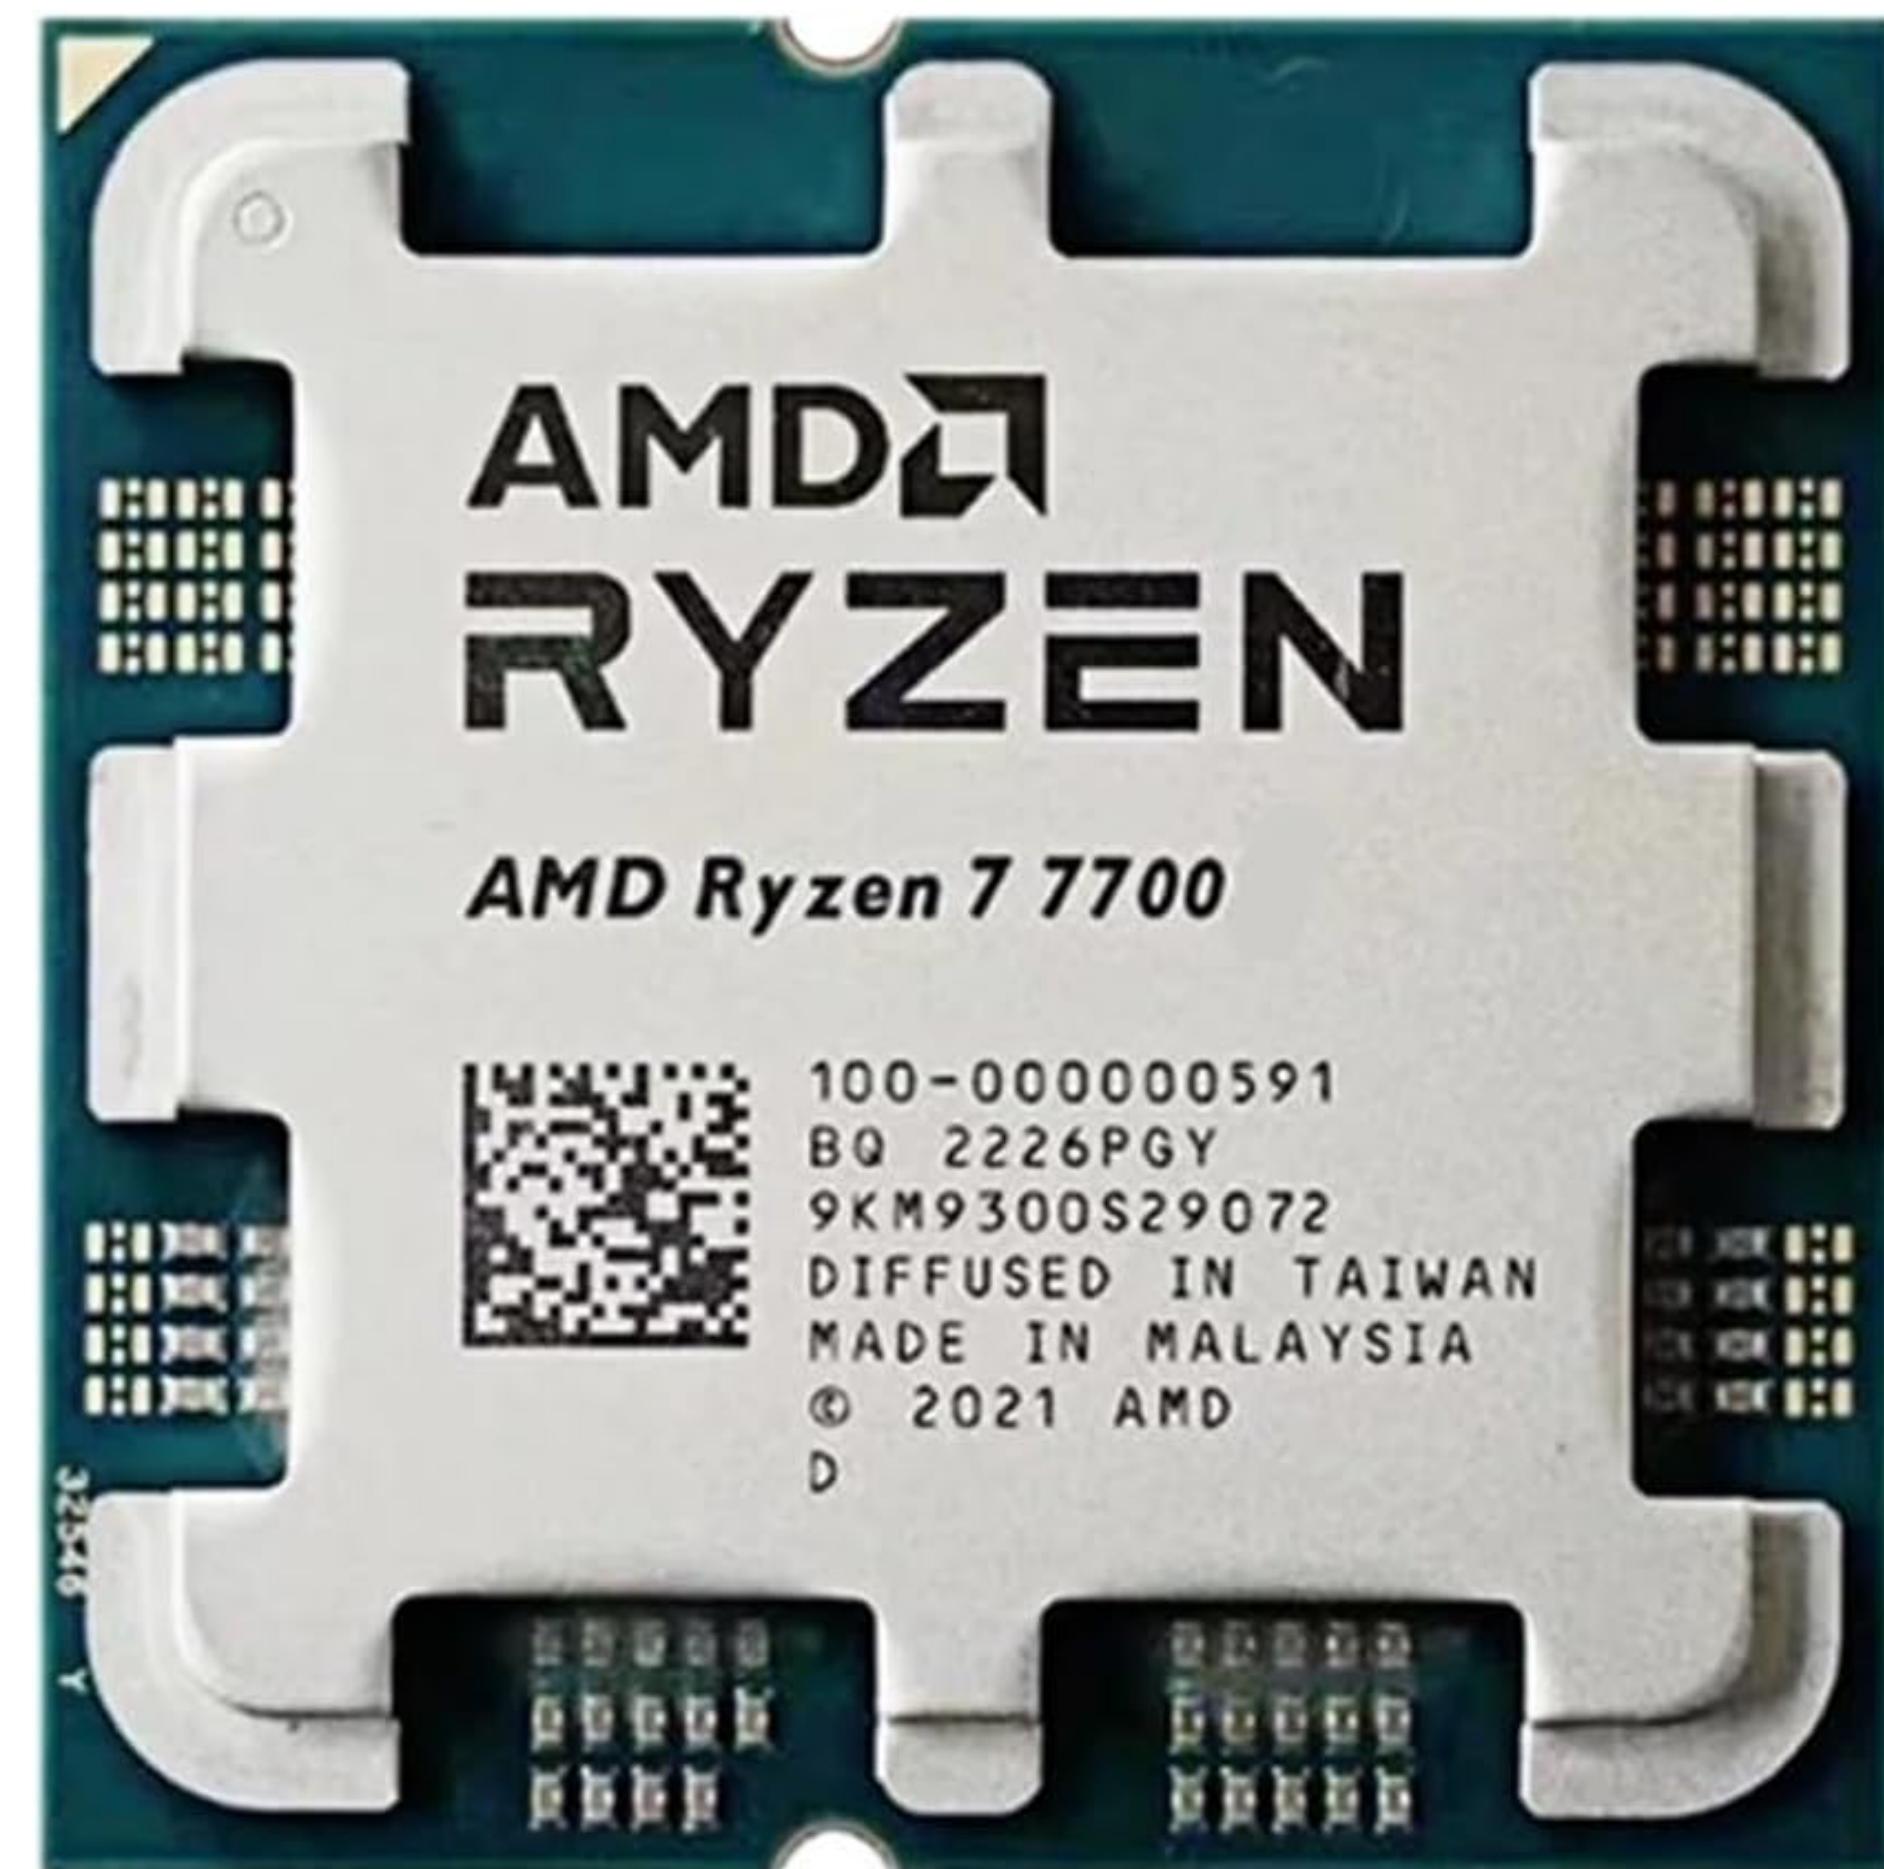
$192 \times 12 = 2,304$ threads

96k 1024-bit registers
= 12 MB

48 MB L2 cache



Our CPU: AMD Ryzen 7 7700



One core of our CPU (AMD Zen 4)

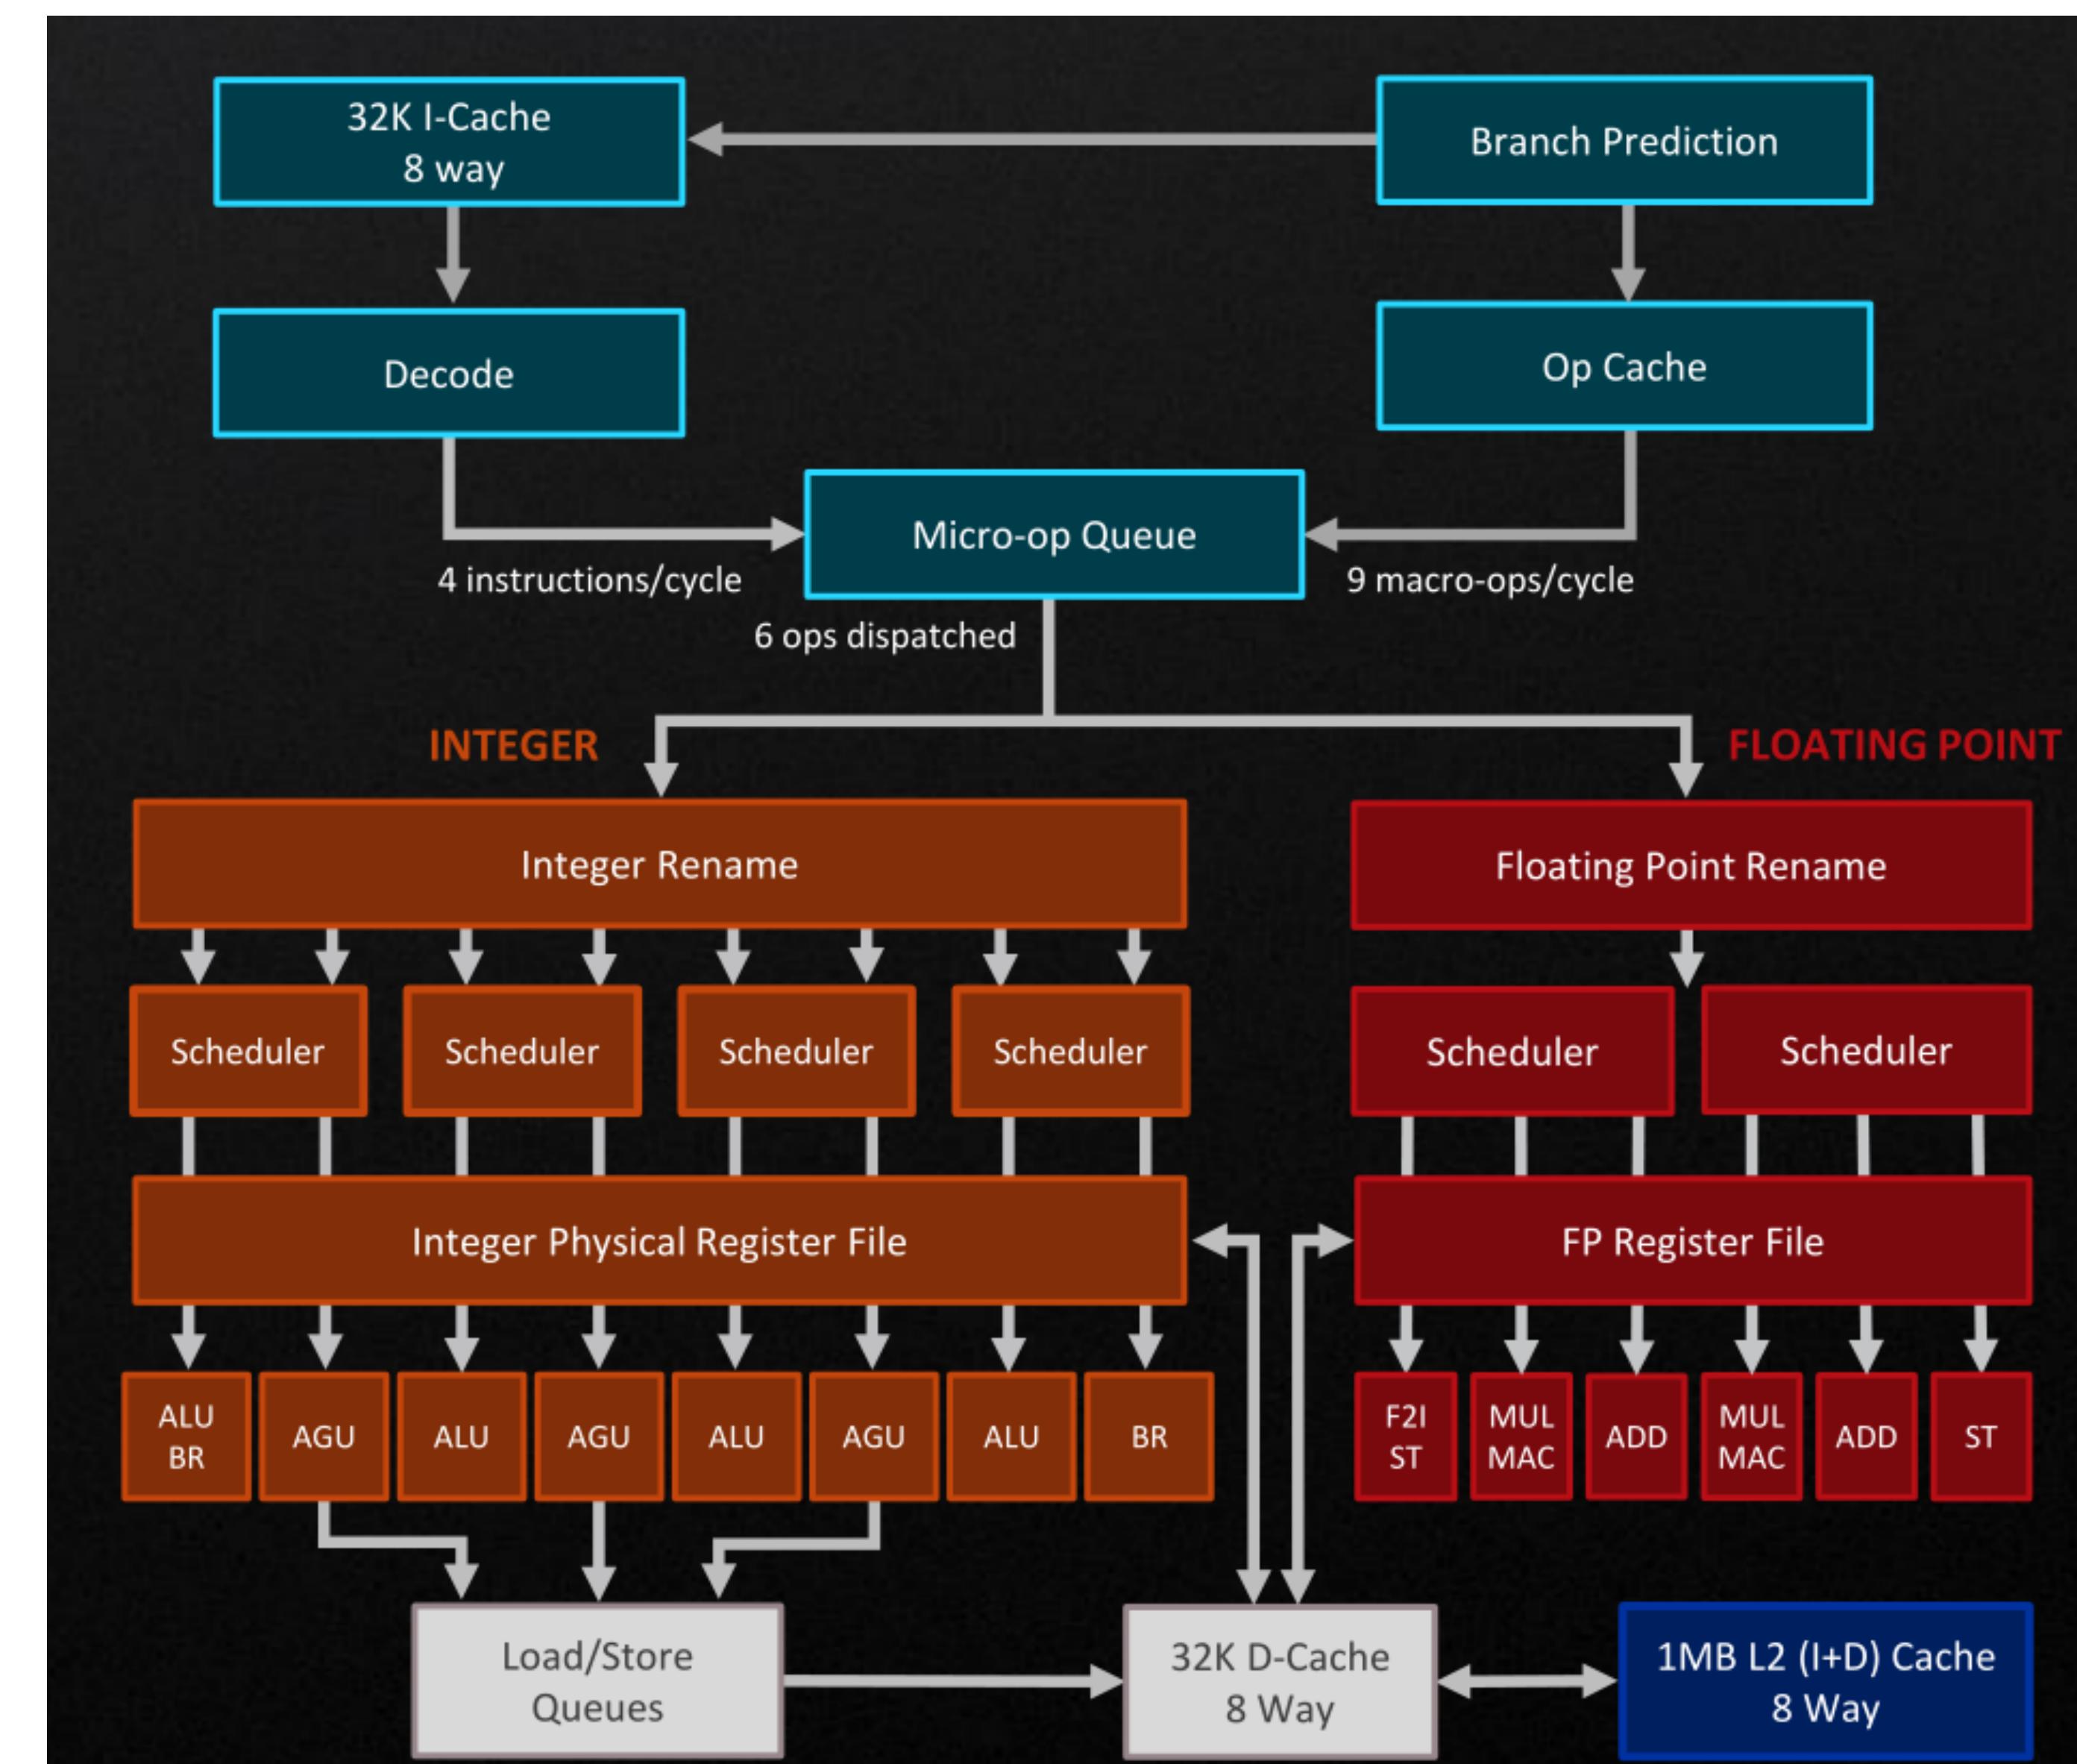
Huge, complex control logic (6-wide, OoO, ...)

Two 256-bit vector ALUs
= 16×32

192 vector registers

1 MB cache

2 “hyperthreads”



Our whole CPU

8 “Zen 4” cores

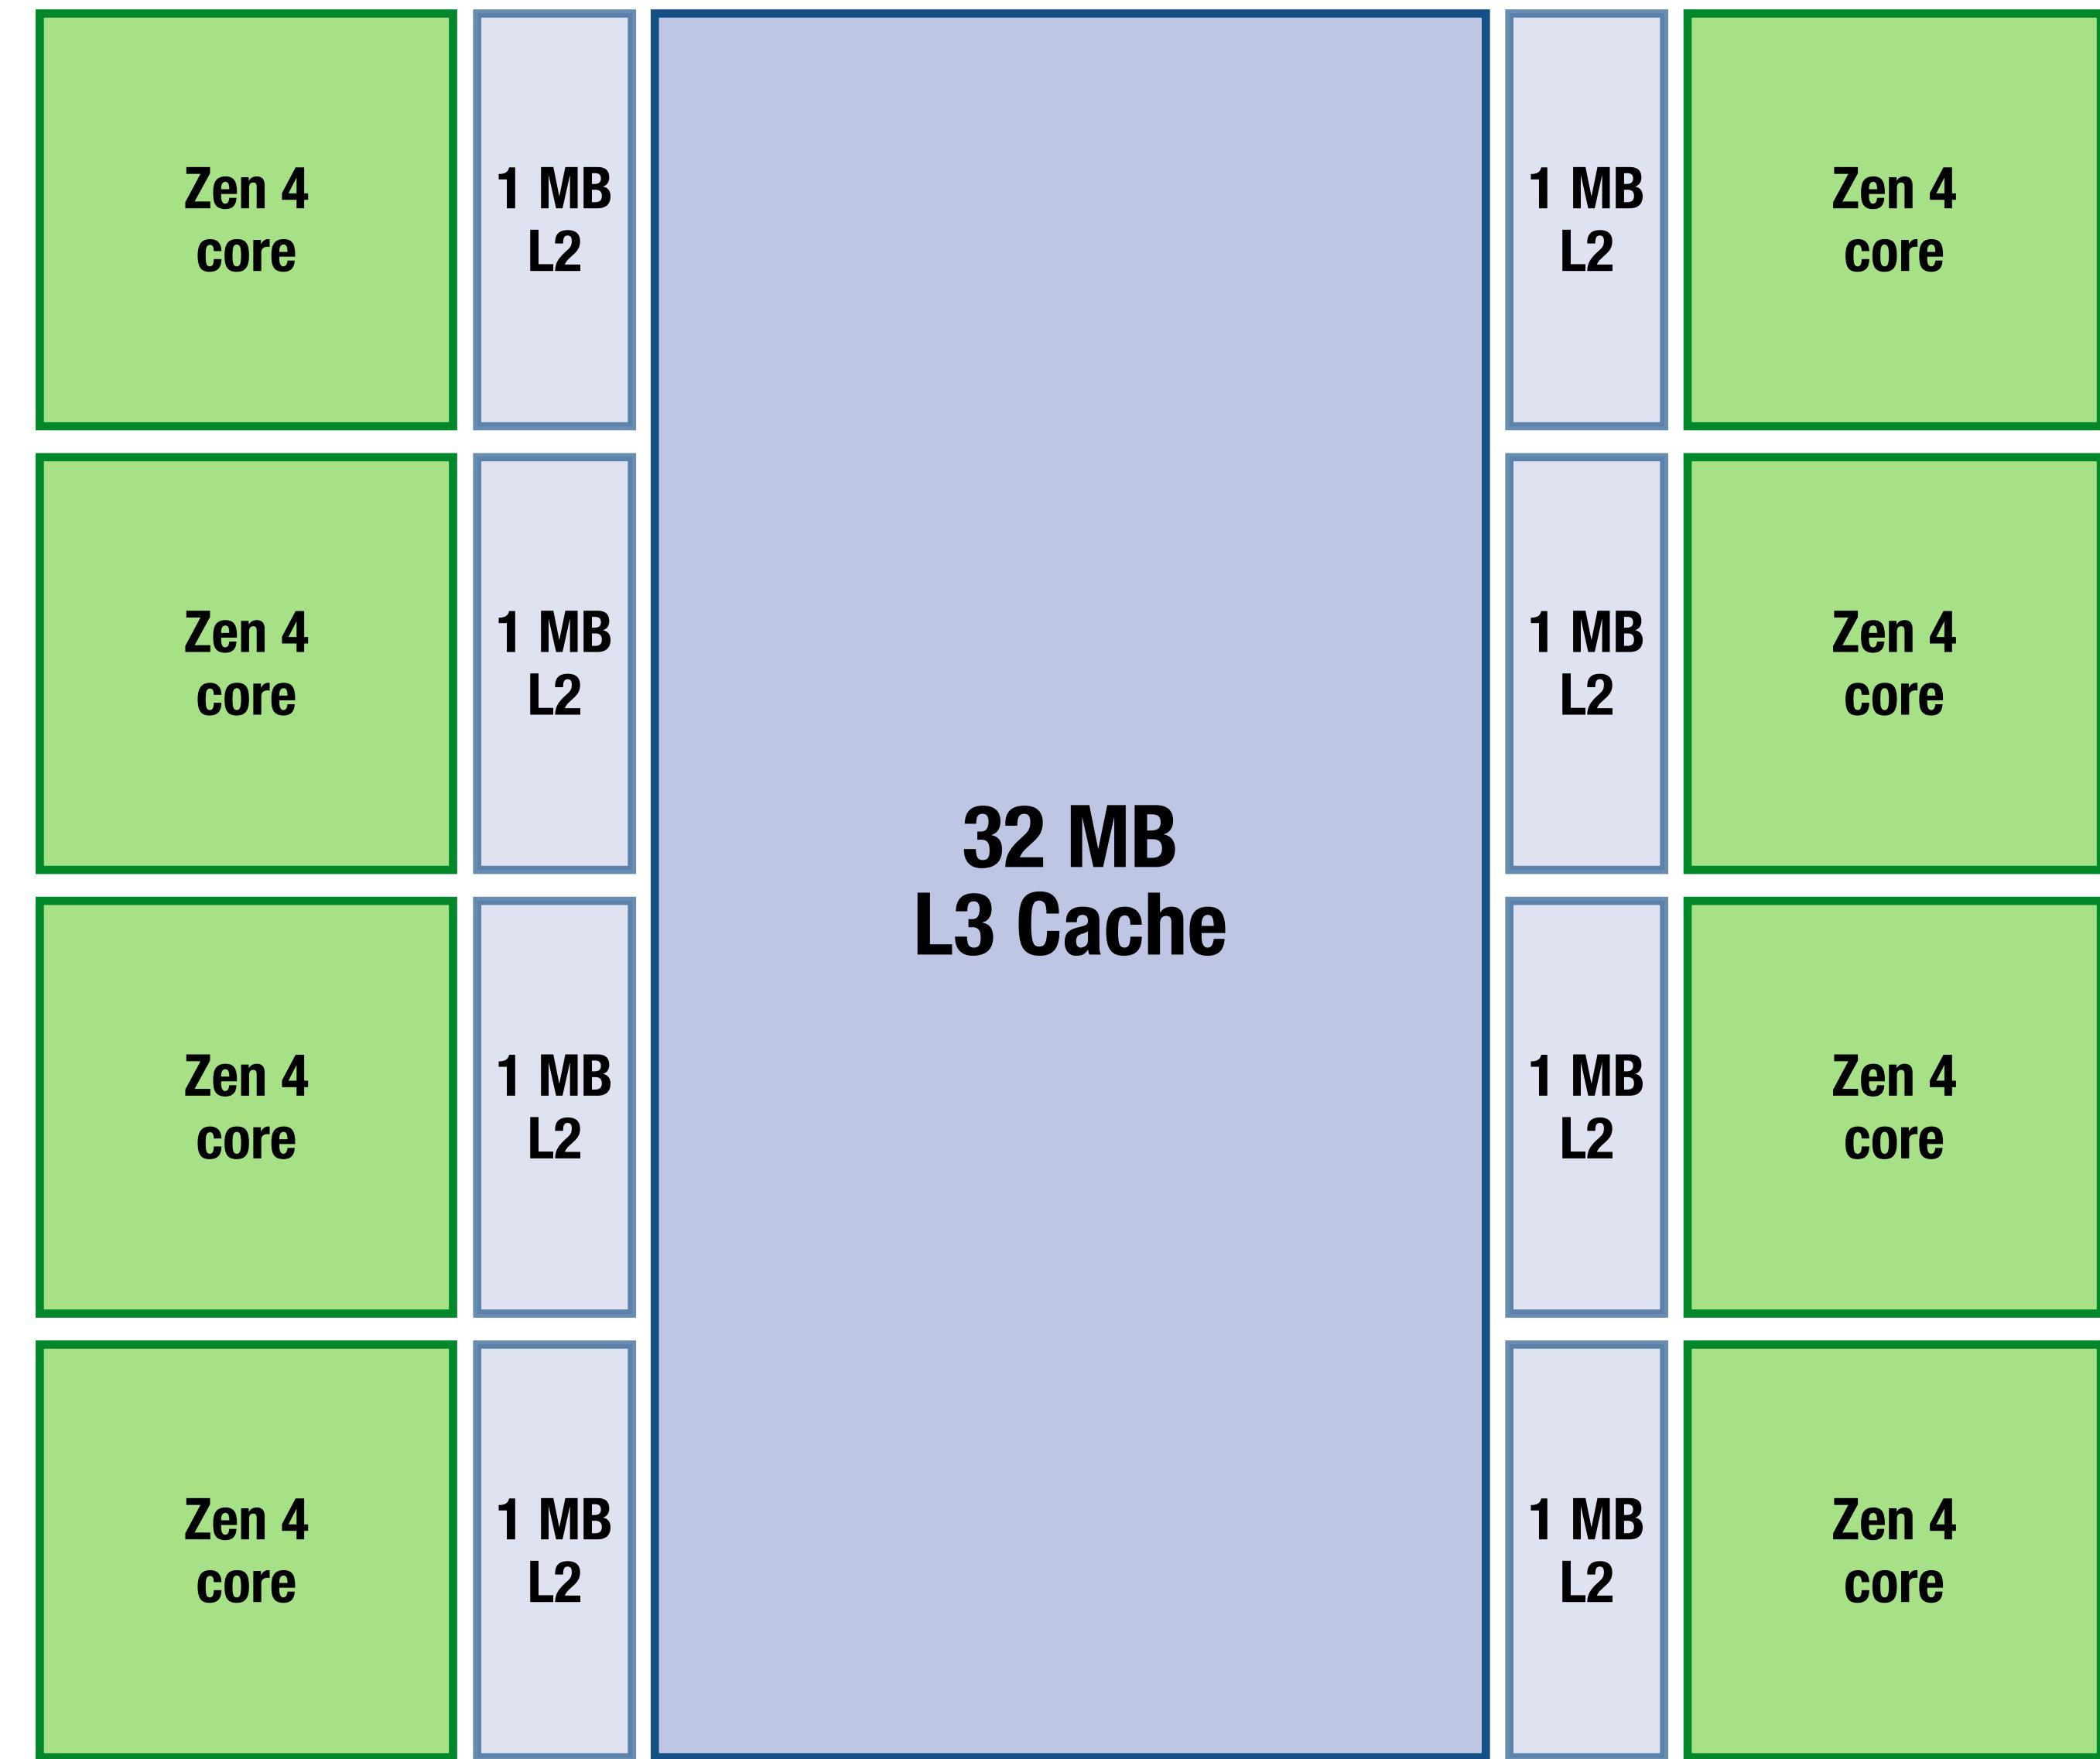
$\times 16 = 128$ exec units

@ 3.8 ~ 5.5 GHz
= 0.97 ~ 1.4 TFLOPS

$8 \times 2 = 16$ threads

1.5k 512-bit registers
= 24 KB

40 MB total L2+L3



Our CPU

8 “Zen 4” cores
 $\times 16 = 128$ exec units
@ 3.8 ~ 5.3 GHz
= 0.97 ~ 1.4 TFLOPS

8 x 2 = 16 threads

1.5k 512-bit registers
= 24 KB

40 MB total L2+L3

Our GPU

48 SMs
 $\times 4 = 192$ cores
 $\times 32 = 6144$ exec. units
@ 2.175 GHz
= 26.7 TFLOPS

192x12 = 2,304 threads

96k 1024-bit registers
= 12 MB

48 MB L2 cache

Our CPU

70 mm²
(CCD only)

70 mm²

/ 8 cores

= 8.75 mm²

6-7x faster
on Mandelbrot

Our GPU

294 mm²
(full die)

~200 mm²

/ 192 cores

= 1.05 mm² / core

8x smaller

**A throughput processor
is still a processor!**

(or actually, many of them)

**But we need to change our
programs to use it efficiently
expose explicit parallelism
within & across instruction streams**