

October 7, 2025

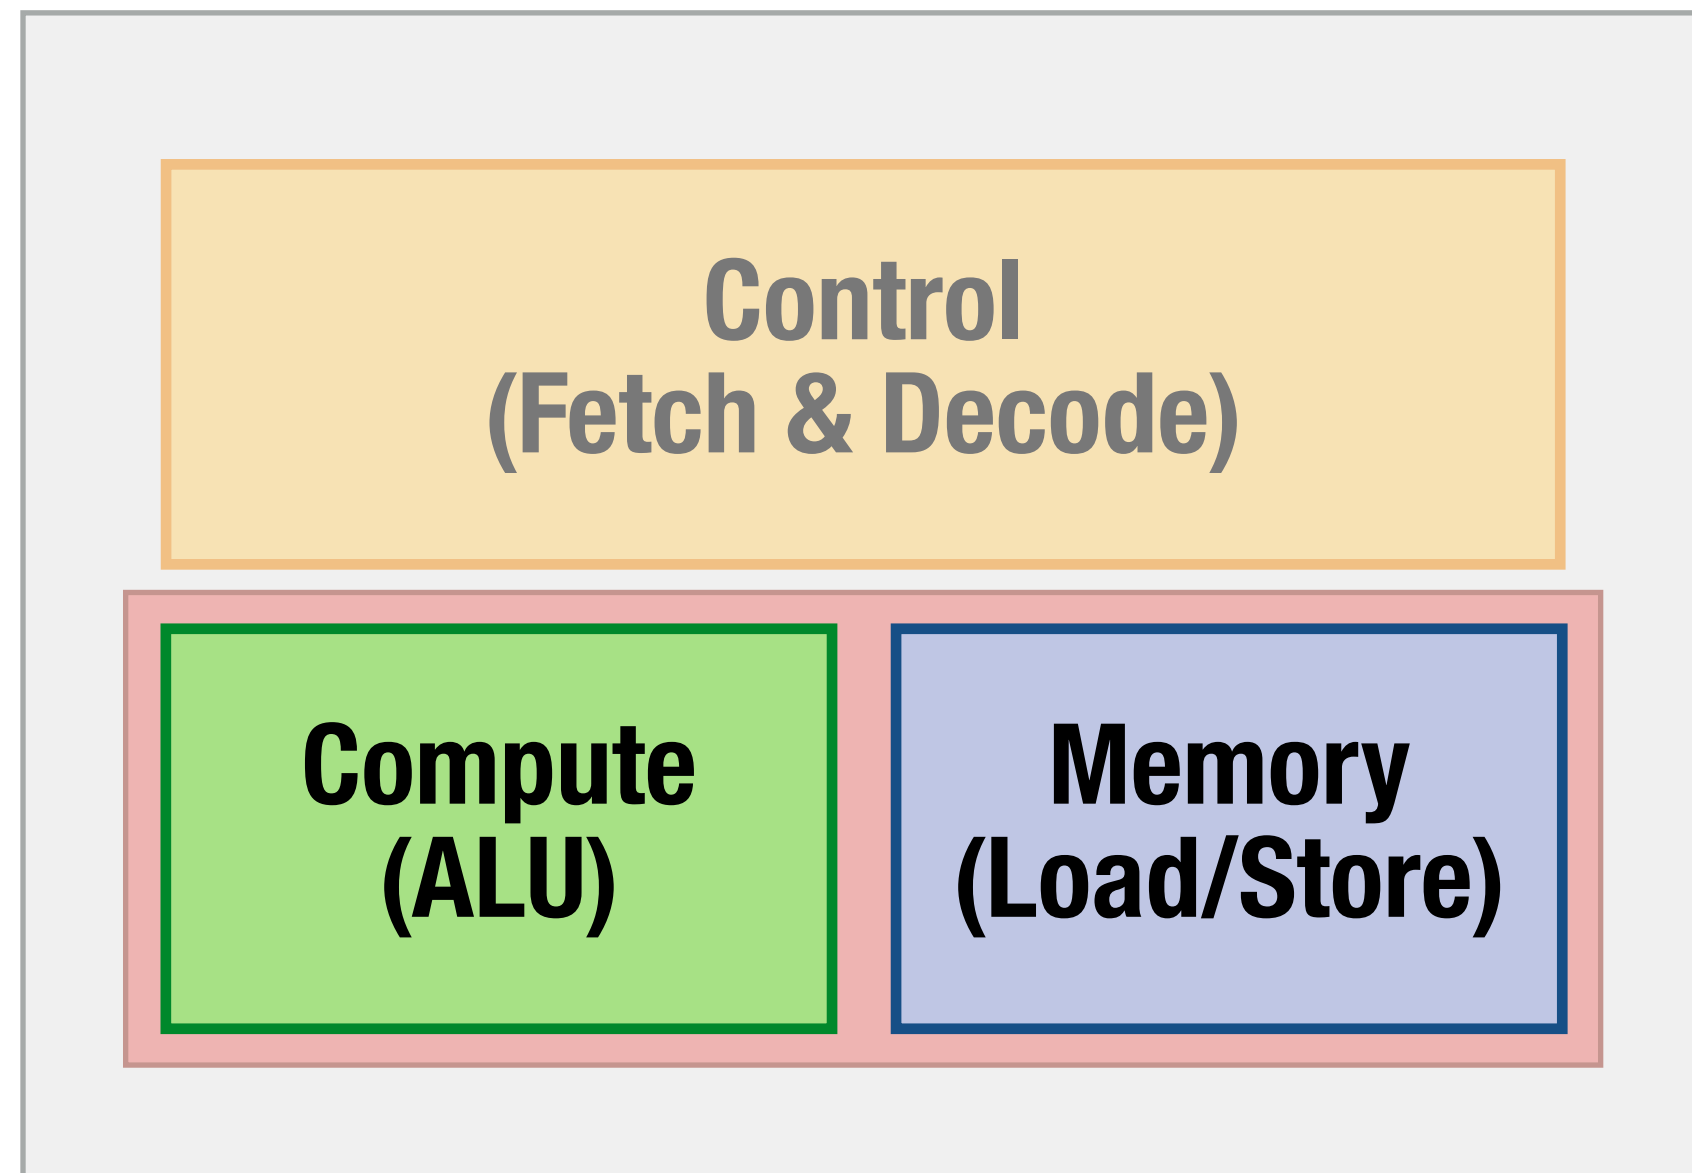
6.S894

Accelerated Computing

**Live Lab 5:
Matmul, part 2**

Jonathan Ragan-Kelley 

Overlapping compute & I/O



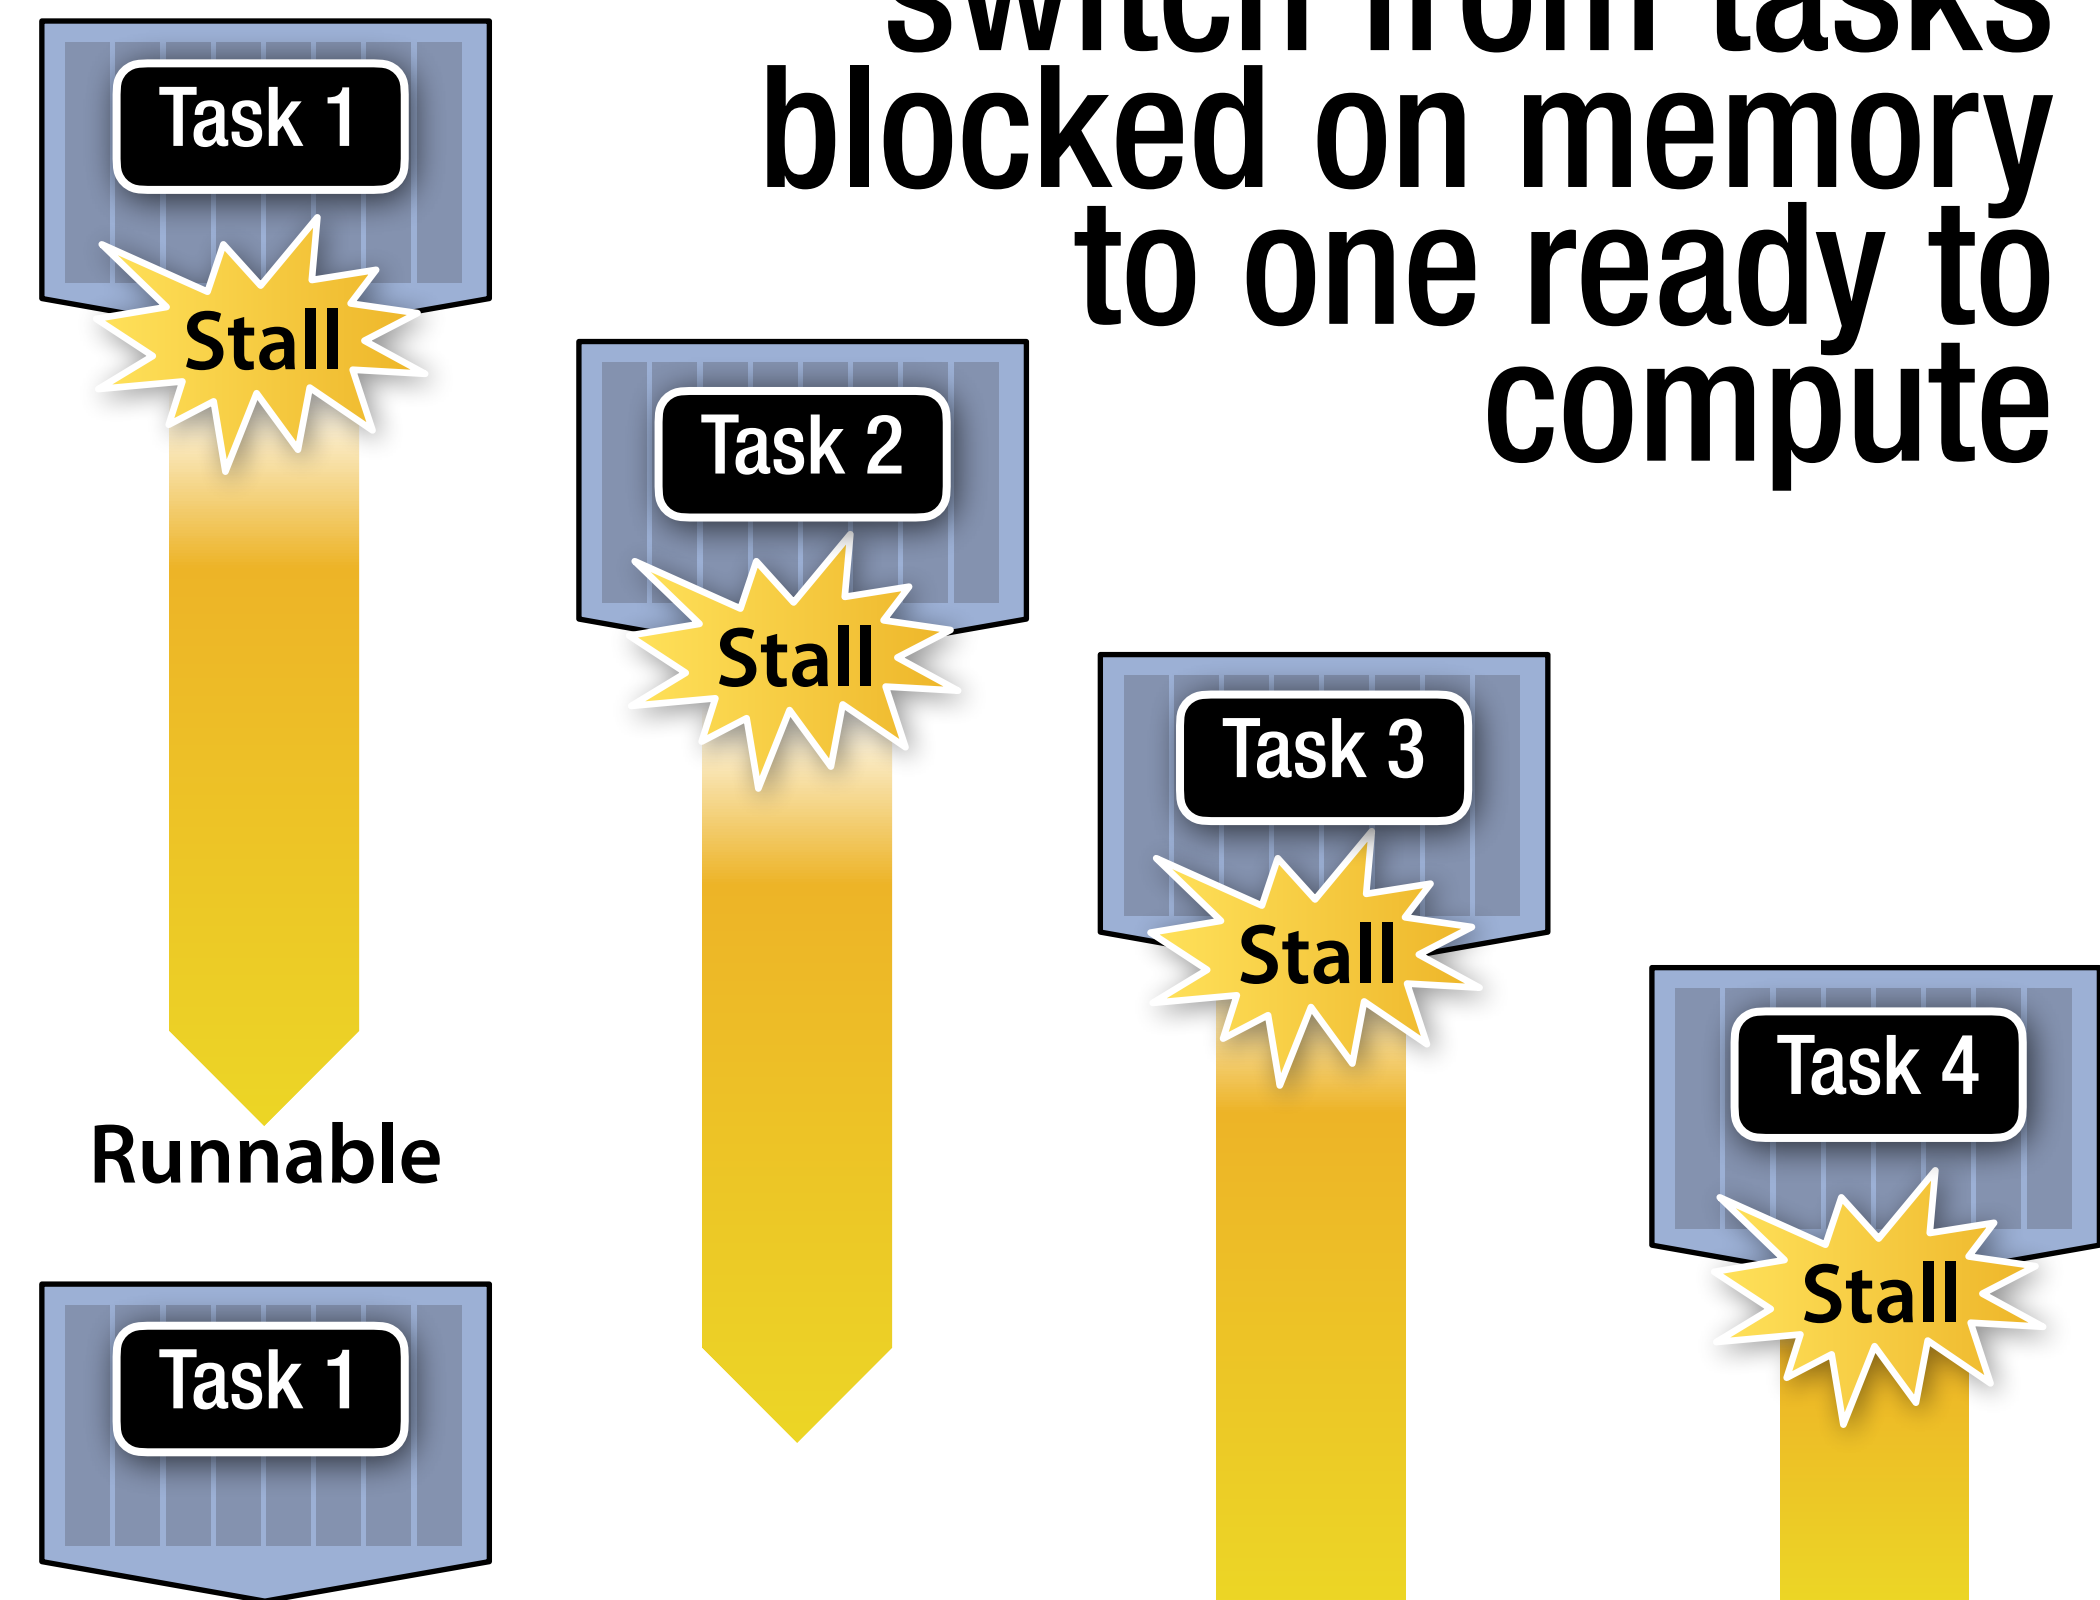
**Goal: fully utilize
both resources**

Problem 1: load / store instructions are asynchronous & long-latency

Solution 1: ILP hoist loads early to avoid blocking

ld	→	ld
fma		ld
ld		ld
fma		fma
ld		fma
fma		fma
...		...

Solution 2: multithreading switch from tasks blocked on memory to one ready to compute



Problem 2: load / store instructions waste **issue slots**

**Solution: bulk load /
store instructions**
e.g., “vectorized” ld / st

ld.f32
ld.f32
ld.f32
ld.f32
fma
fma
fma
fma
...

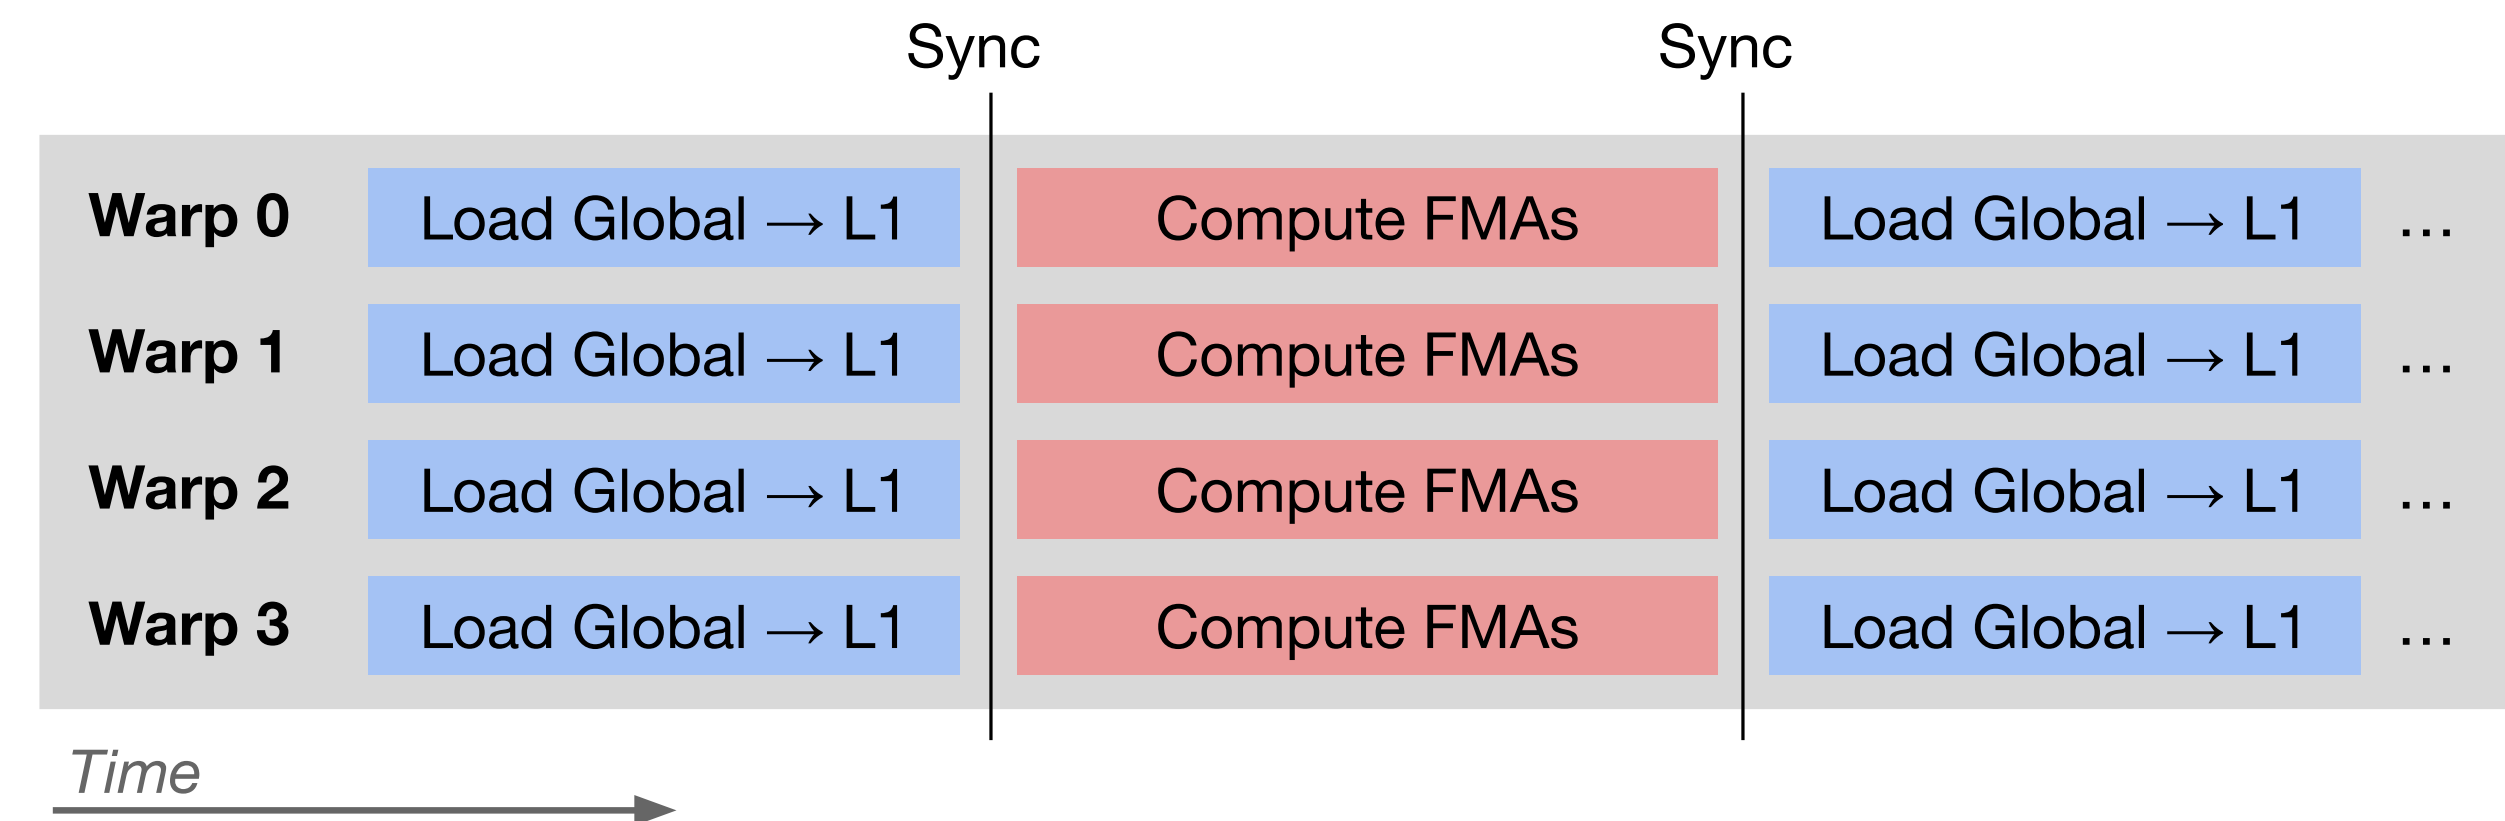


ld.v4.f32
fma
fma
fma
fma
...

Problem 3: overlapping compute with loading to the **scratchpad**

```
foreach tile:  
    // load into scratchpad  
    for i,j:  
        load next A,B → scratch  
    sync  
    // compute!  
    for i,j,k:  
        compute C += A*B  
    sync
```


No Overlapping:



Problem 3: overlapping compute with loading to the **scratchpad**

```
foreach tile:  
    // load into scratchpad  
    for i,j:  
        load next A,B → scratch  
    sync  
    // compute!  
    for i,j,k:  
        compute C += A*B  
    sync
```

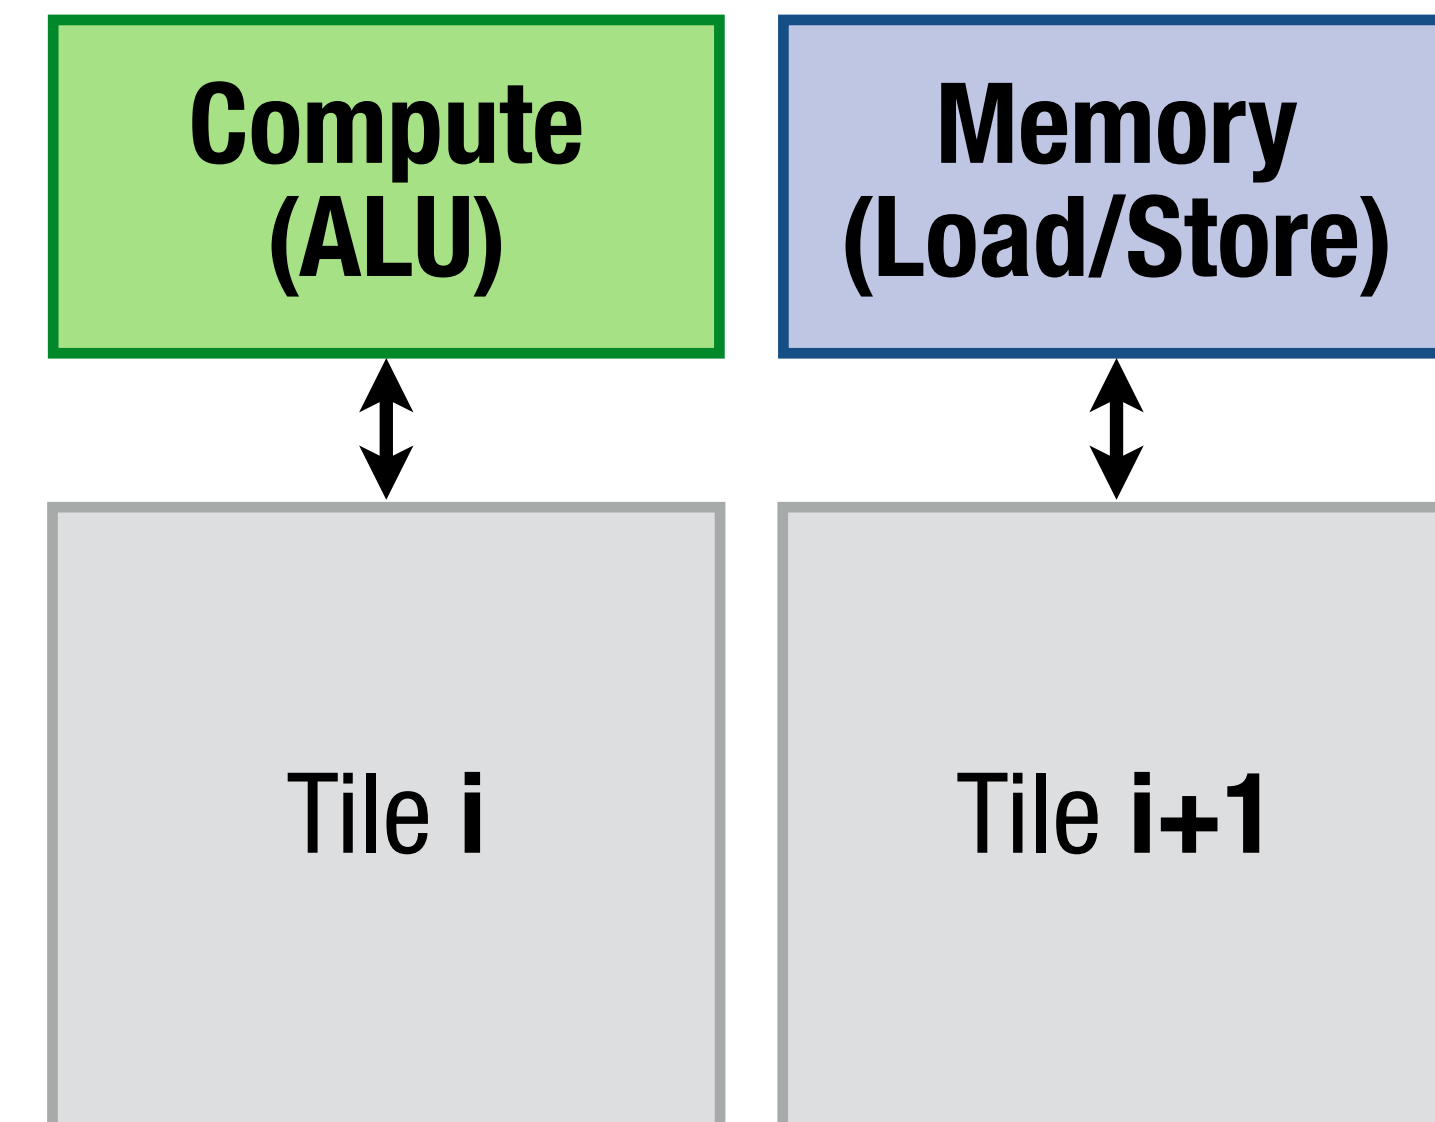
Solution:
asynchronous fetch & double-buffering



Problem 3: overlapping compute with loading to the **scratchpad**


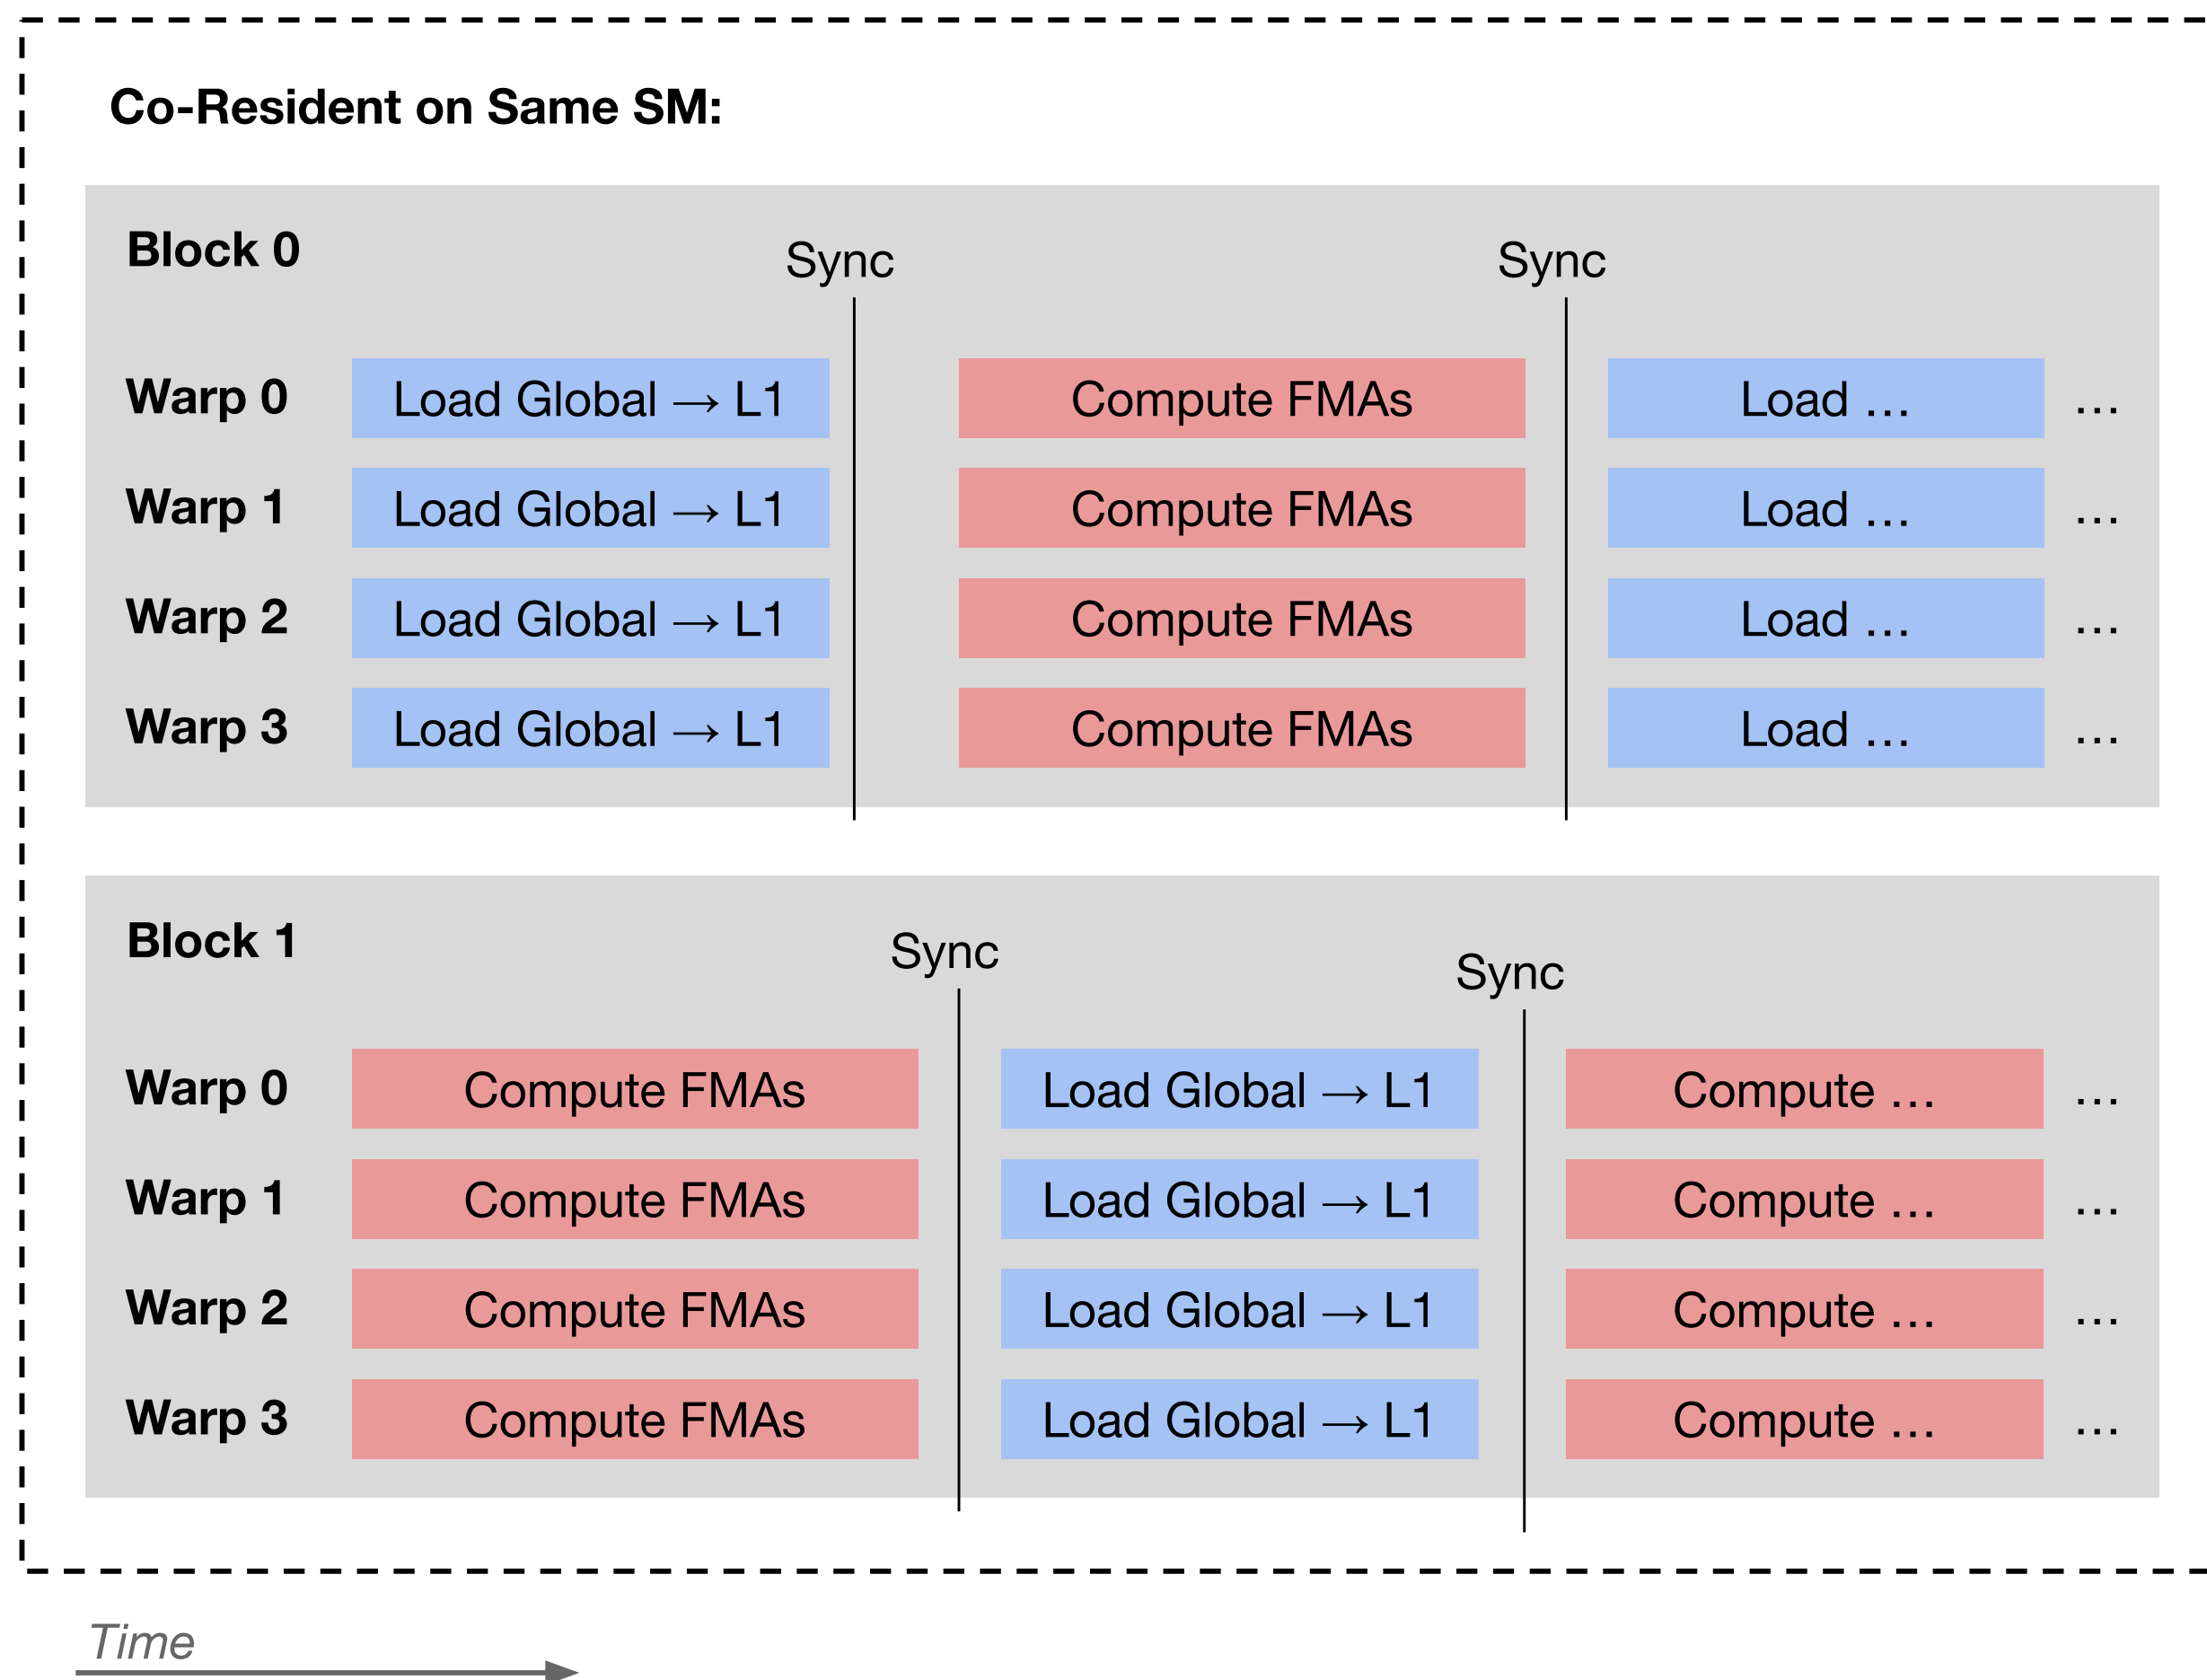
```
foreach tile:  
    // load into scratchpad  
    for i,j:  
        load next A,B → scratch  
    sync  
    // compute!  
    for i,j,k:  
        compute C += A*B  
    sync
```

Solution:
asynchronous fetch & double-buffering



Problem 3: overlapping compute with loading to the **scratchpad**

Solution:
asynchronous fetch &
double-buffering




Problem 3: overlapping compute with loading to the **scratchpad**

Implementation approach:
warp specialization

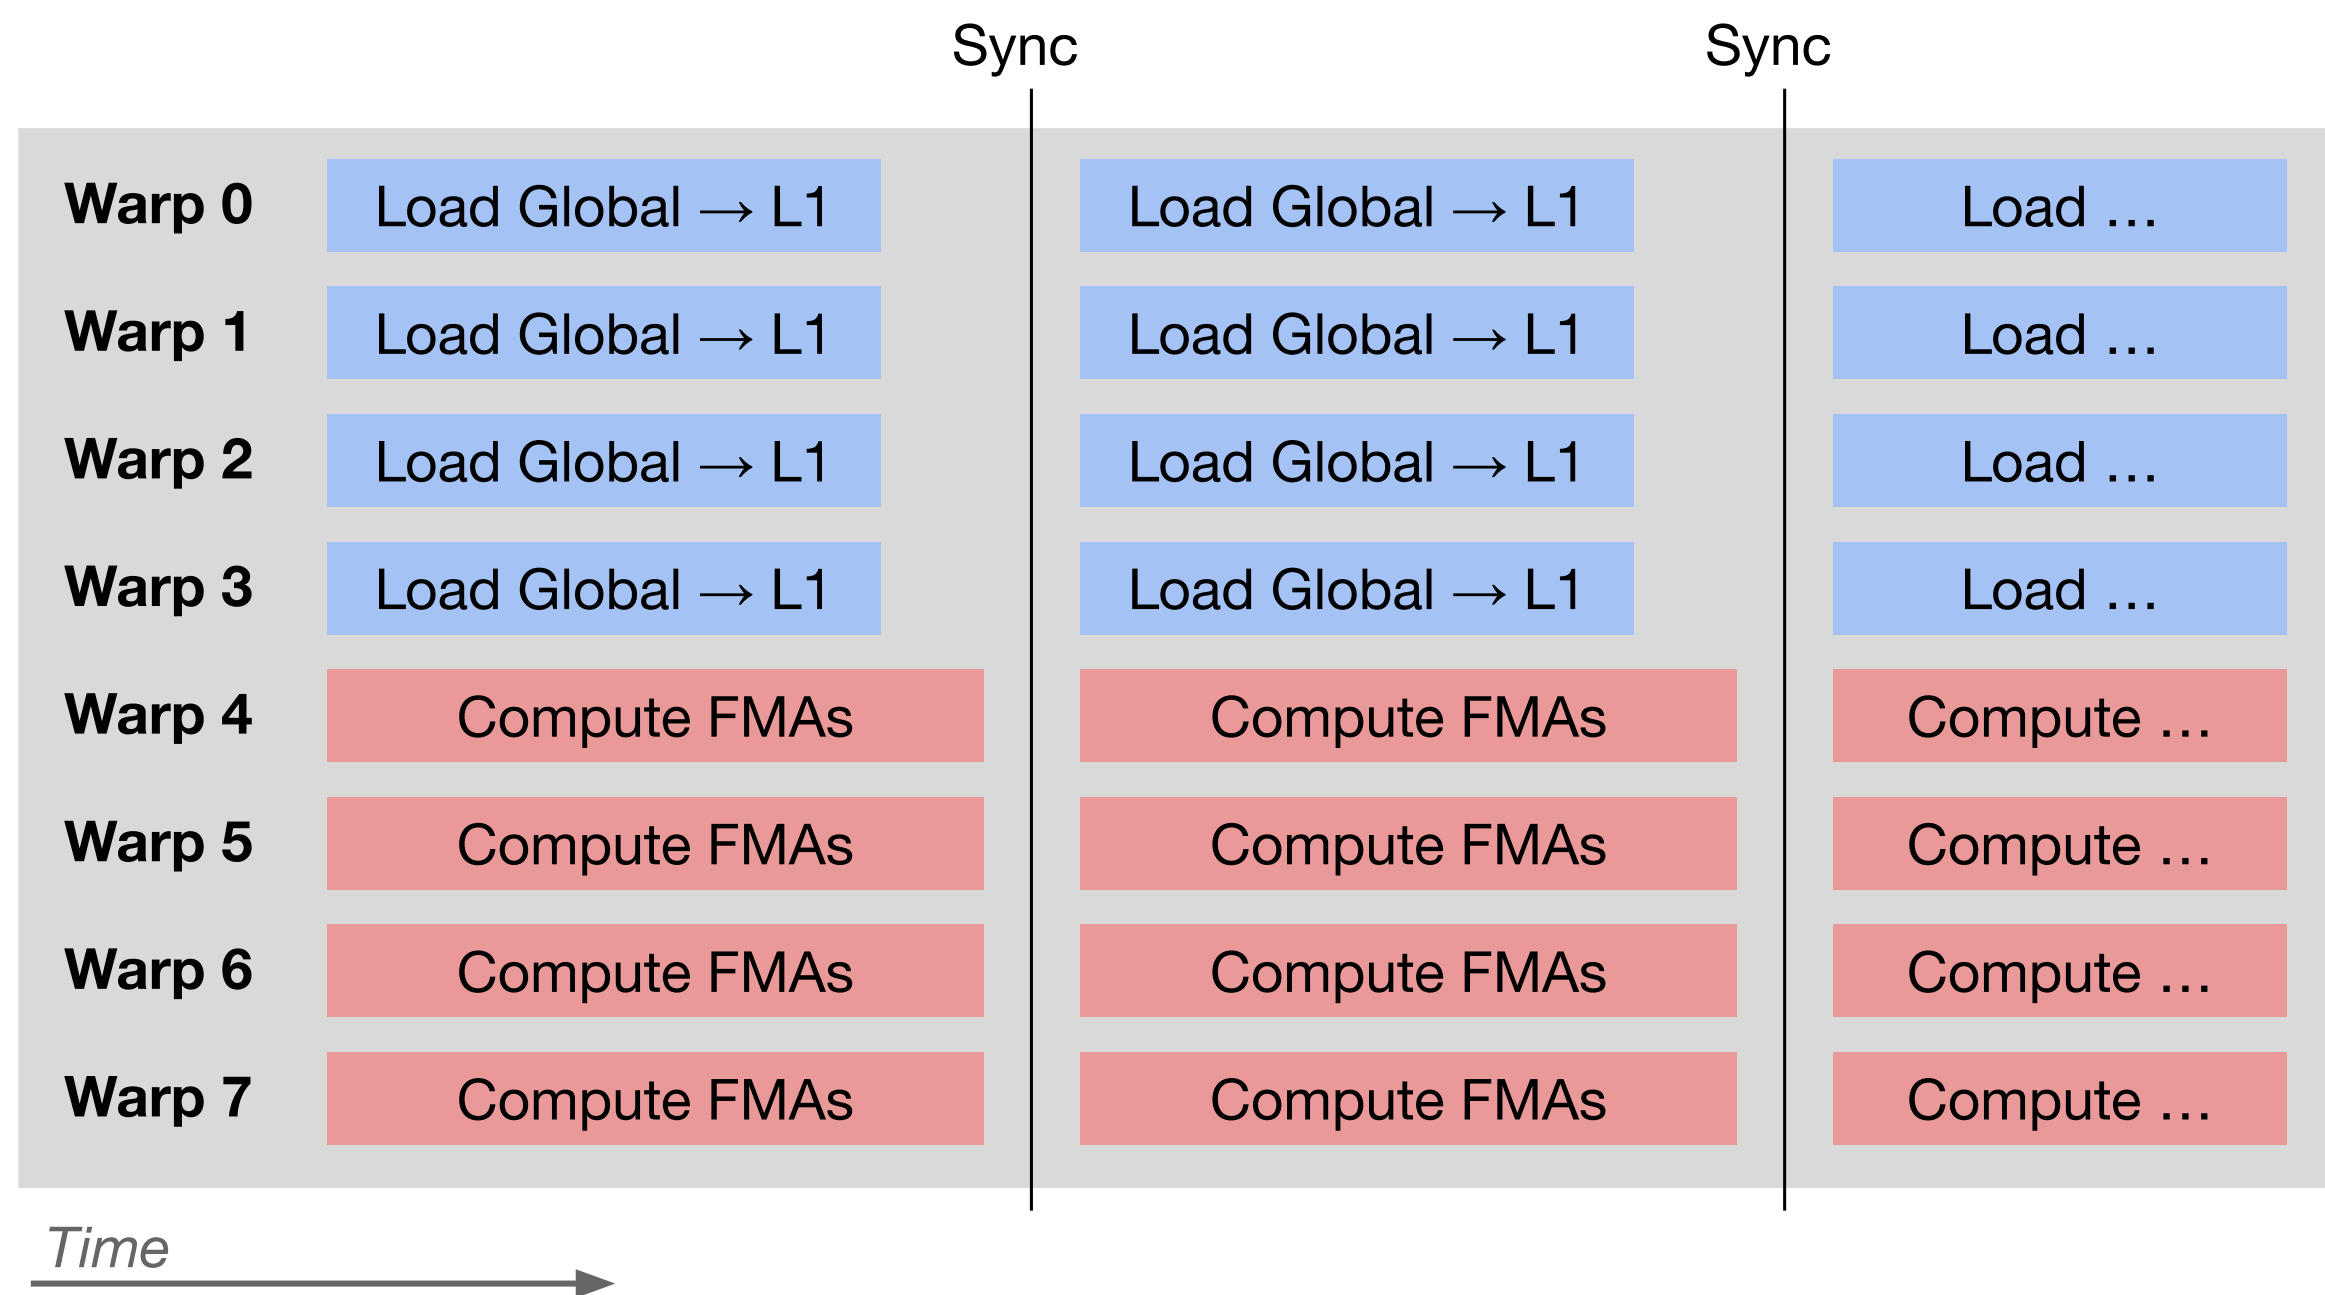
```
if threadIdx.y < 4:  
    // load into scratchpad  
    for i,j:  
        load next A,B → scratch  
else:  
    // compute!  
    for i,j,k:  
        compute C += A*B  
sync & swap buffers...
```

Solution:
asynchronous fetch &
double-buffering

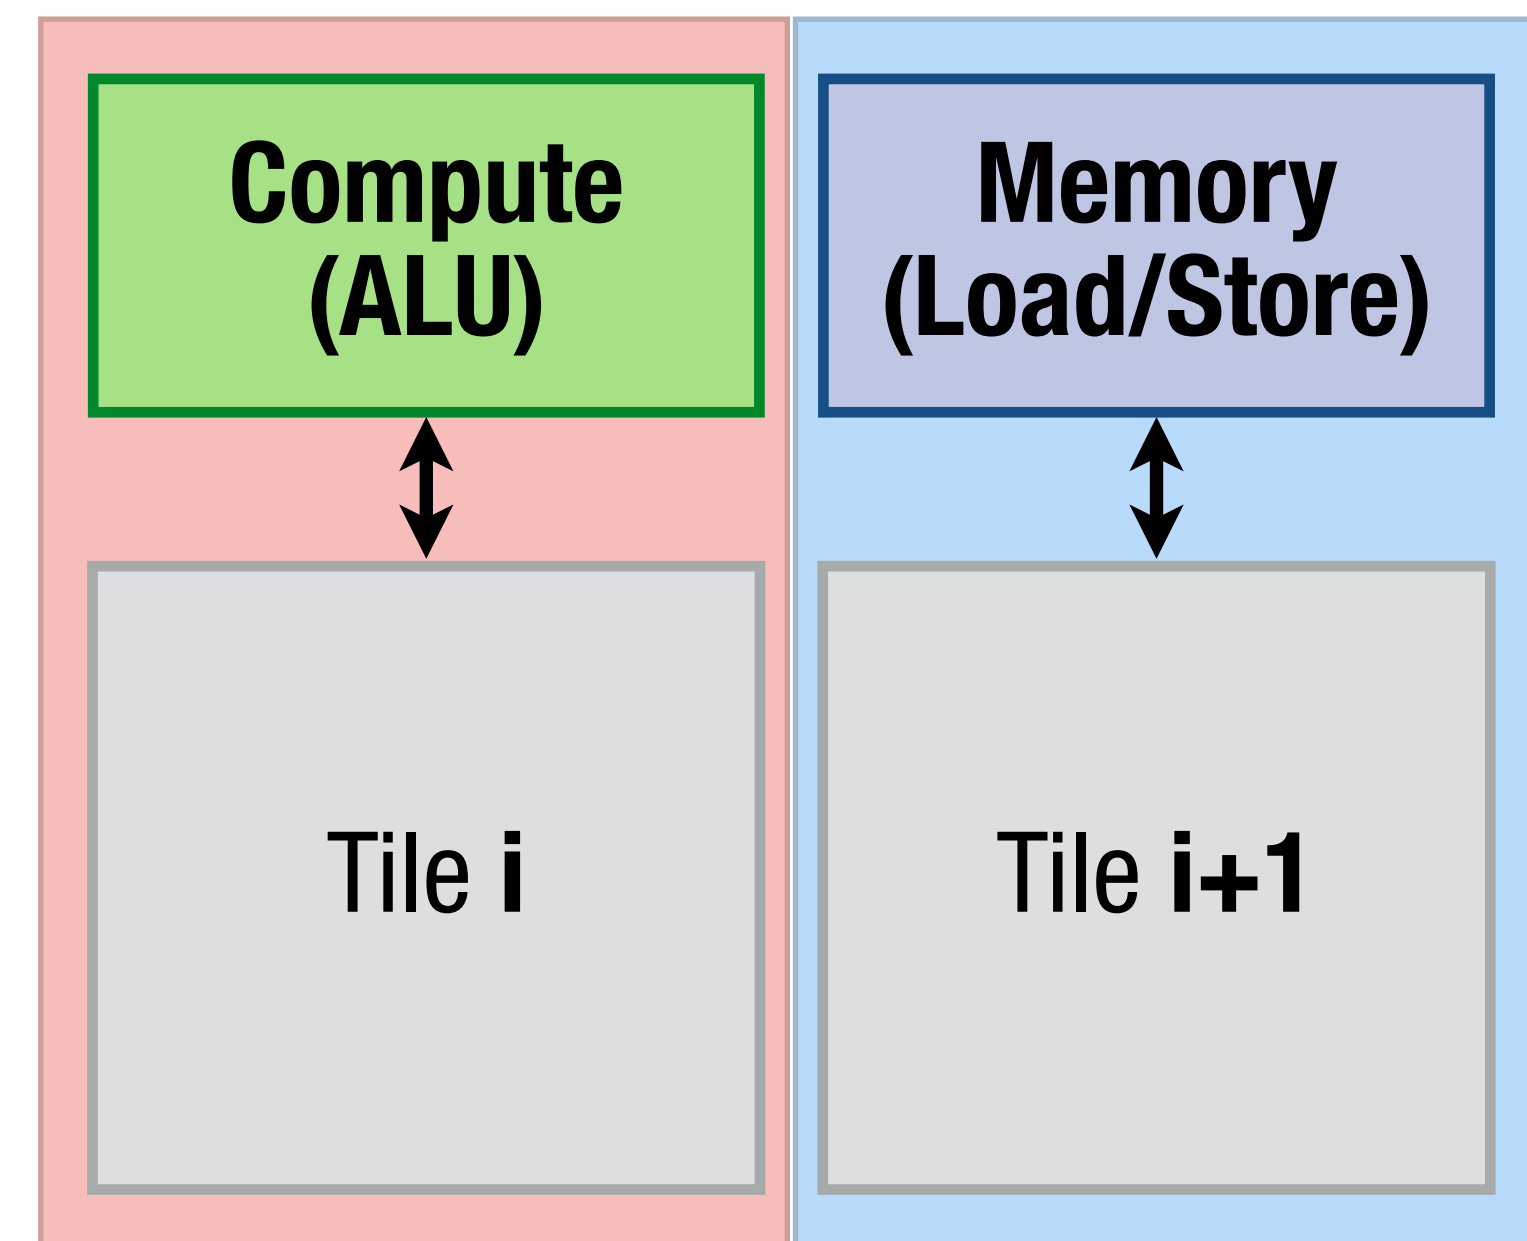


Problem 3: overlapping compute with loading to the **scratchpad**

Warp-Specialized Overlapping:



Solution:
asynchronous fetch &
double-buffering



Problem 3: overlapping compute with loading to the **scratchpad**

Implementation approach:
warp specialization

foreach tile:

```
// load into scratchpad
```

```
for i,j:  
    async load next A,B → scratch
```


```
await previous tile load
```

```
// compute!
```

```
for i,j,k:  
    compute C += A*B
```

```
sync & swap buffers...
```

Solution:
asynchronous fetch &
double-buffering



Problem 4: load to scratch wastes issue slots, register file space & bandwidth

```
ld.global r1, [G]  
st.shared [S], r1
```

...

```
ld.shared r2, [S]
```

```
cp.async.shared.global \  
[S], [G], 16
```

...

```
cp.async.wait_all  
ld.shared r2, [S]
```

**Alternative implementation:
async memcpy instructions**

Questions?