

October 2, 2025

**6.S894**

# **Accelerated Computing**

## **Lecture 5: Memory, Overlapping Compute & I/O**

Jonathan Ragan-Kelley 

# L1 SRAM

128 KB per-SM

(  $\times 48$  SMs = 6 MB )

128 bytes / cycle / SM

↳ 1 warp-wide ld/st  
(Per-core: 1 every 4 cycles)

$\times 48$  SMs  $\times 2.18$ GHz = 13.4 TB/s



# L1 SRAM: also used as **explicit scratchpad**




**Each block (SM) only sees  
its own scratchpad**

# Tradeoffs: scratchpads vs. caches

- + No tag / lookup overhead
- + Predictable / controllable
- + No need for coherence
- + Read  $>1$  line / cycle
- Programming effort
- Software depends on size


# How is the L1 SRAM built?

How can it deliver 128 B / cycle? **Parallelism!**



**gather**


$32 \times 4 \text{ Bytes / cycle}$




$32 \times 4 \text{ Bytes / cycle}$

**scatter**

# How can we handle **gather/scatter** in the **L1 SRAM**?

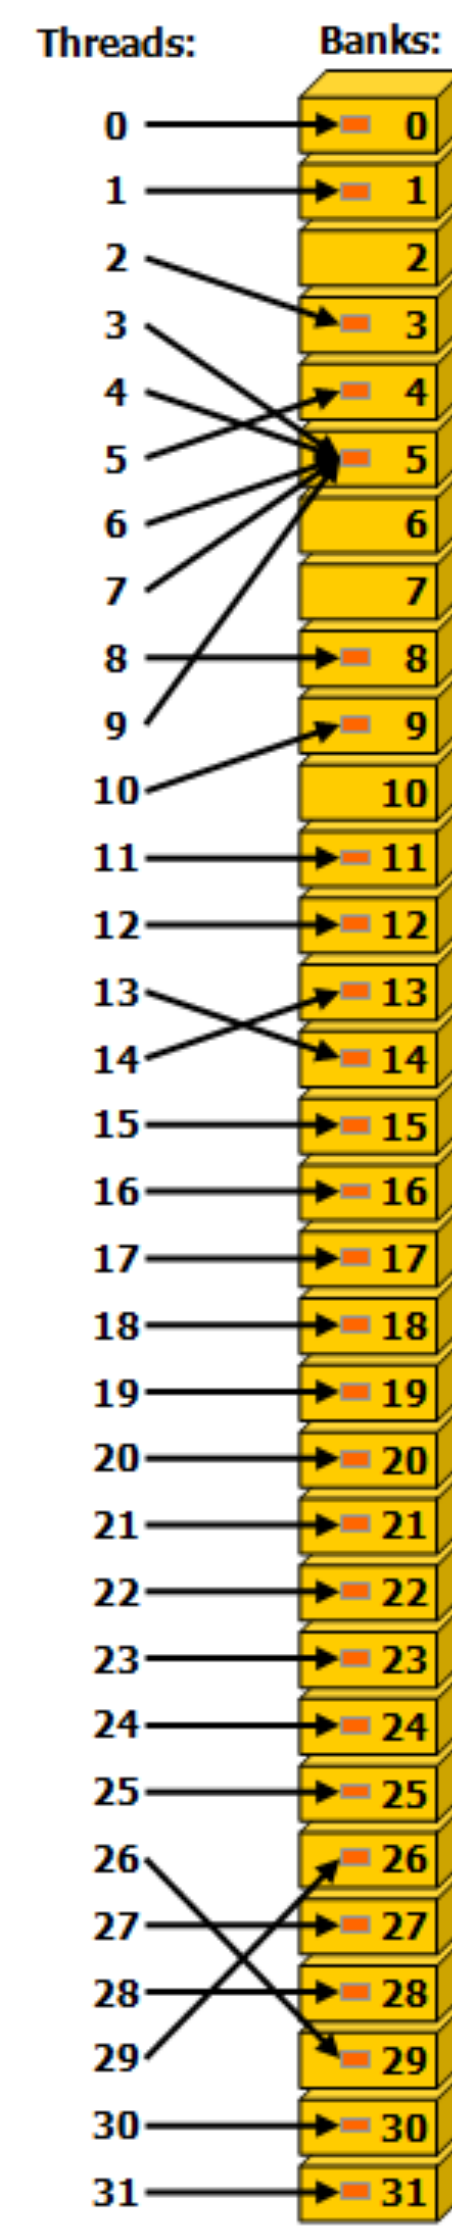







How can we handle **gather/scatter**  
in the **L1 SRAM**? All-to-all **crossbar**!






# L1 bank conflicts



1 load /  
150 ops


360  
GB/sec



high **bandwidth**,  
limited **capacity**  
high clocks &  
wide interface

1 load /  
20 ops


2.5  
TB/sec



**aggregate** large  
transactions for  
DRAM  
**streaming**  
access  
large-scale **reuse**

1 load /  
4 ops

13.4  
TB/sec



shared per-SM,  
**not coherent**  
high **bandwidth**  
via **banking**  
cache or  
**scratchpad**



1 load /  
150 ops

360  
GB/sec

1 load /  
20 ops






2.5  
TB/sec

1 load /  
4 ops

13.4  
TB/sec

3 loads + 1 store /  
1 op

>100  
TB/sec




high **bandwidth**,  
limited **capacity**  
high clocks &  
wide interface

**aggregate** large  
transactions for  
DRAM  
**streaming**  
access  
large-scale **reuse**

shared per-SM,  
**not coherent**  
high **bandwidth**  
via **banking**  
cache or  
**scratchpad**

# Register file: parallel banks per-lane

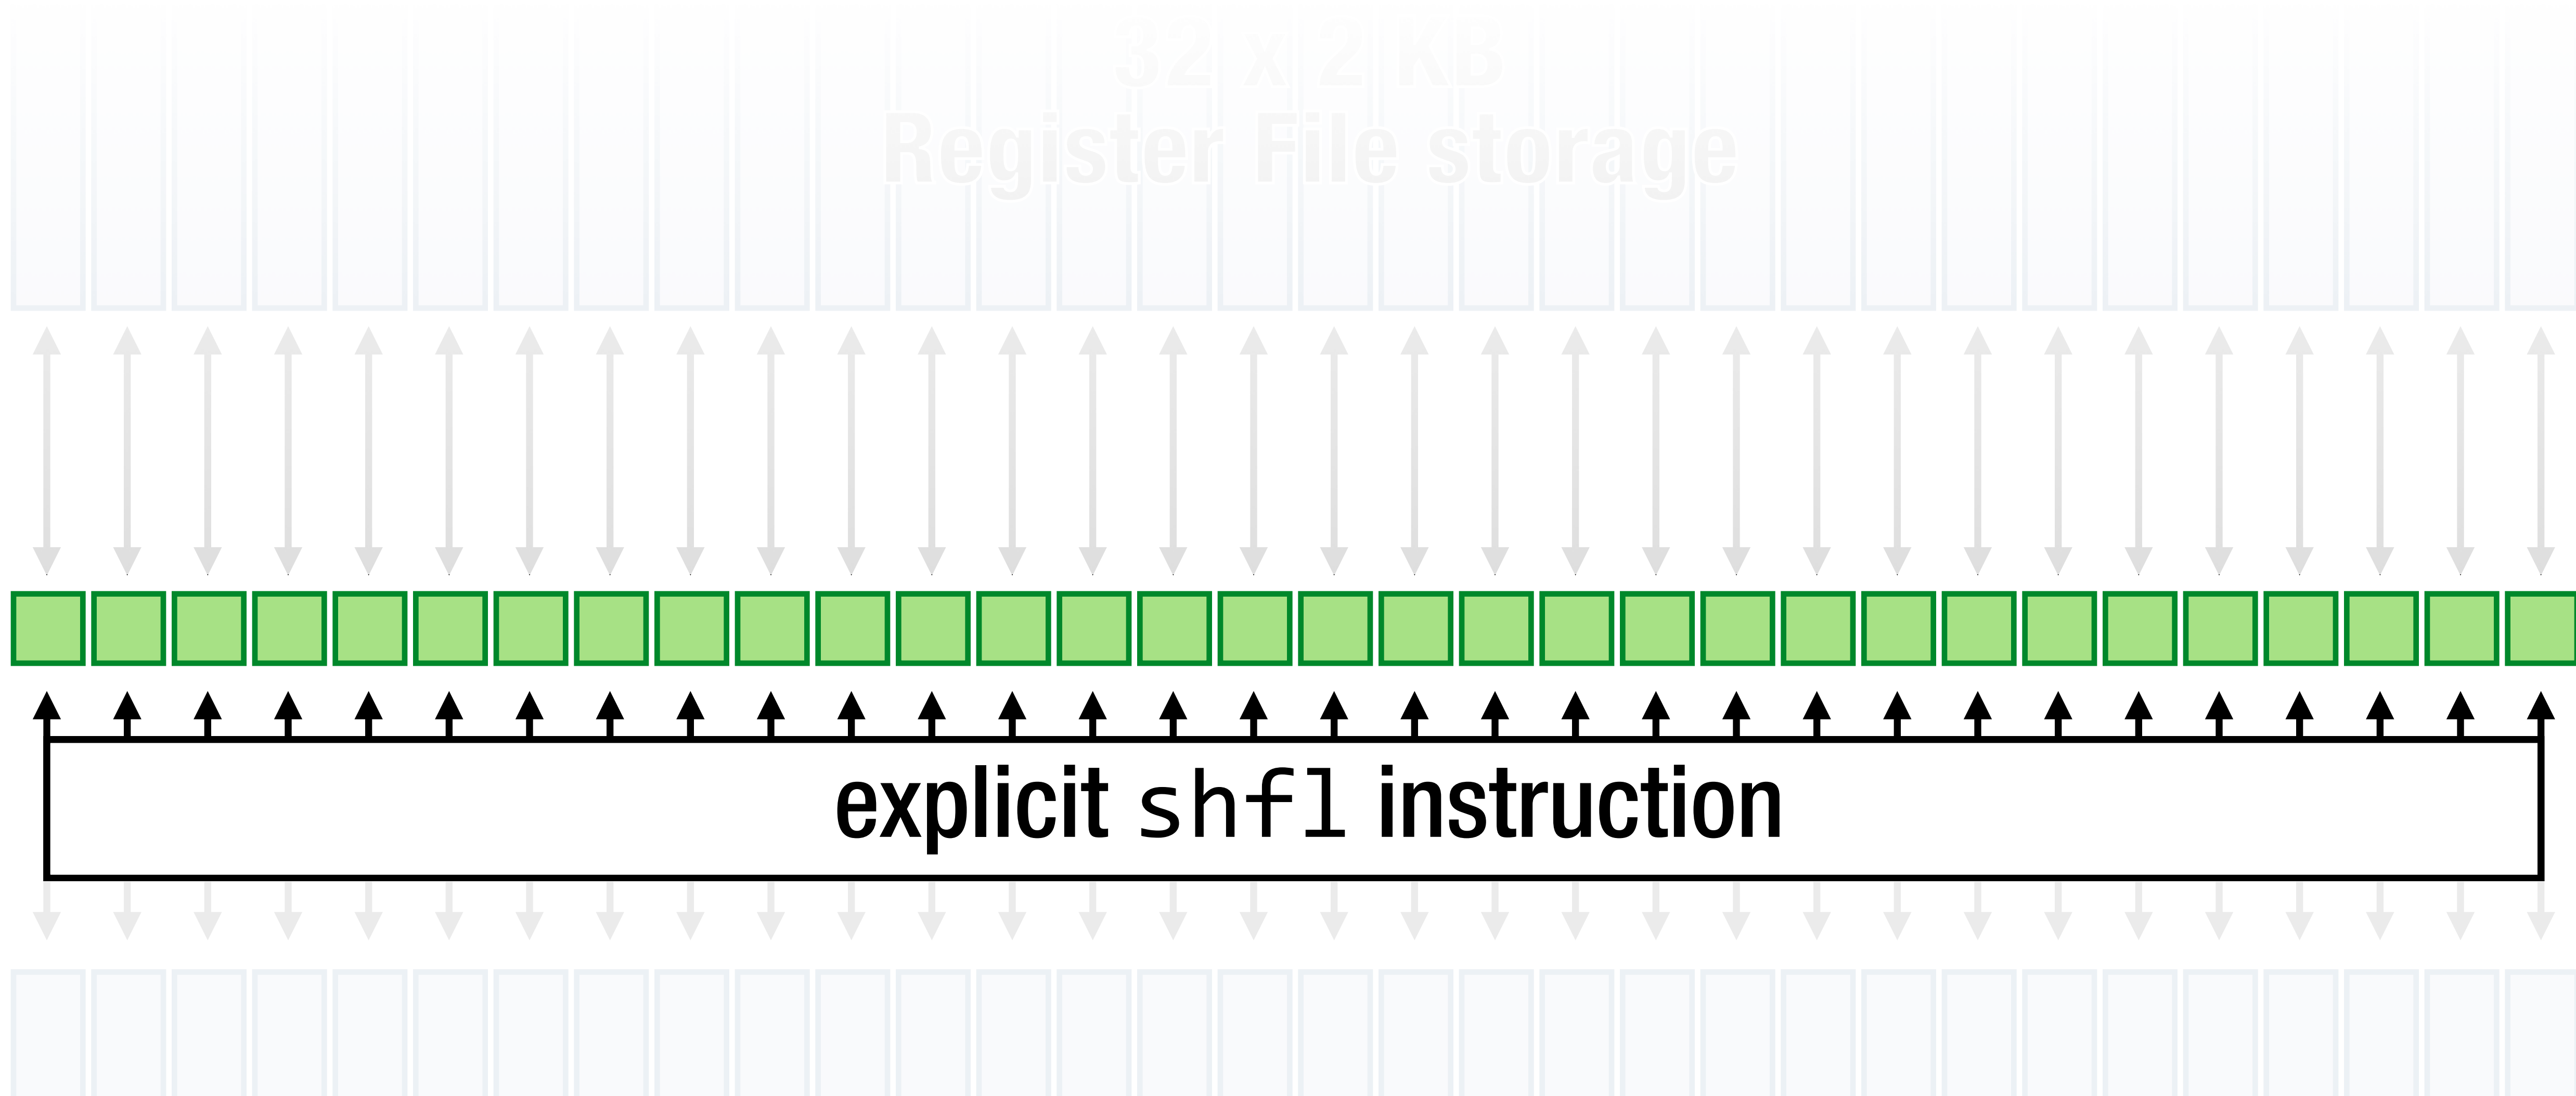


no  $r[3]$  (or  $r[i]$ ), only  $r3$

explicit `shf1` instruction



Data can be **shuffled between lanes**  
using the **L1 crossbar**



1 load /  
150 ops

1 load /  
20 ops

1 load /  
4 ops





3 loads + 1 store /  
1 op

360  
GB/sec

2.5  
TB/sec

13.4  
TB/sec

>100  
TB/sec



high **bandwidth**,  
limited **capacity**  
high clocks &  
wide interface

**aggregate** large  
transactions for  
DRAM  
**streaming**  
access  
large-scale **reuse**


shared per-SM,  
**not coherent**  
high **bandwidth**  
via **banking**  
cache or  
**scratchpad**

**banked** per-lane  
no dynamic  
**indexing**,  
either across or  
within lanes  
“infinite” BW

# **Overlapping compute & I/O**

# **Key facets of a processor:**

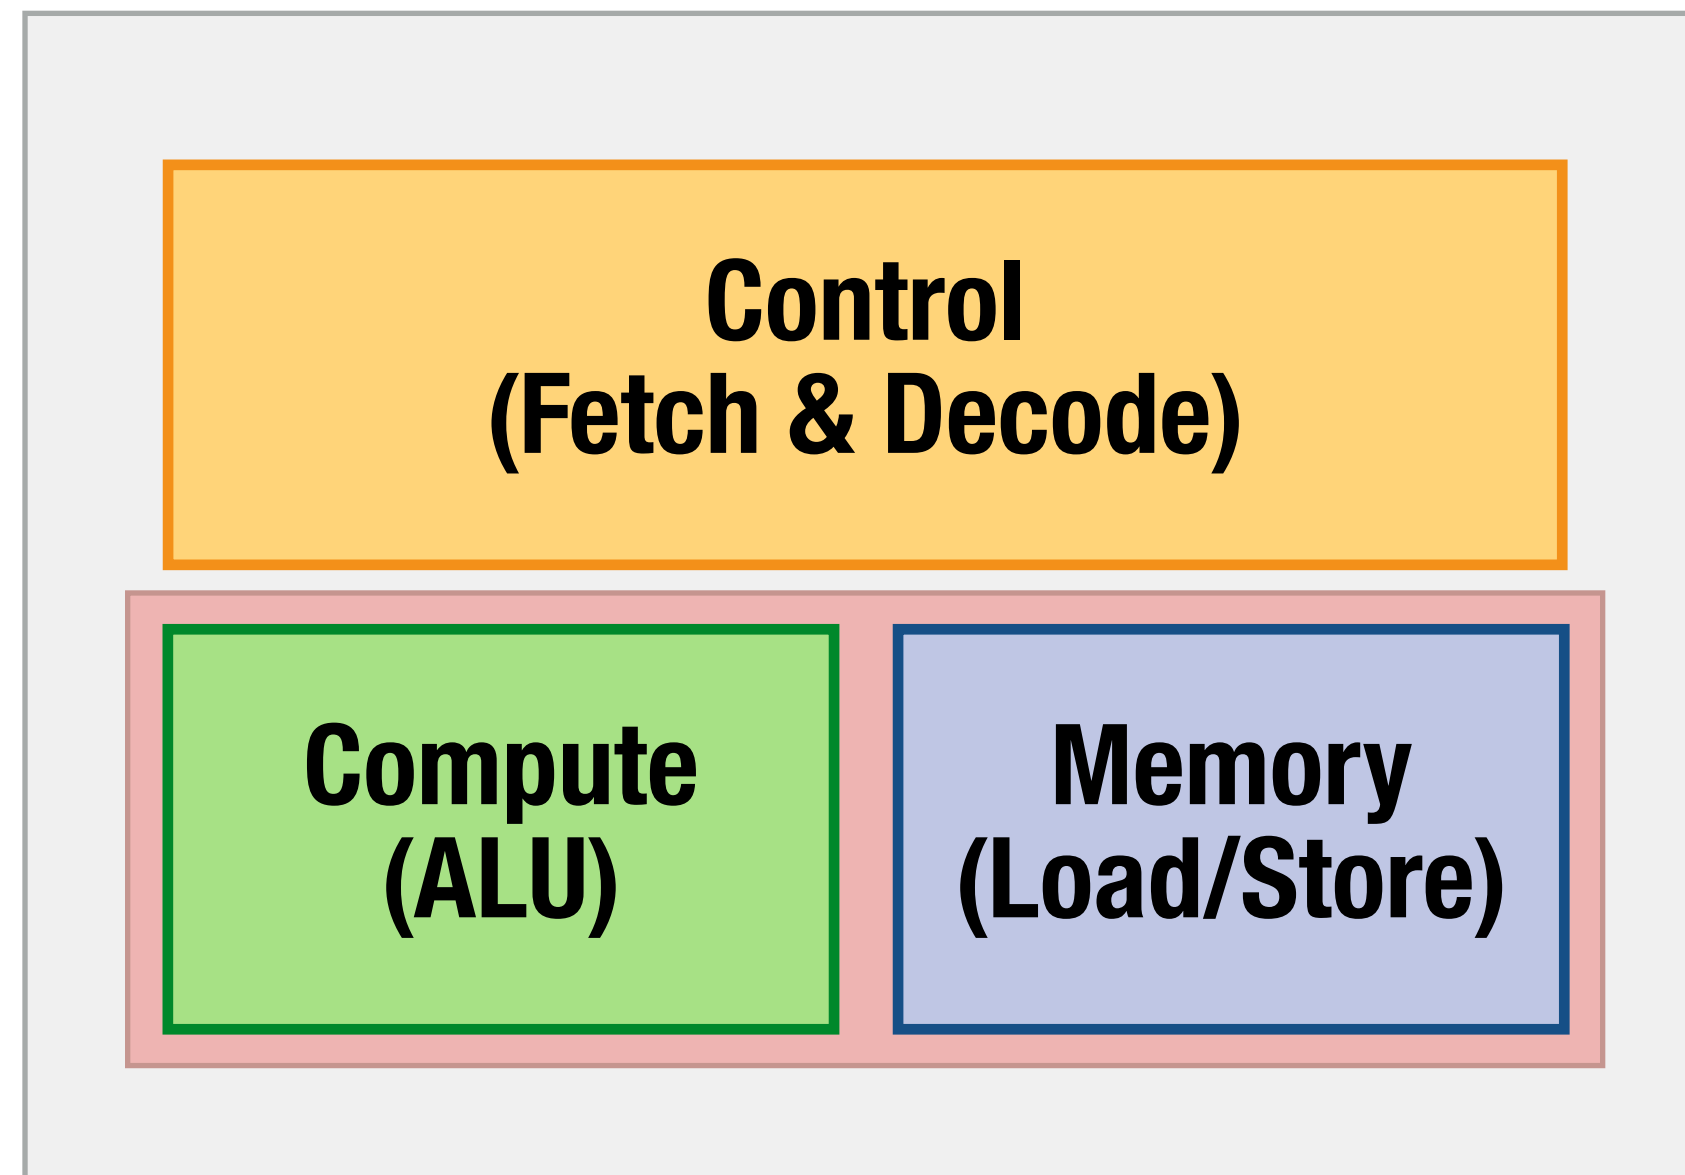
## **Control, Compute, Memory**



**Goal: fully utilize  
both resources**



# Use compute & memory **in parallel**

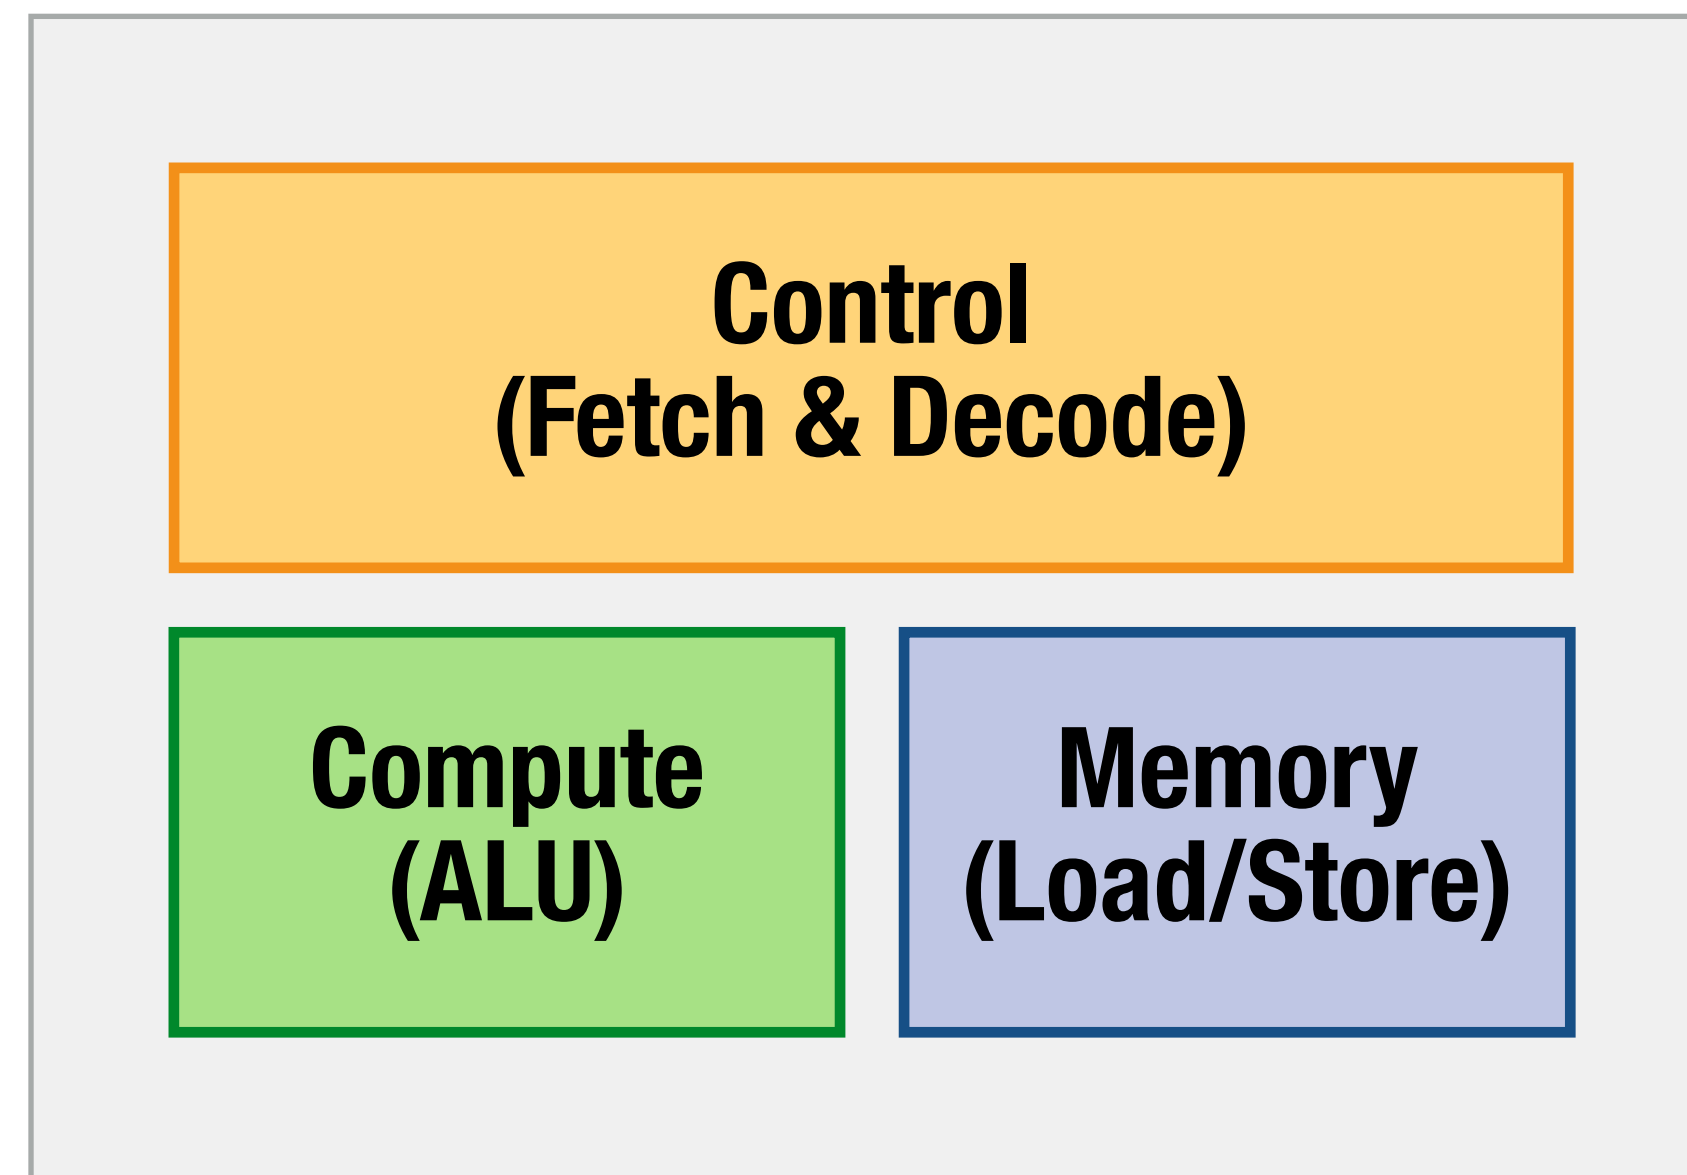


**Goal: fully utilize  
both resources**

**Approach 1: “CPU-style”**  
**wide issue, out-of-order**  
parallelism **within** instruction stream

**Approach 2: “GPU-style”**  
**multithreading**  
parallelism **across** instruction streams

# Both strategies work for common workloads



ld	r3, mem[r0+r2]
ld	r4, mem[r1+r2]
mul	ld r3, mem[r0+r2]
add	ld r4, mem[r1+r2]
add	mul ld r3, mem[r0+r2]
blt	add ld r4, mem[r1+r2]
st	add mul r3, r3, r4
...	blt add r5, r5, r3
	st addi r2, r2, 4
	blt r2, \$400, LOOP
	st addi r6, r5, 4


# What about workloads like **matrix multiplication**?

**lots of  
reuse**



core loop

**load a  
bunch**



**compute  
a bunch**


```
// load into scratchpad  
for i,j:  
    load A,B global → scratch
```

for each microtile:

```
// load into registers  
for i,k:  
    load A scratch → reg  
for k,j:  
    load B scratch → reg
```

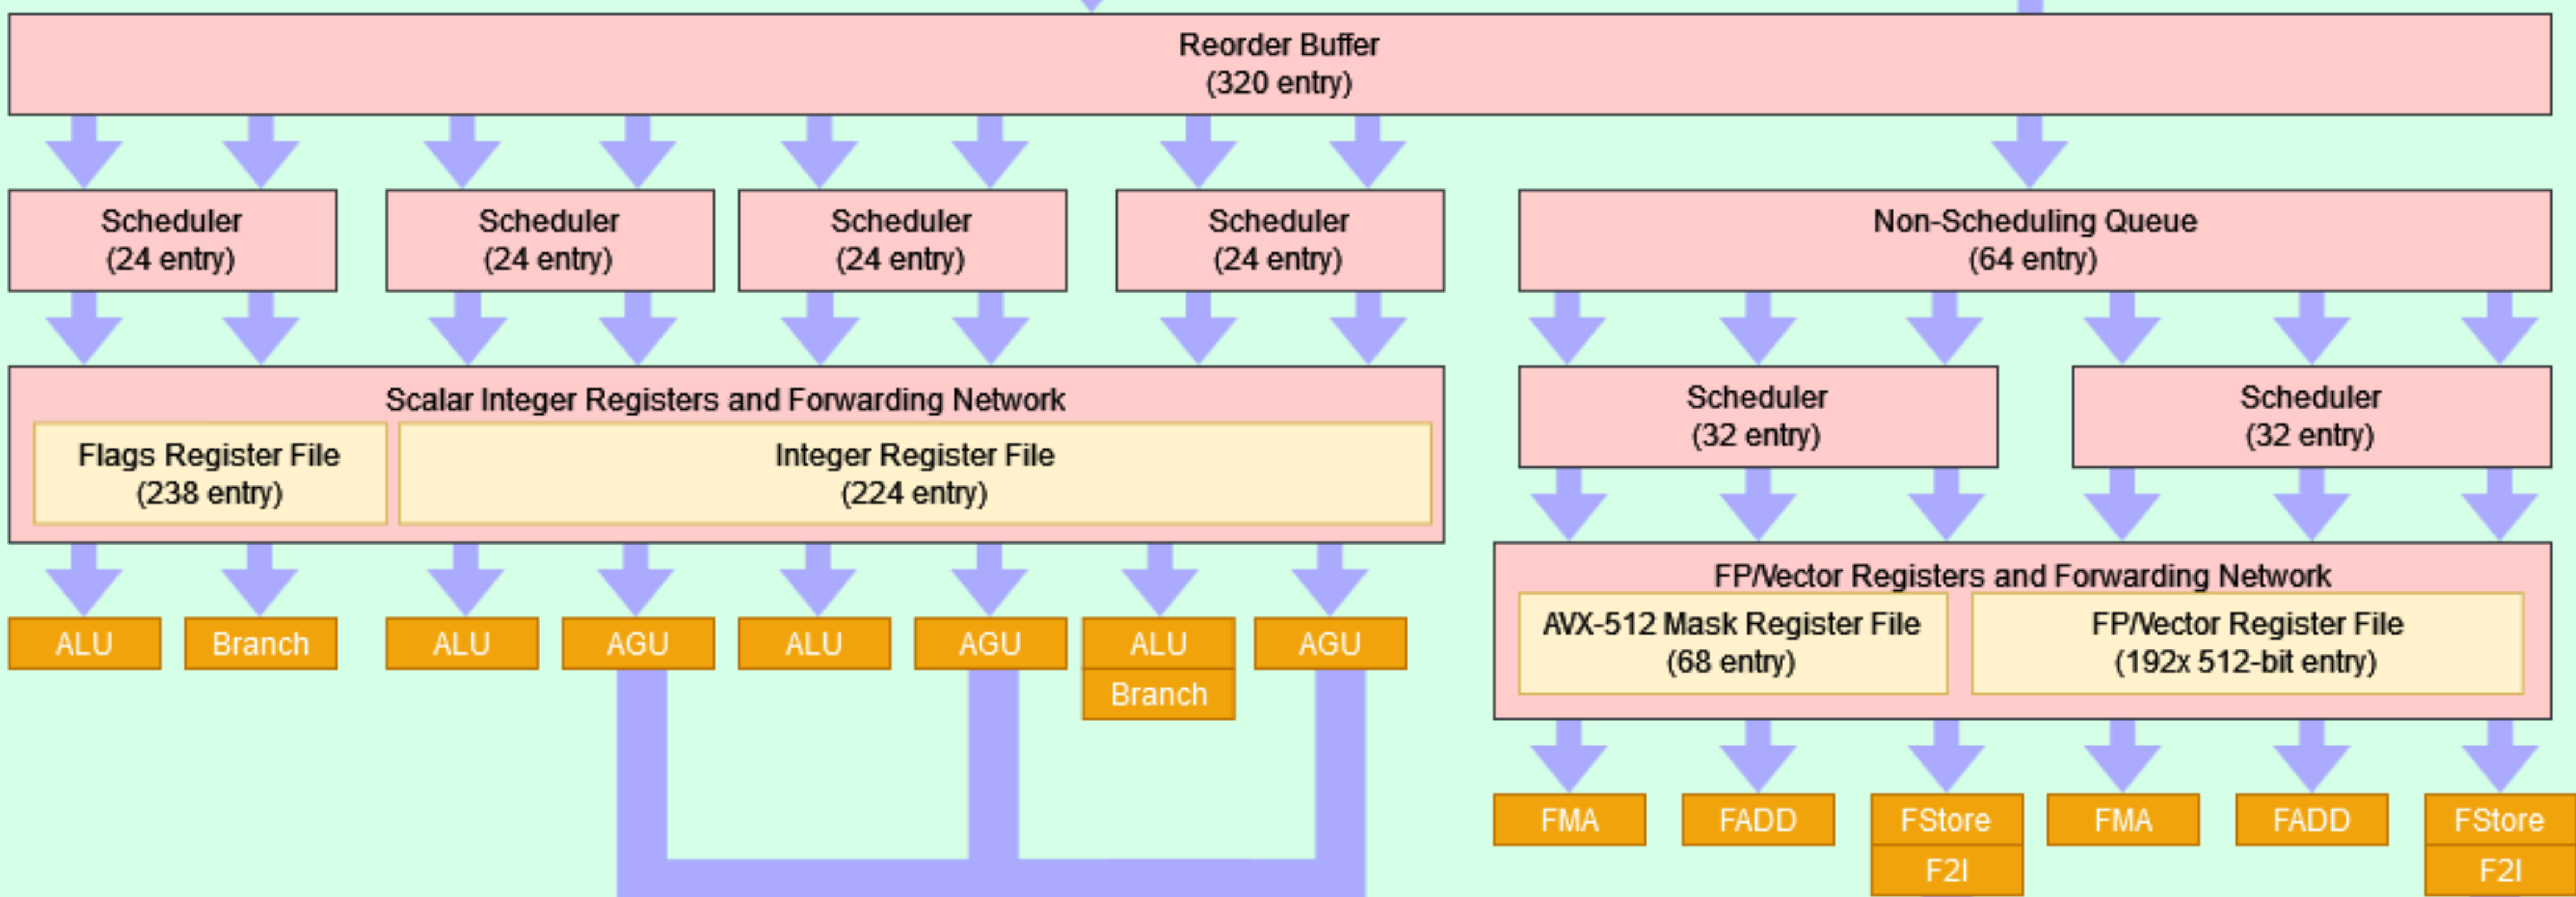
```
// compute!  
for i,j,k:  
    compute C += A*B
```

Front End



6 Micro-Ops

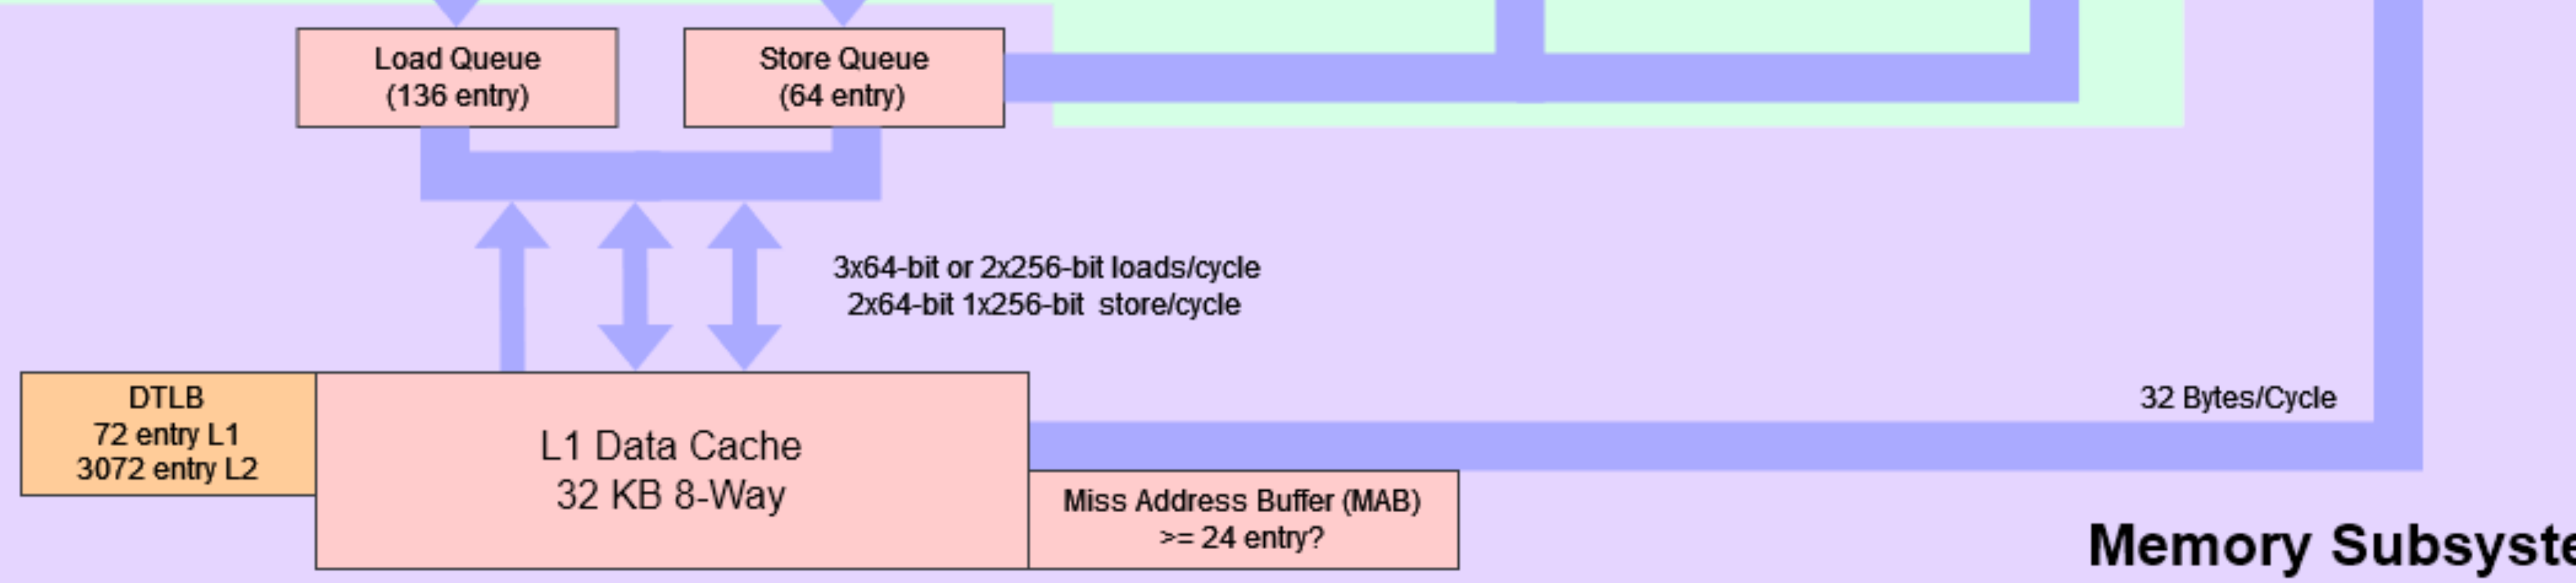
8 Micro-Op Retire



Execution Engine

Zen 4 core


diagram by Chips & Cheese



Memory Subsystem

# What about a **throughput processor** (GPU)?

- ▶ single issue
- ▶ in-order
- ▶ “RISC”



more manual &  
explicit management  
of **overlapping**

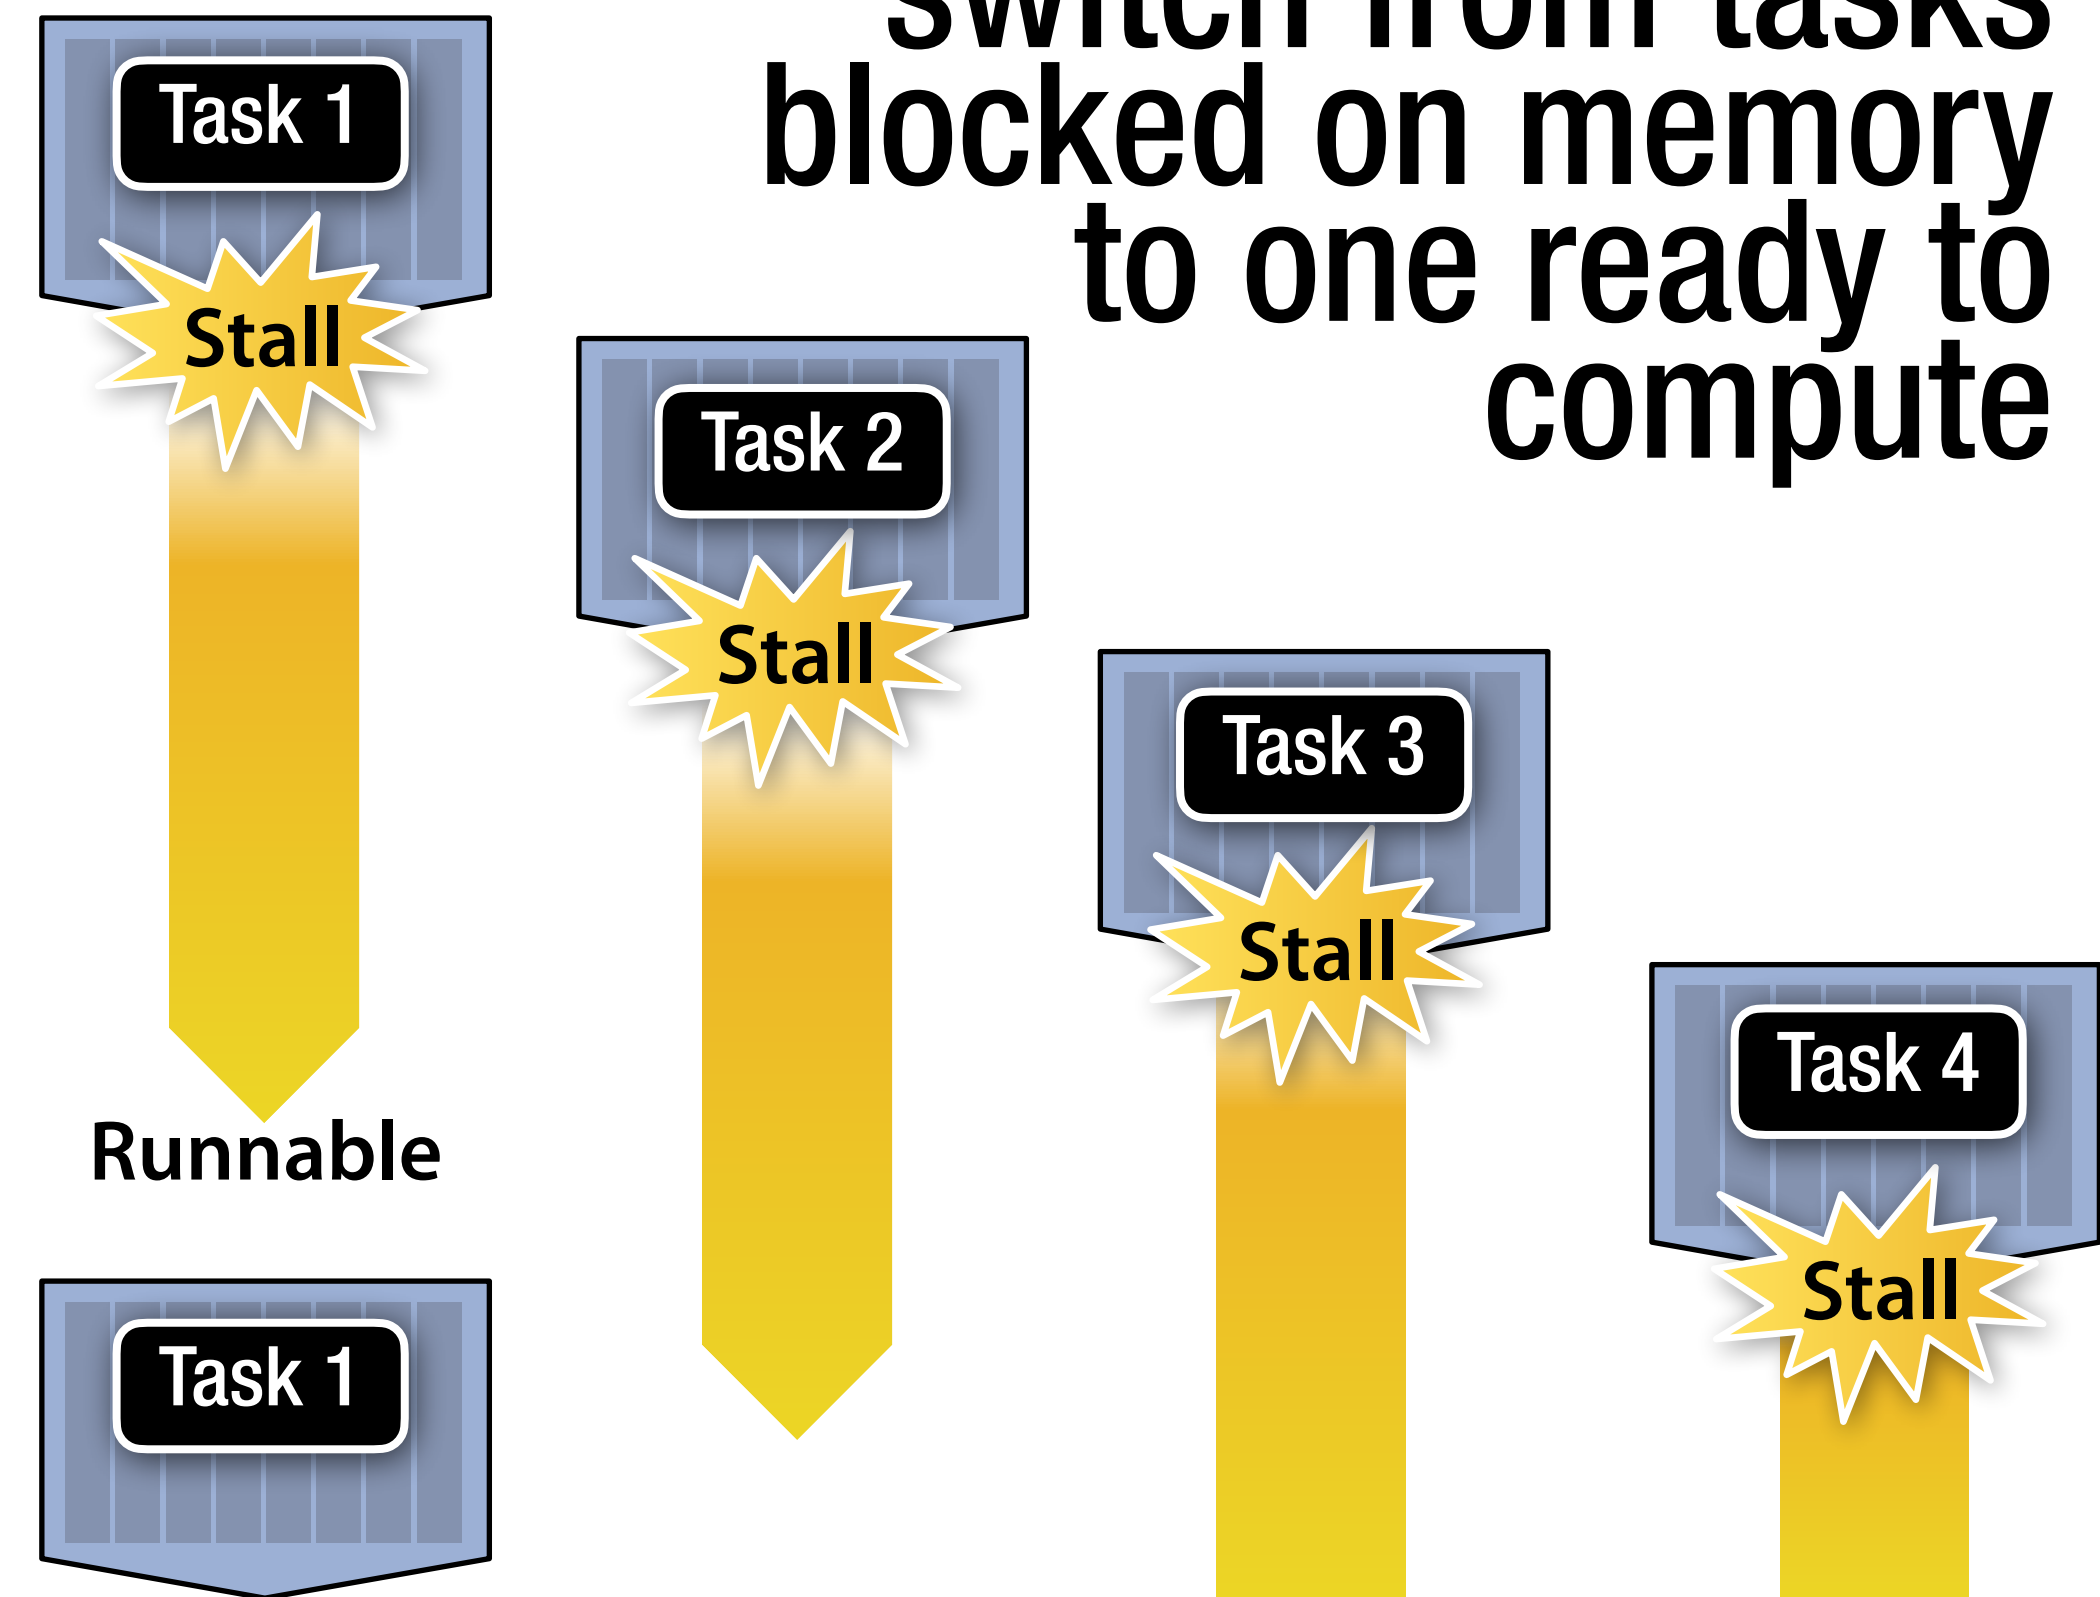
↳ recurs at **many  
levels**

# Problem 1: load / store instructions are asynchronous & long-latency

## Solution 1: ILP hoist loads early to avoid blocking

ld	→	ld
fma		ld
ld		ld
fma		fma
ld		fma
fma		fma
...		...

## Solution 2: multithreading switch from tasks blocked on memory to one ready to compute



# Problem 2: load / store instructions waste **issue slots**

**Solution: bulk load /  
store instructions**  
e.g., “vectorized” ld / st

ld.f32  
ld.f32  
ld.f32  
ld.f32  
fma  
fma  
fma  
fma  
...

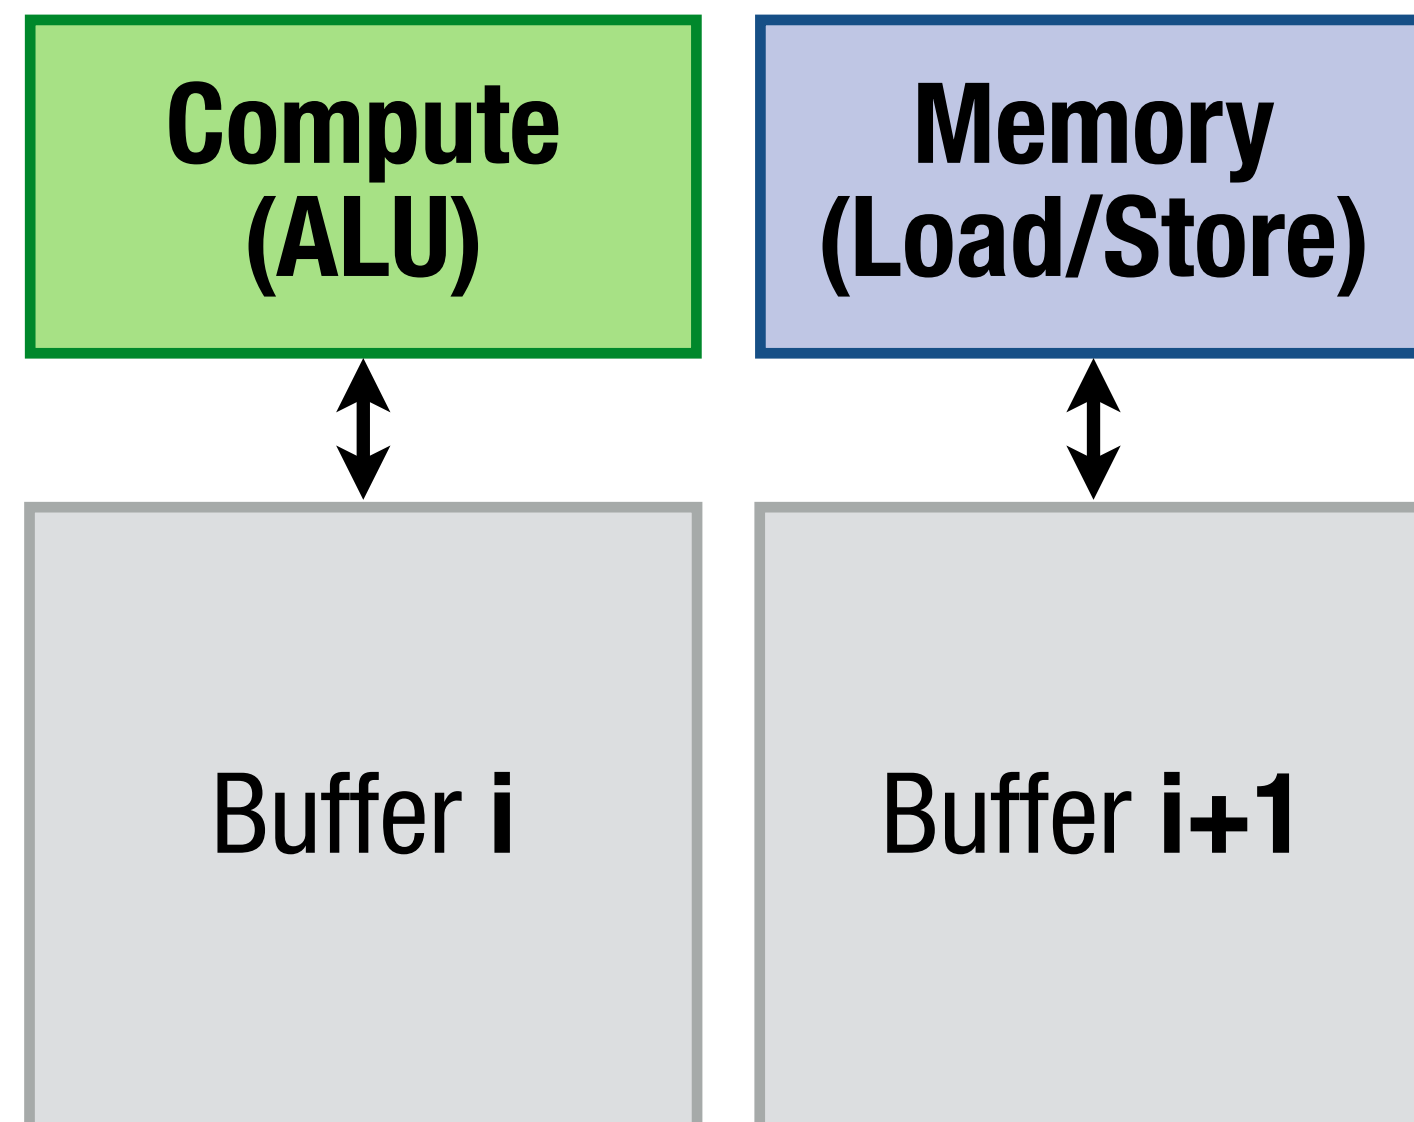


ld.v4.f32  
fma  
fma  
fma  
fma  
...



# Problem 3: overlapping compute with loading to the **scratchpad**

**Solution:**  
asynchronous fetch &  
**double-buffering**



**Implementation approach:**  
warp specialization

```
if threadIdx.y < 4:  
    // load into scratchpad  
    for i,j:  
        load next A,B → scratch  
else:  
    // compute!  
    for i,j,k:  
        compute C += A*B  
sync & swap buffers...
```